

Yesterday, my program worked.  
Today, it does not. Why?

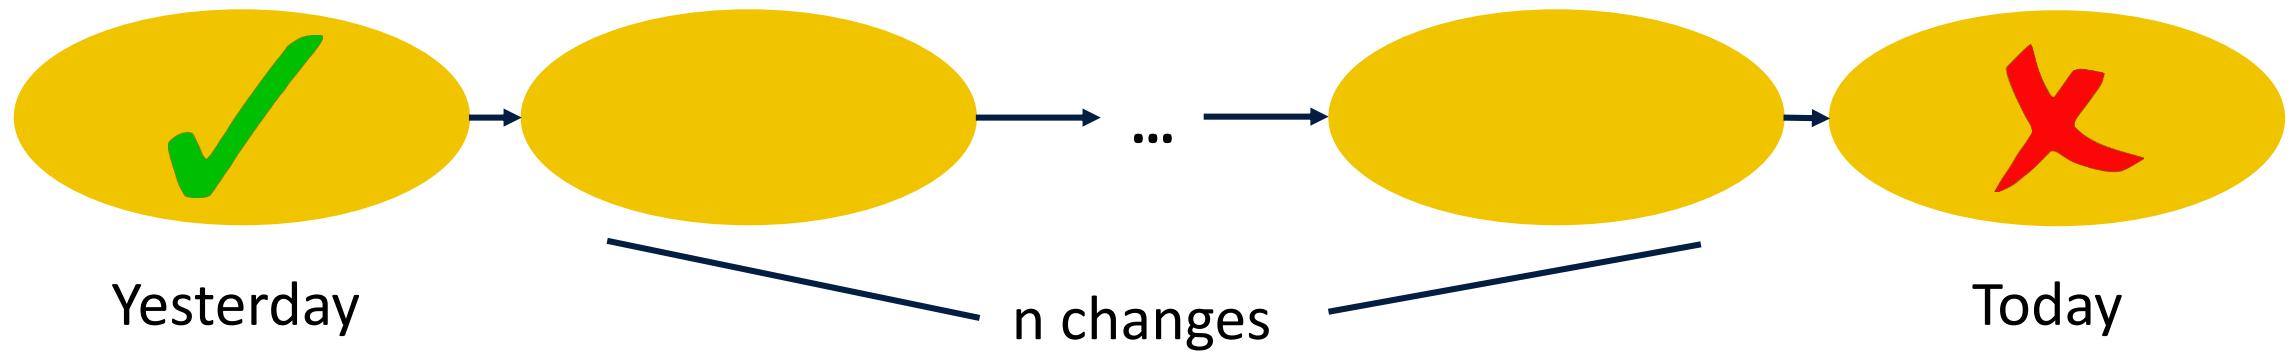
(Delta Debugging)



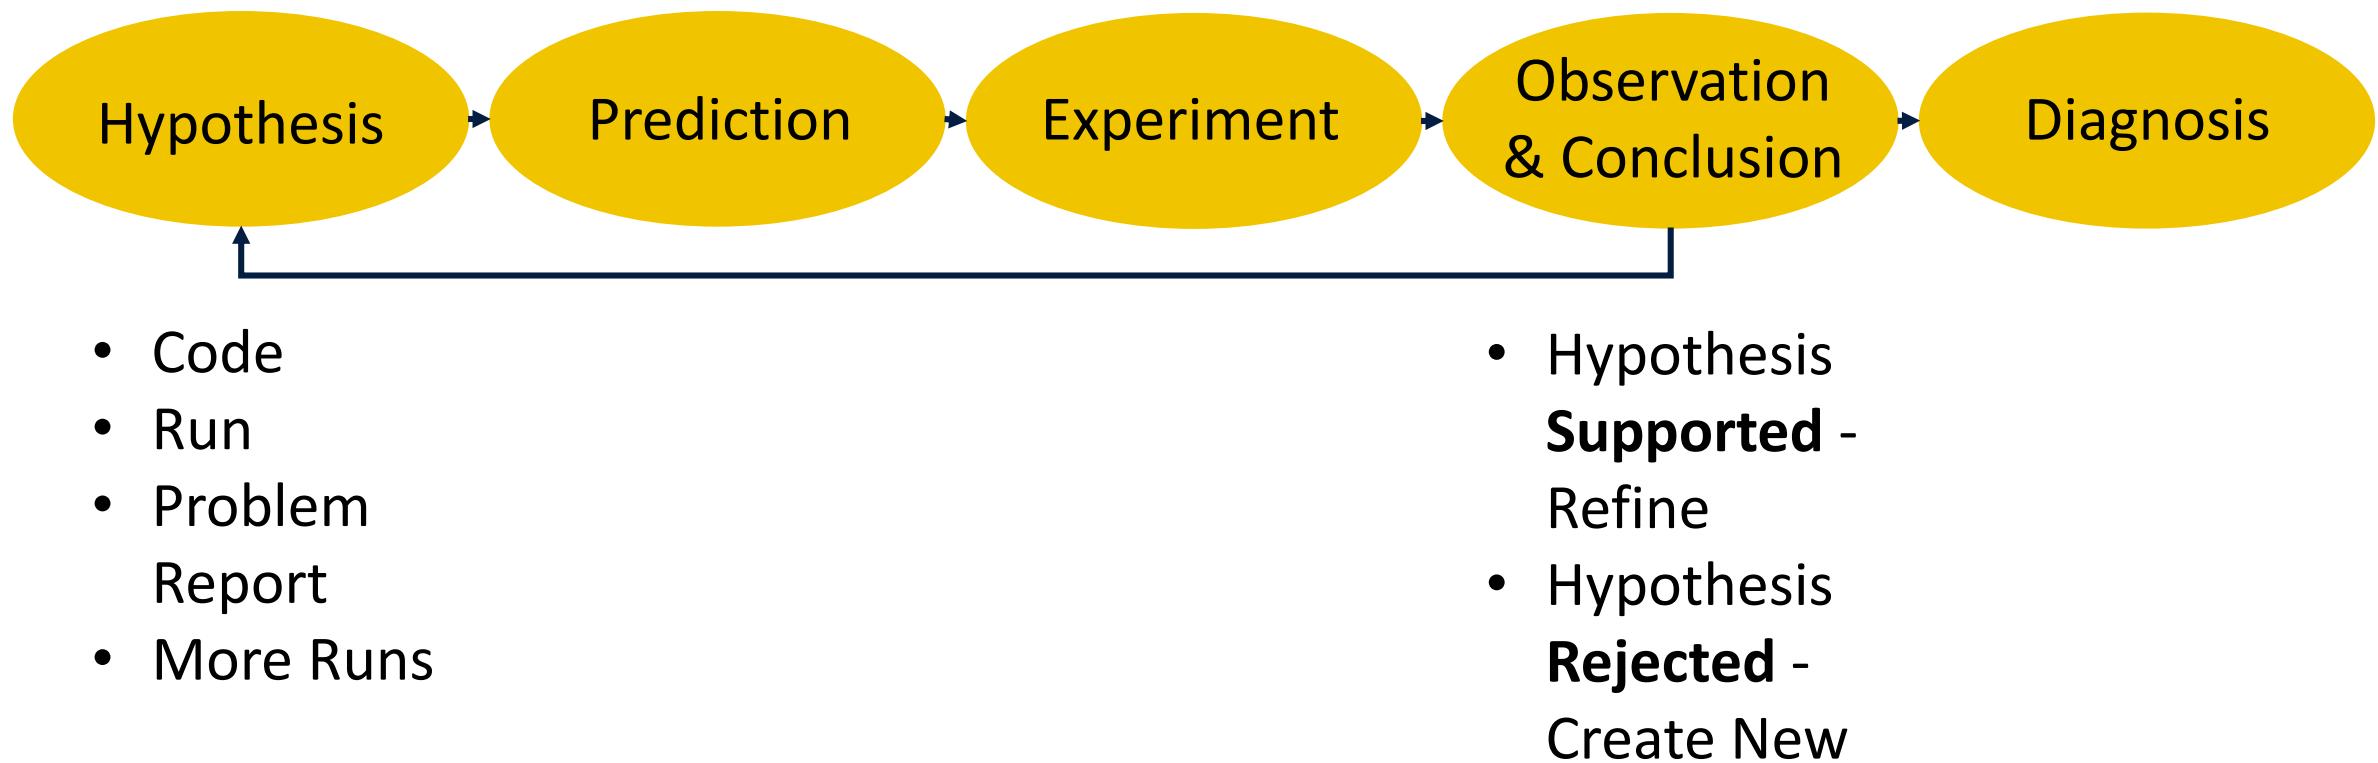
Andreas Zeller  
Presented by Amani Alkayyali  
October 29, 2018  
EECS 582

# Delta Debugging

- Differences determine cause for program behavior
- Slower than conventional debugging
- Automatic debugging

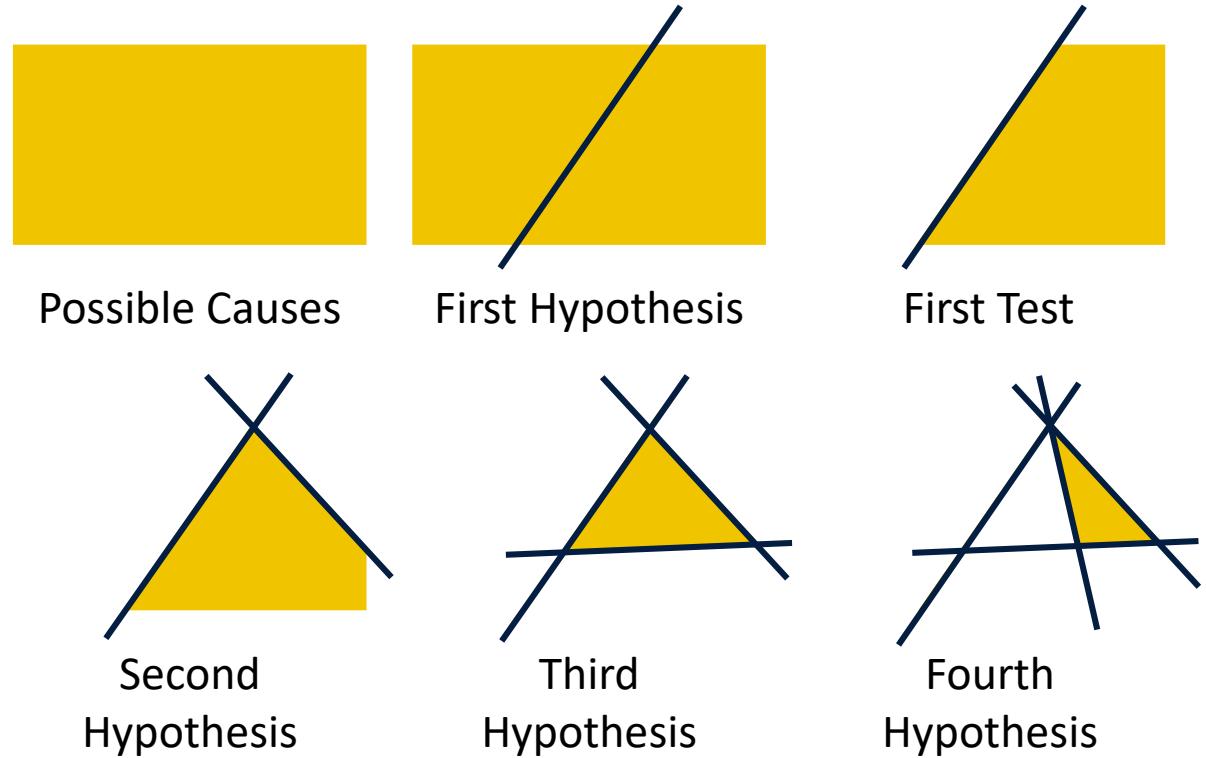


# Scientific Method



# Delta Debugging

- Technique, not a tool
- Similar to unit testing
- Recursive
- Divide and conquer algorithm



# Previous Techniques

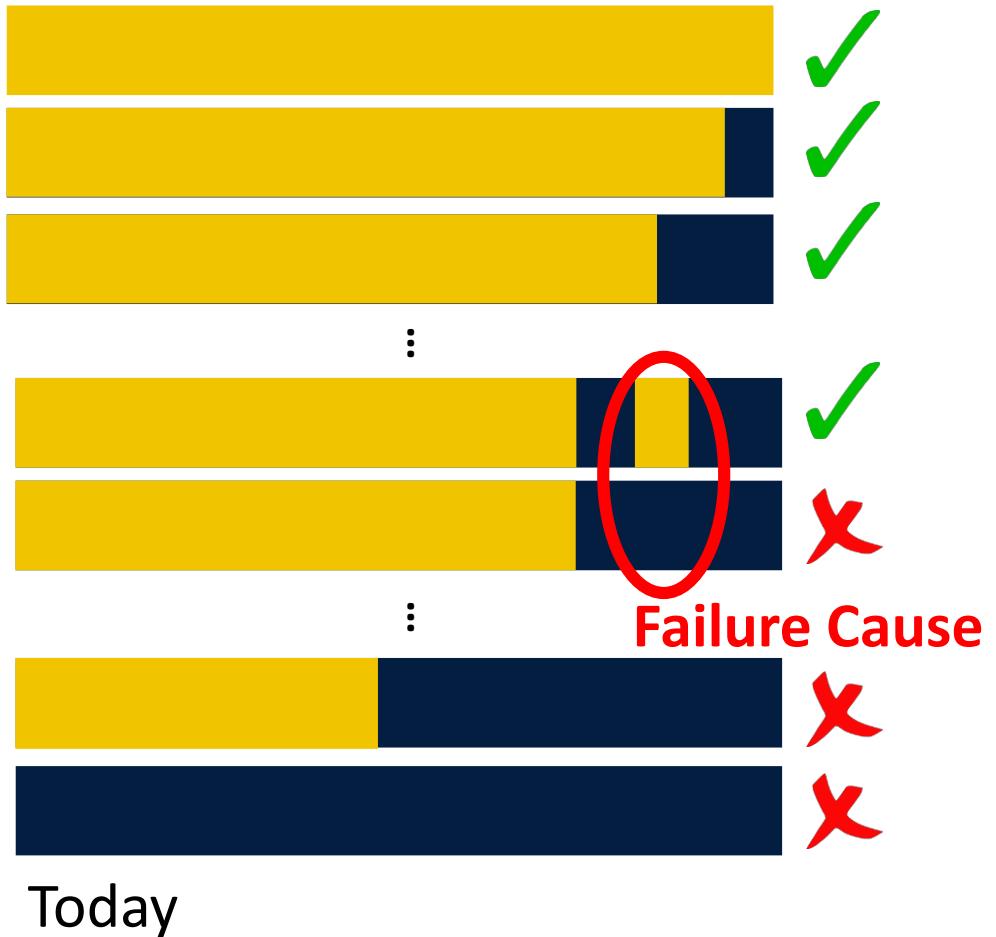
- Linear search – not sufficient
  - Interference
  - Inconsistency
  - Granularity
- Debugging vs. isolating



# Goal

- Minimal failure-inducing set
  - Least required subset to cause failure

Yesterday



# Binary Search

- Three outcomes:
  - Pass
  - Fail
  - Unresolved



# Configuration Properties

Certain properties allow for more efficiency

- **Monotony:** once failed, cannot un-fail
- **Unambiguity:** failure caused by one subset
- **Consistency:** repeatable
- Four scenarios



# Scenario 1:

- Hold all properties
- Possible Results:
  - Found in C1
  - Found in C2
  - Interference – both pass

Assume  $I = \{1, 2, 3, 4, 5, 6, 7, 8\}$

Configuration	Results
1 2 3 4 5 6 <b>7</b> 8	<b>FAIL</b>
1 2 3 4	PASS
5 6 <b>7</b> 8	<b>FAIL</b>
5 6	PASS
7 8	<b>FAIL</b>
7	<b>FAIL</b>



# Scenario 1:

- Hold all properties
- Possible Results:
  - Found in C1
  - Found in C2
  - Interference – both pass

Assume  $I = \{1, 2, 3, 4, 5, 6, 7, 8\}$

Configuration	Results
1 2 <b>3</b> 4 <b>5</b> 6 <b>7</b> 8	<b>FAIL</b>
1 2 <b>3</b> 4	PASS
1 2 <b>3</b> 4 <b>5</b> 6 <b>7</b> 8	PASS
1 2 <b>3</b> 4 <b>5</b> 6	PASS
1 2 <b>3</b> 4	PASS
1 2 <b>3</b> 4 <b>5</b> 6 <b>7</b>	<b>FAIL</b>
1 2 <b>3</b> 4 <b>5</b> <b>7</b>	<b>FAIL</b>
1 2 <b>3</b> 4 <b>5</b> <b>7</b>	<b>FAIL</b>
1 2 <b>5</b> <b>7</b>	PASS
<b>3</b> 4 <b>5</b> <b>7</b>	<b>FAIL</b>
<b>3</b> <b>5</b> <b>7</b>	<b>FAIL</b>



# Scenarios 2 & 3

## Scenario 2

- Configuration is ambiguous
- Find one, rerun, find another until complete

## Scenario 3

- Not monotone
- “undoing” changes that cause pass



# Scenario 4

- Inconsistency
- Causes:
  - Integration failure
  - Construction failure
  - Execution failure
- Possible Results:
  - Found
  - Interference
  - Preference
  - Try again



# Scenario 4

Preference

Configuration

1 2 3 4 5 6 7 8

1 2 3 4

5 6 7 8

1 2 5 6 7 8

Results

FAIL

?

PASS

PASS

Try Again

Configuration

1 2 3 4

5 6 7 8

1 2

3 4

Results

?

?

?

?

...



# Scenario 4

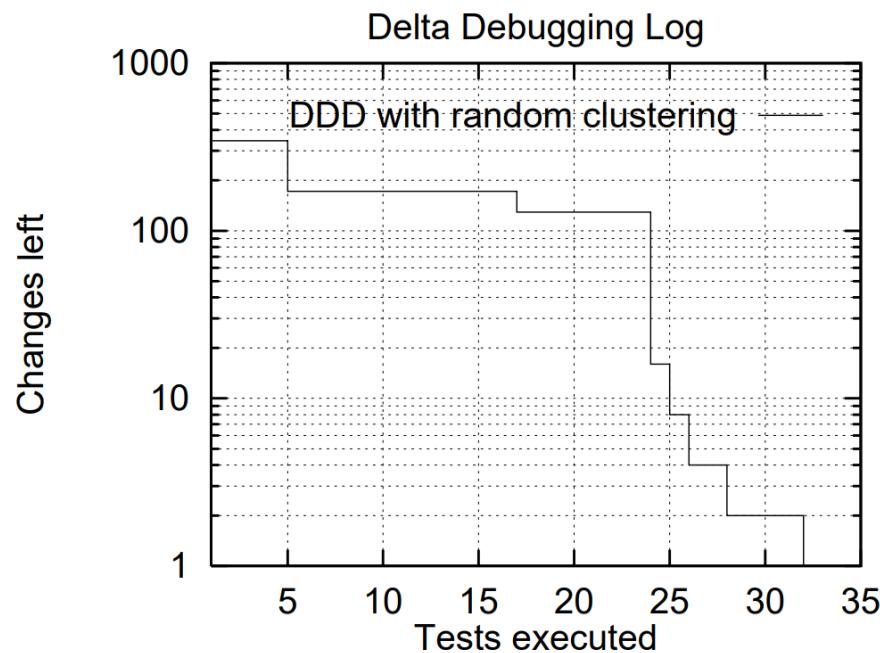
- Avoiding inconsistency
  - Grouping related changes
    - Location
    - Syntactic
  - Predicting test outcomes

Configuration	1	2	3	4	5	6	7	8	Results
									PASS
									Predict ?
	1	2	3	4	5	6			PASS
	1	2	3	4			7	8	Predict ?
	1	2	3	4	5	6	7	8	FAIL

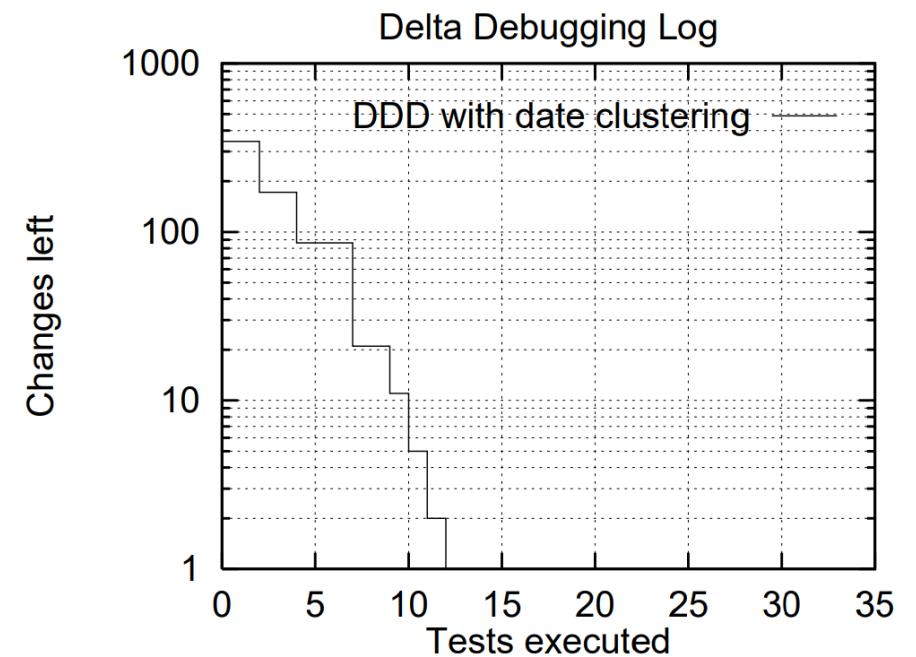


# Case Study 1

Random Clustering

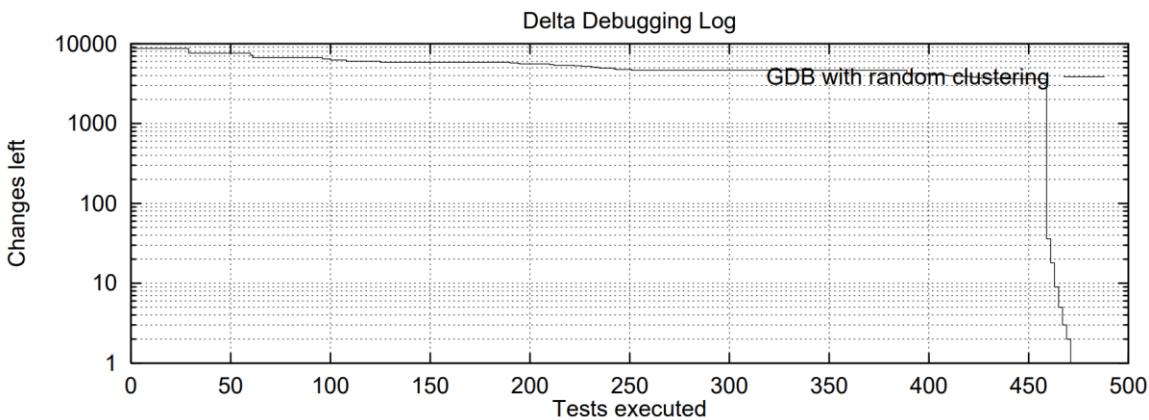


Date Clustering

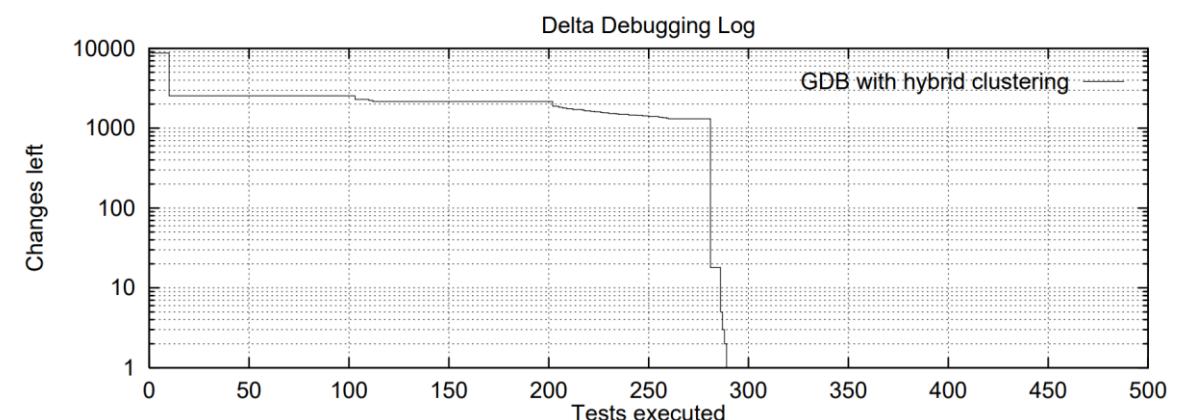


# Case Study 2

## Random Clustering



## Hybrid Clustering



# Using Magpie for request extraction and workload modeling

Paul Barham, Austin Donnelly, Rebecca Isaacs, Richard Mortier

Presented by; Won Park

# Motivation

- Ability to understand complex system behavior is useful
  - Performance analysis, debugging, etc.
- Still open questions / limitations in current work

# What is Magpie?

- Toolchain for automatically extracting system's workload under realistic operating conditions
- Provides ability to capture control path / resource demands of application requests
- Magpie operates both online / offline
  - Intention when designing was online

# Features of Magpie

- Collects events directly generated from kernel, middleware, application
- Ability to measure request's resource demands while discarding scheduling artefacts
  - OS multitasking / timesharing

# In short...

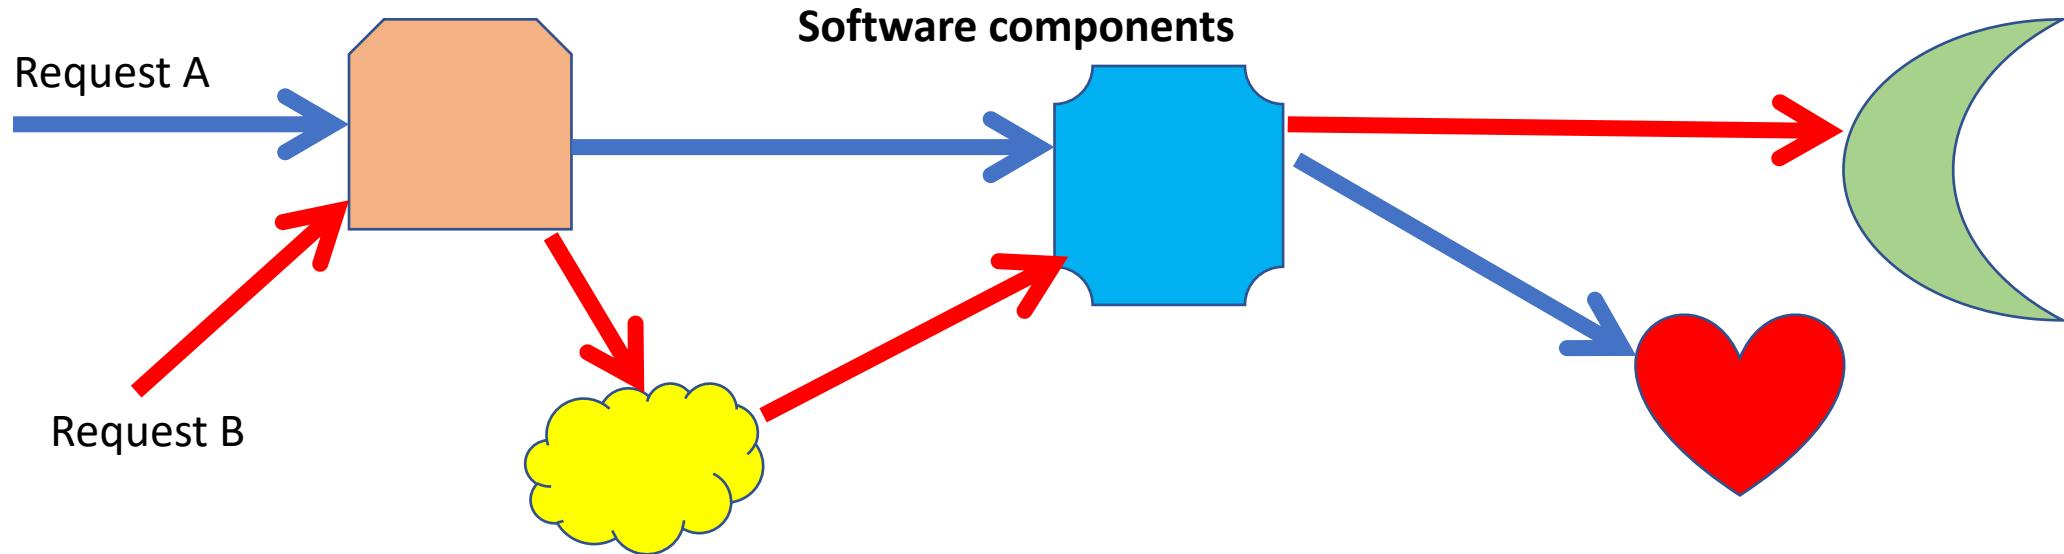
- Magpie provides picture of:
  - How request could have been serviced
  - How it actually was serviced
- Models are quicker + easier to understand
- Causes of problems often revealed by comparing Magpie trace to expected behavior

# In this talk..

- Background / Instrumentation
- Steps in the Toolchain
- Validation / Overhead Analysis
- Evaluation

# Background

- Workload of a system: different requests that take different paths



- Functional progress of request + resource consumption at every stage recorded

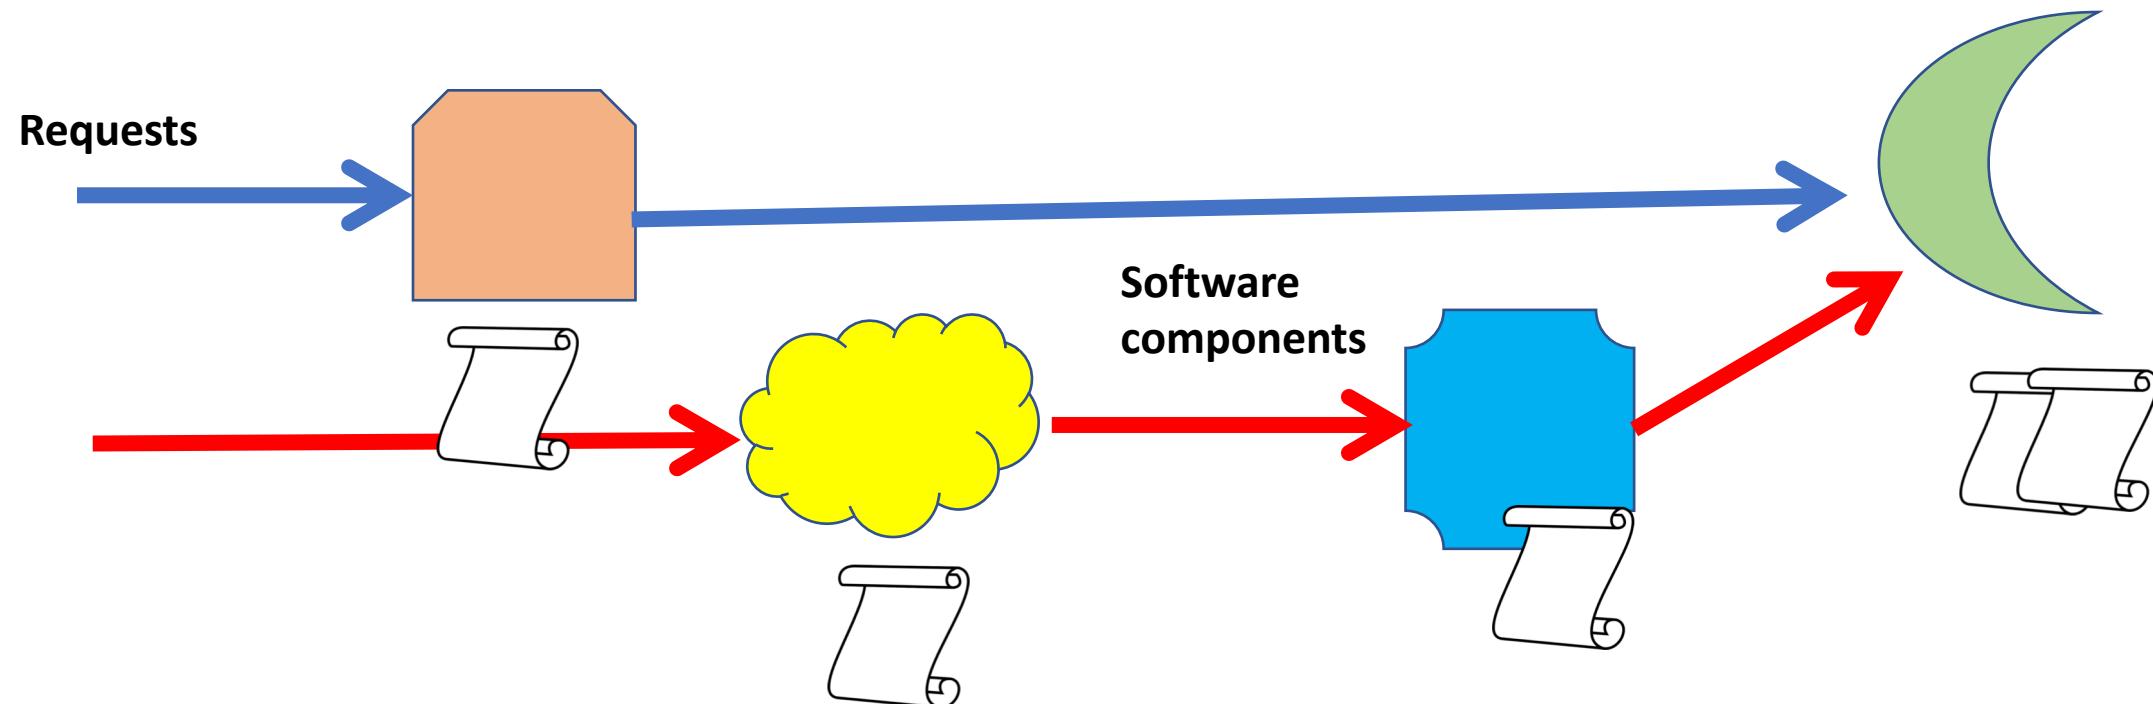
# Instrumentation Overview

- Requirement: accurate accounting of resource usage
- Logical consequence: high precision timestamps needed
- Magpie's solution: Processor cycle counter

# Event Tracing for Windows (ETW)

- Low overhead event logging infrastructure
- ETW event:
  - Timestamp
  - Provider / Event ID
  - Zero or more typed attributes

# Event Trace Production



Ordered but likely  
interleaved

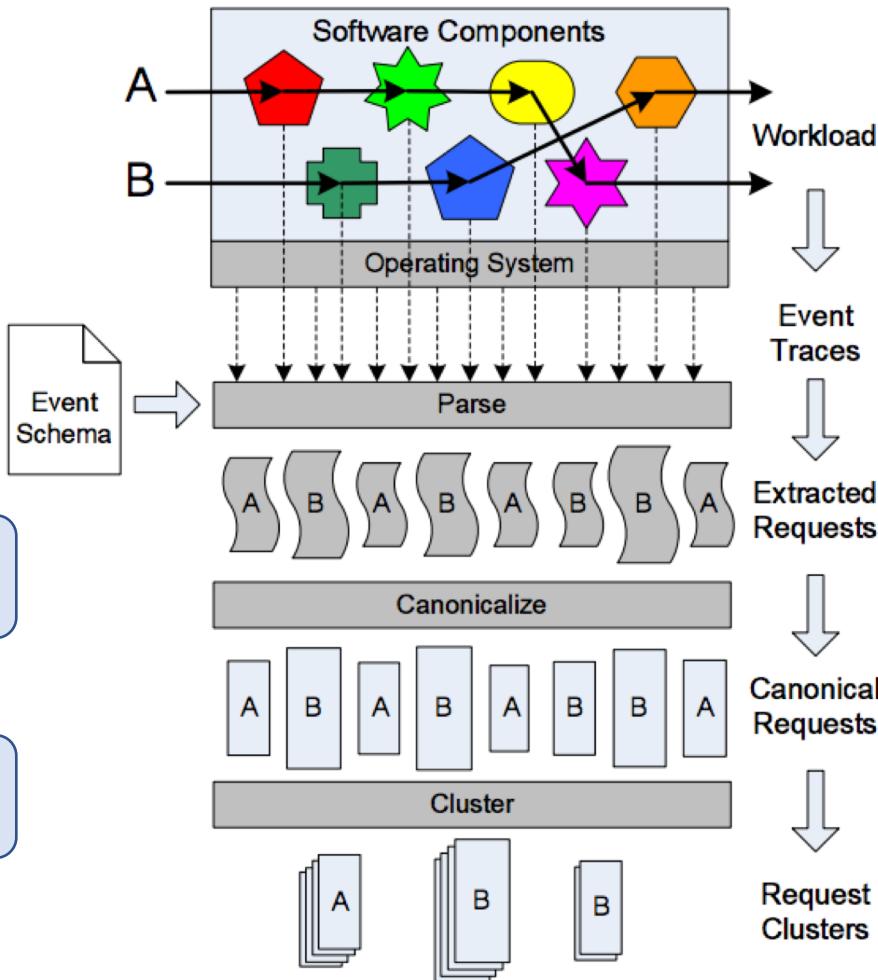
# In this talk..

- Background / Instrumentation
- Steps in the Toolchain
- Validation / Overhead Analysis
- Evaluation

# Toolchain Overview

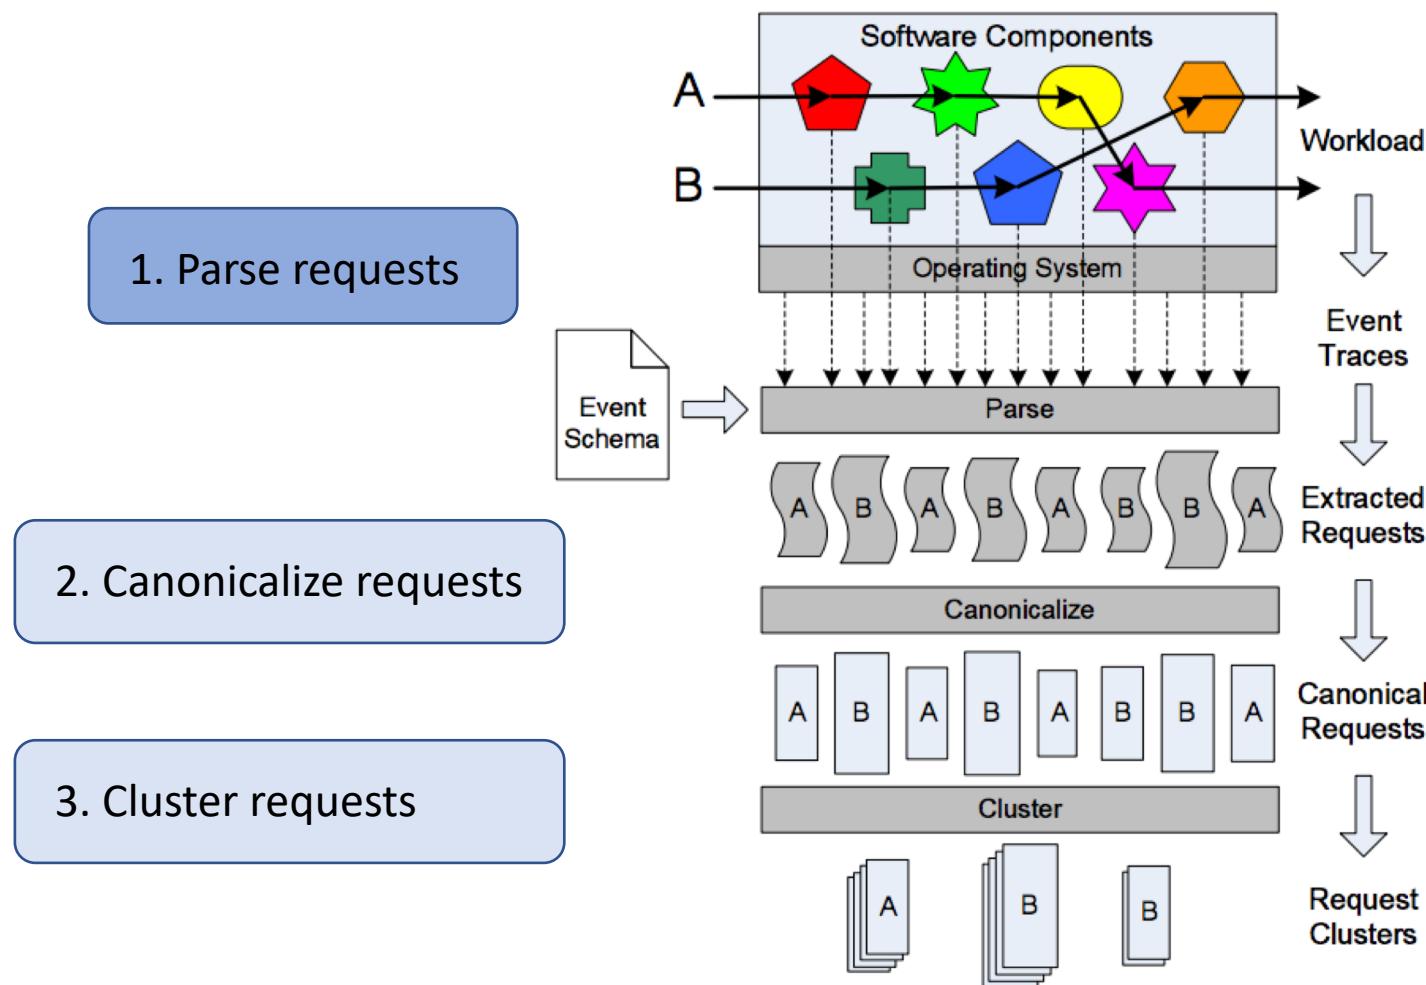
TODO: during practice talk, think about where to insert graphics talkinga bout the software componetns

1. Parse requests



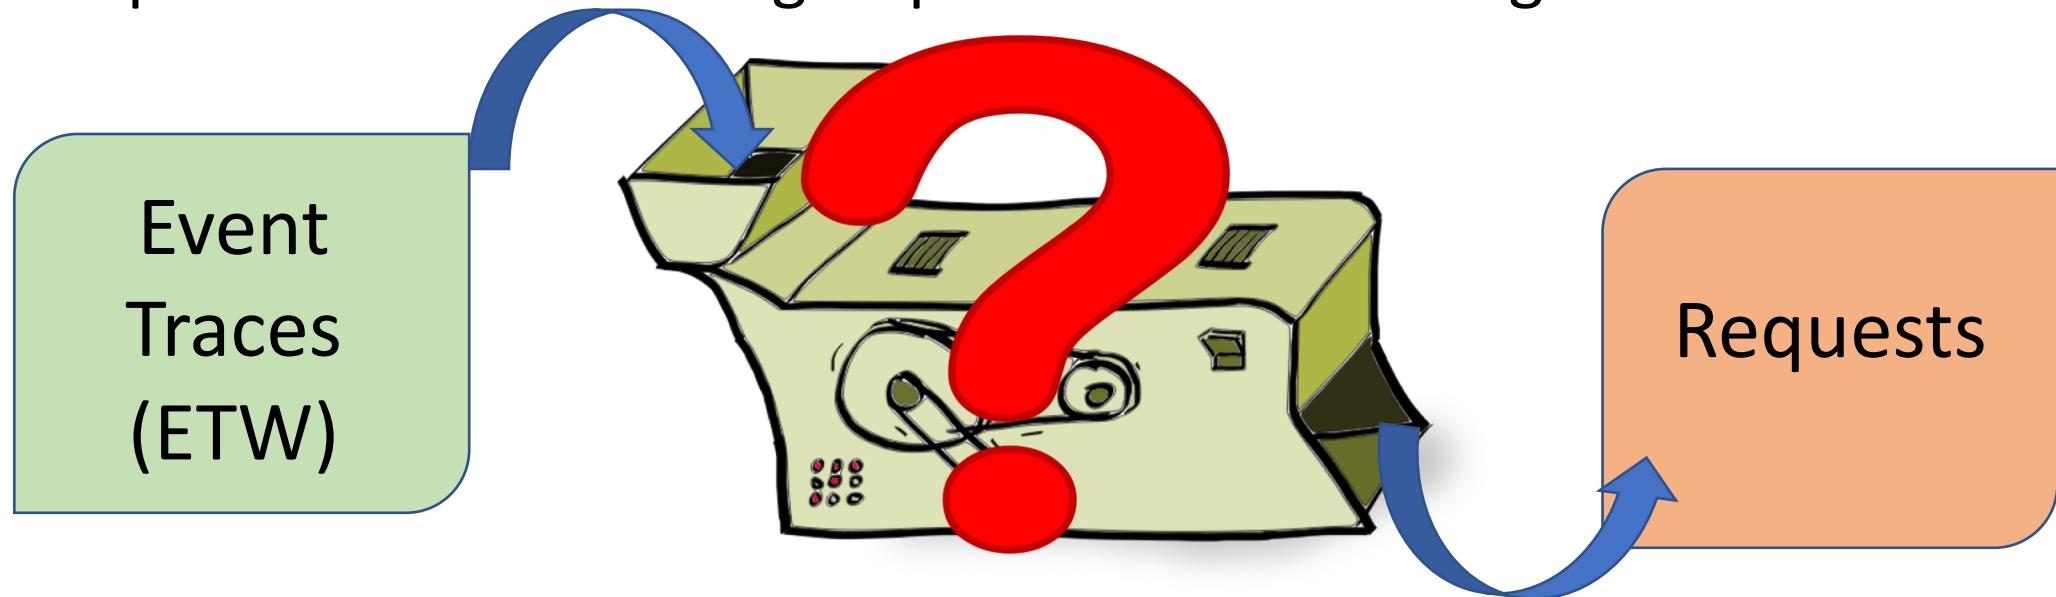
2. Canonicalize requests

3. Cluster requests



# Request Parser

- Responsible for extracting requests from event logs



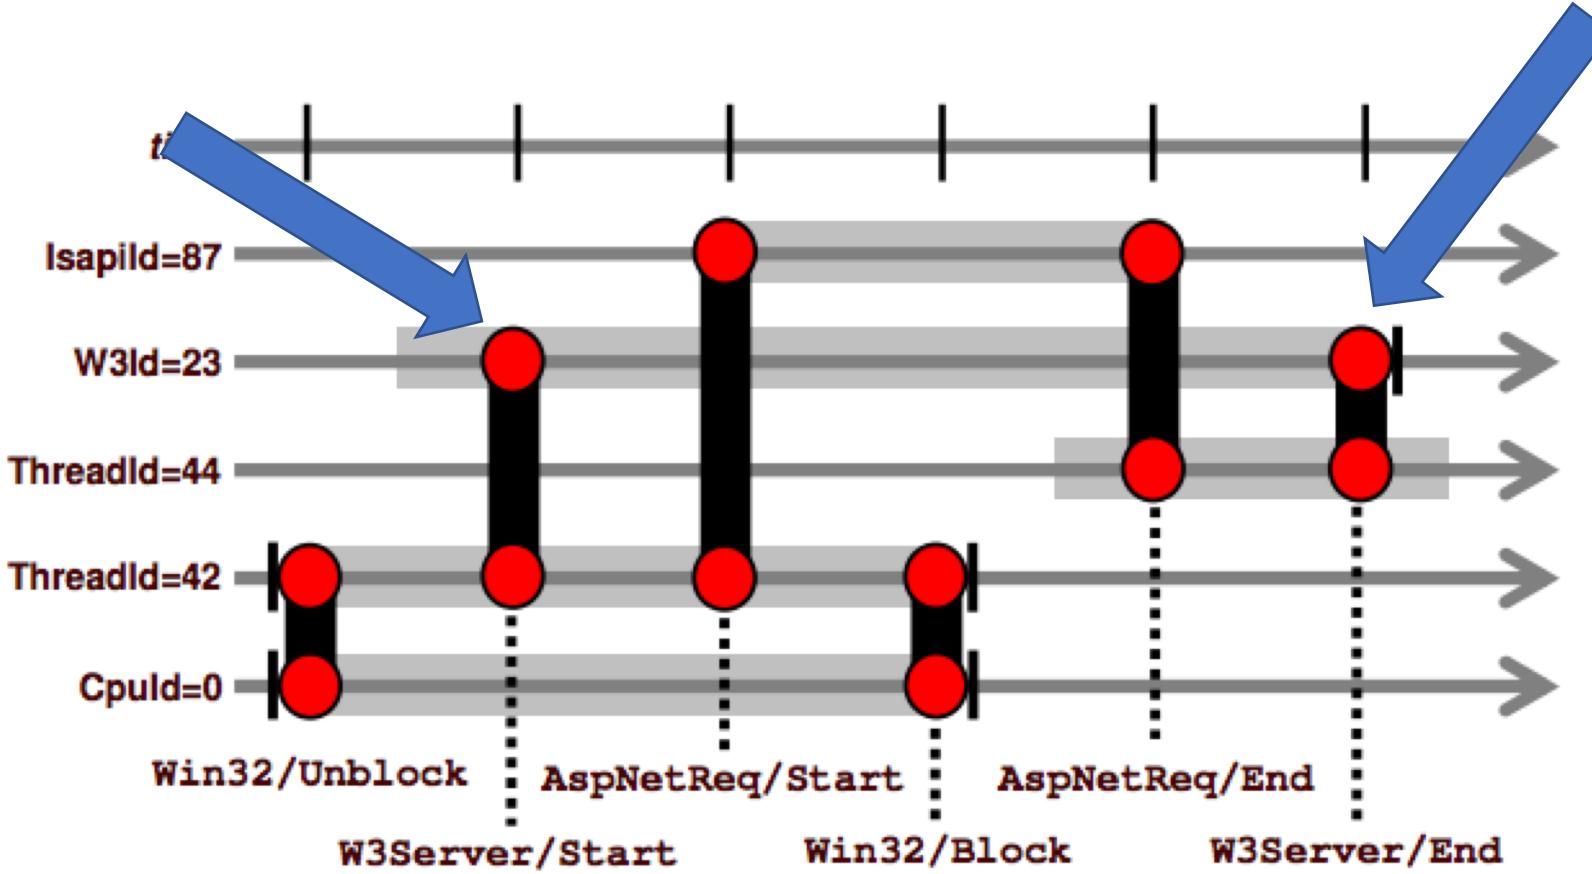
- Speculatively builds up sets of related events

# Event Schema

- Specifies which attribute connects events to other events
- Application specific
  - Request parser itself makes no built-in assumptions

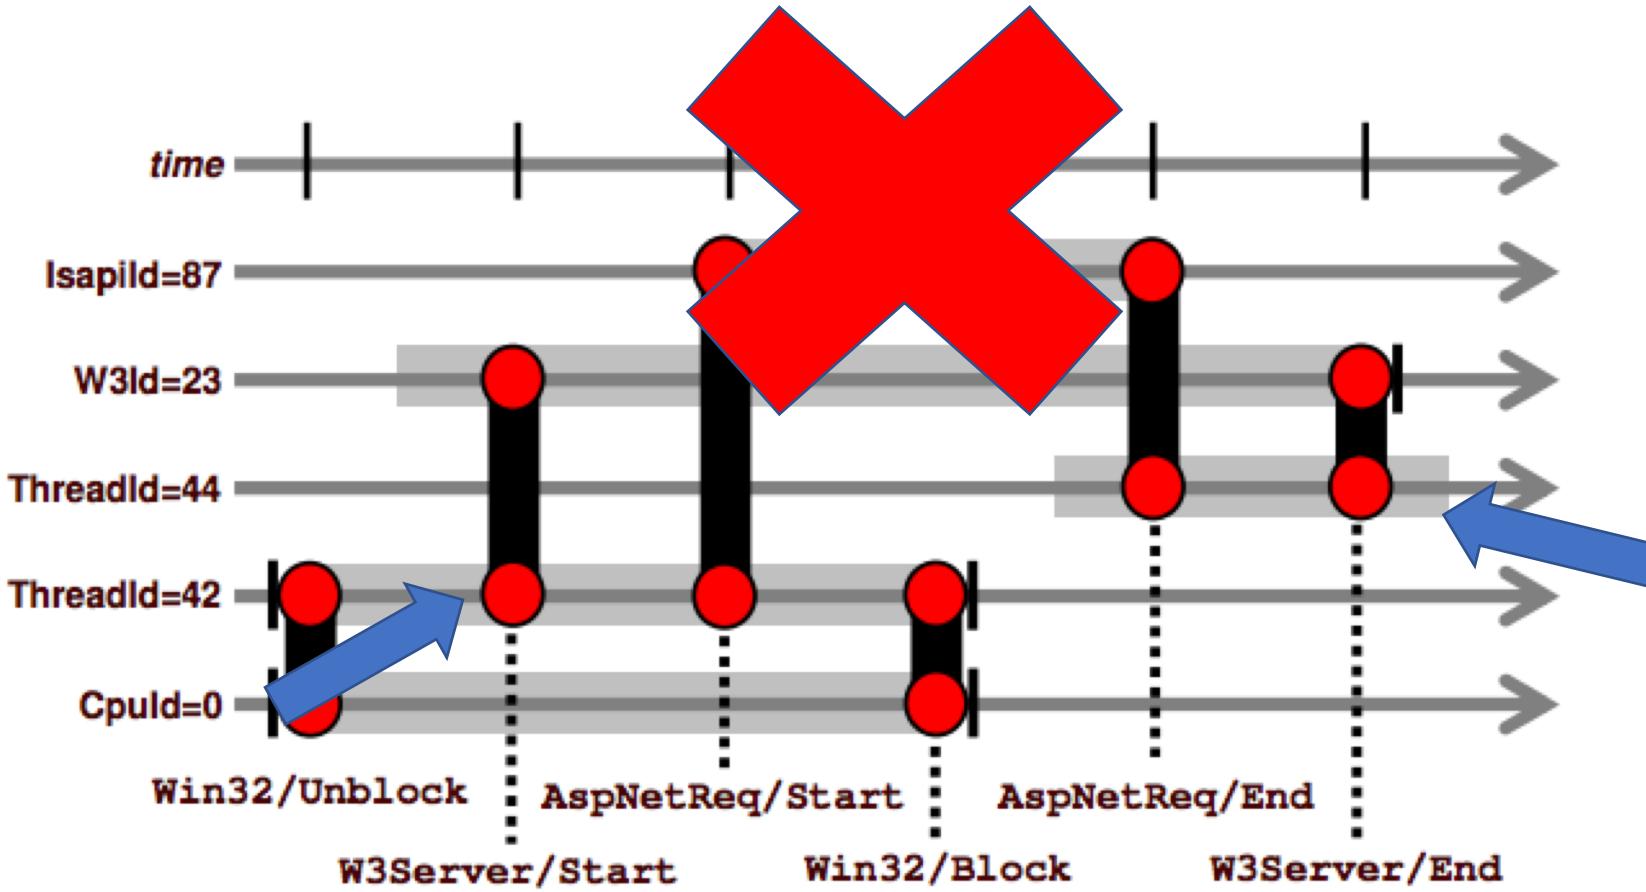
```
Event ("W3Server", "Start");
Attribute ("TId", BIND_BASIC, 0);
Attribute ("W3Id", BIND_BASIC, 0);
```

```
Event ("W3Server", "End");
Attribute ("TId", BIND_BASIC, 0);
Attribute ("W3Id", BIND_STOP, 0);
```



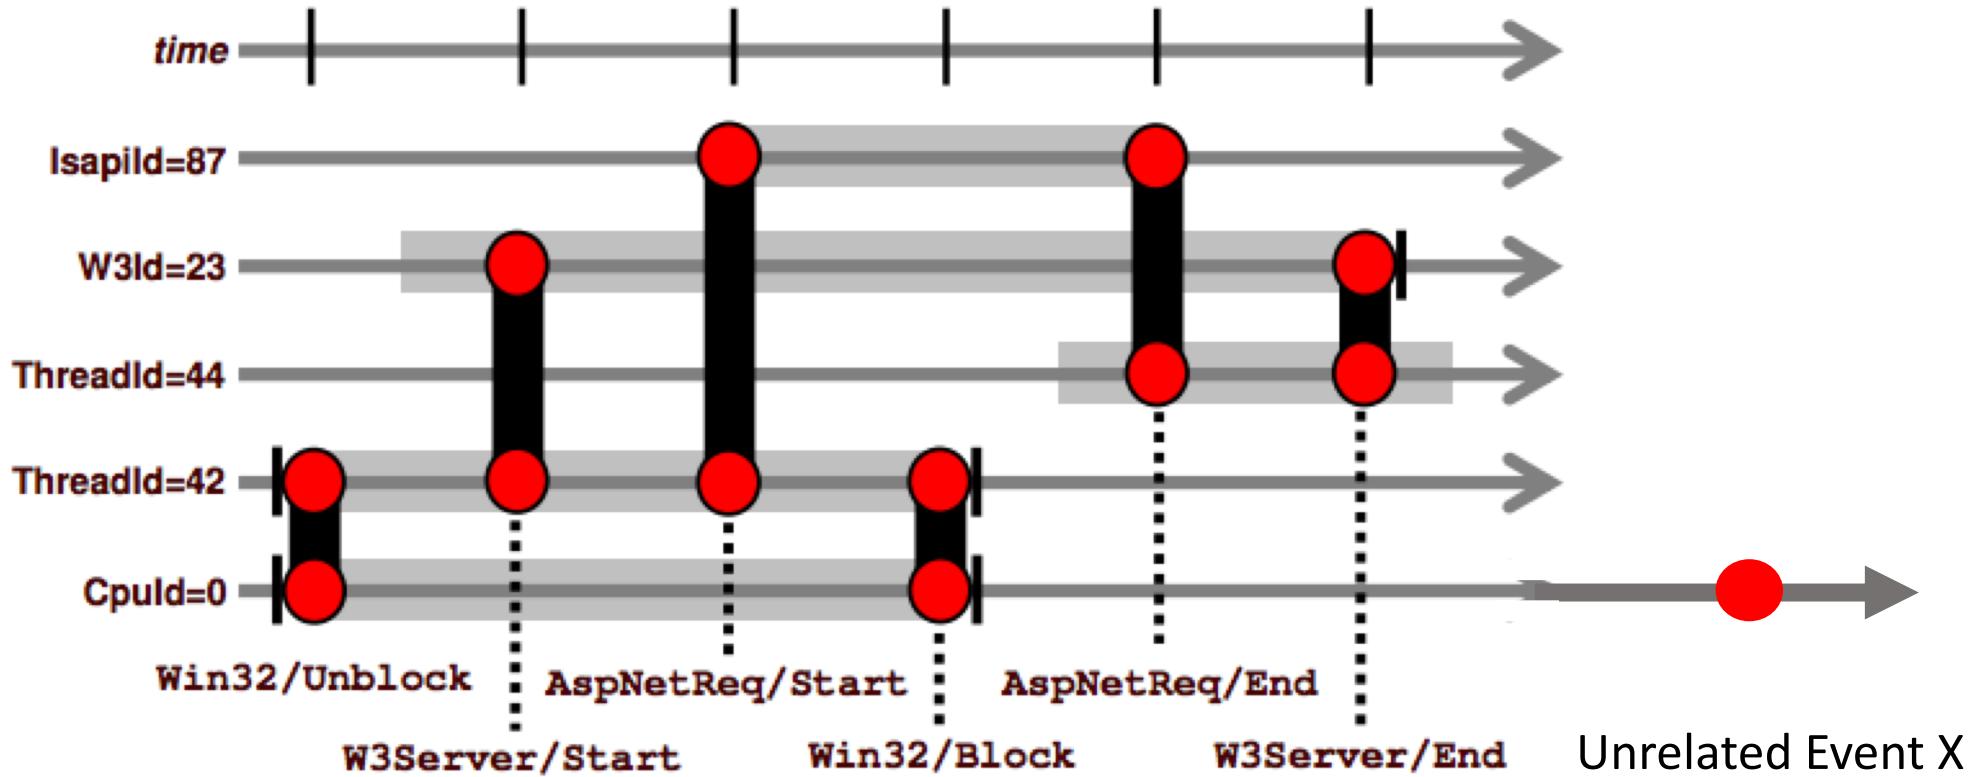
Event ("W3Server", "Start");  
 Attribute ("TId", BIND\_BASIC, 0);  
 Attribute ("W3Id", BIND\_BASIC, 0);

Event ("W3Server", "End");  
 Attribute ("TId", BIND\_BASIC, 0);  
 Attribute ("W3Id", BIND\_STOP, 0);



→ Event ("W3Server", "Start");  
 Attribute ("TId", BIND\_BASIC, 0);  
 Attribute ("W3Id", BIND\_BASIC, 0);

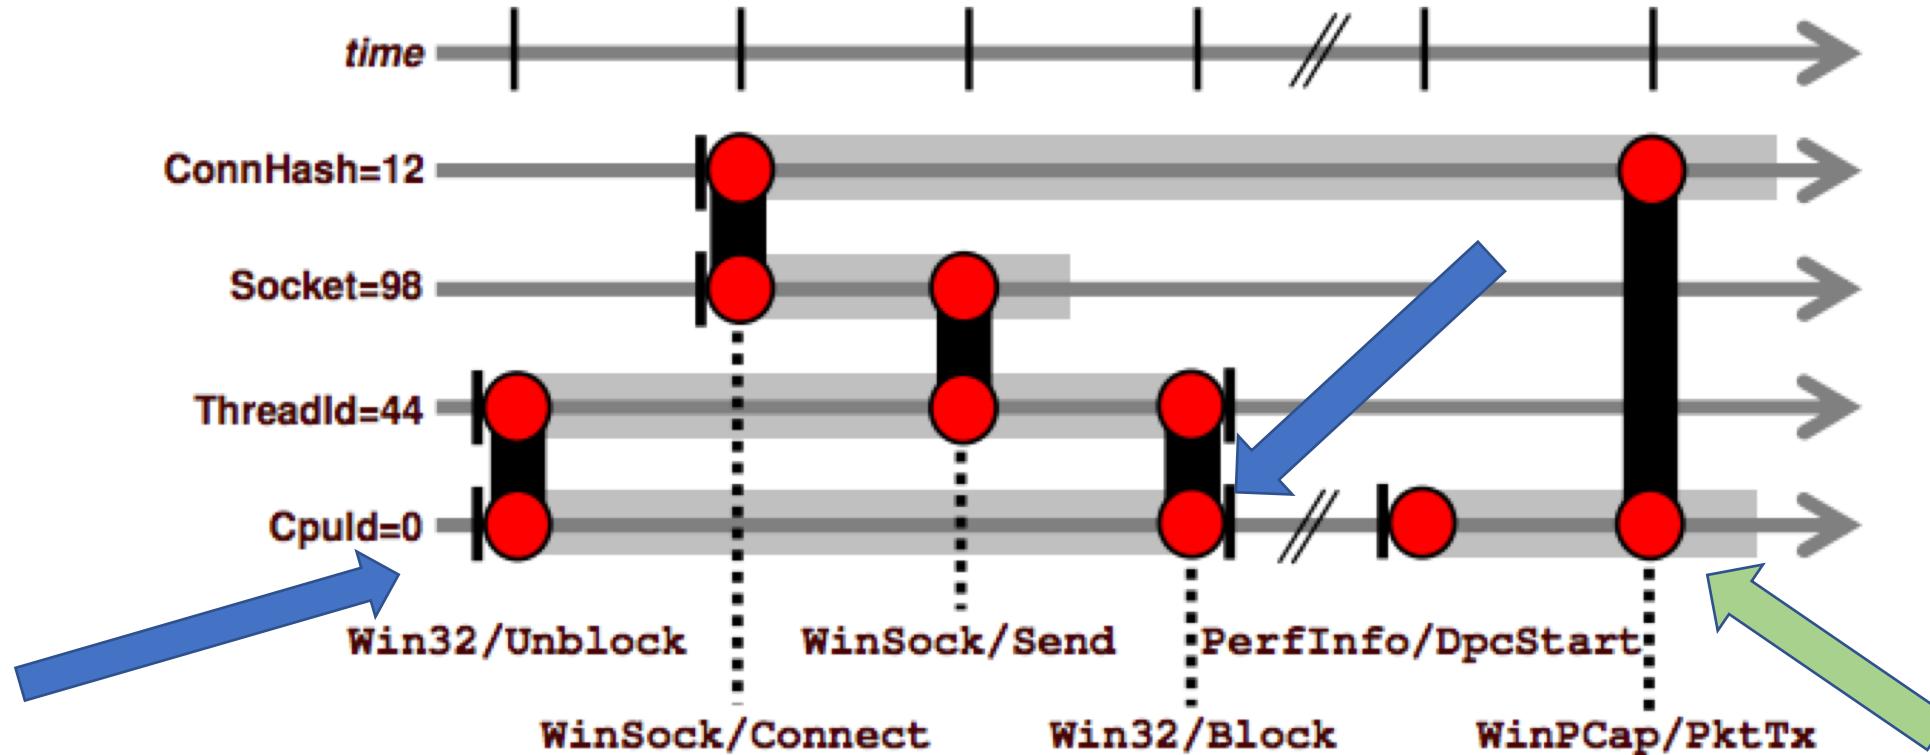
→ Event ("W3Server", "End");  
 Attribute ("TId", BIND\_BASIC, 0);  
 Attribute ("W3Id", BIND\_STOP, 0);



What prevents X from being joined with other unrelated events?

# Temporal Joins

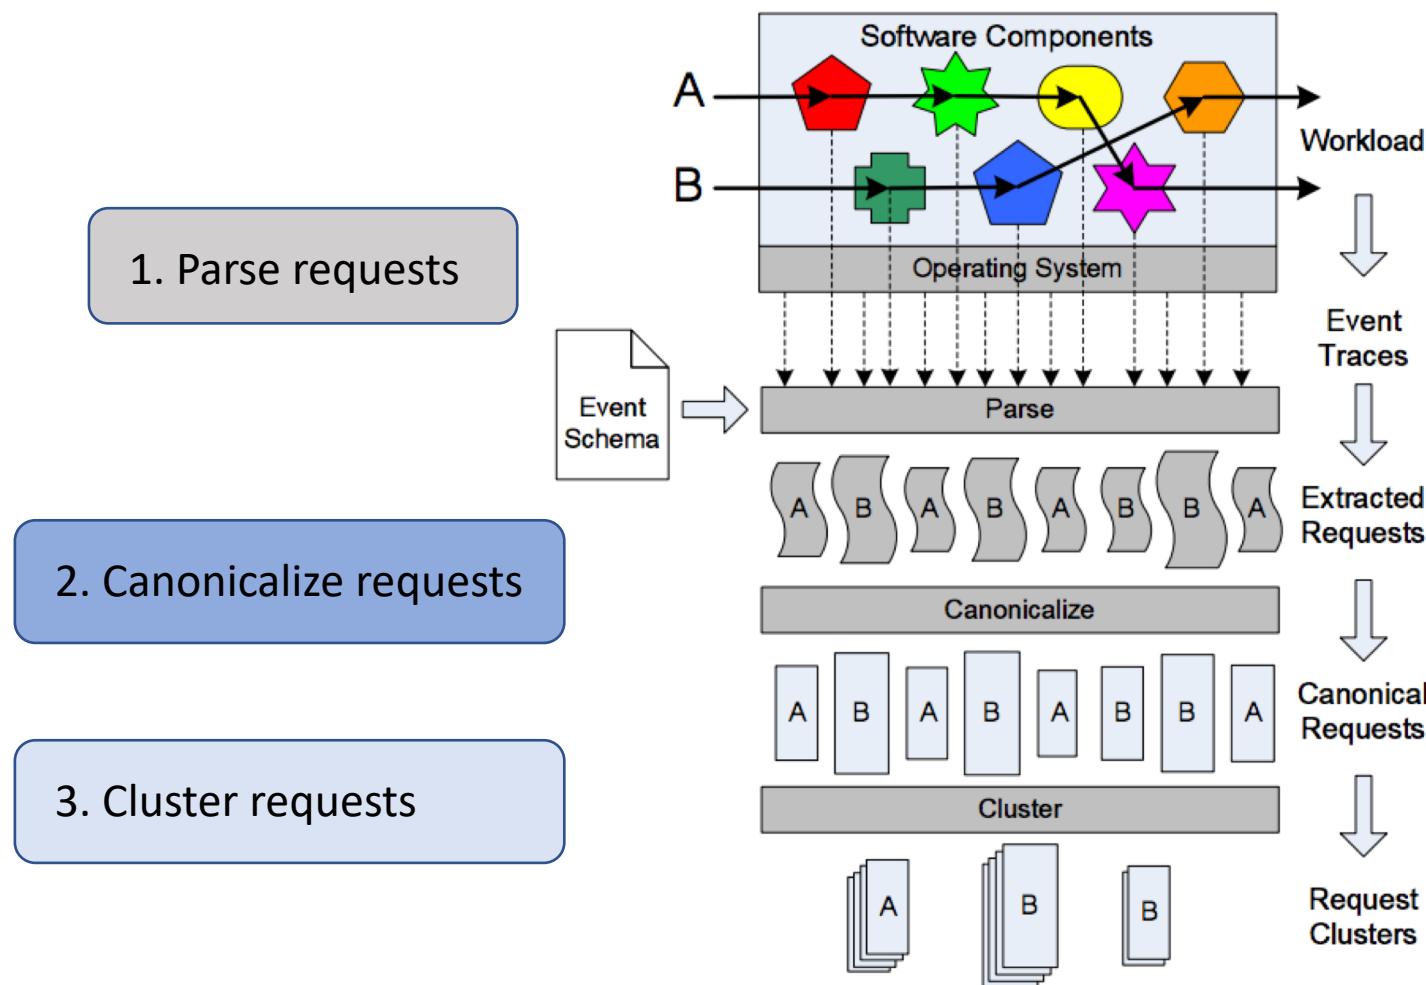
- Valid interval: period in which we know one resource is working exclusively on one request
- During a valid interval: events joined together as usual
- Once valid interval closed: no more events added to event via the specified pair
- This data is built into schema



Event ("Win32", "Unblock");  
 Attribute ("CpuID", BIND\_START, 0);  
 Attribute ("TId", BIND\_START, 0);

Event ("Win32", "Block");  
 Attribute ("CpuID", BIND\_STOP, 0);  
 Attribute ("TId", BIND\_STOP, 0);

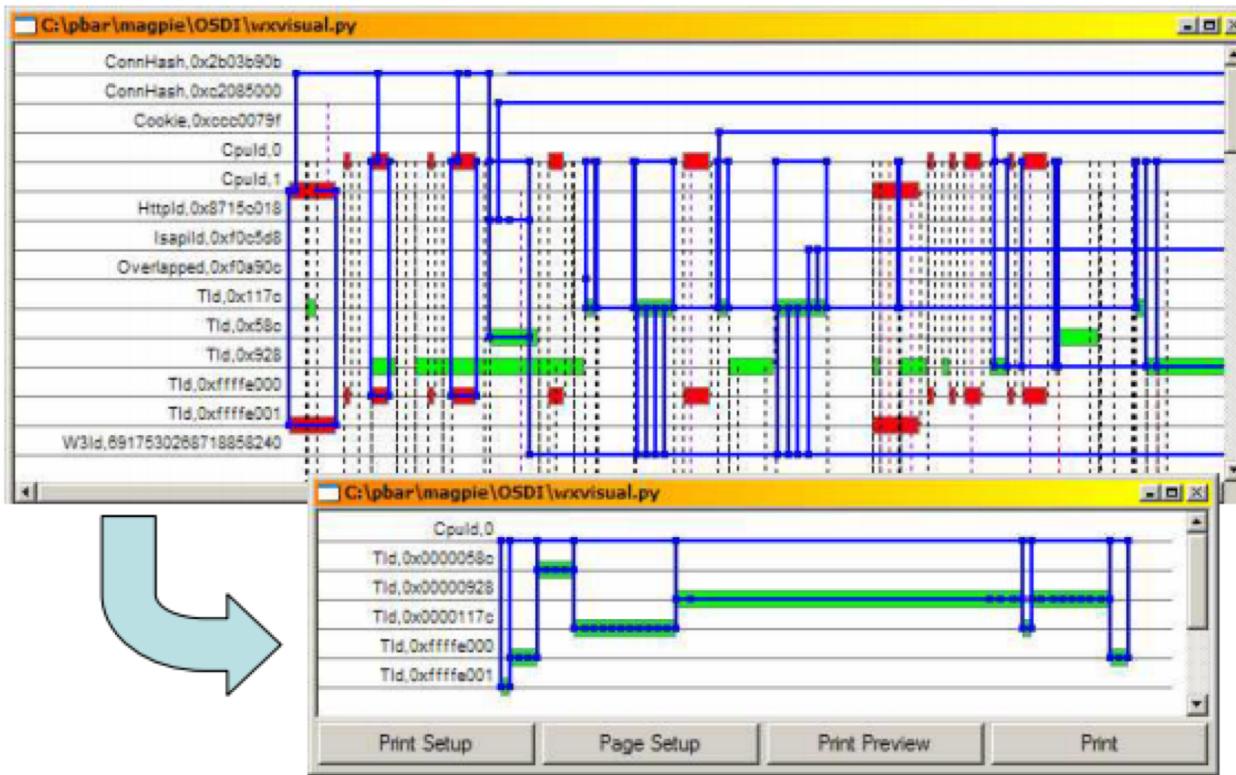
New events not  
 added via that  
 attribute pair



# Canonicalization

- Problem: we may want to know requests as sequence of demands
  - Not about how it was actually serviced
- Use cause ordering annotations to create canonical version of each request

“Clean out  
the request”



# How to handle synchronization?

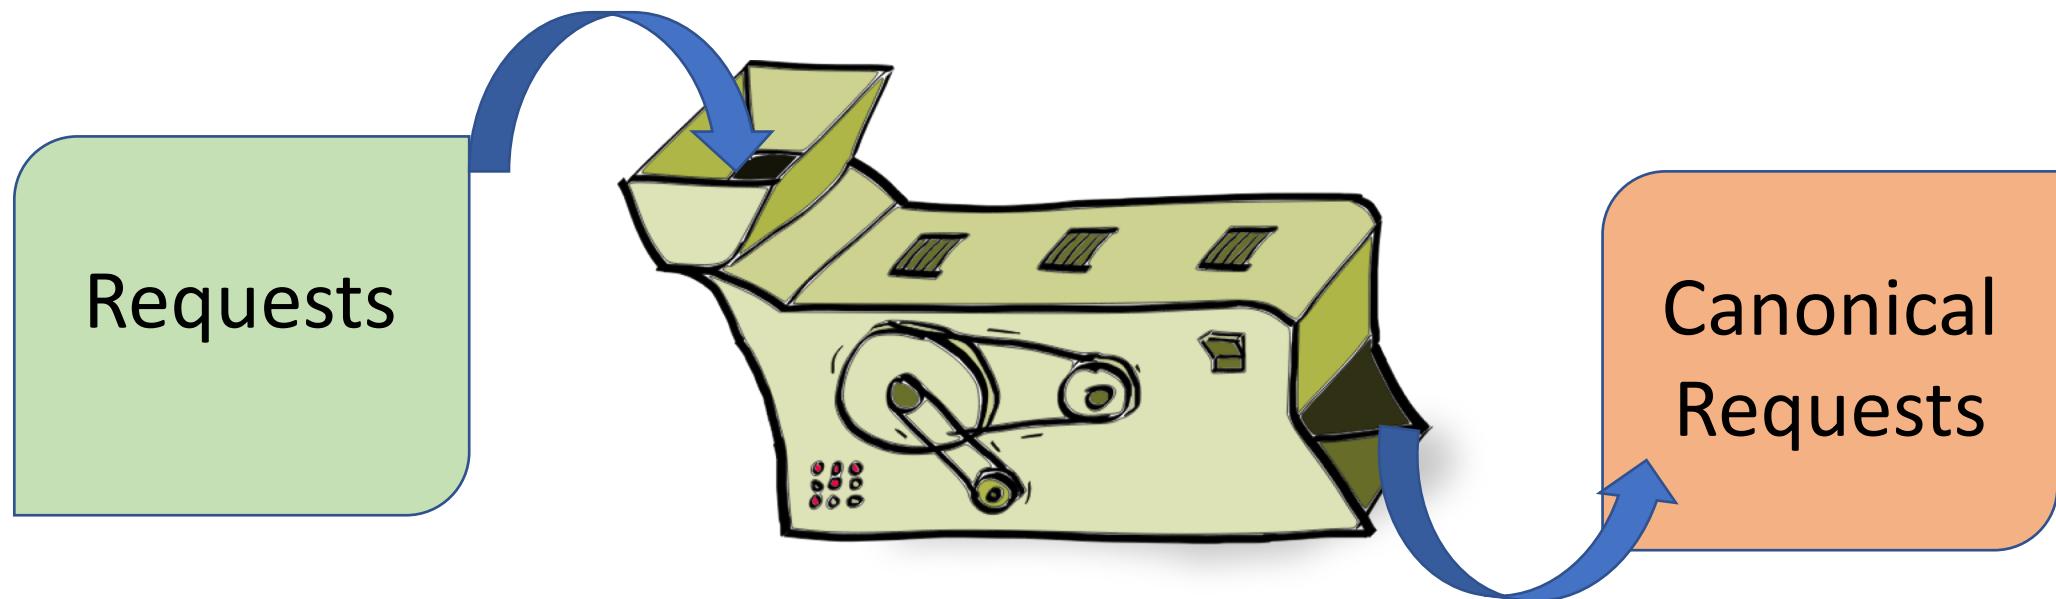
- Need to recover threading structure within request
  - E.g. identify when thread blocks or causes another to unblock
- Magpie has mechanism to explicitly insert causal dependencies into parsed event graphs

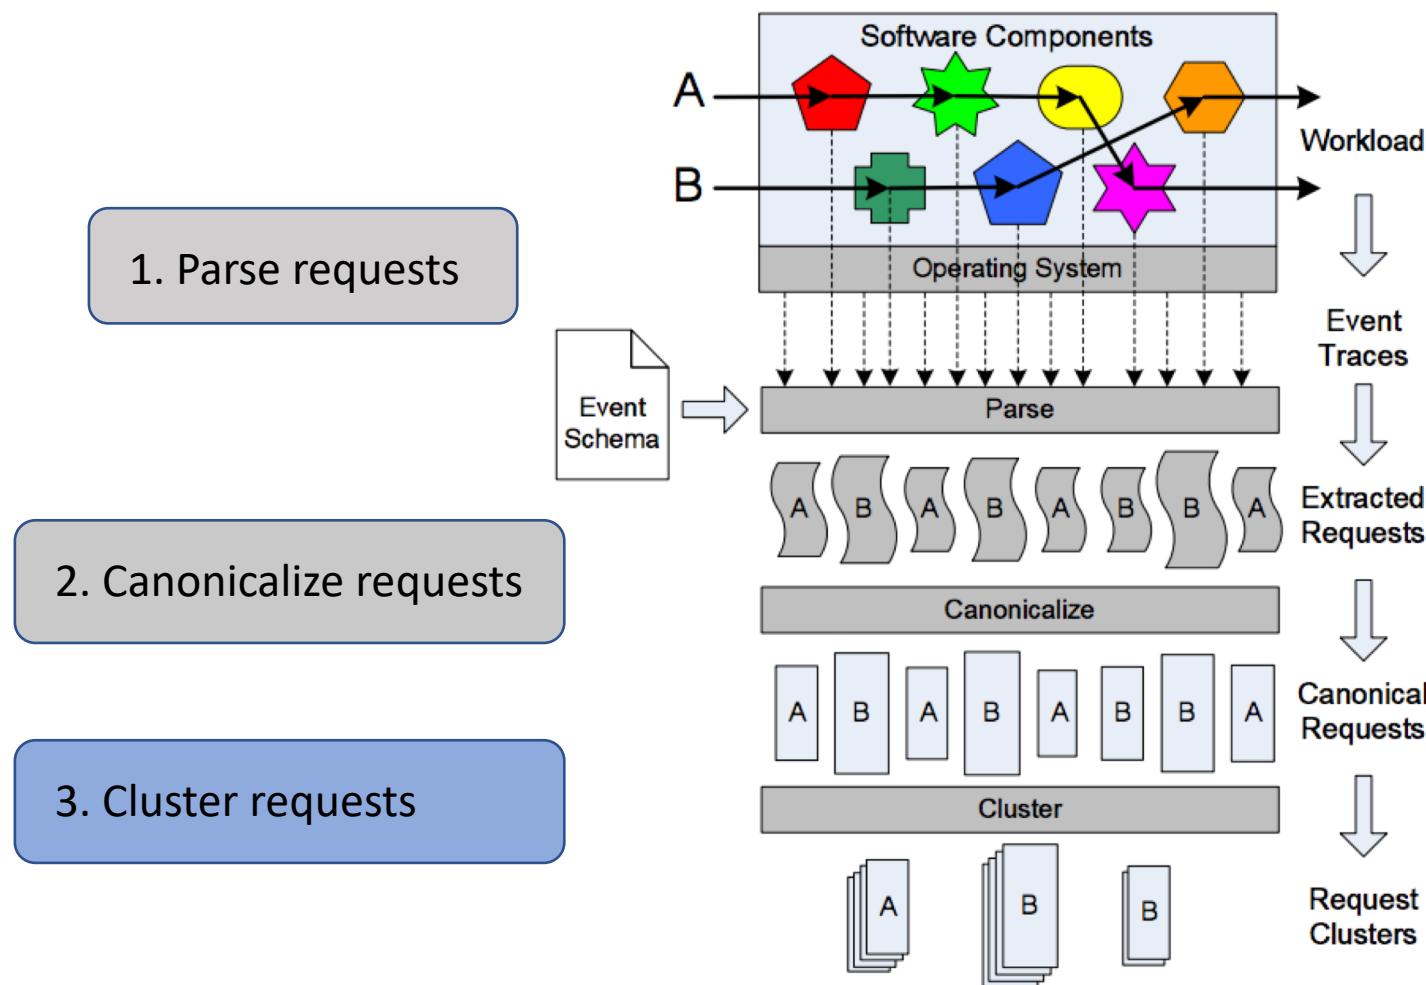
# How to handle requests on multiple machines?

- Run request parser once for each machine
- Request fragments canonicalized
- Run offline tool to combine canonical fragments by connecting synchronization events

# In short...

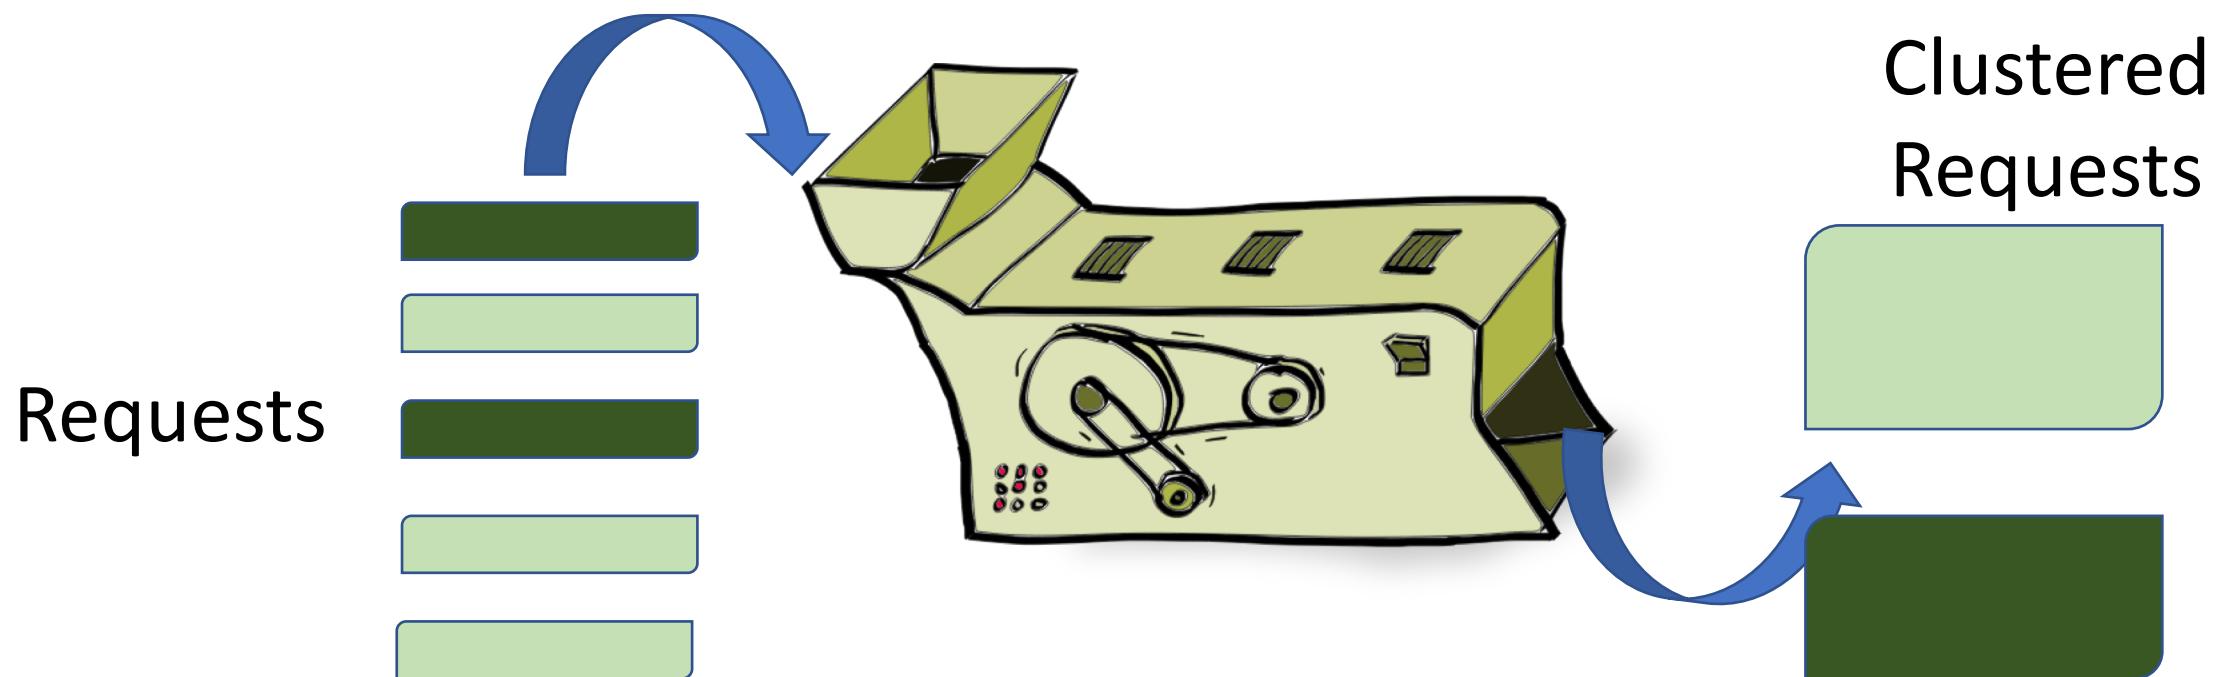
- Create canonicalized version of requests
  - Concatenate resource demands between synchronization points





# Behavioral Clustering

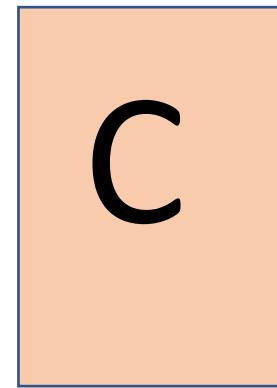
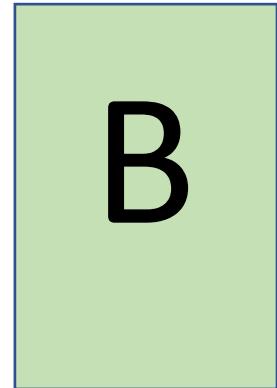
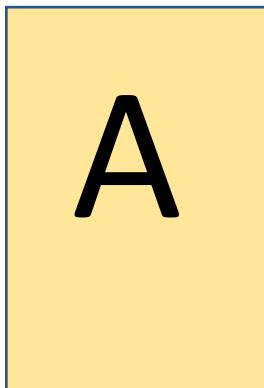
- Groups together requests with similar behavior
- Use single incremental clustering algorithm



# Behavioral Clustering Steps

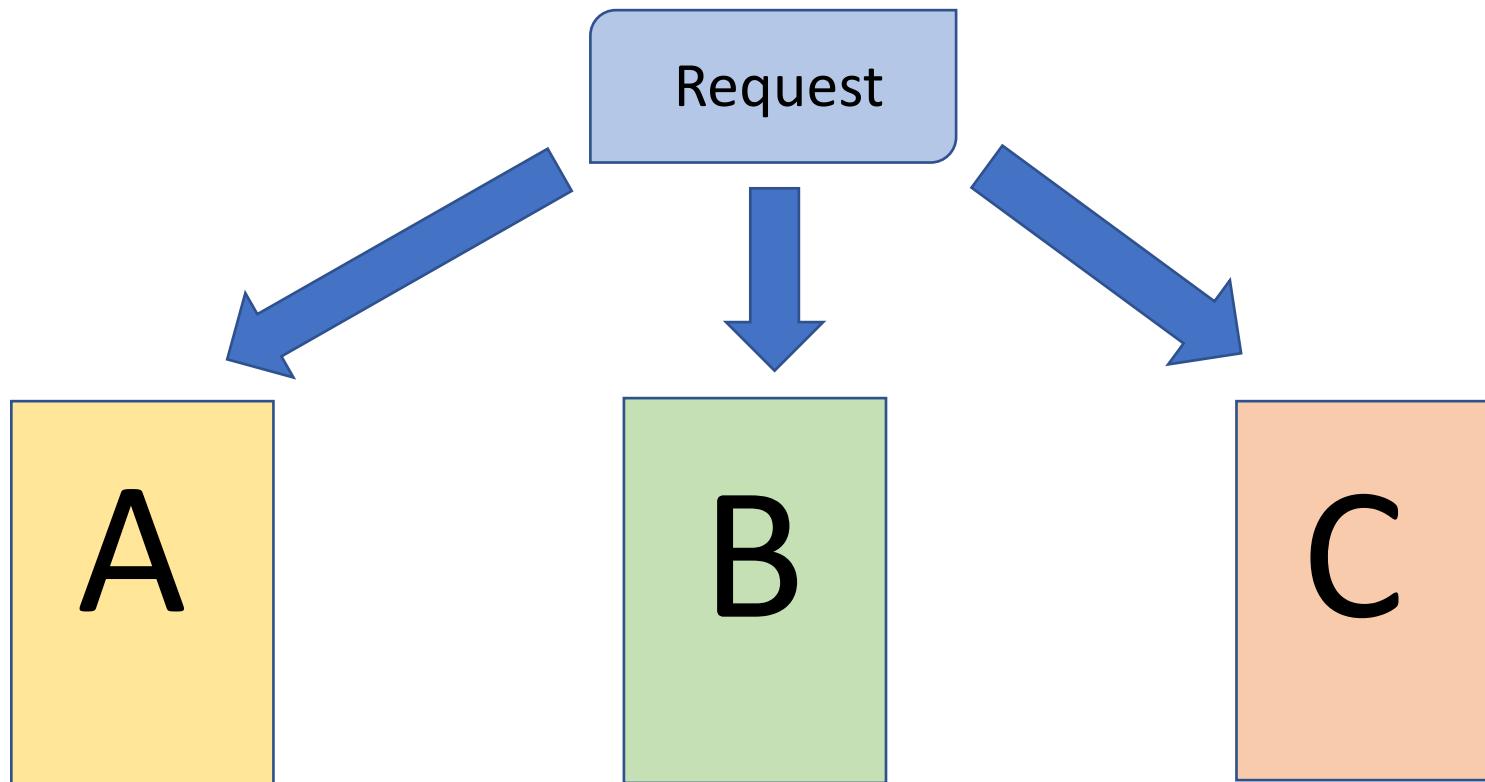
Clusterer maintains set of active clusters

Use representative centroid



# Behavioral Clustering Steps

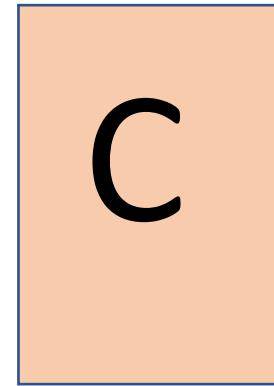
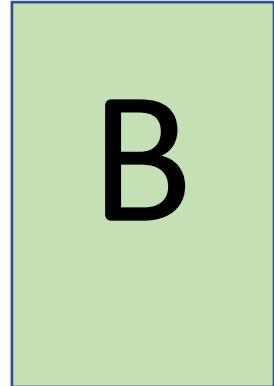
For each request, compute string-edit-distance to centroid(s)

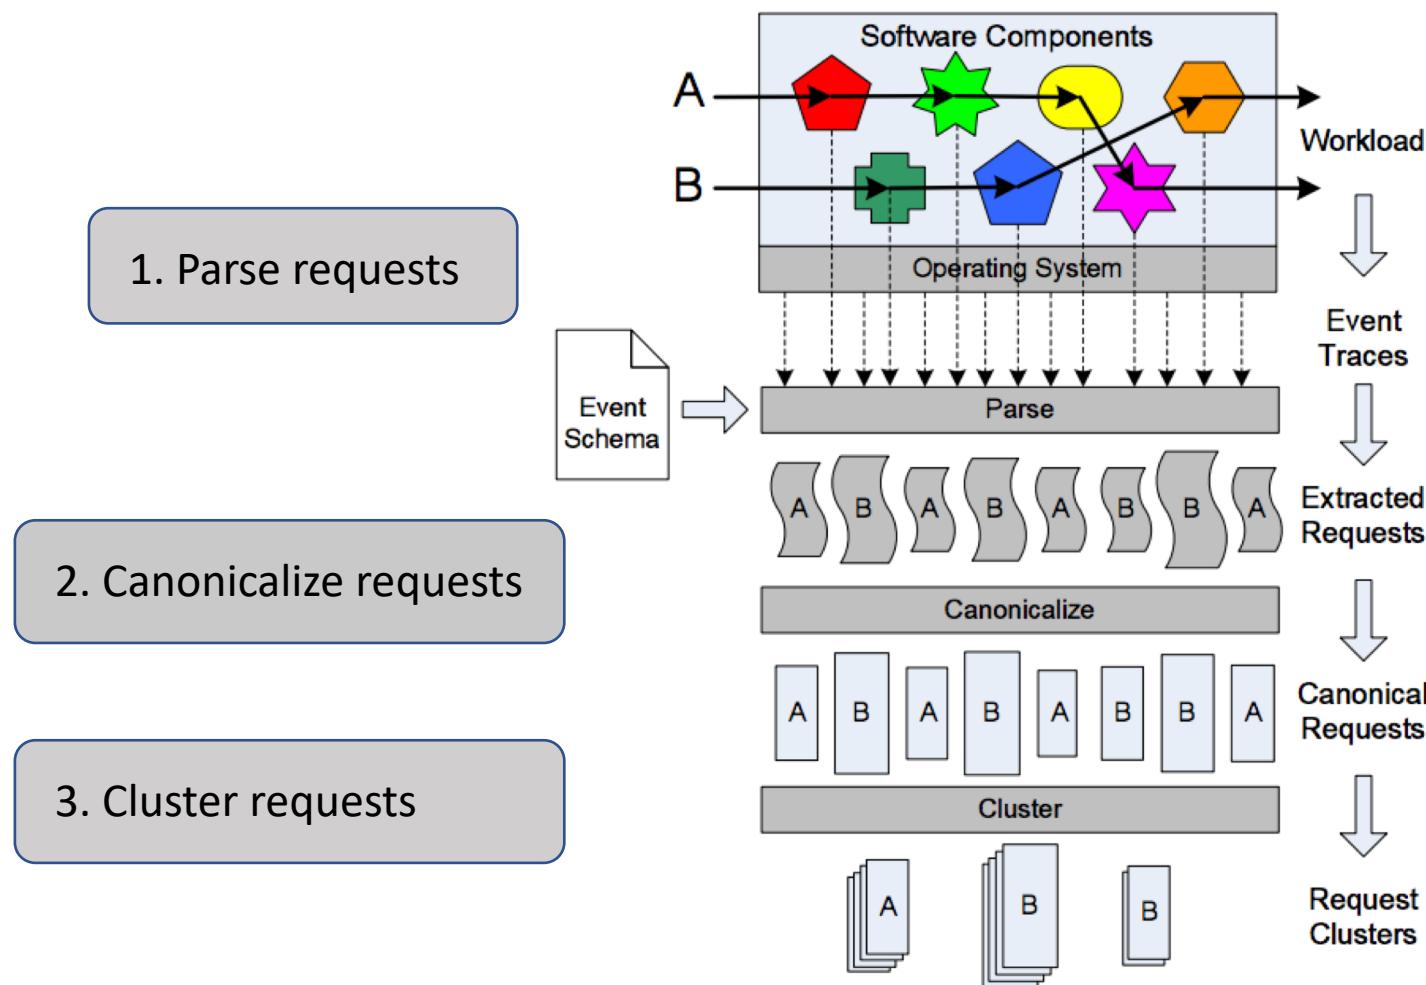


# Behavioral Clustering Steps

Add to cluster with smallest distance or create new cluster

Request





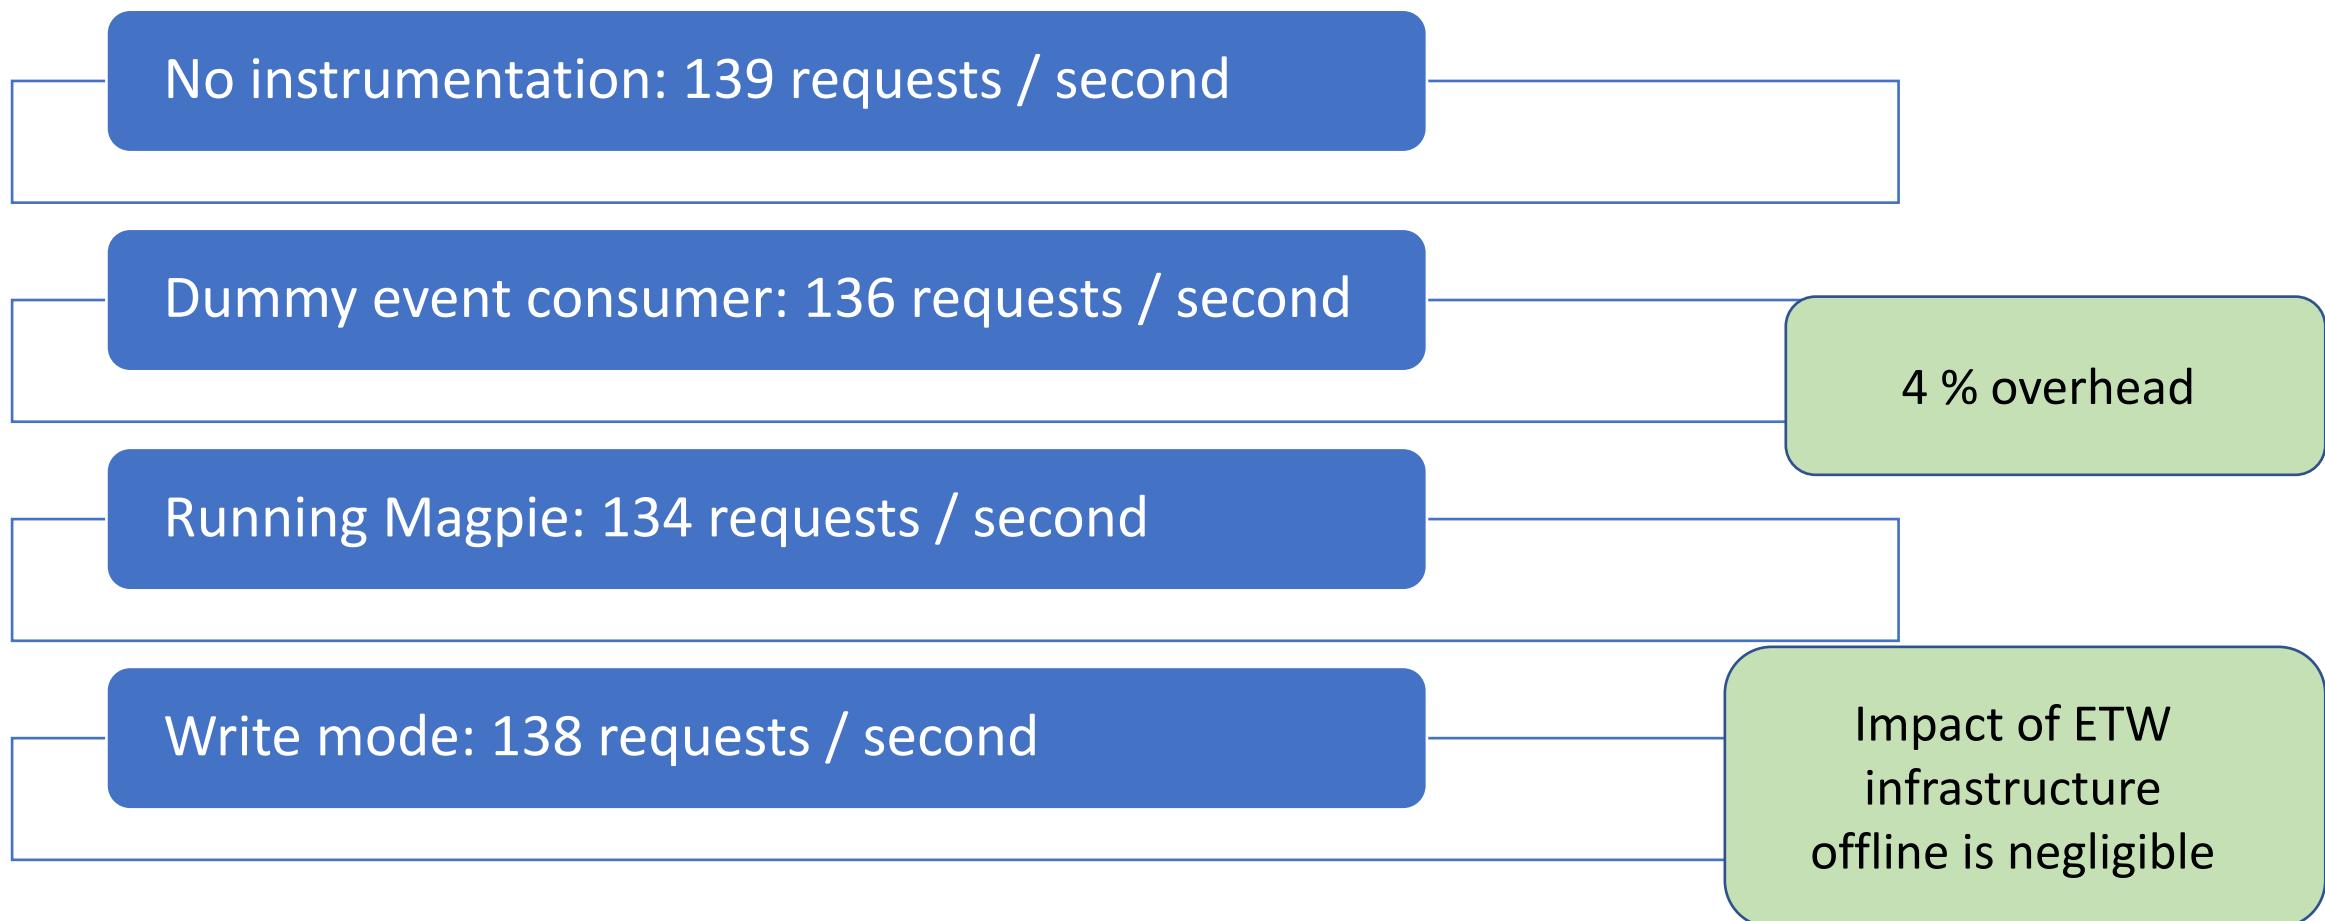
# In this talk..

- Background / Instrumentation
- Steps in the Toolchain
- Validation / Overhead Analysis
- Evaluation

# Overhead

- Assessed performance impact of event tracing / parsing using web stress-test
- CPU bound workload that generated active HTML content that saturated CPU
- Counted # HTTP requests served over two minute interval

# Overhead



# Validation

- Evaluated accuracy of Magpie's models using traces from synthetic workload
- Calibrated to check workload output was same as input

# Validation: Processor Time

- Used two types of requests:

Type	URL	Resource
A	longspin.aspx	1 thread spin CPU 30ms
B	shortspin.aspx	1 thread spin CPU 15ms

- Combination of:
  - Serialized vs. Concurrent (from 5 clients)
  - Single Variety or Mixed

Exp no.	Workload / # clients		Single?	Acct CPU %	Clusters found	Dia.	Min. Sep.	Model Error
1.1	500 A	x1	Single	96.6%	499	0.6	130	2.4%
1.2	Accnt CPU%: fraction of CPU consumed by relevant process			94.0%	1	0	130	
1.3				96.6%	500	0.4	0.0	
1.4	500 B	x5			49	1.4	21.7	
1.5	500 (A/B)	x1			1	1.4	21.7	
1.6	500 (A/B)	x5			1	1.4	13.1	
					500	0.0	0.0	
					266A	1.1	8.5	
					234B	1.1	8.5	
					244A	1.1	8.5	
					254B	1.1	8.5	
					1A	1.1	8.5	
					1B	0.0	33.2	
						0.0	103	2.0%

Exp no.	Workload / # clients		Single?	Accnt CPU %	Clusters found	Dia.	Min. Sep.	Model Error
1.1	500 A	x1	Single	96.6%	499	0.6	130	2.4%
					1	0	130	
1.2	500 B	x1	Single	94.0%	500	0.4	0.0	2.1%
1.3	500 A	x5	Conc.	96.6%	498	1.2	21.7	3.2%
					1	0.0	21.7	
					1	0.0	13.1	
1.4	500 B	x5	Conc.	94.6%	500	1.2	0.0	2.8%
1.5	500 (A/B)	x1	Mixed	95.9%	266A	1.1	8.5	0.9%
					234B	1.7	8.5	
1.6	500 (A/B)	x5	Mixed	95.9%	244A	1.2	8.0	2.0%
					254B	1.2	8.0	
					1A	0.0	33.2	
					1B	0.0	103	

Exp no.	Workload / # clients		Single?	Accnt CPU %	Clusters found	Dia.	Min. Sep.	Model Error
1.1	500 A	x1	Single	96.6%	499	0.6	130	<b>2.4%</b>
1.2	500 B	x1	Single	94.1%	499	0.6	130	<b>2.1%</b>
1.3	500 A	x5	Conc.	96.1%	499	0.6	130	<b>3.2%</b>
1.4	500 B	x5	Conc.	94.1%	499	0.6	130	<b>2.8%</b>
1.5	500 (A/B)	x1	Mixed	95.1%	499	0.6	130	<b>0.9%</b>
1.6	500 (A/B)	x5	Mixed	95.1%	499	0.6	130	<b>2.0%</b>
					1A	0.0	33.2	
					1B	0.0	103	

In all cases, cluster size / representatives can be used to represent workload to within 3.2%

# Magpie: anomaly detection

Exp no.	Workload / # clients		Single?	Accnt CPU %	Clusters found	Dia.	Min. Sep.	Model Error
1.1	500 A	x1	Single	96.6%	499	0.6	130	2.4%
					1	0	130	
1.2	500 B	x1	Single	94.0%	500	0.4	0.0	2.1%
1.3	500 A	x5	Conc.	96.6%	498	1.2	21.7	3.2%
					1	0.0	21.7	
					1	0.0	13.1	
1.4	500 B	x5	Conc.	94.6%	500	1.2	0.0	2.8%
1.5	500	x1	Mixed	95.9%	266A	1.1	8.5	0.9%
Discovered that driver was calling some function with proper args				234B	1.7	8.5		
				244A	1.2	8.0		
				254B	1.2	8.0		
				1A	0.0	33.2		
				1B	0.0	103		

# Concurrency

How well is Magpie able to capture differences in concurrency structure?

Type	URL	Resource
A	longspin.aspx	1 thread spin CPU 30ms
B	shortspin.aspx	1 thread spin CPU 15ms
E	parspin.aspx	2 threads spin CPU 15ms concurrently

Exp No	Workload/Clients		Accnt CPU %	Clusters Found	Dia.	Min. Sep	Model Error
3.3	500 (B/E)	x1	95.8%	267E	0.5	8.7	0.6%
				232B	0.4	8.8	
				1E	0.0	34.7	
3.5	500 (A/B/E)	x1	96.1%	172E	0.5	8.9	3.0%
				168A	0.6	8.1	
				160B	0.5	8.1	
3.6	500 (A/B/E)	x5	96.2%	187E	0.9	0.0	0.3%
				152B	0.6	8.6	
				160A	1.2	8.6	
				1A	0.0	130	

Exp No	Workload/Clients		Accnt CPU %	Clusters Found	Dia.	Min. Sep	Model Error
3.3	500 (B/E)	x1	95.8%	267E	0.5	8.7	0.6%
				232B	0.4	8.8	
				1E	0.0	34.7	
3.5	500 (A/B/E)	x1	96.1%	172E	0.5	8.9	3.0%
				168A	0.6	8.1	
				160B	0.5	8.1	
3.6	500 (A/B/E)	x5	96.2%	187E	0.9	0.0	0.3%
				152B	0.6	8.6	
				160A	1.2	8.6	
				1A	0.0	130	

Distance metric / clustering algo capable of separating requests that differ only in internal concurrency structure

# In this talk..

- Background / Instrumentation
- Steps in the Toolchain
- Validation / Overhead Analysis
- Evaluation

# Small two-tier website: Duwamish

- Setup: 2 machine system running Duwamish bookstore
  - Packed with generated data in accordance with volume / size

Exp no.	Workload	Accnt CPU %	Clusters Found	Max Dia	Min Sep
5.1	Book (B)	82.8%	WEB: 404 + 92 + 4	0.8	3.6
		96.1%	SQL: 500	0.5	0.0
			E2E: 408 + 92	1.2	5.1
5.2	Logon (L)	86.5%	WEB: 492 + 2 + 1	1.3	7.3
5.3	Categories (C)	95.2%	WEB: 445 + 26 + 29	2.3	14.9
		97.3%	SQL: 1499 + 1	1.3	16.4
			E2E: 444 + 26 + 29 + 1	2.6	17.9
5.4	Mixed (B/L/C)	91.7%	WEB: (138L + 1B) + 220B + 61C+ 21C + 24C + 10 * 25C	4.9	5.9
		96.7%	SQL: (141C <sub>1</sub> + 122 C <sub>2</sub> + 141 C <sub>3</sub> + 221B) + 9 C <sub>2</sub> + 10C <sub>2</sub>	2.3	19.7
			E2E:(138L + 1B) + 220B + 55C + 47CF + 14C + 10 * 25C	5.1	8.2

Exp no.	Workload	Accnt CPU %	Clusters Found	Max Dia	Min Sep	
5.1	Book (B)	82.8%	WEB: 404 + 92 + 4	0.8	3.6	
		96.1%	SQL: 500	0.5	0.0	
			E2E: 408 + 92	1.2	5.1	
5.2	Logon (L)	86.5%	WEB: 492 + 2 + 1	1.3	7.3	
5.3	Categories (C)	95.2%	WEB: 445 + 26 + 29	2.3	14.9	
		97.3%	SQL: 1499 + 1	1.3	16.4	
			E2E: 444 + 26 + 29 + 1	2.6	17.9	
		1.7%	WEB: (138L + 1B) + 220B + 61C+ 21C + 24C + 10 * 25C	4.9	5.9	
		6.7%	SQL: (141C <sub>1</sub> + 122 C <sub>2</sub> + 141 C <sub>3</sub> + 221B) + 9 C <sub>2</sub> + 10C <sub>2</sub>	2.3	19.7	
			E2E:(138L + 1B) + 220B + 55C + 47CF + 14C + 10 * 25C	5.1	8.2	

Groups for Book / Logon  
much tighter than those  
for Categories

Workload model based  
on observed behavior >  
workload model based on  
URL measurements

# Conclusion

- Magpie is a toolchain that takes events generated by OS, middleware, application components
- Able to output a model of how requests were serviced and how it could have been serviced
- Validated approach against synthetic workloads

# Questions?

