

The Locality Principle

Peter J. Denning

Presented by: Matthew Perez

Background History

- Atlas Computer system (1949)
 - Storage: RAM and Disk
- Manual page transfers (overlays)
 - Programmers could utilize RAM and Disk
- Innovation: Virtual Memory (1959)
 - OS managed pages loaded on RAM



Atlas Computer System

Challenges with Virtual Memory

- Issue 1: Translating addresses to locations
- Issue 2: Replacing loaded pages

Issues with Virtual Memory

- Issue 1: Translating addresses to locations
 - Solution: High-speed cache (TLB)
- Issue 2: Replacing loaded pages

Issues with Virtual Memory


- Issue 1: Translating addresses to locations
 - Solution: High-speed cache (TLB)
- Issue 2: Replacing loaded pages
 - Solution: Learning algorithms

Page Faults

- Page faults -> 10,000x slower than CPU instruction cycle
- Industry invested in replacement algorithms
 - Bad algorithm = expensive in long run
 - IBM Watson found LRU to be best (1966)


Thrashing Problem

- Multiprogramming -> thrashing
- Too many programs = page faults
- Working set = Memory needed for program
- Solution: Limit multiprogramming



Open Question: Thrashing

- How should we process programs in queue?
 - Let smaller programs go first?
 - How to ensure big programs get processed?




Locality Principle

- Memory demand represents sequence of locality sets
- Programs reference addresses close in proximity and time
 - Spatial Locality = Nearby addresses
 - Temporal Locality = Recently referenced
- Beneficial for quickly loading memory

Locality Visualization



Pre-planned



Non-planned

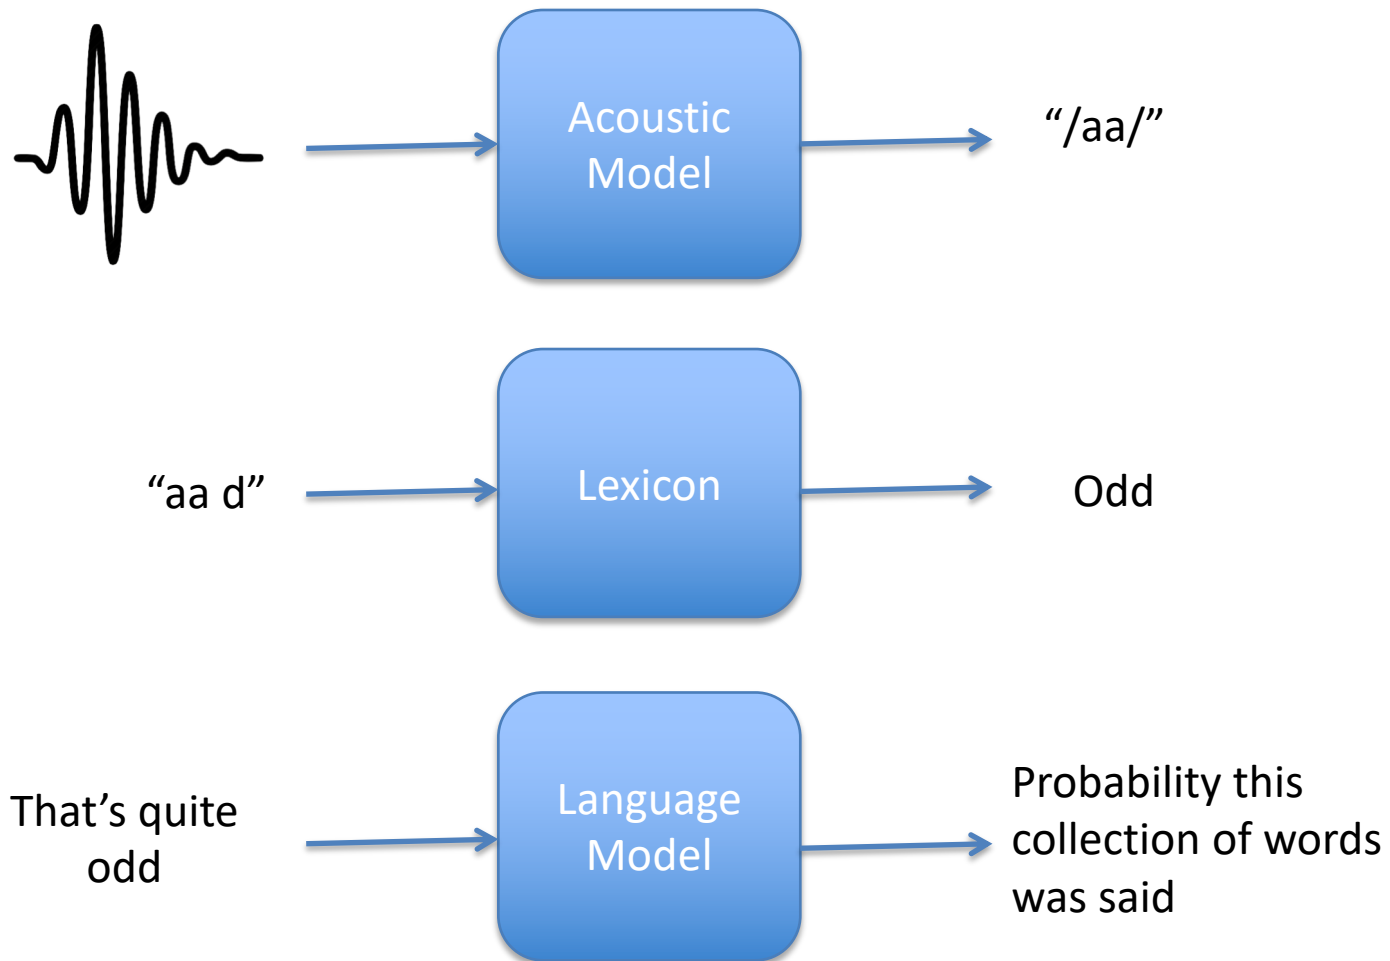
Locality Principle Examples

- Loading data caches
- Web browsers – recent web pages
- Video streaming – reduce distance to server
- Search engines – relevant queries

Personal Example

- Use locality principle to infer context?
- Speech recognition uses context for improving classification

Speech Recognition



ASR Example

- Just Acoustic Model and Lexicon (no context):
“He pored the waiter into a cap”
- Causes:
 - High speaker variability (accents, gender, word selection)
 - Noisy environment

ASR Example

- Acoustic Model and Lexicon produce:
“He poured the ____ into a cup”
- Top acoustic model and lexicon choices:
 - Waiter
 - Otter
 - Water

Language Model
decides most
probable choice

ASR Example

- Acoustic Model and Lexicon produce:
“He poured the ____ into a cup”
- Top acoustic model and lexicon choices:
 - Waiter
 - Otter
 - Water

Discussion Questions

- Strengths
- Weaknesses
- Open Questions

Discussion Questions

- Contrast with Exokernel. Virtual Memory gives OS more control rather than programmer (overlays). What are some of the pros/cons behind this approach?
- How to schedule/manage programs in the stall queue?

Discussion Questions

- Analyze replacement algorithms: Random, FIFO, LRU, empty block. Which is best for small block sizes? Large block sizes?
- What about dynamically allocating different sized memory?



Compiler Optimizations for Improving Data Locality

Steve Carr Kathryn S. McKinley Chau-Wen Tseng

Presented by: Wenjia He
Department: CSE



Motivation

- Fast processor & slow memory
- Cache: based on data locality
- Before:
 - Scientific programmers' effort
 - Machine-dependent programming
- Now:
 - Compiler strategy — reducing number of cache line references in loop nests

Motivation

Example — Matrix Multiply

Given $N \times N$ matrices A and B , calculate $C = A * B$

Loop nest:

Do $J = 1, N$

Do $K = 1, N$

Do $I = 1, N$

$C(I, J) = C(I, J) + A(I, K) * B(K, J)$



Background

- Data reuse
 - Temporal reuse
 - Spatial reuse
- Reference group
 - Group-temporal reuse
 - Group-spatial reuse

Background

- Cost model

$$\text{RefCost} = \begin{cases} 1 & \text{loop invariant} \\ \text{trip}/(\text{cls}/\text{stride}) & \text{consecutive} \\ \text{trip} & \text{non-consecutive} \end{cases}$$
$$\text{LoopCost} = \sum_{k=1}^m \text{RefCost}_k \prod_{i \neq l} \text{trip}_i$$



Key Insights and Design Details

- Compound loop transformations
 - Loop permutation
 - Loop reversal
 - Loop fusion
 - Loop distribution



Key Insights and Design Details

- Compound loop transformations

- Loop permutation

- Loop reversal

- Loop fusion

- Loop distribution



Key Insights and Design Details

■ Loop permutation

Steps:

- Calculate LoopCost for each loop
- Rank loops in order
- Or find the nearest legal order

Complexity

- Sorting: $O(n \log n)$
- Finding legal order: $O(n^2)$

Key Insights and Design Details

- Loop permutation

Example

Do J = 1, N

Do K = 1, N

Do I = 1, N

$C(I, J) = C(I, J) + A(I, K) * B(K, J)$

Key Insights and Design Details

- Loop permutation
Example

References	J	K	I
C(I, J)	$n \cdot n^2$	$1 \cdot n^2$	$1/4 \cdot n \cdot n^2$
A(I, K)	$1 \cdot n^2$	$n \cdot n^2$	$1/4 \cdot n \cdot n^2$
B(K, J)	$n \cdot n^2$	$1/4 \cdot n \cdot n^2$	$1 \cdot n^2$
Total	$2n^3 + n^2$	$5/4 \cdot n^3 + n^2$	$1/2 \cdot n^3 + n^2$



Key Insights and Design Details

- Compound loop transformations
 - Loop permutation
 - Loop reversal
 - Loop fusion
 - Loop distribution



Key Insights and Design Details

- Loop reversal

- Reverse the order of iteration

- May help loop permutation



Key Insights and Design Details

- Compound loop transformations
 - Loop permutation
 - Loop reversal
 - Loop fusion
 - Loop distribution



Key Insights and Design Details

■ Loop fusion

Combine multiple loop nests to one loop nest

Benefit

- Create a permutable loop nest
- Improve temporal locality

Key Insights and Design Details

- Loop fusion

Example

Do I = 2, N

Do K = 1, N

$X(I, K) = X(I-1, K) * A(I, K) / B(I-1, K)$

Do K = 1, N

$B(I, K) = A(I, K) * A(I, K) / B(I-1, K)$

Key Insights and Design Details

- Loop fusion

Example

Do I = 2, N

Do K = 1, N

$X(I, K) = X(I-1, K) * A(I, K) / B(I-1, K)$

$B(I, K) = A(I, K) * A(I, K) / B(I-1, K)$

Key Insights and Design Details

- Loop fusion

Example

Do K = 1, N

Do I = 2, N

$$X(I, K) = X(I-1, K) * A(I, K) / B(I-1, K)$$

$$B(I, K) = A(I, K) * A(I, K) / B(I-1, K)$$



Key Insights and Design Details

- Compound loop transformations
 - Loop permutation
 - Loop reversal
 - Loop fusion
 - Loop distribution



Key Insights and Design Details

- Loop distribution

Divide a single loop into multiple loops

Benefit: create permutable loop nests

Key Insights and Design Details

- Loop distribution

Example: Cholesky factorization

Do K = 1, N

$A(K, K) = \text{SQRT}(A(K, K))$

Do I = K+1, N

$A(I, K) = A(I, K) / A(K, K)$

Do J = K+1, I

$A(I, J) = A(I, J) - A(I, K) * A(J, K)$

Key Insights and Design Details

- Loop distribution

Example: Cholesky factorization

Do K = 1, N

$A(K, K) = \text{SQRT}(A(K, K))$

Do I = K+1, N

$A(I, K) = A(I, K) / A(K, K)$

Do I = K+1, N

Do J = K+1, I

$A(I, J) = A(I, J) - A(I, K) * A(J, K)$

Key Insights and Design Details

- Loop distribution

Example: Cholesky factorization

Do K = 1, N

$A(K, K) = \text{SQRT}(A(K, K))$

Do I = K+1, N

$A(I, K) = A(I, K) / A(K, K)$

Do J = K, N

Do I = J+1, N

$A(I, J+1) = A(I, J+1) - A(I, K) * A(J+1, K)$



Key Insights and Design Details

- **Compound loop transformations**

- Loop permutation

- Loop reversal

- Loop fusion

- Loop distribution



Key Insights and Design Details

- Compound loop transformations
 - Do loop permutation for each loop nest
 - Try loop fusion and do loop permutation
 - Try loop distribution if helping permutation
 - Fuse all loop nests

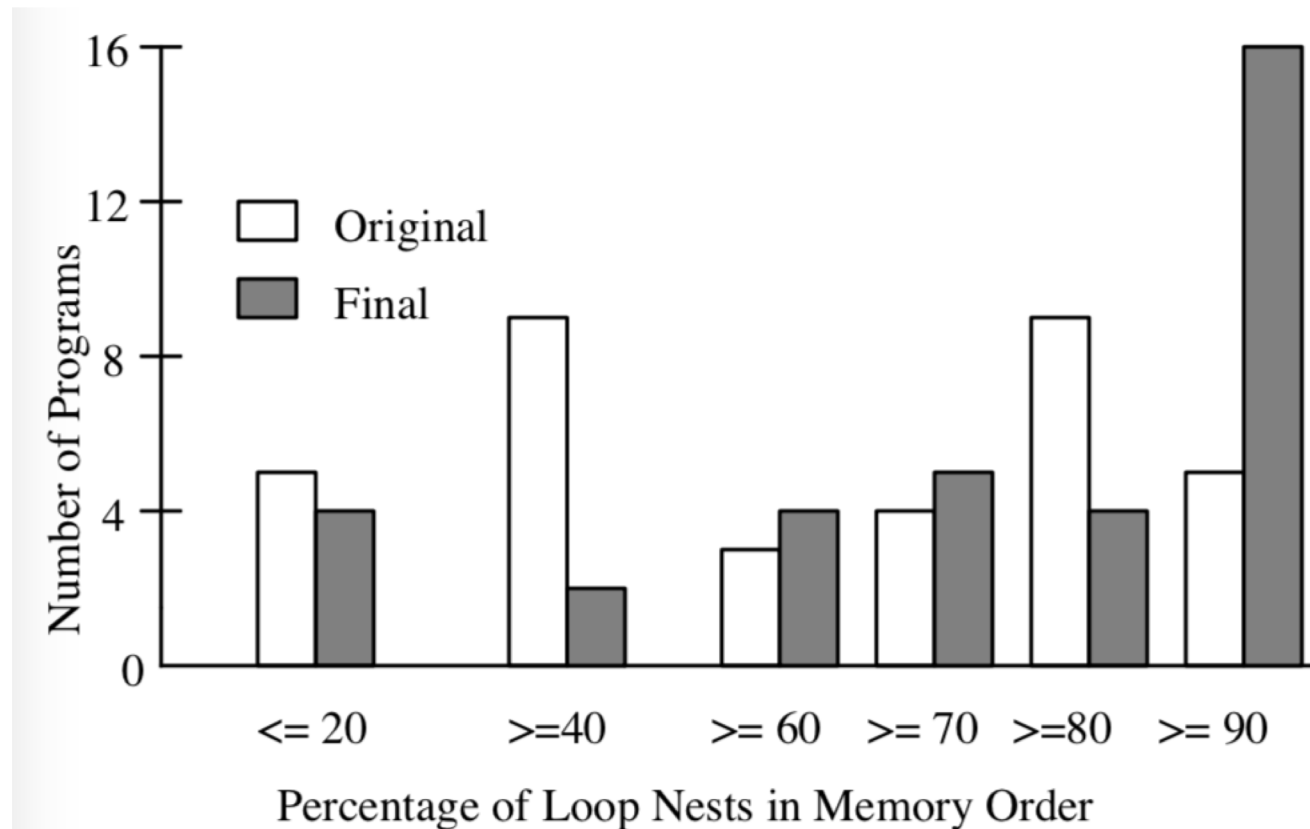


Key Results

- Analyze results in:
 - Number of programs achieving memory order
 - Speed
 - Cache hit rate

Key Results

- Number of programs achieving memory order



Key Results

■ Speed

Program	Original	Transformed	Speedup
<i>Perfect Benchmarks</i>			
arc2d	410.13	190.69	2.15
dyfesm	25.42	25.37	1.00
flo52	62.06	61.62	1.01
<i>SPEC Benchmarks</i>			
dnasa7 (btrix)	36.18	30.27	1.20
dnasa7 (emit)	16.46	16.39	1.00
dnasa7 (gmtry)	155.30	17.89	8.68
dnasa7(vpenta)	149.68	115.62	1.29
<i>NAS Benchmarks</i>			
applu	146.61	149.49	0.98
appsp	361.43	337.84	1.07
<i>Misc Programs</i>			
simple	963.20	850.18	1.13
linpackd	159.04	157.48	1.01
wave	445.94	414.60	1.08

Key Results

■ Cache hit rate

Program	Optimized Procedures				Whole Program			
	Cache 1		Cache 2		Cache 1		Cache 2	
	Orig	Final	Orig	Final	Orig	Final	Orig	Final
<i>Perfect Benchmarks</i>								
adm	100	100	97.7	97.8	99.95	99.95	98.48	98.58
arc2d	89.0	98.5	68.3	91.9	95.30	98.66	88.58	93.61
bdna	100	100	100	100	99.45	99.45	97.32	97.32
dyfesm	100	100	100	100	99.98	99.97	97.02	96.95
flo52	99.6	99.6	96.7	96.3	98.77	98.77	93.84	93.80
mdg	100	100	87.4	87.4	— —	— —	— —	— —
mg3d	98.8	99.7	95.3	98.7	— —	— —	— —	— —
ocean	100	100	93.0	92.8	99.36	99.36	93.71	93.72
qcd	100	100	100	100	99.83	99.83	98.85	98.79
spec77	100	100	100	100	99.28	99.28	93.79	93.78
track	100	100	100	100	99.81	99.81	97.49	97.54
trfd	99.9	99.9	93.7	93.7	99.92	99.92	96.43	96.40



Key Results

- Analyze results in:
 - Number of programs achieving memory order
 - Speed
 - Cache hit rate
- Prove effect of loop fusion and distribution



Strengths

- Introduce loop fusion & distribution & reversal
 - Improve performance of loop permutation
 - Permit imperfect loop nests
- Proper reference group formulation



Weaknesses

- Limitation in uniprocessor system
- Limitation in rectangular and triangular nests



Open Questions

- Loop tiling
- Loop fusion's defect



Discussion



Thank you!