

Development of the Domain Name System -> 1988

Authors: Paul V. Mockapetris & Kevin J. Dunlap

Presented by Steve Sprecher

Problem

- IP addresses are hard to remember
- Keep track of hostname -> address mapping
- Centrally controlled
- Definitely not scalable
- Current domain name systems are Limited or host specific



Design Requirements & Constraints

- DNS same or more info as HOSTS.TXT
- Distributed maintenance
- No obvious size limits
- Work cross Internet(s)
- Tolerable performance
- Must provide extensible services and work everywhere (to justify cost)
- Do not force an OS or organizational style



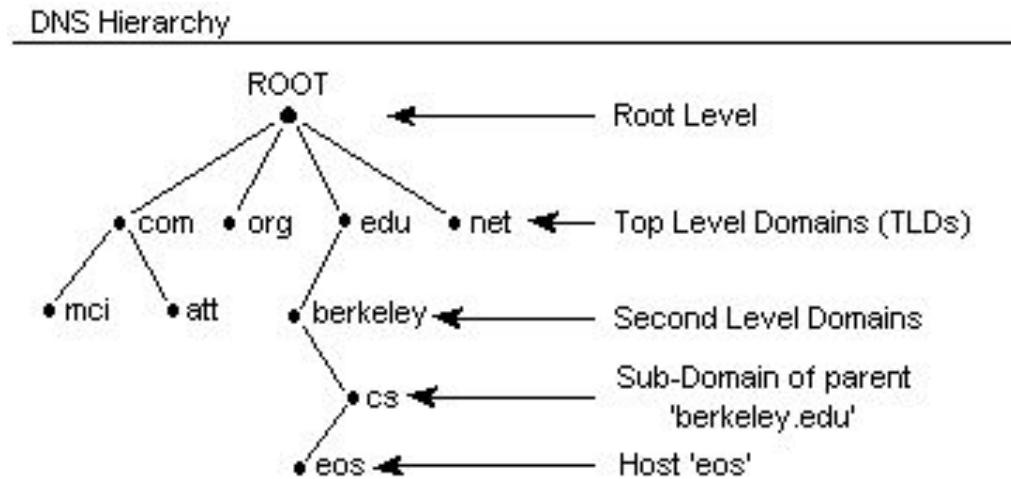
Key Insights and Design Details

- Hierarchical
- Distributed
- Cached



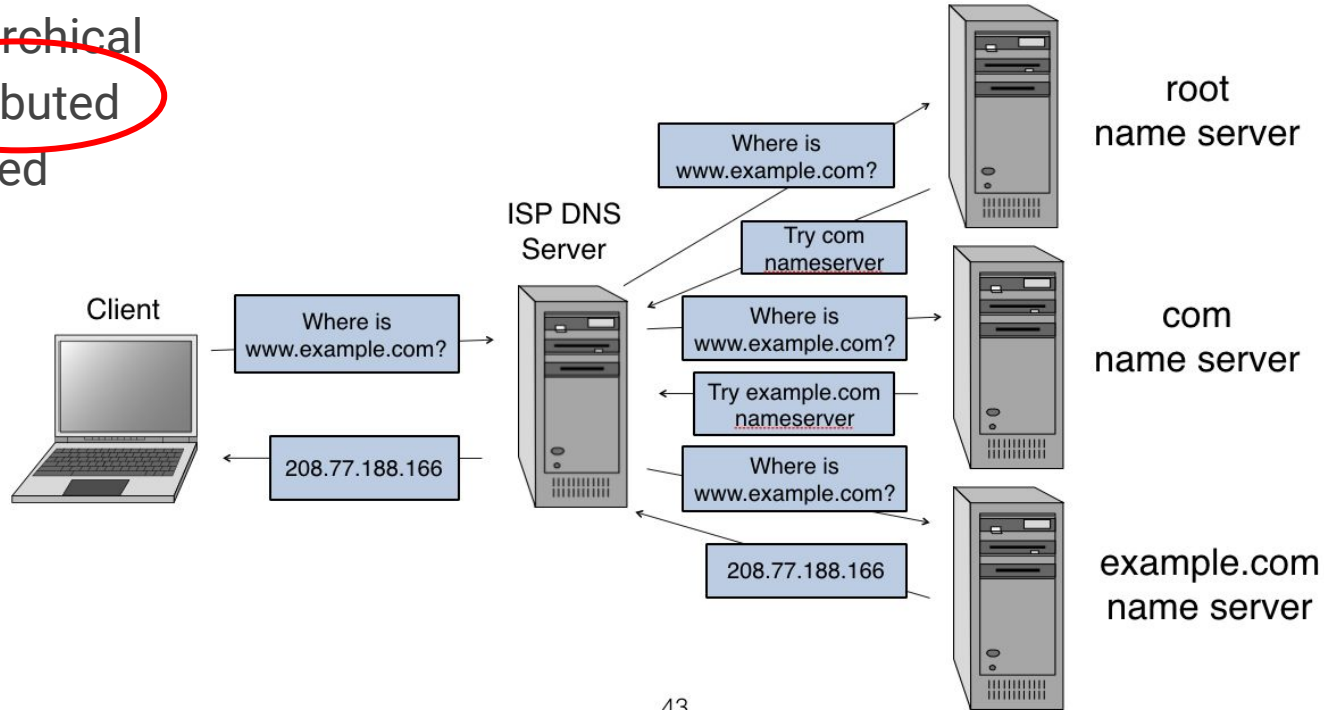
Key Insights and Design Details

- Hierarchical
- Distributed
- Cached



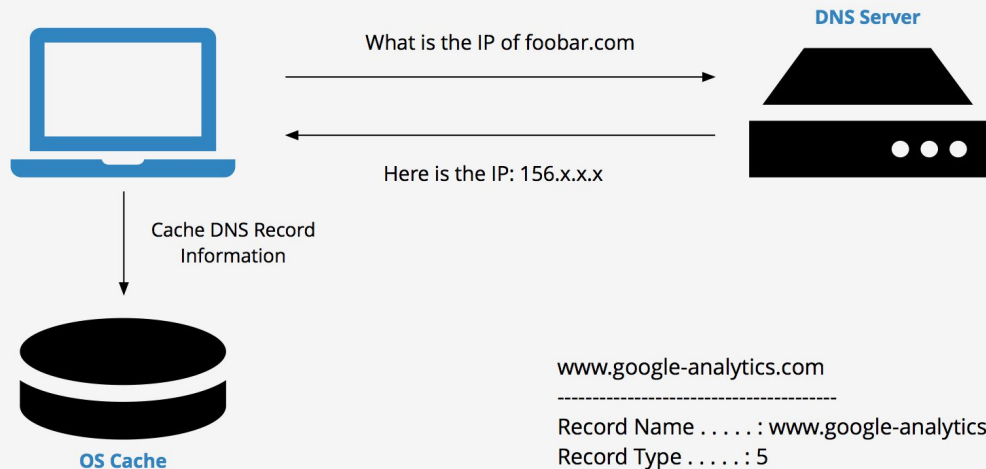
Key Insights and Design Details

- Hierarchical
- **Distributed**
- Cached



Key Insights and Design Details

- Hierarchical
- Distributed
- **Cached**



DNS Cache

www.google-analytics.com

Record Name : www.google-analytics.com

Record Type : 5

Time To Live : 104

Data Length : 4

Section : Answer

CNAME Record : www-google-analytics.l.google.com

Results / Key Insights

- The form and content of information was well understood
 - Nope!
- Network performance wasn't great
- Measurements were very hard to do
- Expected negative responses to be rare
 - Needed to add caching for this as well



Results / Lessons Learned

- Successes

- Variable depth hierarchy (self organizing, different size orgs, easy encapsulation)
- Organizational structuring of names
- Datagram access
- Additional section processing
- Caching!
- Mail address cooperation

- Shortcomings

- Type and class growth (expected huge demand, got almost none)
- Transition wasn't easy
- Distribution of control \neq Distribution of expertise / responsibility

Types of Records

...

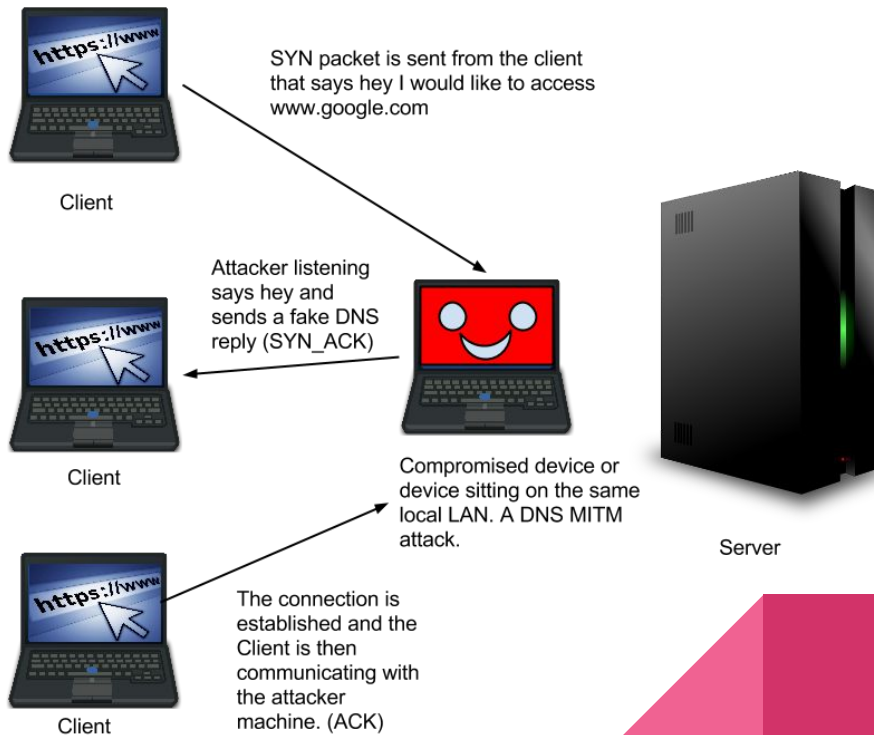


Shorthand	Resource Record	Description
A	Address Mapping records	IPv4
AAAA	IP Version 6 Address records	IPv6
CNAME	Canonical Name records	The CNAME record specifies a domain name that has to be queried in order to resolve the original DNS query
HINFO	Host Information records	General information about a host - not typically used on public servers
ISDN	Integrated Services Digital Network records	Telephone number
MX	Mail exchanger record	Mail exchange server for a DNS domain name
NS	Name Server records	The NS record specifies an authoritative name server for given host
PTR	Reverse-lookup Pointer records	Domain name -> IP address
SOA	Start of Authority records	The record specifies core information about a DNS zone
TXT	Text records	The text record can hold arbitrary non-formatted text string. Typically, the record is used by Sender Policy Framework (SPF) to prevent fake emails to appear to be sent by you

DNS & Security

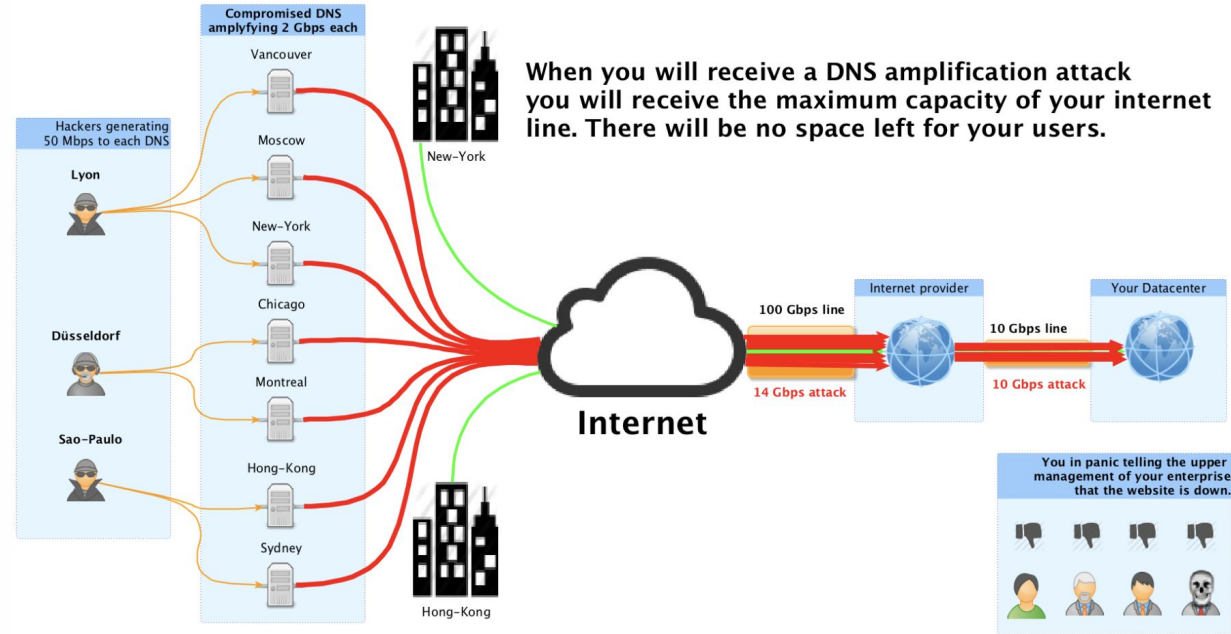
- DNS Spoofing

DNS spoofing TCP Three Way Handshake (SYN, SYN-ACK, ACK)




DNS & Security

- DNS amplification for DDoS attacks



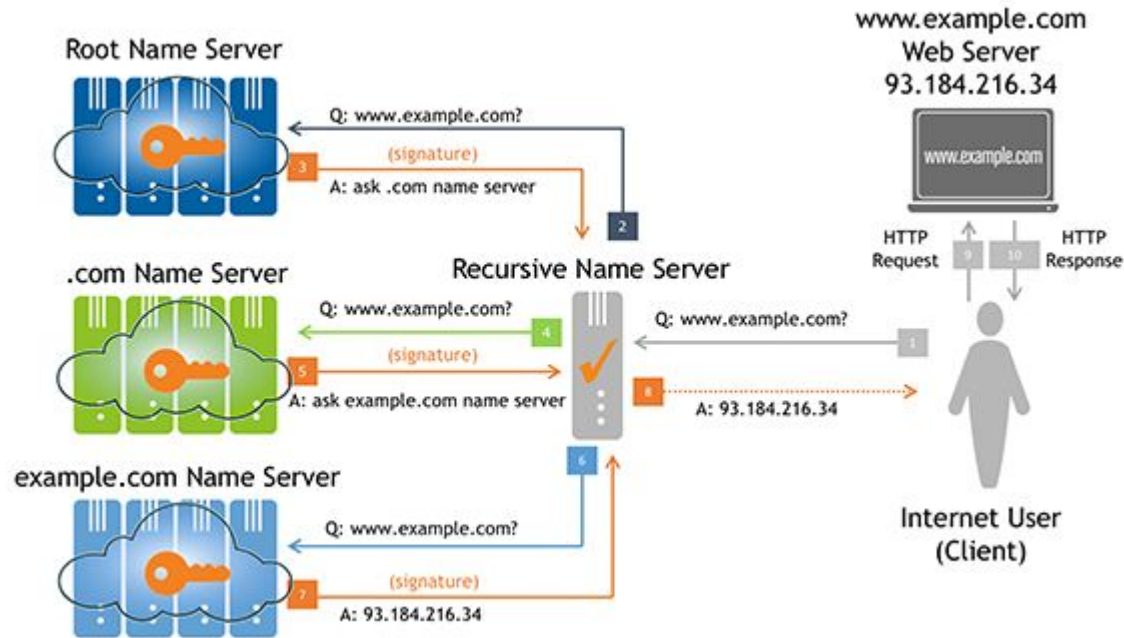
DNS & Security

- DNS cache poisoning



DNS & Security

- DNSSEC



All the Info


Standards[\[edit\]](#)

The Domain Name System is defined by [Request for Comments](#) (RFC) documents published by the [Internet Engineering Task Force](#) ([Internet standards](#)). The following is a list of RFCs that define the DNS protocol.

- [RFC 1034](#), *Domain Names - Concepts and Facilities*
- [RFC 1035](#), *Domain Names - Implementation and Specification*
- [RFC 1123](#), *Requirements for Internet Hosts—Application and Support*
- [RFC 1995](#), *Incremental Zone Transfer in DNS*
- [RFC 1996](#), *A Mechanism for Prompt Notification of Zone Changes (DNS NOTIFY)*
- [RFC 2136](#), *Dynamic Updates in the domain name system (DNS UPDATE)*
- [RFC 2181](#), *Clarifications to the DNS Specification*
- [RFC 2308](#), *Negative Caching of DNS Queries (DNS NCACHE)*
- [RFC 2672](#), *Non-Terminal DNS Name Redirection*
- [RFC 2845](#), *Secret Key Transaction Authentication for DNS (TSIG)*
- [RFC 3225](#), *Indicating Resolver Support of DNSSEC*
- [RFC 3226](#), *DNSSEC and IPv6 A6 aware server/resolver message size requirements*
- [RFC 3596](#), *DNS Extensions to Support IP Version 6*
- [RFC 3597](#), *Handling of Unknown DNS Resource Record (RR) Types*

All the Info


Standards[\[edit\]](#)

- ...
 - [RFC 4343](#), *Domain Name System (DNS) Case Insensitivity Clarification*
 - [RFC 4592](#), *The Role of Wildcards in the Domain Name System*
 - [RFC 4635](#), *HMAC SHA TSIG Algorithm Identifiers*
 - [RFC 5001](#), *DNS Name Server Identifier (NSID) Option*
 - [RFC 5452](#), *Measures for Making DNS More Resilient against Forged Answers*
 - [RFC 5890](#), *Internationalized Domain Names for Applications (IDNA): Definitions and Document Framework*
 - [RFC 5891](#), *Internationalized Domain Names in Applications (IDNA): Protocol*
 - [RFC 5892](#), *The Unicode Code Points and Internationalized Domain Names for Applications (IDNA)*
 - [RFC 5893](#), *Right-to-Left Scripts for Internationalized Domain Names for Applications (IDNA)*
 - [RFC 6891](#), *Extension Mechanisms for DNS (EDNS0)*
 - [RFC 7766](#), *DNS Transport over TCP - Implementation Requirements*
- 

All the Info

Standards

Security

- [RFC 4033](#), *DNS Security Introduction and Requirements*
 - [RFC 4034](#), *Resource Records for the DNS Security Extensions*
 - [RFC 4035](#), *Protocol Modifications for the DNS Security Extensions*
 - [RFC 4509](#), *Use of SHA-256 in DNSSEC Delegation Signer (DS) Resource Records*
 - [RFC 4470](#), *Minimally Covering NSEC Records and DNSSEC On-line Signing*
 - [RFC 5011](#), *Automated Updates of DNS Security (DNSSEC) Trust Anchors*
 - [RFC 5155](#), *DNS Security (DNSSEC) Hashed Authenticated Denial of Existence*
 - [RFC 5702](#), *Use of SHA-2 Algorithms with RSA in DNSKEY and RRSIG Resource Records for DNSSEC*
 - [RFC 5910](#), *Domain Name System (DNS) Security Extensions Mapping for the Extensible Provisioning Protocol (EPP)*
 - [RFC 5933](#), *Use of GOST Signature Algorithms in DNSKEY and RRSIG Resource Records for DNSSEC*
 - [RFC 7858](#), *Specification for DNS over Transport Layer Security (TLS)*
- 

All the Info

Experimental

- [RFC 1183](#), *New DNS RR Definitions*

Best Current Practices

- [RFC 2182](#), *Selection and Operation of Secondary DNS Servers* (BCP 16)
- [RFC 2317](#), *Classless IN-ADDR.ARPA delegation* (BCP 20)
- [RFC 5625](#), *DNS Proxy Implementation Guidelines* (BCP 152)
- [RFC 6895](#), *Domain Name System (DNS) IANA Considerations* (BCP 42)
- [RFC 7720](#), *DNS Root Name Service Protocol and Deployment Requirements* (BCP 40)

Informational

These RFCs are advisory in nature, but may provide useful information despite defining neither a standard or BCP. ([RFC 1796](#))

- [RFC 1178](#), *Choosing a Name for Your Computer* (FYI 5)
- [RFC 1591](#), *Domain Name System Structure and Delegation*
- [RFC 1912](#), *Common DNS Operational and Configuration Errors*
- [RFC 2100](#), *The Naming of Hosts*


All the Info

Informational...

- [RFC 3696](#), *Application Techniques for Checking and Transformation of Names*
- [RFC 4892](#), *Requirements for a Mechanism Identifying a Name Server Instance*
- [RFC 5894](#), *Internationalized Domain Names for Applications (IDNA):Background, Explanation, and Rationale*
- [RFC 5895](#), *Mapping Characters for Internationalized Domain Names in Applications (IDNA) 2008*
- [RFC 7626](#), *DNS Privacy Considerations*
- [RFC 7706](#), *Decreasing Access Time to Root Servers by Running One on Loopback*

Unknown

These RFCs have an official status of [Unknown](#), but due to their age are not clearly labeled as such.

- [RFC 920](#), *Domain Requirements* – Specified original top-level domains
 - [RFC 1032](#), *Domain Administrators Guide*
 - [RFC 1033](#), *Domain Administrators Operations Guide*
 - [RFC 1101](#), *DNS Encodings of Network Names and Other Types*
- 

Designing a Global Name Service

Butler W. Lampson

Published in the proceedings of the 1986 Conference on Principles of Distributed Computing

Presented by
Ramakrishnan Sundara Raman

What is the problem?

In a large distributed computing system which continuously evolves and grows, there needs to be an efficient decentralized system for managing resource location and authentication.

Name Service

- Think about the old telephone system
- In simple terms, a name service maps names to (a set of) values
- Most prominent example today: DNS
- The name service design in this work focuses on mapping names to a set of labeled properties in a distributed system.

Design Considerations

- Large size
- Long lifetime
- High availability
- Fault isolation
- Tolerance of mistrust

Best Design: Hierarchical System

Two Levels

Client level:


- Hierarchical names and values
- Operations for reading and updating
- No idea about replication

Administrative level:

- All copies are visible
- Mechanisms for synchronizing
- copies

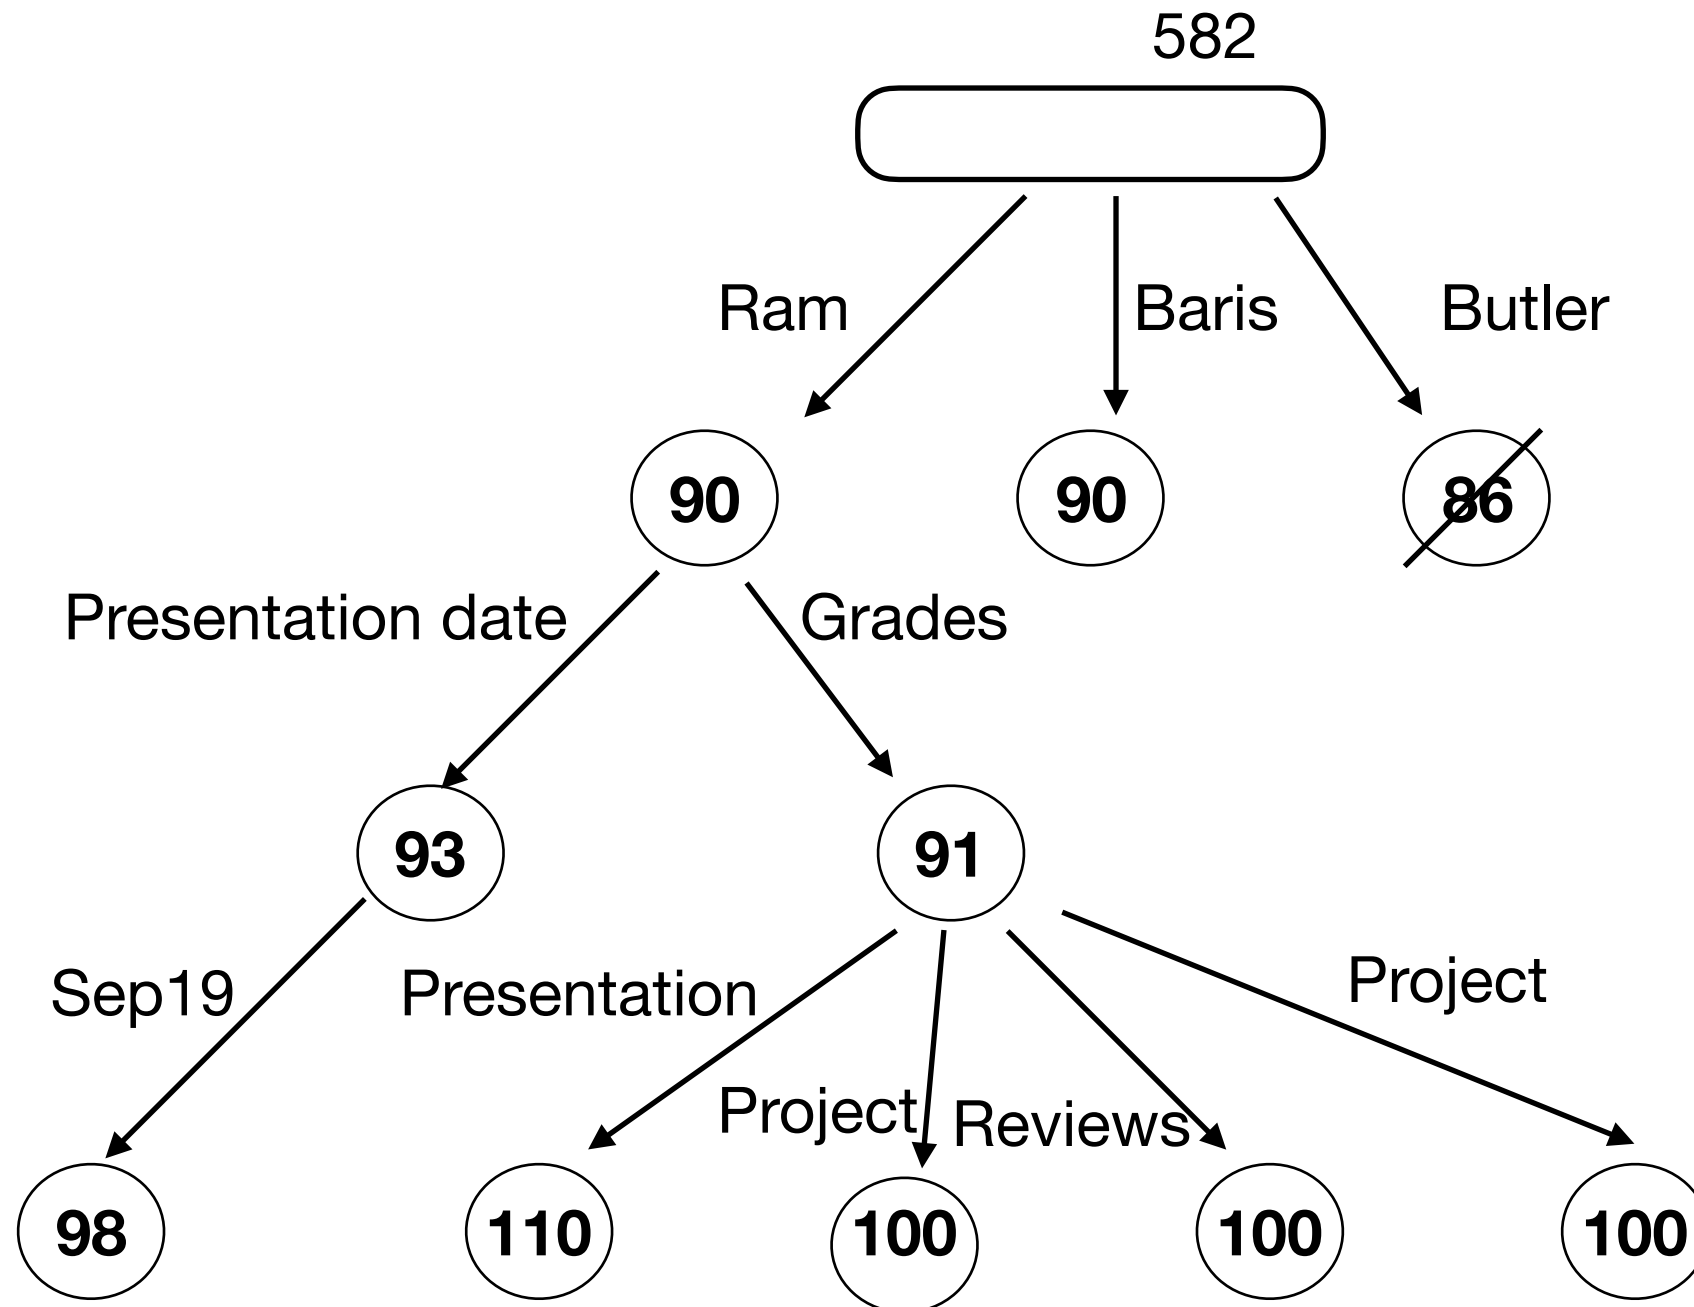
Client Level

Directory Tree



- Directory structure much like filesystems
- Full Name (FN) of directory 582 - Fall2018/EECS/582

Values in a directory

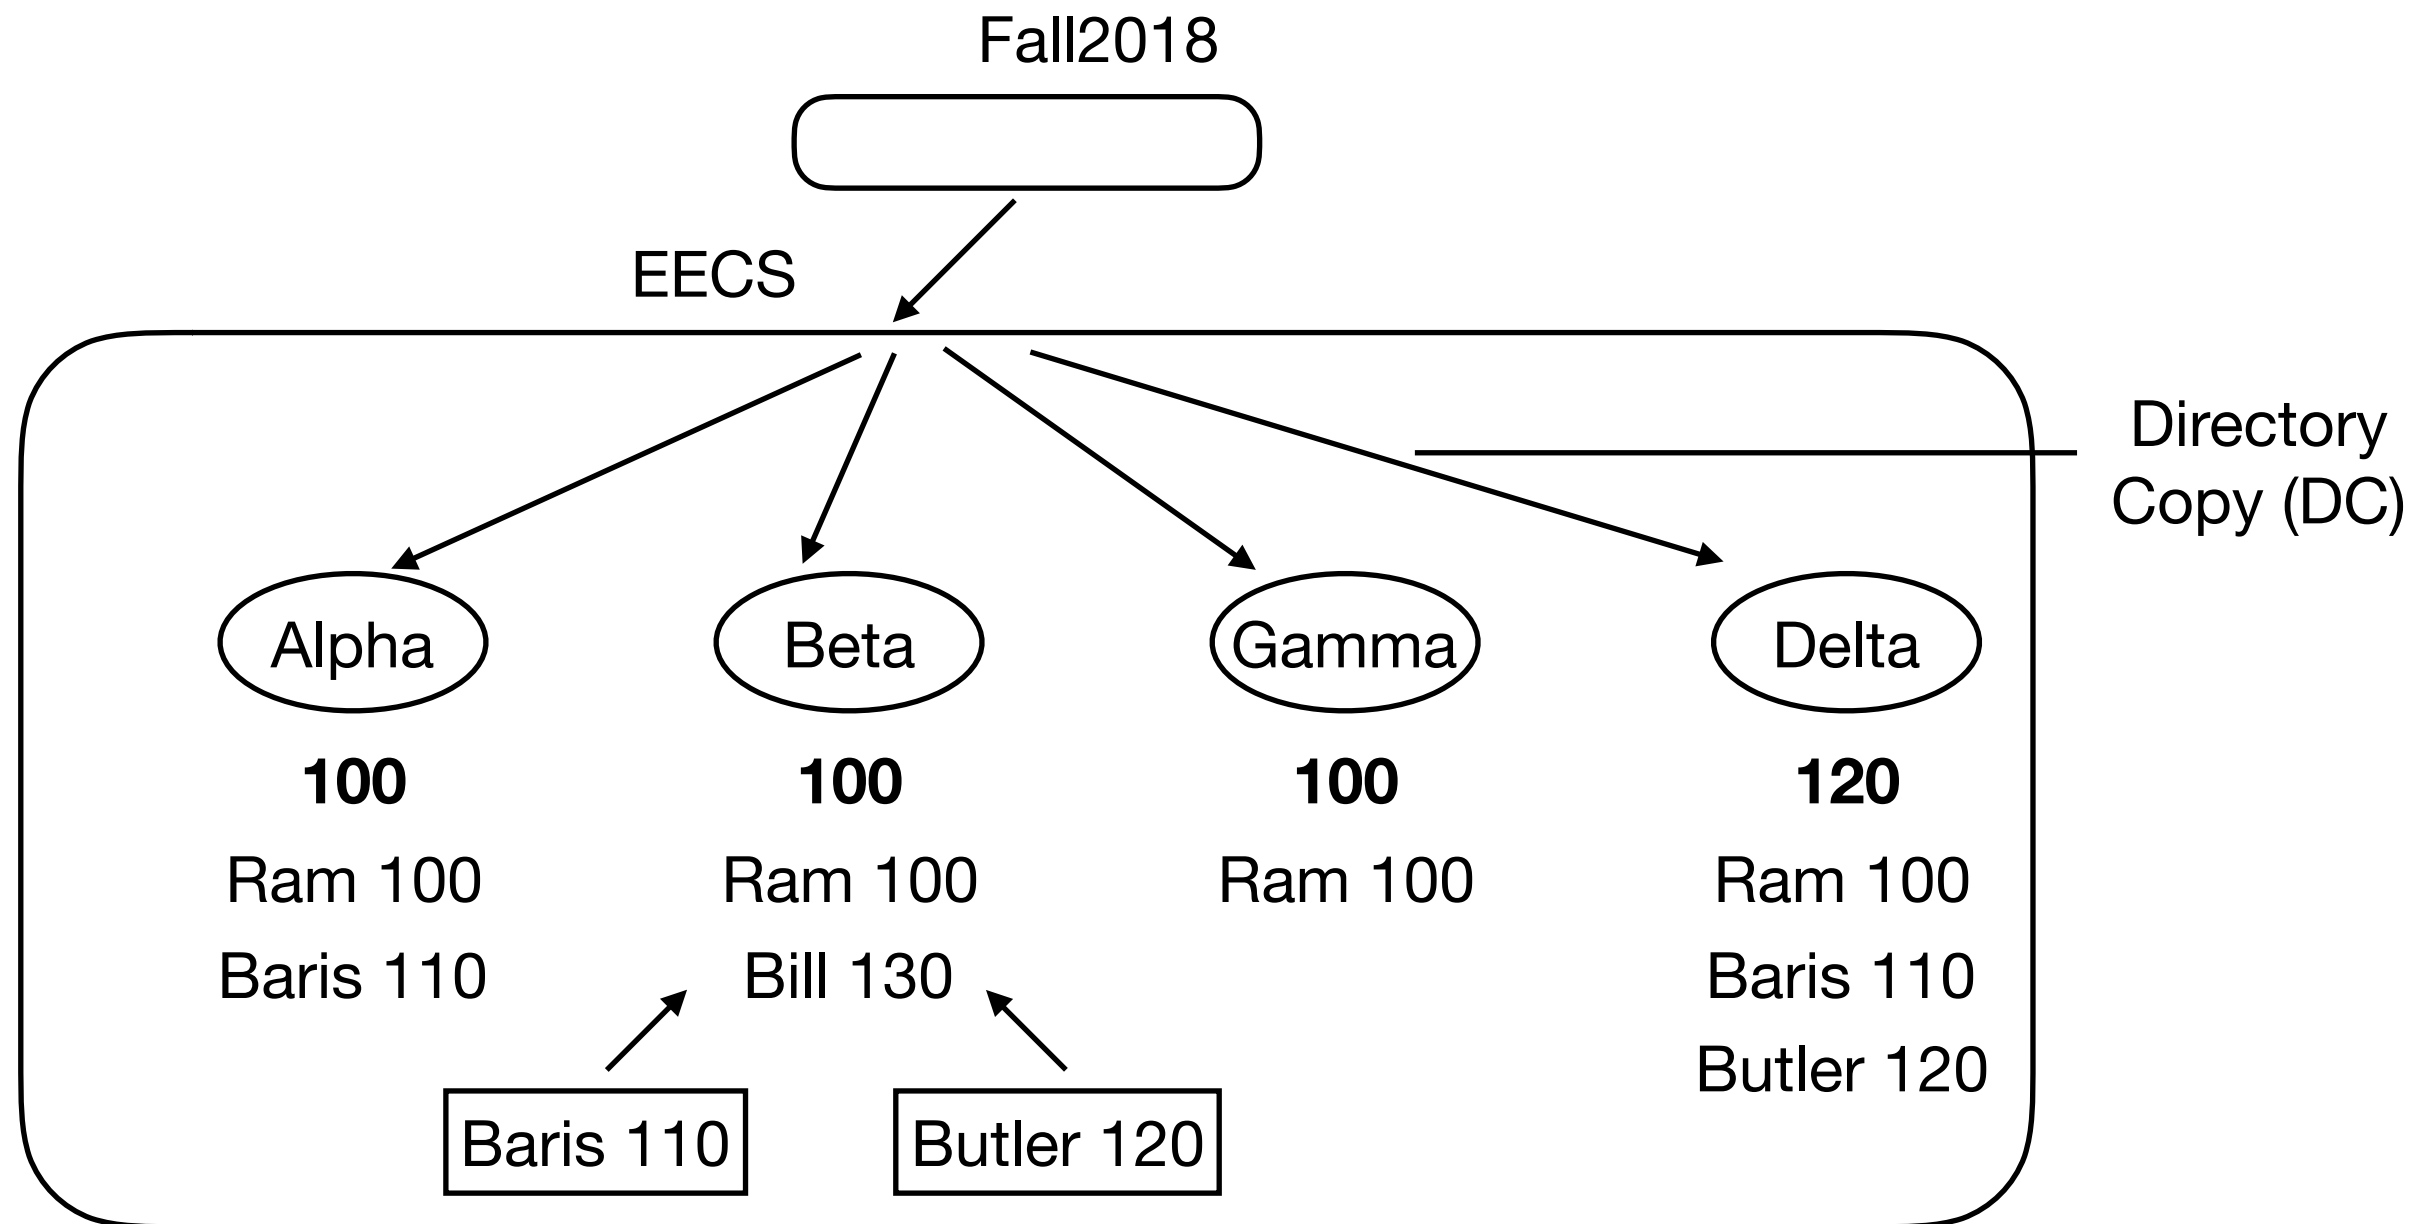


Access Control and Authentication

- Triplet (*Fall2018/EECS/582/Baris*, *Ram*/*, write) gives principal Baris write access to the subtree which is the value of *Ram*
- Authentication is based on secure channel between caller of operation and implementor
- Authentication can be extended with a certificate

Administrative Level

Directory copies

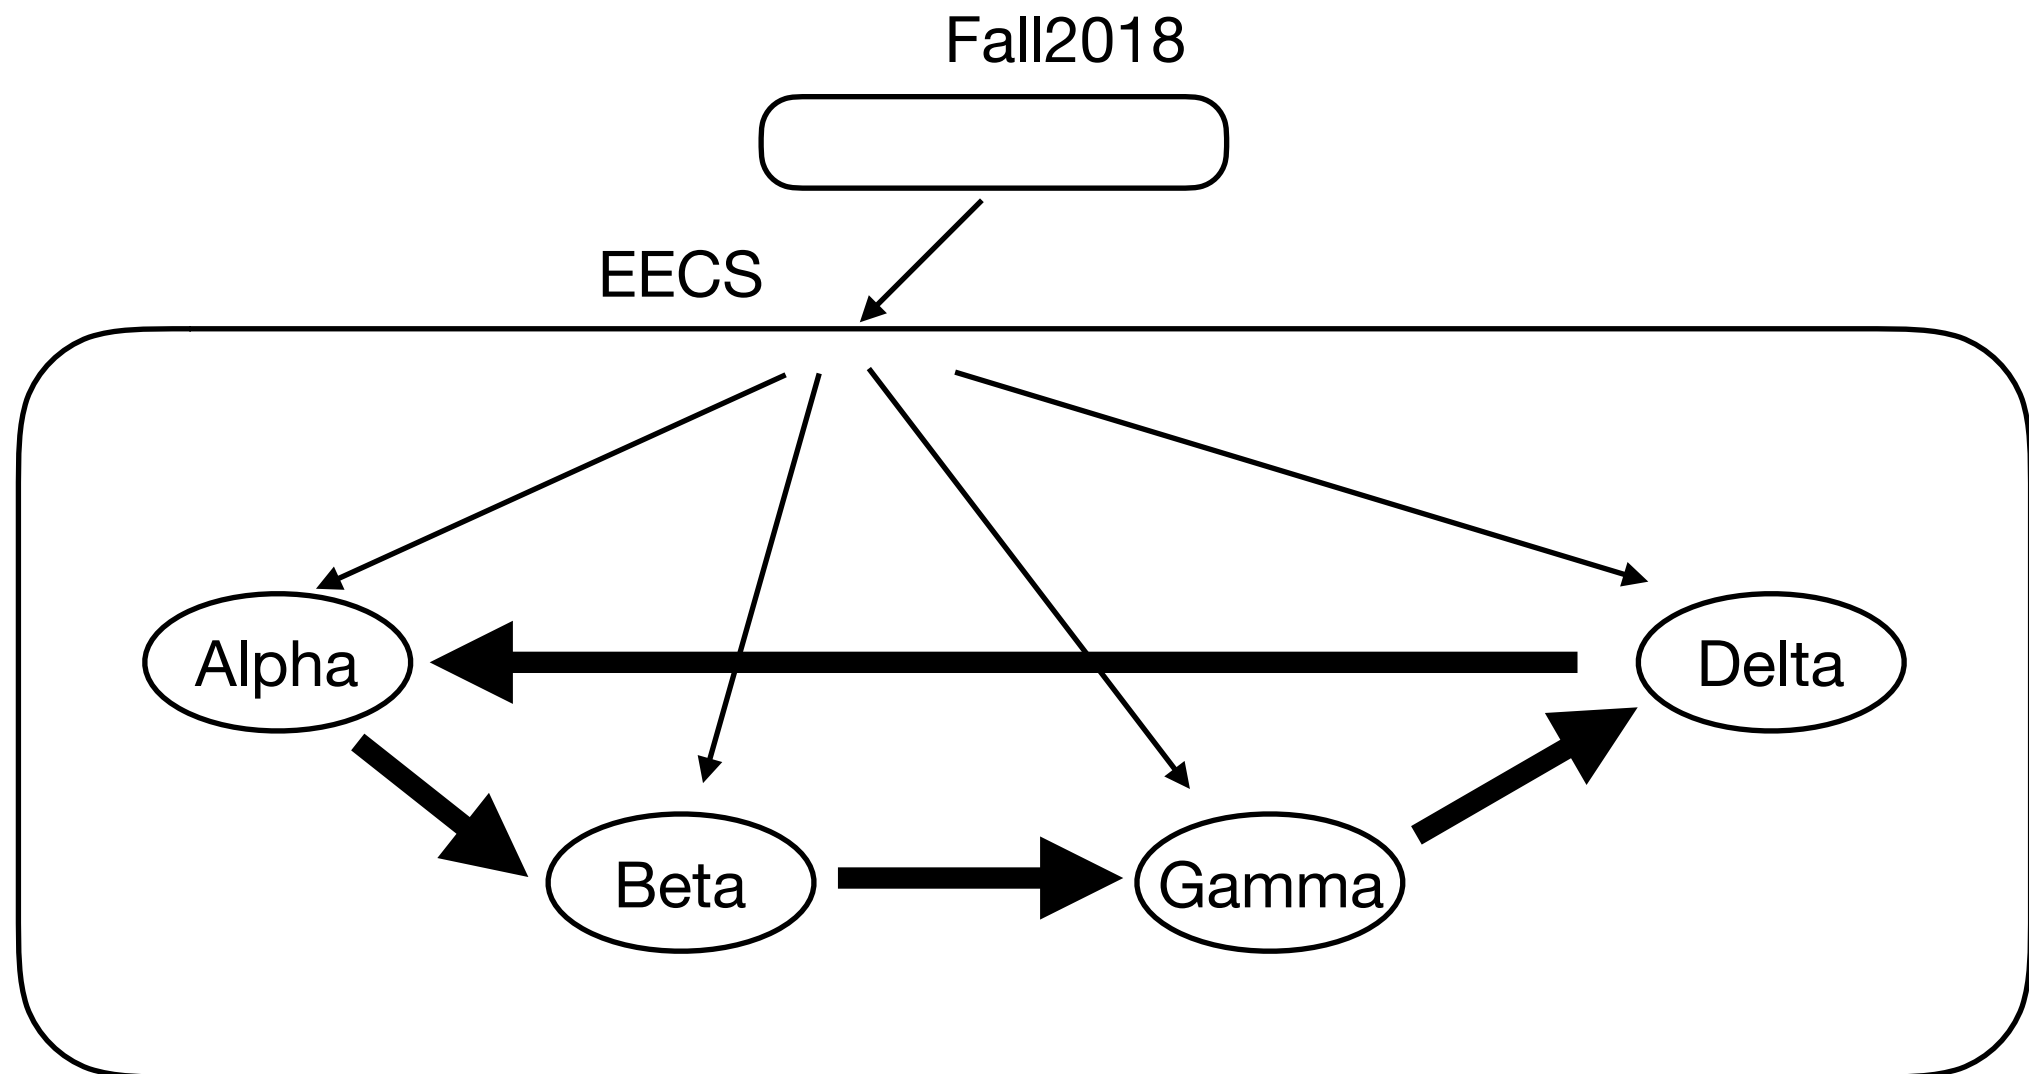


- DR now includes DI and DC.

Sweep updates

- Four updates with timestamps 100, 110, 120 and 130
- lastSweep of alpha, beta and gamma is 100, and lastSweep of delta is 120
- Update originates at one DC at sweepTS -> visits every other DC and collects all updates -> Write updates back to each DC
- All updates collected till sweepTS, sets lastSweep of all DCs to sweepTS
- Updates also possible through direct messages from one DC to another but this is not reliable


Sweep Ring



The Name Space


Growth of name space

- Inside each directory, names can be generated without considering any other directory.
- Combining exiting name services?



Restructuring

Fall 2018



EECS


LSE

#13

#14

If LSE came under EECS

Fall 2018



EECS

LSE


#13

#20/DEC/LSE

LSE

#14

Caching



#10/EECS = #13
Valid until 20 Sep 2018

#13/LSE = #14
Valid until 17 Oct 2018

#10/EECS/LSE = #14
Valid until 20 Sep 2018

The name service itself and authentication are also candidates for caching

The name service interface

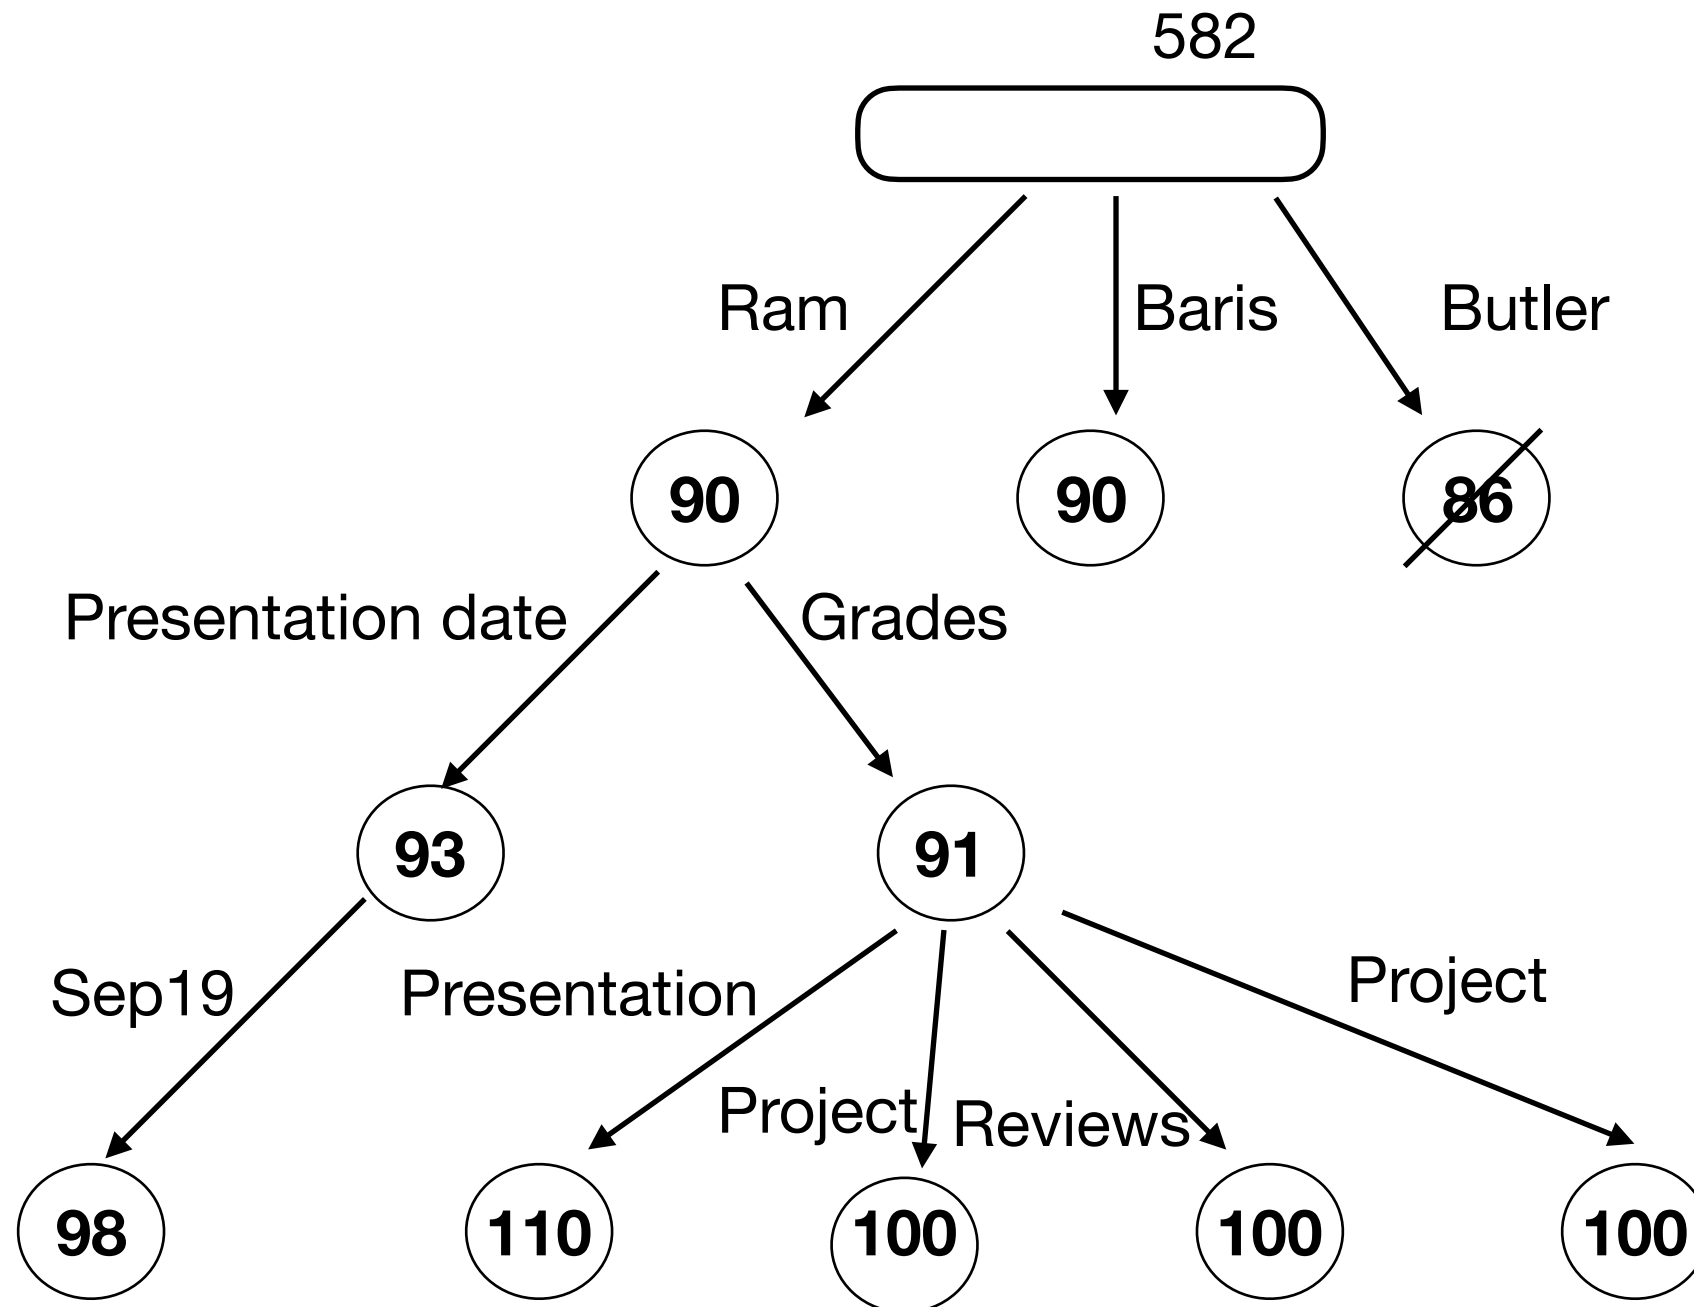
- The interface is based on remote procedure calls
- Example interfaces for
 - Reading values - GetValue(VD) -> Tree
 - Directories - NewD(SN,FN) -> DI
 - Directory Copies - Sweep(FN,SN) -> time-stamp
 - Servers - NewServer() -> Server Address

Specifications

Name Service Specifications

- Model of computation - Interleaving of atomic actions in separate processes
- Collection of predicates
- Two kinds - Invariants and post-conditions
- Predicates can be both variables in the program and auxiliary variables.
- DB - Function DI -> directory - Defines current content of database

Values and Updates



Values and Updates


- A value v is modified if an update operation names a node by giving path p (including TS) and says whether it is present or absent
- Update operation:
 1. Find arcs and nodes in v for longest prefix of $p - v'$
 2. If prefix is all of p and operation = present, then do nothing. Otherwise, replace v' with absent leaf node
 3. Otherwise, next element of $p = (l, ts)$
 4. If there is an arc labeled l from v' to a node with timestamp $> ts$, do nothing. Otherwise, remove subtree from v' labeled by l , and add subtree defined by p
- This specification is total, commutative and idempotent, so the order of applying updates doesn't matter.

Directories (D)

- D1 - The result of a read operation on directory d with identifier $d.di$ is determined by the state of d after S set of updates are applied on D where S includes
 1. All updates with timestamps less than $d.lastSweep$
 2. Arbitrary set of updates with timestamp greater than $d.lastSweep$
- D2 - Looking up FN $di/n1/.../nk$ yields a directory if either di is the root and each n_j is defined in $di/n1/..n(j-1)$ or di is in the root's well-known table and $di'/.../ni'/n1/...nk$ satisfies the previous condition

Directories (D)

- Back-references (BRs) are child-parent arcs in the tree and is the primary invariant in ensuring the structure of the tree, while the DRs are secondary



- D3 - The D's defined in DB form a tree rooted in root whose arcs are DRs that are the reverse of the BR backpointers.
- D4 - Each DR is pointed to by a BR with a longer TX.

Directory Copies (DC)

- DC1 - The database value of a directory $DB(di) = \text{Union}(\text{Updates in all DCs})$
- DC2 - There is always at least one complete ring in the set of DCs for a directory
- DC3 - There are never two complete non-intersecting rings

Servers (S)

- Can reference server names through directory references. This might create infinite loops. Server invariants solve this problem by using a boolean inSN in each directory. A directory cannot be part of a server name unless its inSN is true.
- S1 - All d's from root to an entry that stores a server address has d.inSN = true
- S2 - If d.inSN is true, then either d is a root, or a copy of d is stored on a server with a name shorter than any direct name of d
- S3 - Every server s either stores the root, or s.up is a shorter name of another server, and s stores a copy of the directory for s.up

Conclusion

Formal specifications for functional requirements of a name service that supports growth and evolution of the name space.

Discussion

Strengths

- Decentralized management -> High replication -> large size and high availability
- Hierarchical design supports growth, change in structure and fault isolation
- Authentication and access control mechanisms are easy to implement
- Simple interface
- Exact specifications for the name service

Weaknesses

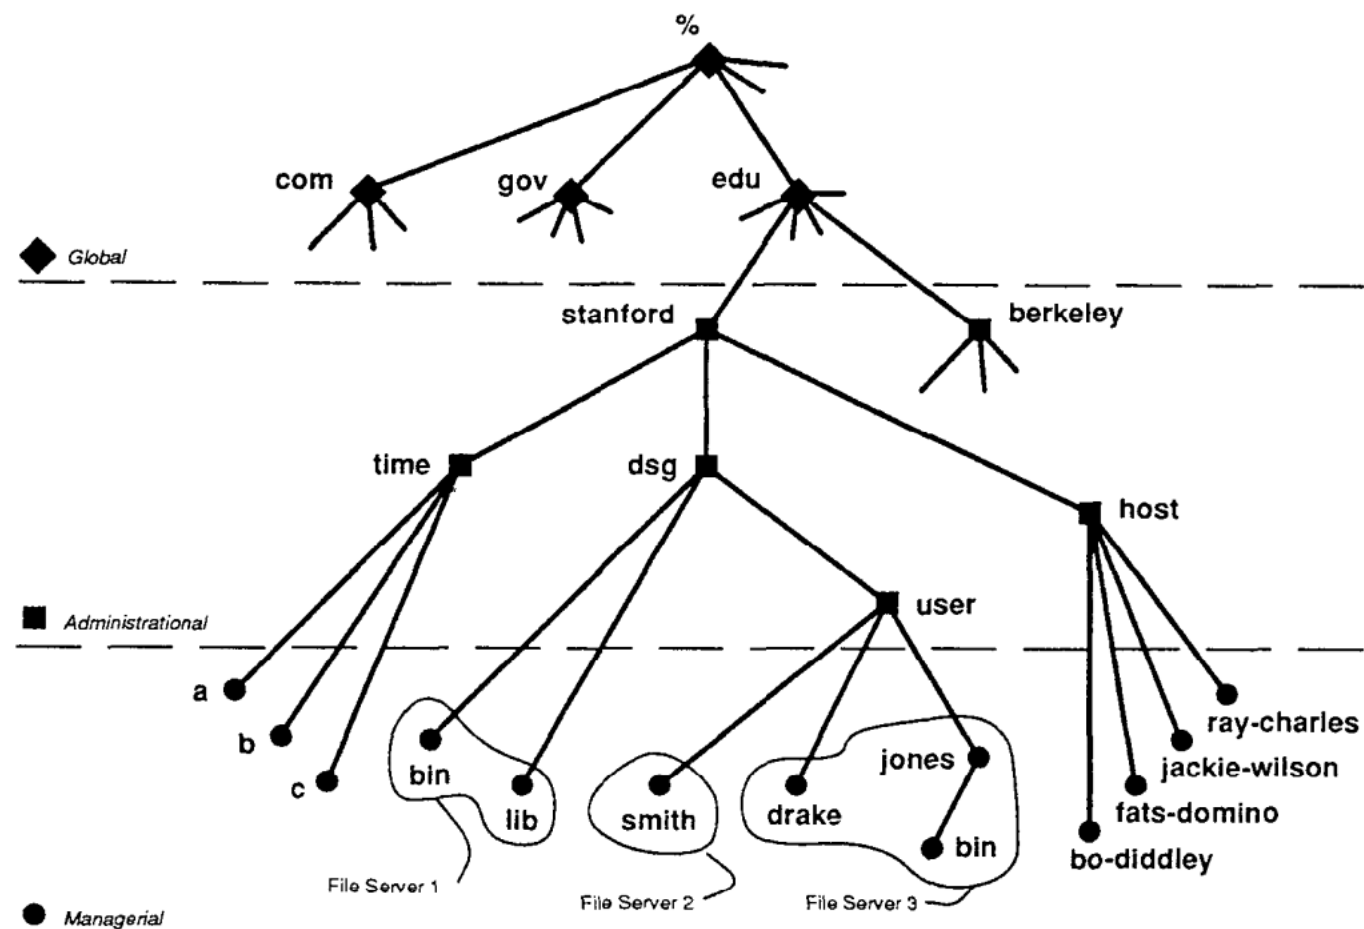
- Frequent changes may cause database to be outdated
- Cost of update propagation might go up as size increases
- Security issues
- Reforming rings (starting a new epoch) needs an administrator, so cannot be fully automated

Open questions

- Why can't direct message updates be made reliable?
- How do current name services like the DNS compare to this design?
- Can this design be practically implemented for efficient name lookup?
- Space and computation requirements on the servers?
How many must there be to enable fast lookup?

More on name services

“Decentralizing a Global Naming Services for Improved Performance and Fault Tolerance,” (Cheriton & Mann, 1989)



More on name services

- P. Mockapetris and K. J. Dunlap. 1988. **Development of the domain name system**. In Symposium proceedings on Communications architectures and protocols (SIGCOMM '88), Vinton Cerf (Ed.). ACM, New York, NY, USA, 123-133. DOI=<http://dx.doi.org/10.1145/52324.5233>
- More recently, there has been work on developing name services designed for high mobility and security.
- A. Venkataramani *et al.*, "**Design requirements of a global name service for a mobility-centric, trustworthy internetwork**," *2013 Fifth International Conference on Communication Systems and Networks (COMSNETS)*, Bangalore, 2013, pp. 1-9.