

AutoFDO: Automatic Feedback-Directed Optimization for Warehouse-Scale Applications

Dehao Chen, David Xinliang Li, Tipp Moseley

Presenter: Yifan Zhao

FDO: Feedback-directed Optimization

- Steps:
 - Compile with instrumentation
 - Run a benchmark to generate *representative profile*
 - Recompile with the profile
- Representative profile requires representative input
- Many other problems with Google datacenter

Problems

- Binaries change rapidly between releases
 - Best if tolerating stale profiles
- Programs deal with sensitive data
 - Profile could expose everything to an attack
- Instrumentation overhead triggers timeouts on servers
 - Instrumented code runs differently from “release” code
 - Lead to even less representative profiles

Solution and Key Insights

- Hardware sampling
 - Instruction frequency
 - Branch taken frequency

Line	Offset	Source:	Binary:
-----			-----
#1	#0	foo() {	foo():
#2	#1	if (cond)	0x670: if_stmt.binary;
#3	#2	foo_stmt;	0x675: foo_stmt.binary;
#4	#3	}	
#5			
#6	#0	bar() {	bar():
#7	#1	bar_stmt;	0x690: bar_stmt.binary;
#8	#2	foo();	0x69d: if_stmt.binary;
#9	#3	}	0x6a2: foo_stmt.binary;

Binary Level Profile:

Instruction Address	Sample Count

0x670	50
0x675	10
0x690	200
0x69d	200
0x6a2	100

Solution and Key Insights

- Hardware sampling
- Profile mapping
 - Source location
 - Extended source location (ESL)
 - ESL on the tree

Binary:

foo():

0x670: if_stmt.binary;

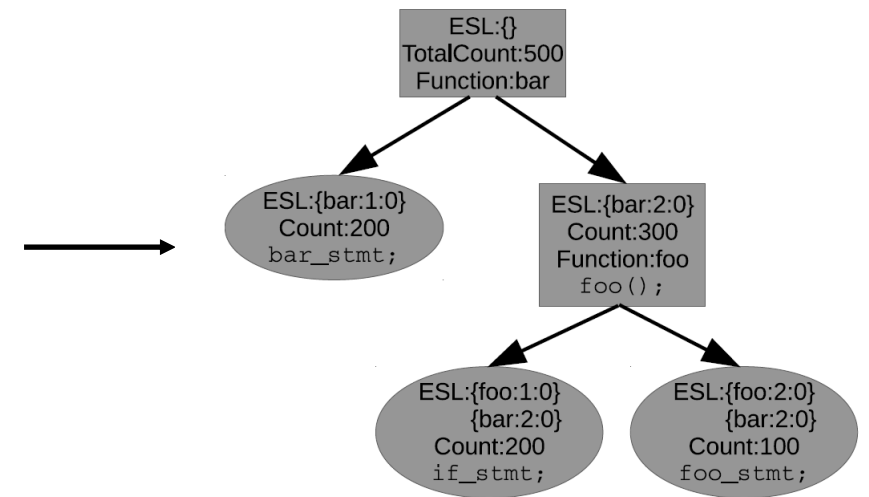

0x675: foo_stmt.binary;

bar():

0x690: bar_stmt.binary;

0x69d: if_stmt.binary;

0x6a2: foo_stmt.binary;



Solution and Key Insights


- Hardware sampling
- Profile mapping
- Top-first profile collection
 - Get a 10-second profile from 10% of machines each day
 - Low coverage in general -- 0.001% of observable cycles
 - Enough coverage for top workloads

Evaluation

- AutoFDO, vs FDO
 - Achieve >90% of the FDO speedup
 - Gap mostly due to inaccurate debug info (not sampling)
 - Mostly happen to loop nesting


Application	FDO	AutoFDO	Ratio
server	17.61%	15.89%	90.23%
graph1	14.68%	14.04%	95.65%
graph2	7.16%	6.27%	87.50%
machine learning1	8.92%	8.46%	94.85%
machine learning2	7.09%	6.60%	93.06%
encoder	8.63%	3.31%	38.37%
protobuf	16.96%	14.40%	84.94%
artificial intelligence1	10.12%	10.12%	100.00%
artificial intelligence2	13.24%	11.33%	85.61%
data mining	20.48%	15.54%	75.86%
mean	12.40%	10.52%	84.84%

Mostly
nested loops



Evaluation

- AutoFDO, vs FDO
- Iterative AutoFDO
 - Performance varies in some apps
 - Again due to loop nesting
 - Usually not improving beyond 3 iterations
 - After all hot callees are inlined



Evaluation

- AutoFDO, vs FDO
- Iterative AutoFDO
- Stale profiles
 - Line number in function outperforms line number in file
 - Inserting a function does not break the former
 - 50% of the speedup with 6-months old profile

Strengths and Weaknesses

- Design strengths
 - Handling degraded debugging information
 - Integration into compiler
 - Deferring symbolization
 - Support of iterative FDO
- Shortcomings
 - Obscuring bugs in apps (more than a typical -O2 flag)
 - If stability comes before average performance
 - Breaking *release integration*

Open Questions

- AutoFDO is *automated* FDO, not quite *advanced* FDO
 - Issues of traditional FDO remains
 - Security, input dependence, etc.
- AutoFDO samples and transforms text segment
 - How about extending FDO to data transformation
- Cross-module optimization
 - AutoLIPO / ThinLTO