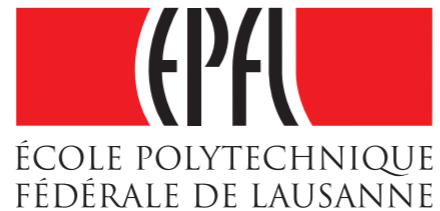


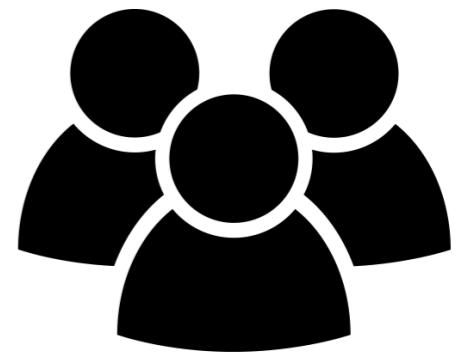
# Failure Sketching: A Technique for Automated Root Cause Diagnosis of In-Production Failures

Baris Kasikci, Benjamin Schubert, Cristiano Pereira,  
Gilles Pokam, George Candea



# Debugging In-Production Software Failures Today

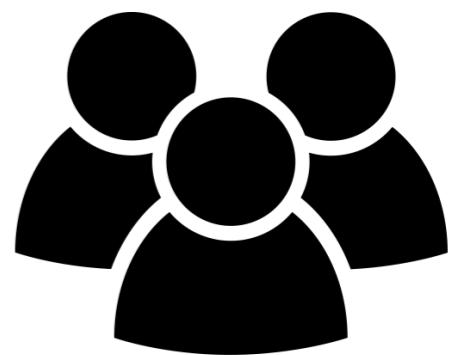
# Debugging In-Production Software Failures Today



# Debugging In-Production Software Failures Today

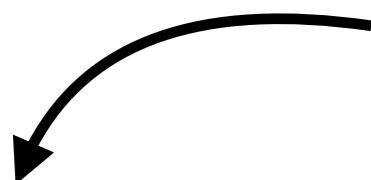


```
#0 0x00007f51abae820b in raise (sig=11) at ../nptl/
sysdeps/unix/sysv/linux/pt-raise.c:37
#1 0x00000000042d289 in ap_buffered_log_writer
(r=0x7f51a40053d0, handle=0x20eeba0,
strs=0x7f51a4003578, strl=0x7f51a40035e8, nelts=14,
len=82) at mod_log_config.c:1368
#2 0x00000000042b10d in config_log_transaction
(r=0x7f51a40053d0, cls=0x20b9d50,
default_format=0x20ee370) at mod_log_config.c:930
#3 0x00000000042aad6 in multi_log_transaction
(r=0x7f51a40053d0) at mod_log_config.c:950
#4 0x00000000046cb2d in ap_run_log_transaction
(r=0x7f51a40053d0) at protocol.c:1563
#5 0x000000000436e81 in ap_process_request
(r=0x7f51a40053d0) at http_request.c:312
#6 0x00000000042e9da in ap_process_http_connection
(c=0x7f519c000b68) at http_core.c:293
#7 0x000000000465cdd in ap_run_process_connection
(c=0x7f519c000b68) at connection.c:85
#8 0x0000000004661f5 in ap_process_connection
(c=0x7f519c000b68, csd=0x7f519c000a20) at
connection.c:211
#9 0x000000000451ba0 in process_socket
(p=0x7f519c0009b8, sock=0x7f519c000a20,
my_child_num=0, my_thread_num=0,
bucket_alloc=0x7f51a4001348) at worker.c:632
#10 0x000000000451221 in worker_thread
(thd=0x210fa90, dummy=0x7f51a40008c0) at worker.c:946
#11 0x00007f51ac87c555 in dummy_worker
(opaque=0x210fa90) at thread.c:127
#12 0x00007f51abae0182 in start_thread
(arg=0x7f51aa8ef700) at pthread_create.c:312
#13 0x00007f51ab80d47d in clone () at ../sysdeps/
unix/sysv/linux/x86_64/clone.S:111
```

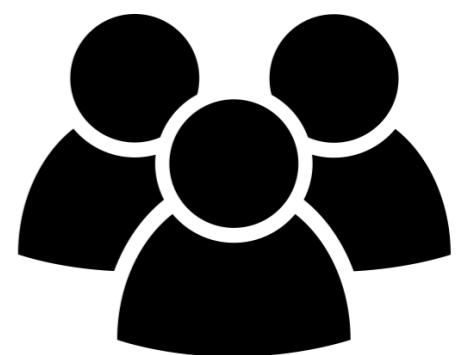


# Debugging In-Production Software Failures Today

Understand root cause



```
#0 0x00007f51abae820b in raise (sig=11) at ../nptl/
sysdeps/unix/sysv/linux/pt-raise.c:37
#1 0x00000000042d289 in ap_buffered_log_writer
(r=0x7f51a40053d0, handle=0x20eeba0,
strs=0x7f51a4003578, strl=0x7f51a40035e8, nelts=14,
len=82) at mod_log_config.c:1368
#2 0x00000000042b10d in config_log_transaction
(r=0x7f51a40053d0, cls=0x20b9d50,
default_format=0x20ee370) at mod_log_config.c:930
#3 0x00000000042aad6 in multi_log_transaction
(r=0x7f51a40053d0) at mod_log_config.c:950
#4 0x00000000046cb2d in ap_run_log_transaction
(r=0x7f51a40053d0) at protocol.c:1563
#5 0x000000000436e81 in ap_process_request
(r=0x7f51a40053d0) at http_request.c:312
#6 0x00000000042e9da in ap_process_http_connection
(c=0x7f519c000b68) at http_core.c:293
#7 0x000000000465cdd in ap_run_process_connection
(c=0x7f519c000b68) at connection.c:85
#8 0x0000000004661f5 in ap_process_connection
(c=0x7f519c000b68, csd=0x7f519c000a20) at
connection.c:211
#9 0x000000000451ba0 in process_socket
(p=0x7f519c0009b8, sock=0x7f519c000a20,
my_child_num=0, my_thread_num=0,
bucket_alloc=0x7f51a4001348) at worker.c:632
#10 0x000000000451221 in worker_thread
(thd=0x210fa90, dummy=0x7f51a40008c0) at worker.c:946
#11 0x00007f51ac87c555 in dummy_worker
(opaque=0x210fa90) at thread.c:127
#12 0x00007f51abae0182 in start_thread
(arg=0x7f51aa8ef700) at pthread_create.c:312
#13 0x00007f51ab80d47d in clone () at ../sysdeps/
unix/sysv/linux/x86_64/clone.S:111
```



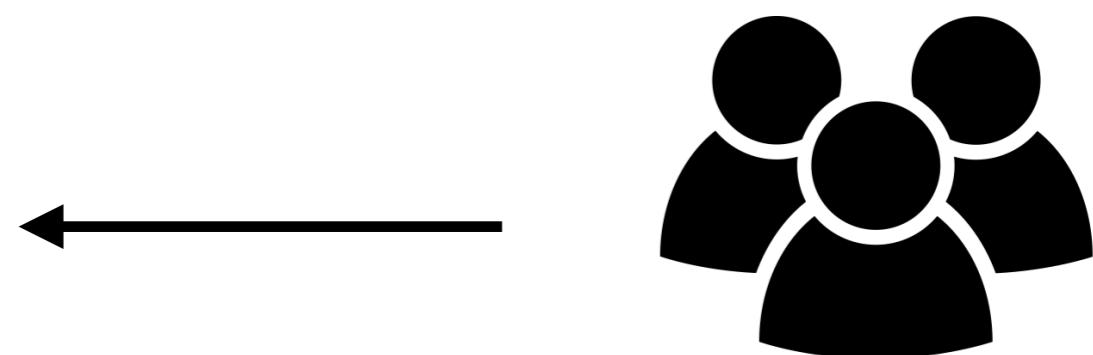
# Debugging In-Production Software Failures Today

Understand root cause



Reproduce the failure

```
#0 0x00007f51abae820b in raise (sig=11) at ../nptl/sysdeps/unix/sysv/linux/pt-raise.c:37
#1 0x00000000042d289 in ap_buffered_log_writer (r=0x7f51a40053d0, handle=0x20eeba0, strrs=0x7f51a4003578, strl=0x7f51a40035e8, nelts=14, len=82) at mod_log_config.c:1368
#2 0x00000000042b10d in config_log_transaction (r=0x7f51a40053d0, cls=0x20b9d50, default_format=0x20eee370) at mod_log_config.c:930
#3 0x00000000042aad6 in multi_log_transaction (r=0x7f51a40053d0) at mod_log_config.c:950
#4 0x00000000046cb2d in ap_run_log_transaction (r=0x7f51a40053d0) at protocol.c:1563
#5 0x000000000436e81 in ap_process_request (r=0x7f51a40053d0) at http_request.c:312
#6 0x00000000042e9da in ap_process_http_connection (c=0x7f519c000b68) at http_core.c:293
#7 0x000000000465cdd in ap_run_process_connection (c=0x7f519c000b68) at connection.c:85
#8 0x0000000004661f5 in ap_process_connection (c=0x7f519c000b68, csd=0x7f519c000a20) at connection.c:211
#9 0x000000000451ba0 in process_socket (p=0x7f519c0009b8, sock=0x7f519c000a20, my_child_num=0, my_thread_num=0, bucket_alloc=0x7f51a4001348) at worker.c:632
#10 0x000000000451221 in worker_thread (thd=0x210fa90, dummy=0x7f51a40008c0) at worker.c:946
#11 0x00007f51ac87c555 in dummy_worker (opaque=0x210fa90) at thread.c:127
#12 0x00007f51abae0182 in start_thread (arg=0x7f51aa8ef700) at pthread_create.c:312
#13 0x00007f51ab80d47d in clone () at ../sysdeps/unix/sysv/linux/x86_64/clone.S:111
```



# Debugging In-Production Software Failures Today

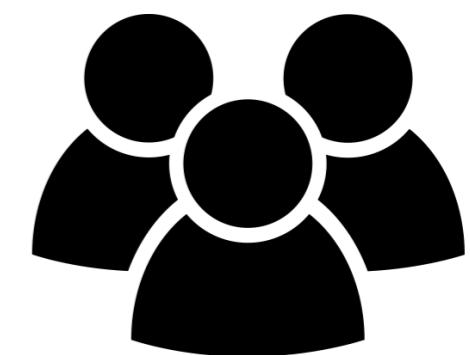
Understand root cause



Reproduce the failure

```
#0 0x00007f51abae820b in raise (sig=11) at ../nptl/sysdeps/unix/sysv/linux/pt-raise.c:37
#1 0x00000000042d289 in ap_buffered_log_writer
(r=0x7f51a40053d0, handle=0x20eeba0,
strs=0x7f51a4003578, strl=0x7f51a40035e8, nelts=14,
len=82) at mod_log_config.c:1368
#2 0x00000000042b10d in config_log_transaction
(r=0x7f51a40053d0, cls=0x20b9d50,
default_format=0x20eee370) at mod_log_config.c:930
#3 0x00000000042aad6 in multi_log_transaction
(r=0x7f51a40053d0) at mod_log_config.c:950
#4 0x00000000046cb2d in ap_run_log_transaction
(r=0x7f51a40053d0) at protocol.c:1563
#5 0x000000000436e81 in ap_process_request
(r=0x7f51a40053d0) at http_request.c:312
#6 0x00000000042e9da in ap_process_http_connection
(c=0x7f519c000b68) at http_core.c:293
#7 0x000000000465cdd in ap_run_process_connection
(c=0x7f519c000b68) at connection.c:85
#8 0x0000000004661f5 in ap_process_connection
(c=0x7f519c000b68, csd=0x7f519c000a20) at connection.c:211
#9 0x000000000451ba0 in process_socket
(p=0x7f519c0009b8, sock=0x7f519c000a20,
my_child_num=0, my_thread_num=0,
bucket_alloc=0x7f51a4001348) at worker.c:632
#10 0x000000000451221 in worker_thread
(thd=0x210fa90, dummy=0x7f51a40008c0) at worker.c:946
#11 0x00007f51ac87c555 in dummy_worker
(opaque=0x210fa90) at thread.c:127
#12 0x00007f51abae0182 in start_thread
(arg=0x7f51aa8ef700) at pthread_create.c:312
#13 0x00007f51ab80d47d in clone () at ../sysdeps/unix/sysv/linux/x86_64/clone.S:111
```

Failures Today



# Related Work

- Collaborative approaches
  - *WER [SOSP'09], CBI [PLDI'05], CCI [OOPSLA'10]*
- Identifying differences of failing and successful runs
  - *Delta debugging [TSE'02], Symbiosis [PLDI'15]*
- Record & replay, checkpointing
  - *ODR [SOSP'09], Triage [SOSP'07]*
- Hardware support
  - *PBI [ASPLOS'13], LBRA/LCRA [ASPLOS'14]*

# Related Work

- Collaborative approaches
  - *WER [SOSP'09], CBI [PLDI'05], CCI [OOPSLA'10]*
- Identifying differences of failing and successful runs
  - *Delta debugging [TSE'02], Symbiosis [PLDI'15]*
- Record & replay, checkpointing
  - *ODR [SOSP'09], Triage [SOSP'07]*
- Hardware support
  - *PBI [ASPLOS'13], LBRA/LCRA [ASPLOS'14]*

# Contributions

# Contributions

Goal: automate the manual detective work of debugging

# Contributions

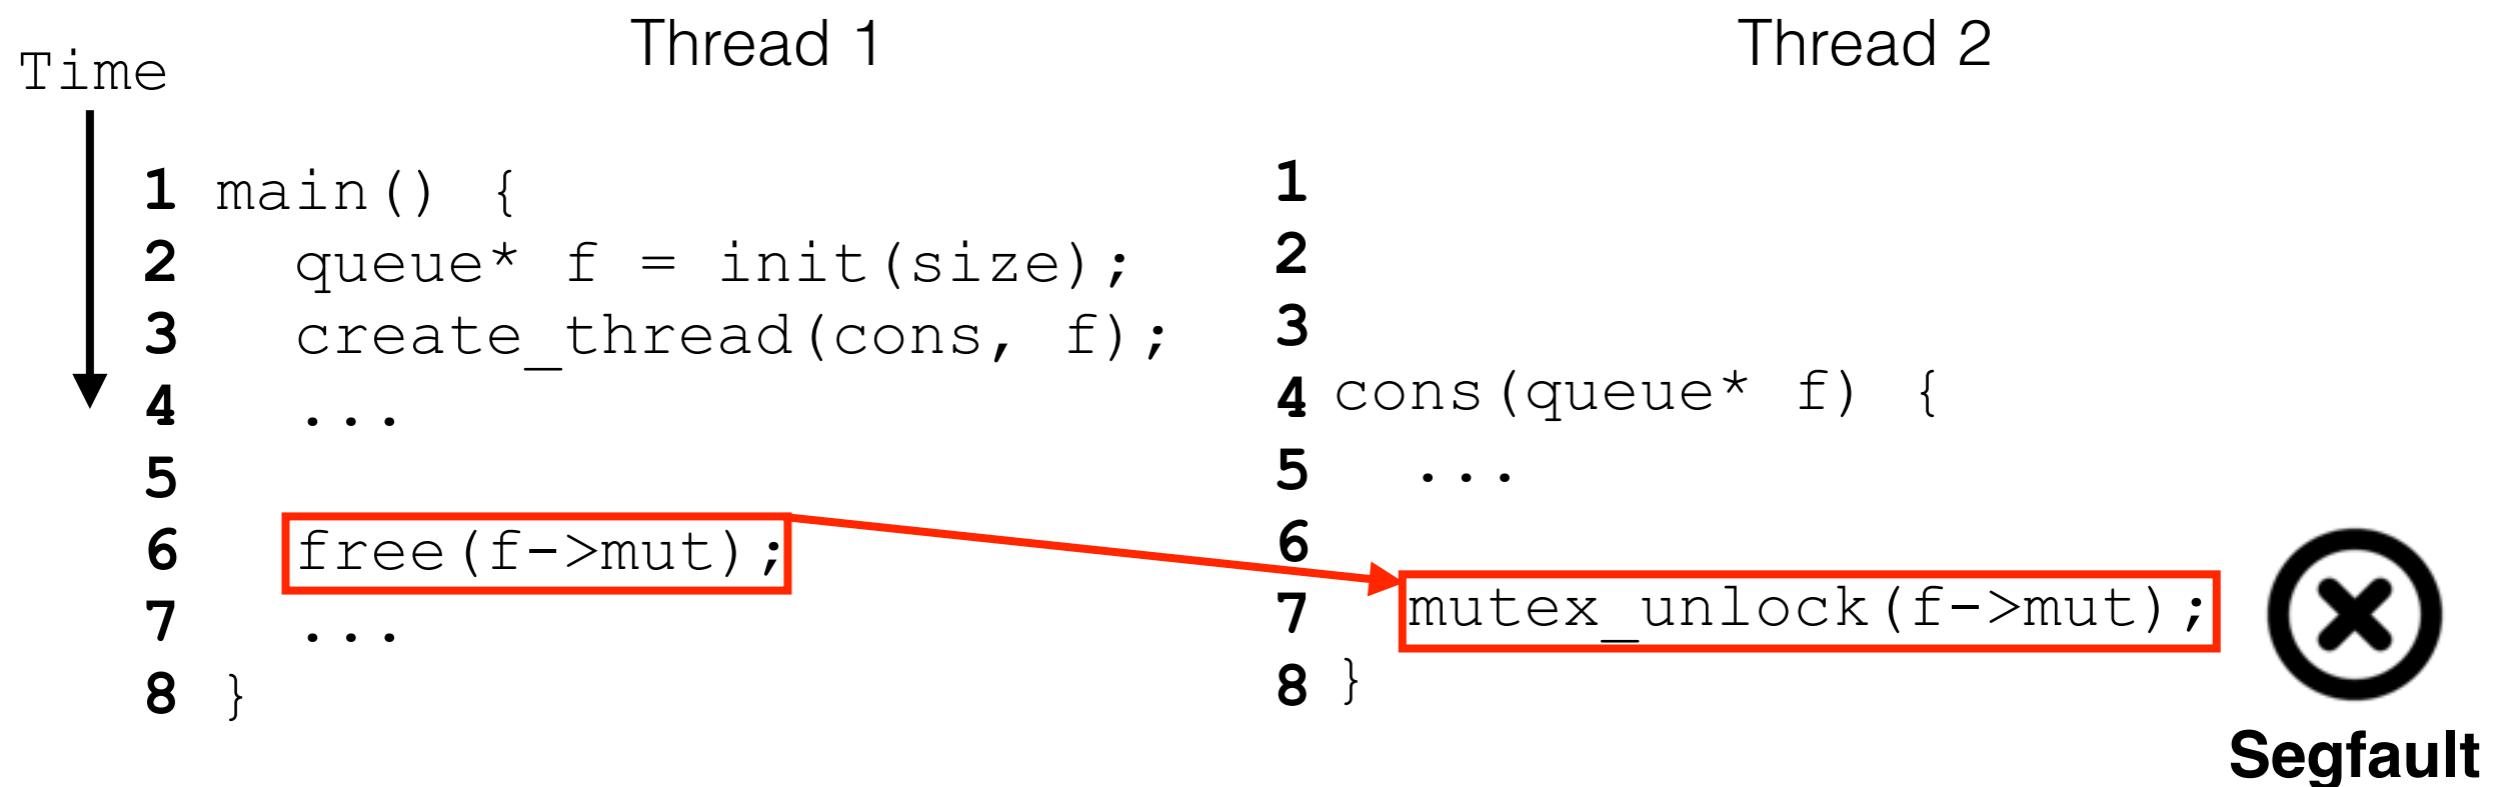
Goal: automate the manual detective work of debugging

## **Failure sketching**

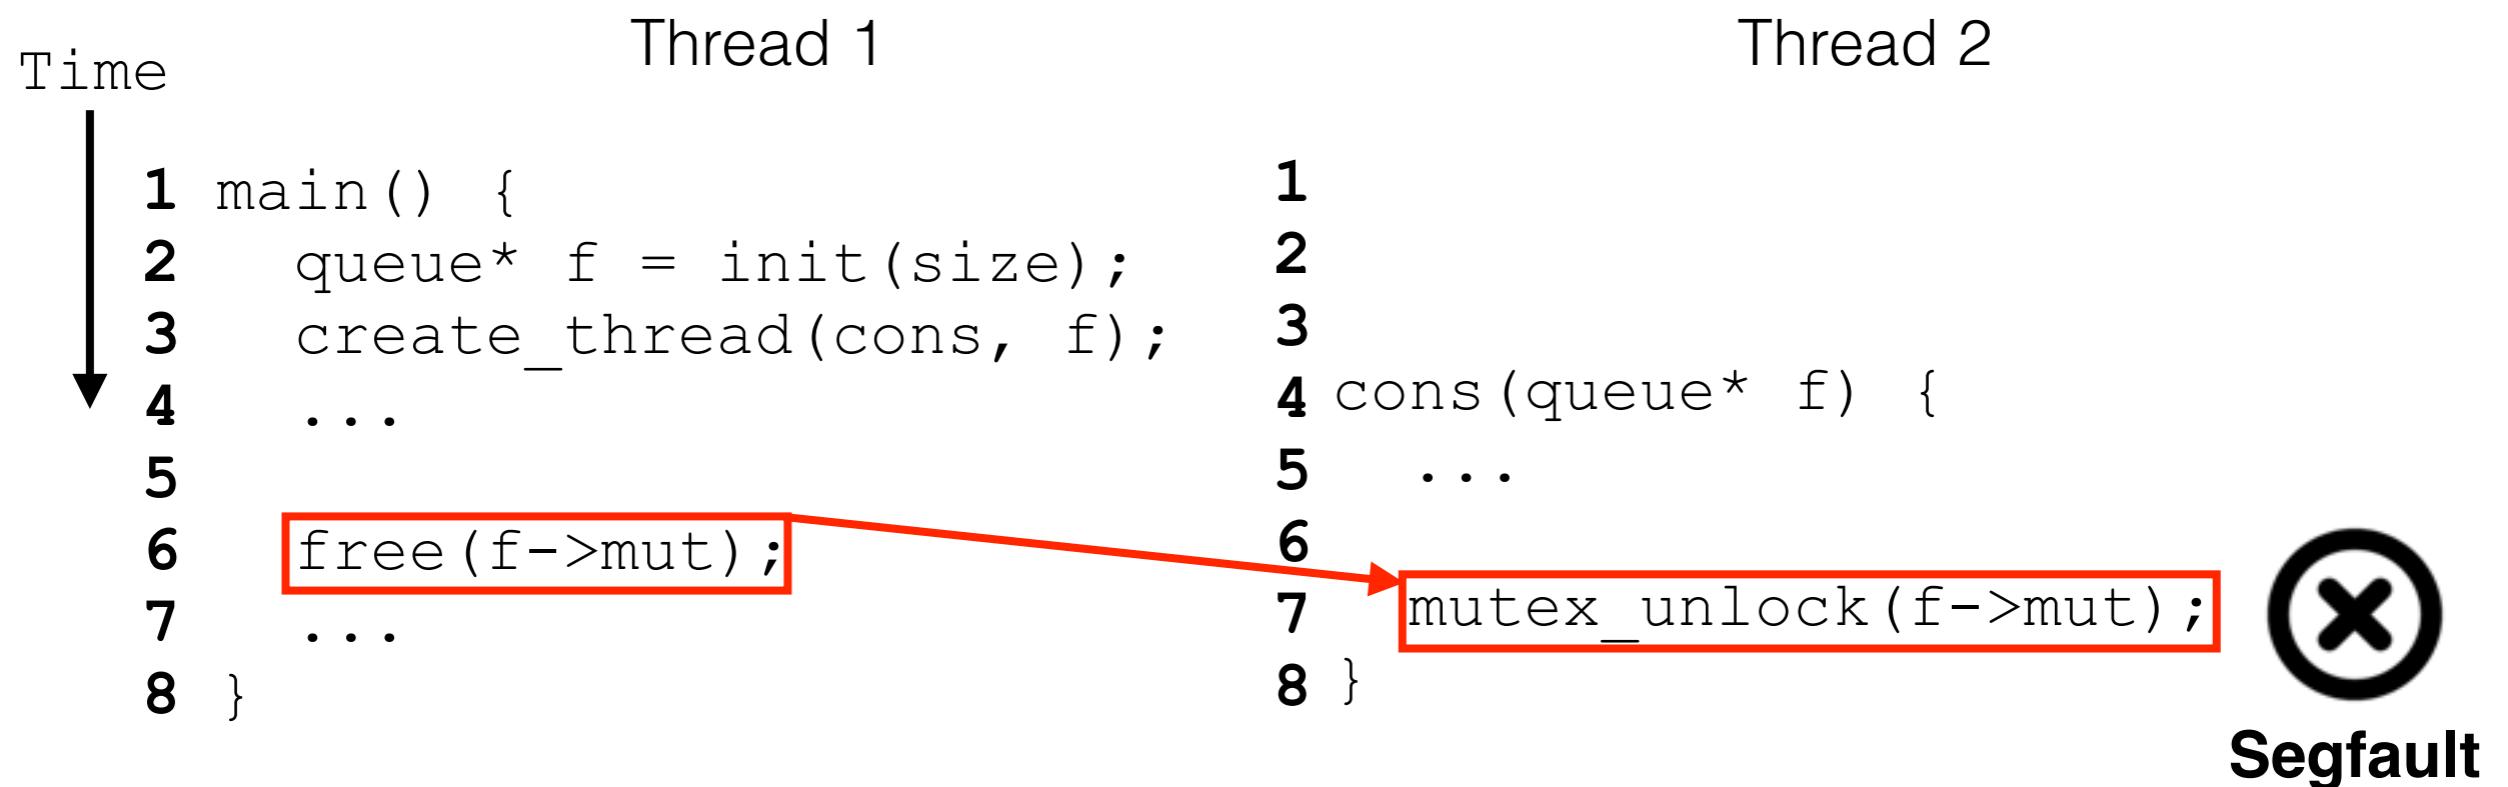
*Complements in-house static analysis with in-production dynamic analysis*

*Automatically and efficiently builds accurate failure sketches that show root causes of failures*

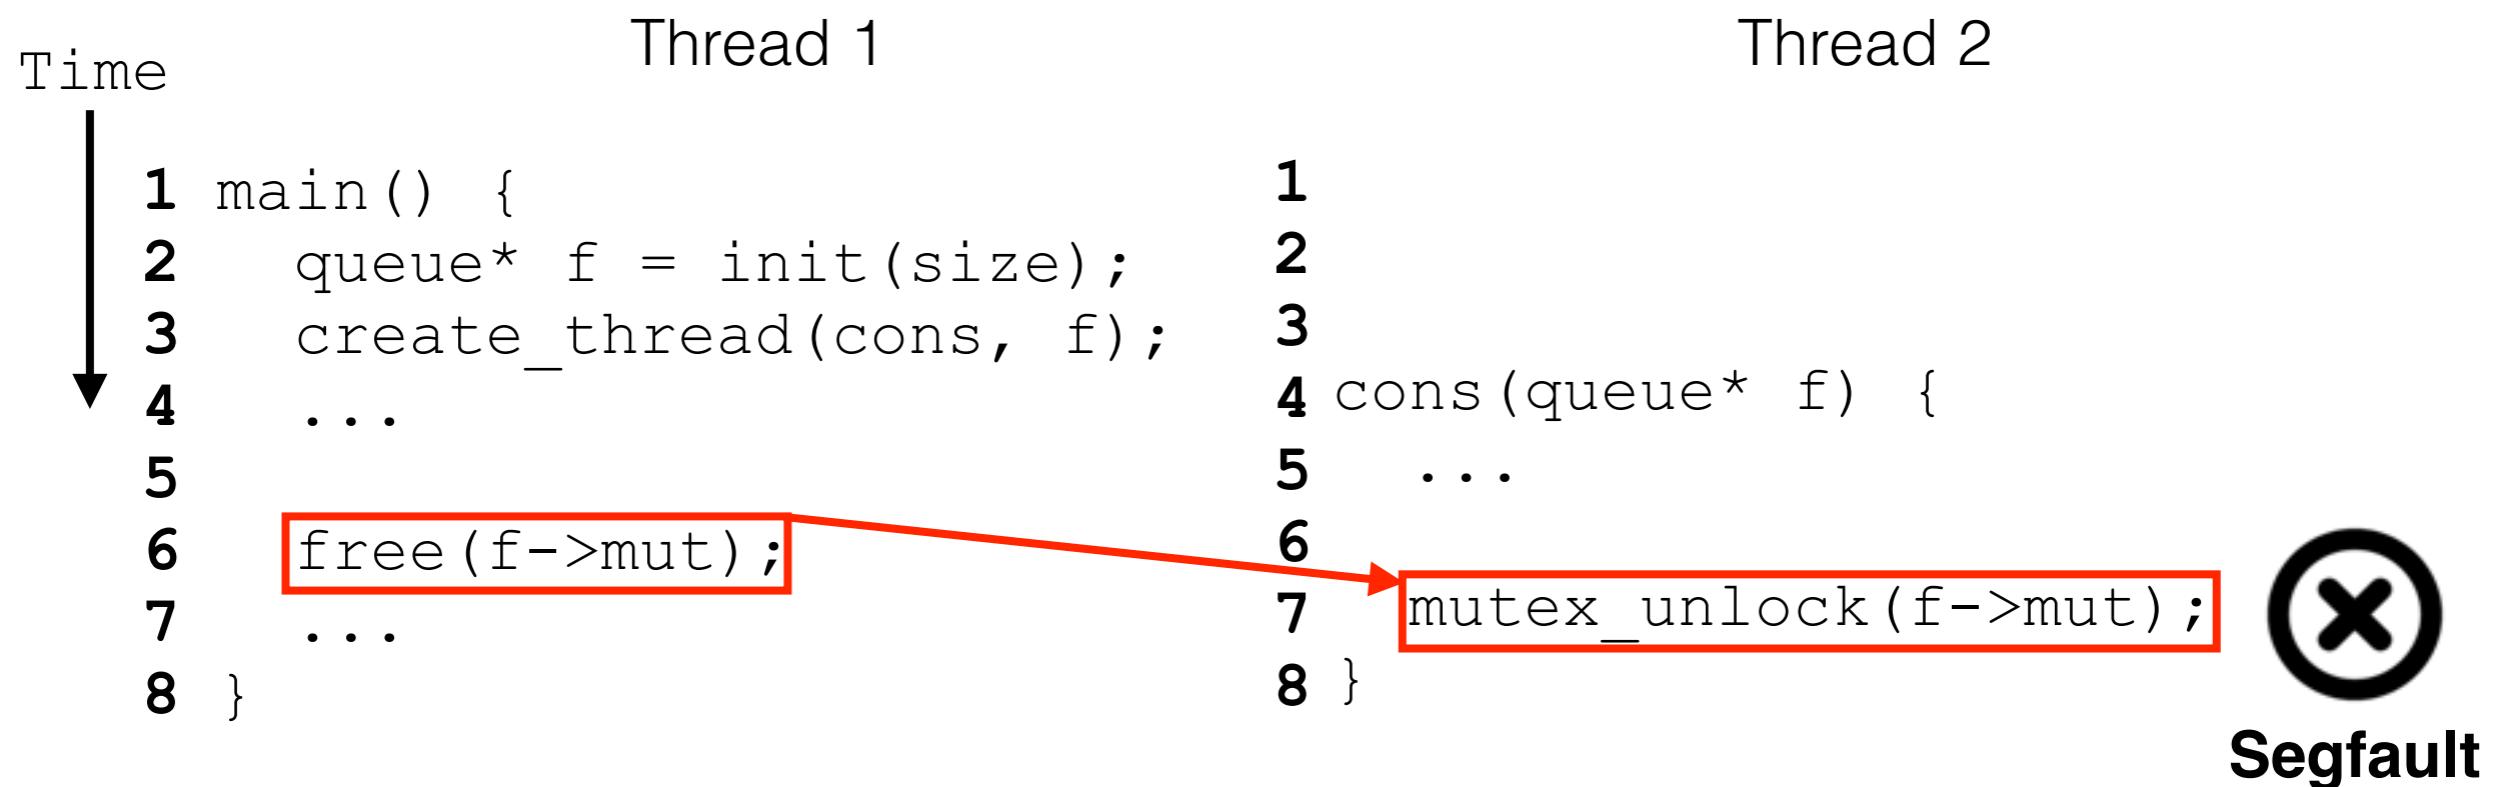
# Failure Sketch



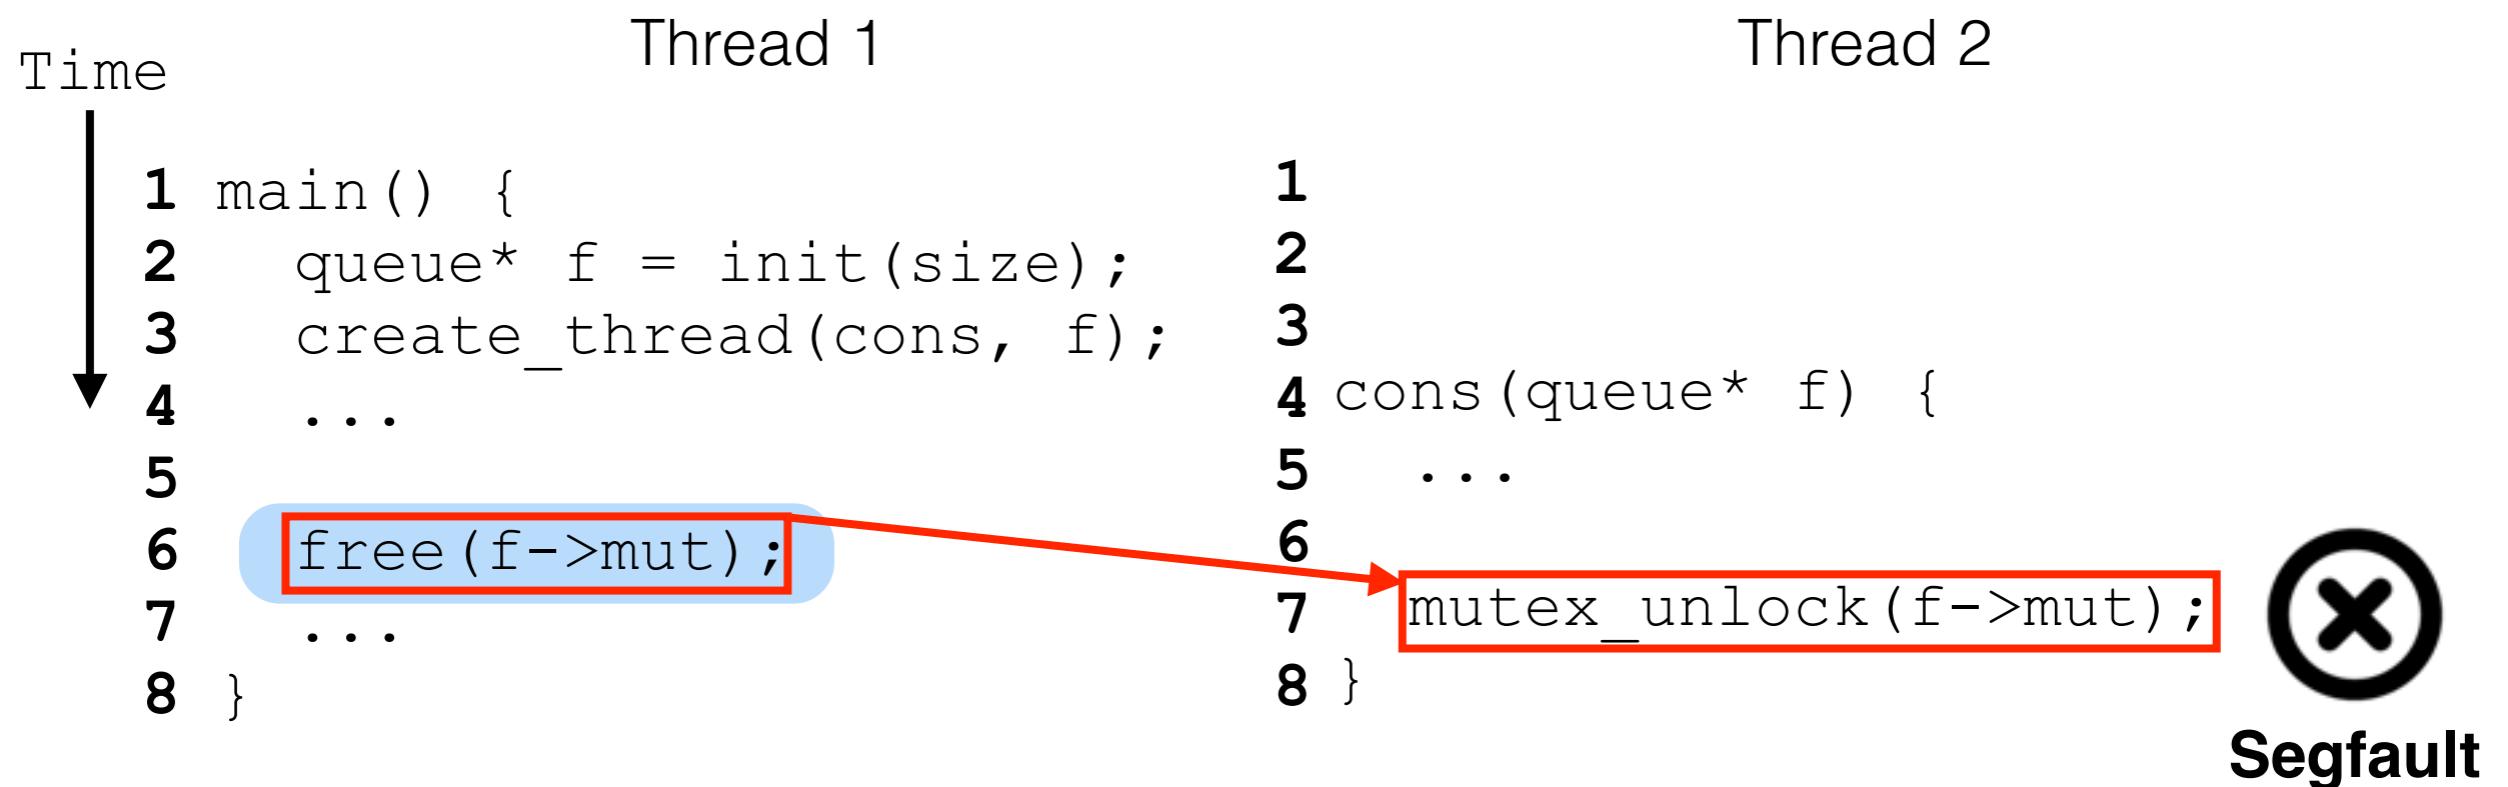
# Failure Sketch



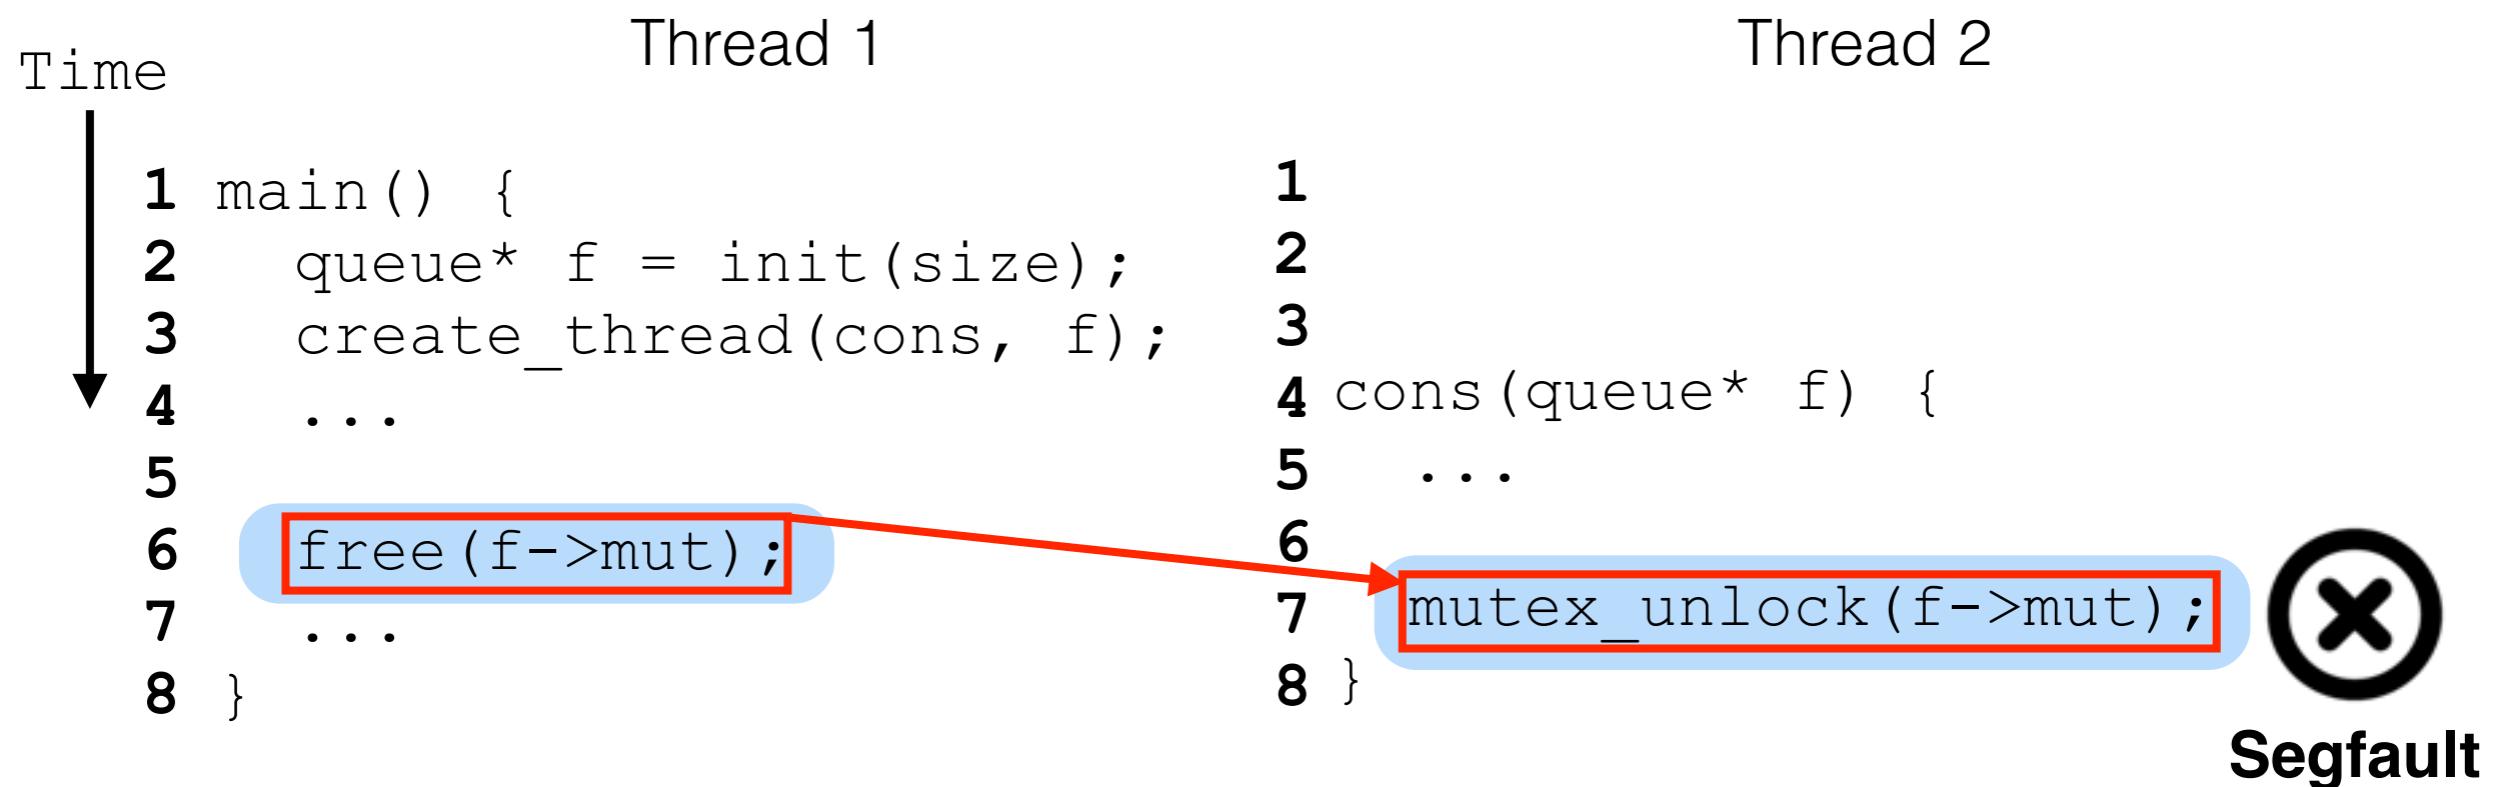
# Failure Sketch



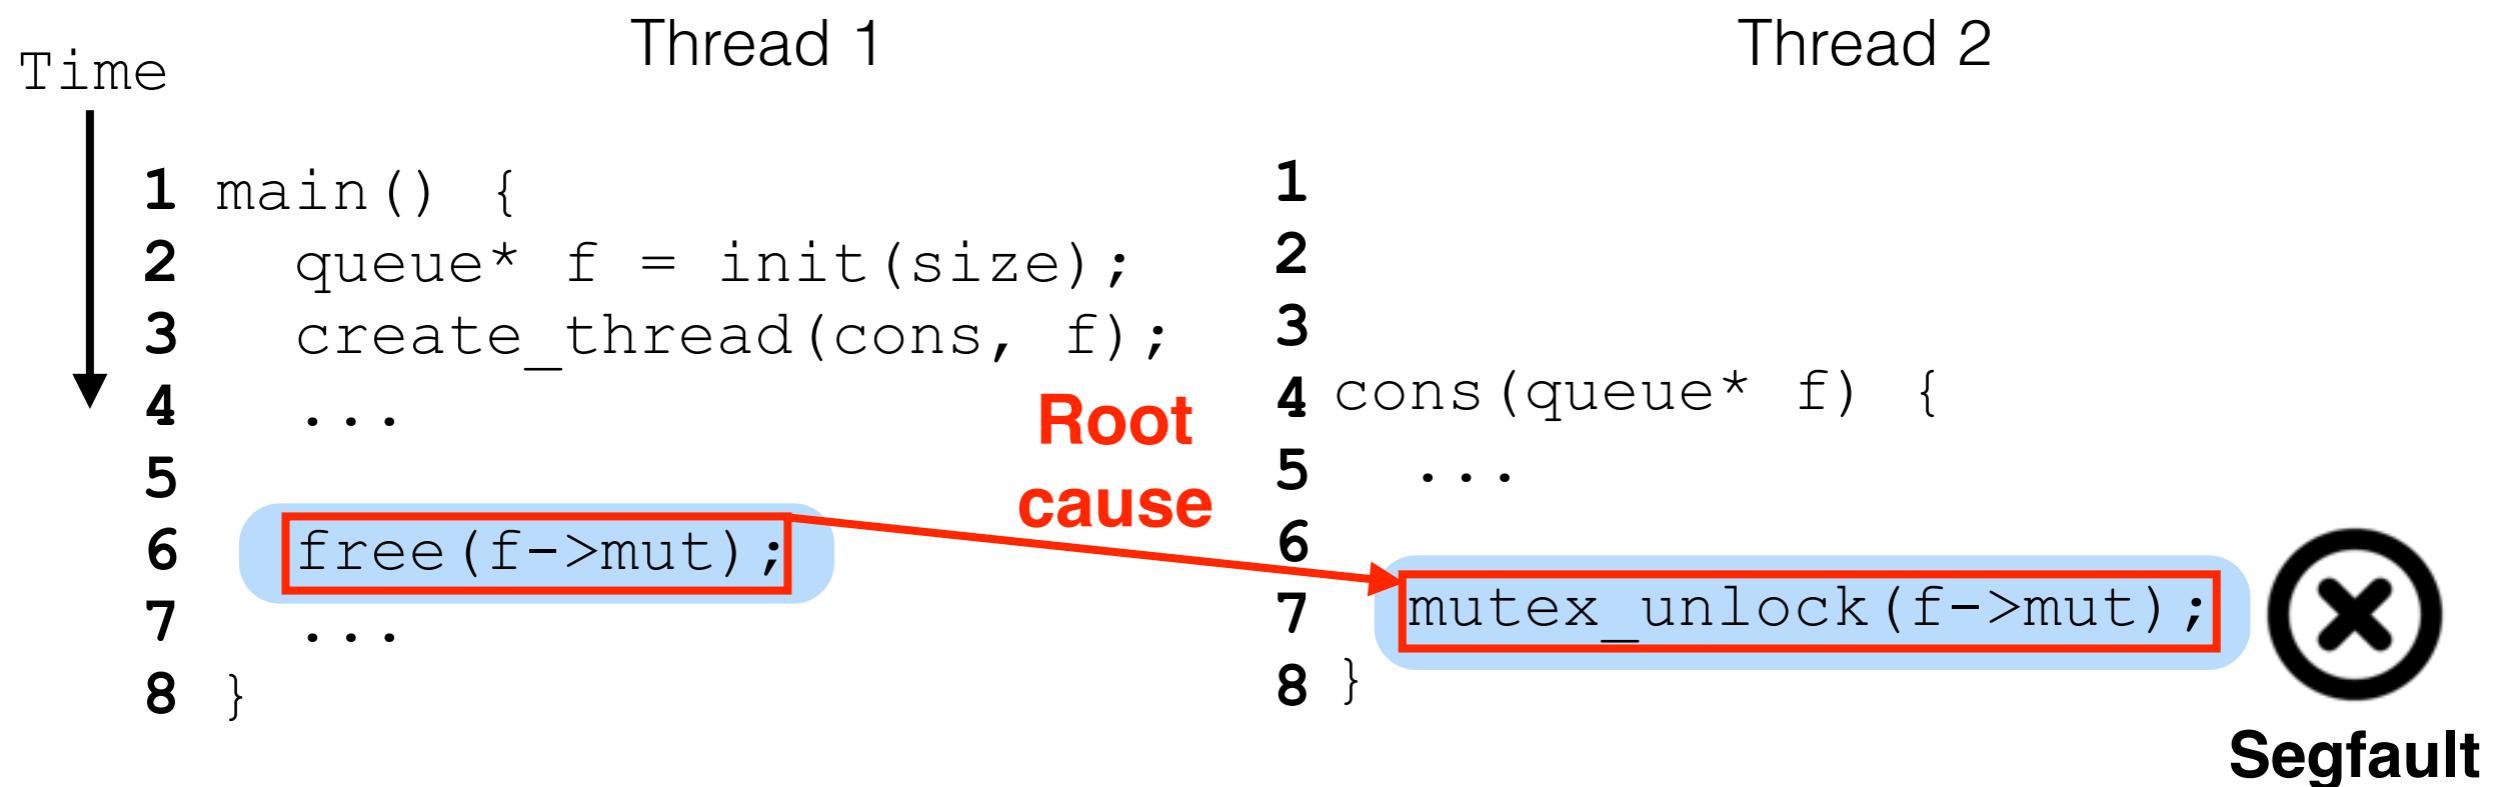
# Failure Sketch



# Failure Sketch



# Failure Sketch

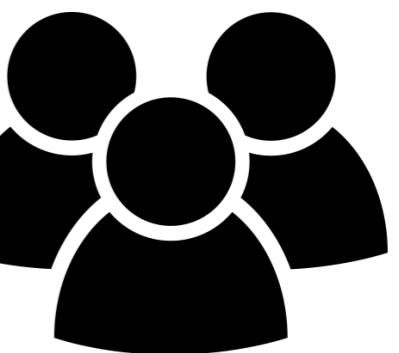


# Failure Sketch Usage Model

Understand  
root cause

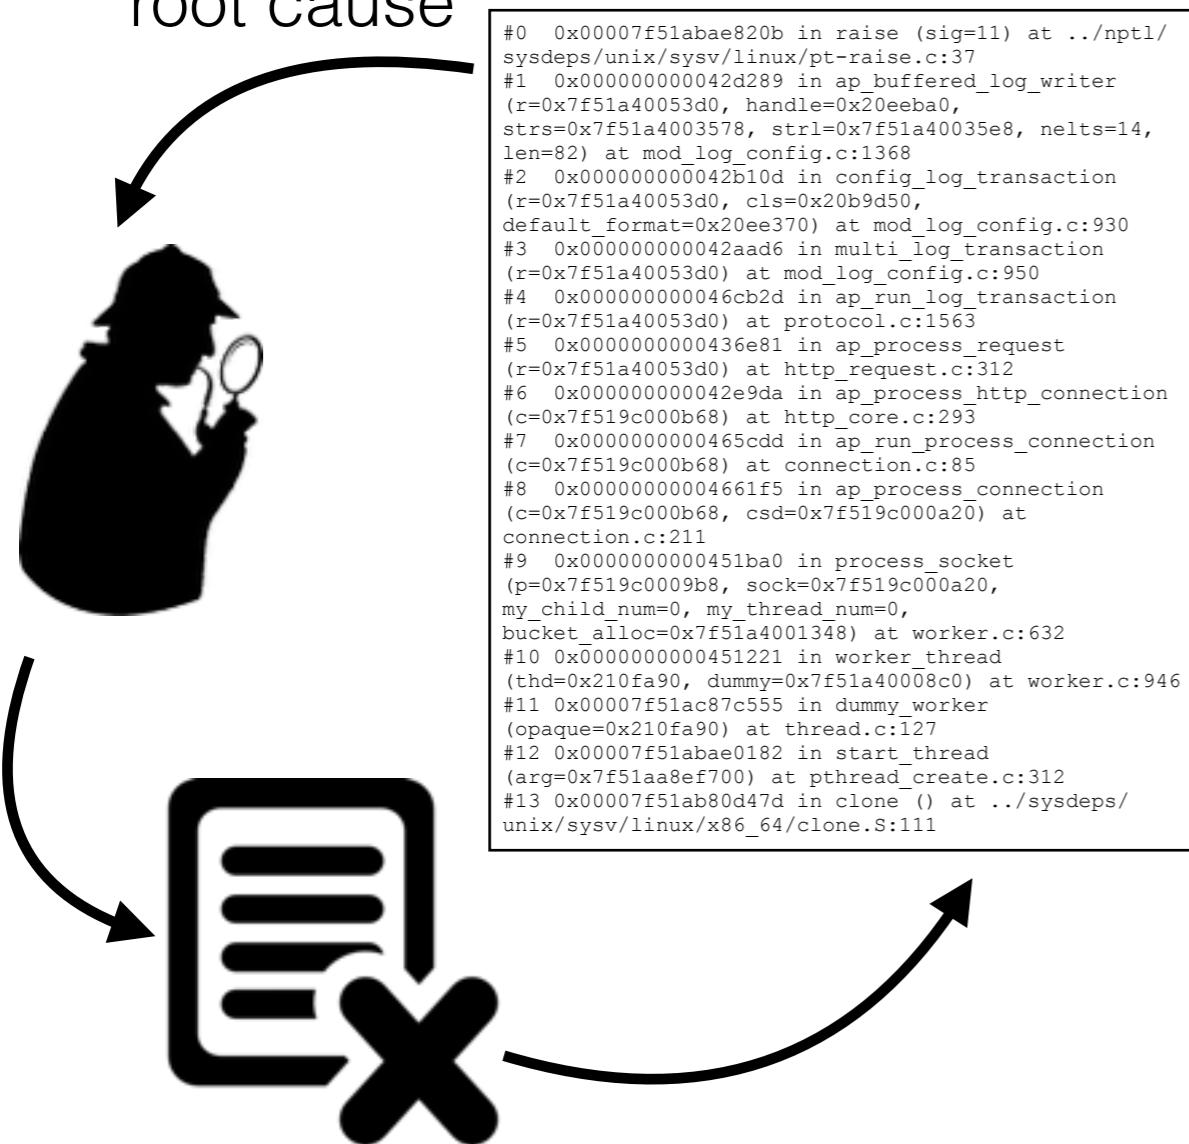


Reproduce  
the failure

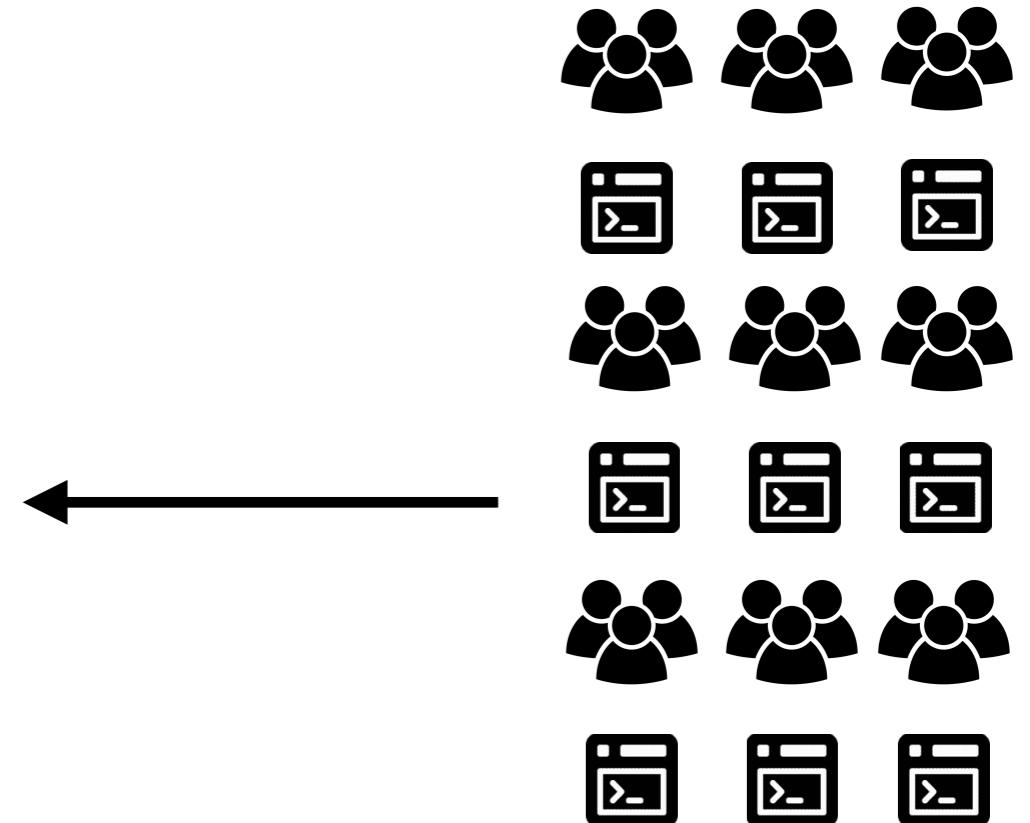


# Failure Sketch Usage Model

Understand  
root cause

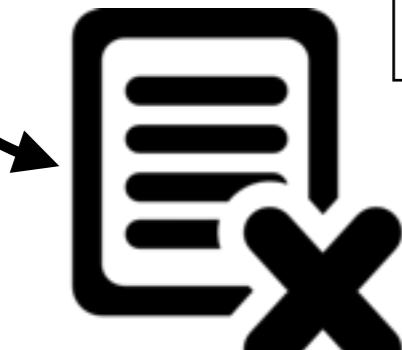


Reproduce  
the failure



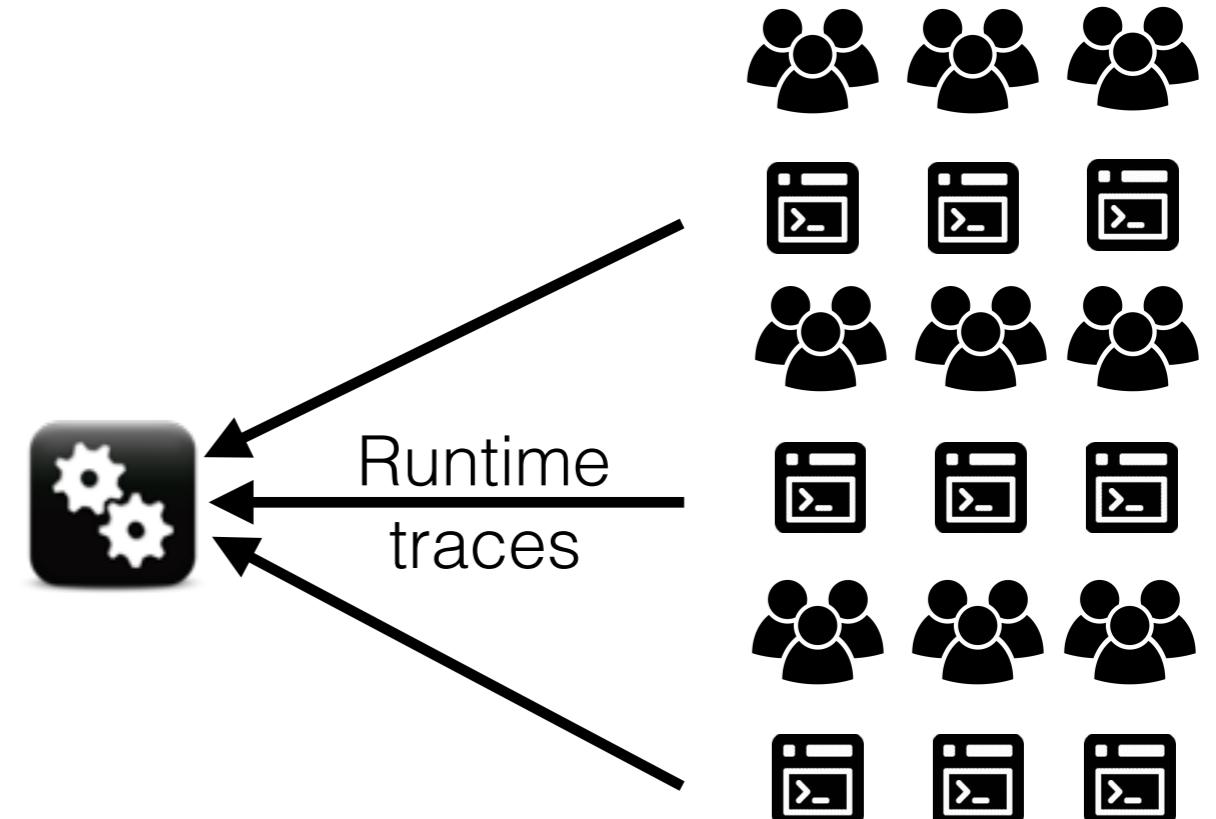
# Failure Sketch Usage Model

Understand  
root cause



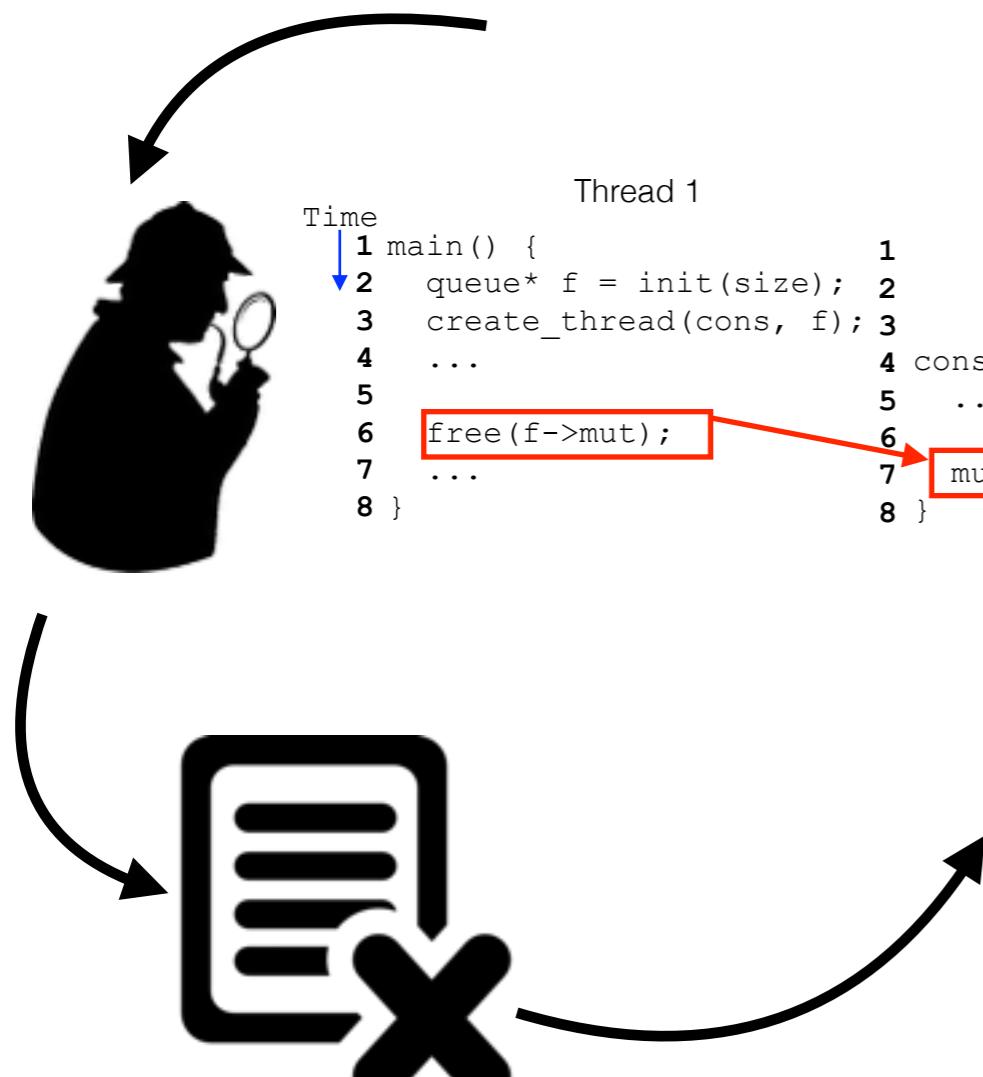
Reproduce  
the failure

```
#0 0x00007f51abae820b in raise (sig=11) at ../nptl/
sysdeps/unix/sysv/linux/pt-raise.c:37
#1 0x00000000042d289 in ap_buffered_log_writer
(r=0x7f51a40053d0, handle=0x20eeba0,
strs=0x7f51a4003578, strl=0x7f51a40035e8, nelts=14,
len=82) at mod_log_config.c:1368
#2 0x000000000042b10d in config_log_transaction
(r=0x7f51a40053d0, cls=0x20b9d50,
default_format=0x20ee370) at mod_log_config.c:930
#3 0x000000000042aad6 in multi_log_transaction
(r=0x7f51a40053d0) at mod_log_config.c:950
#4 0x000000000046cb2d in ap_run_log_transaction
(r=0x7f51a40053d0) at protocol.c:1563
#5 0x0000000000436e81 in ap_process_request
(r=0x7f51a40053d0) at http_request.c:312
#6 0x000000000042e9da in ap_process_http_connection
(c=0x7f519c000b68) at http_core.c:293
#7 0x0000000000465cdd in ap_run_process_connection
(c=0x7f519c000b68) at connection.c:85
#8 0x00000000004661f5 in ap_process_connection
(c=0x7f519c000b68, csd=0x7f519c000a20) at
connection.c:211
#9 0x0000000000451ba0 in process_socket
(p=0x7f519c0009b8, sock=0x7f519c000a20,
my_child_num=0, my_thread_num=0,
bucket_alloc=0x7f51a4001348) at worker.c:632
#10 0x0000000000451221 in worker_thread
(thd=0x210fa90, dummy=0x7f51a40008c0) at worker.c:946
#11 0x00007f51ac87c555 in dummy_worker
(opaque=0x210fa90) at thread.c:127
#12 0x00007f51abae0182 in start_thread
(arg=0x7f51aa8ef700) at pthread_create.c:312
#13 0x00007f51ab80d47d in clone () at ../sysdeps/
unix/sysv/linux/x86_64/clone.S:111
```

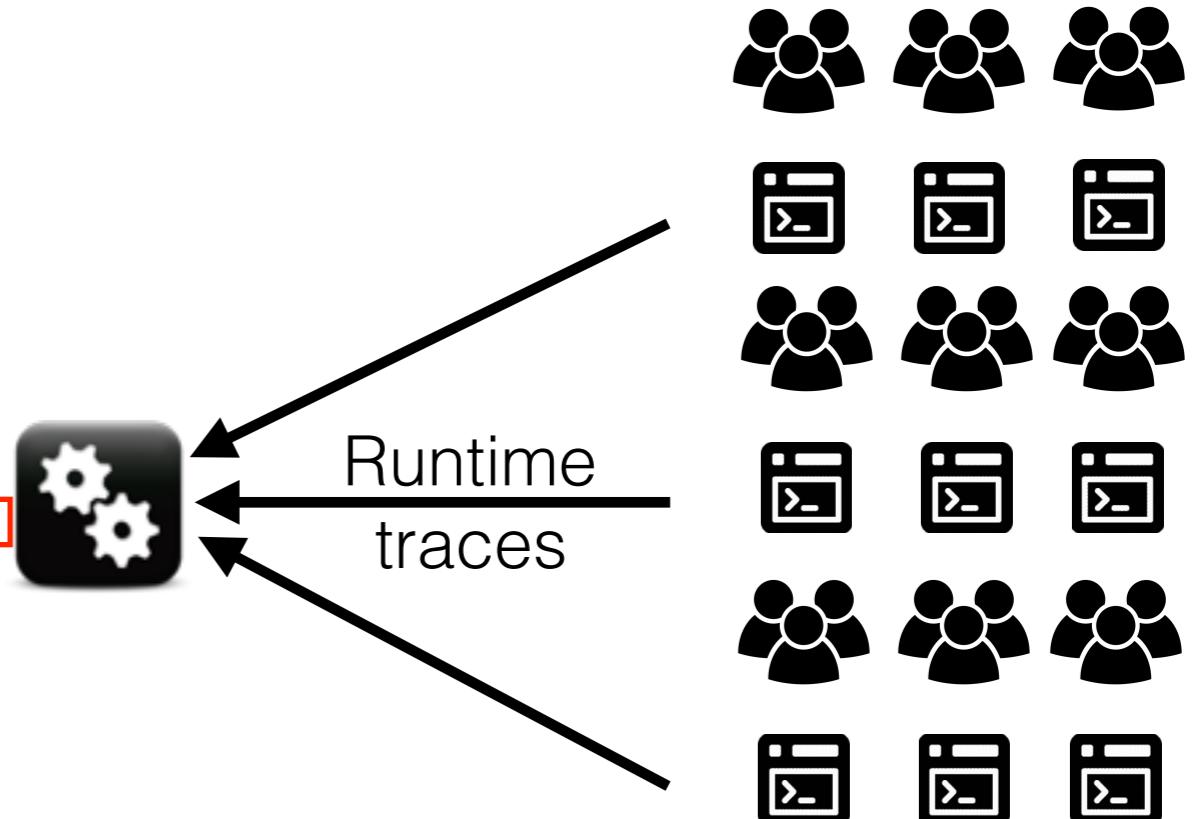


# Failure Sketch Usage Model

Understand  
root cause



Reproduce  
the failure



# Failure Sketch Usage Model

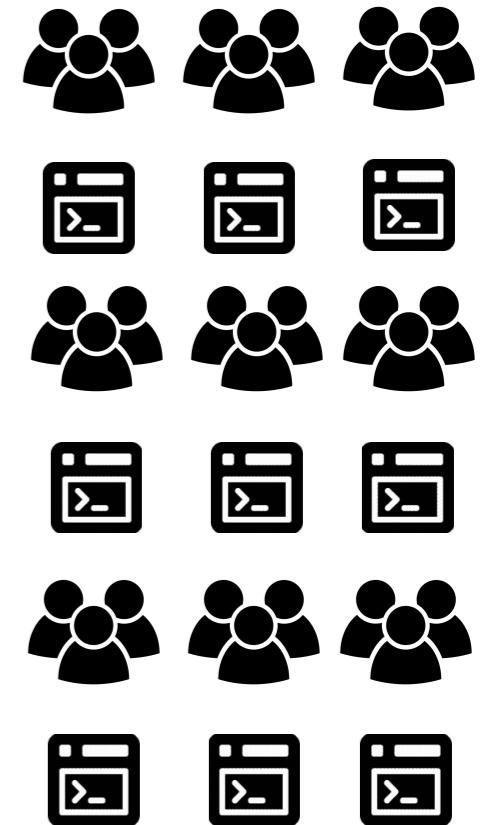


Time ↓

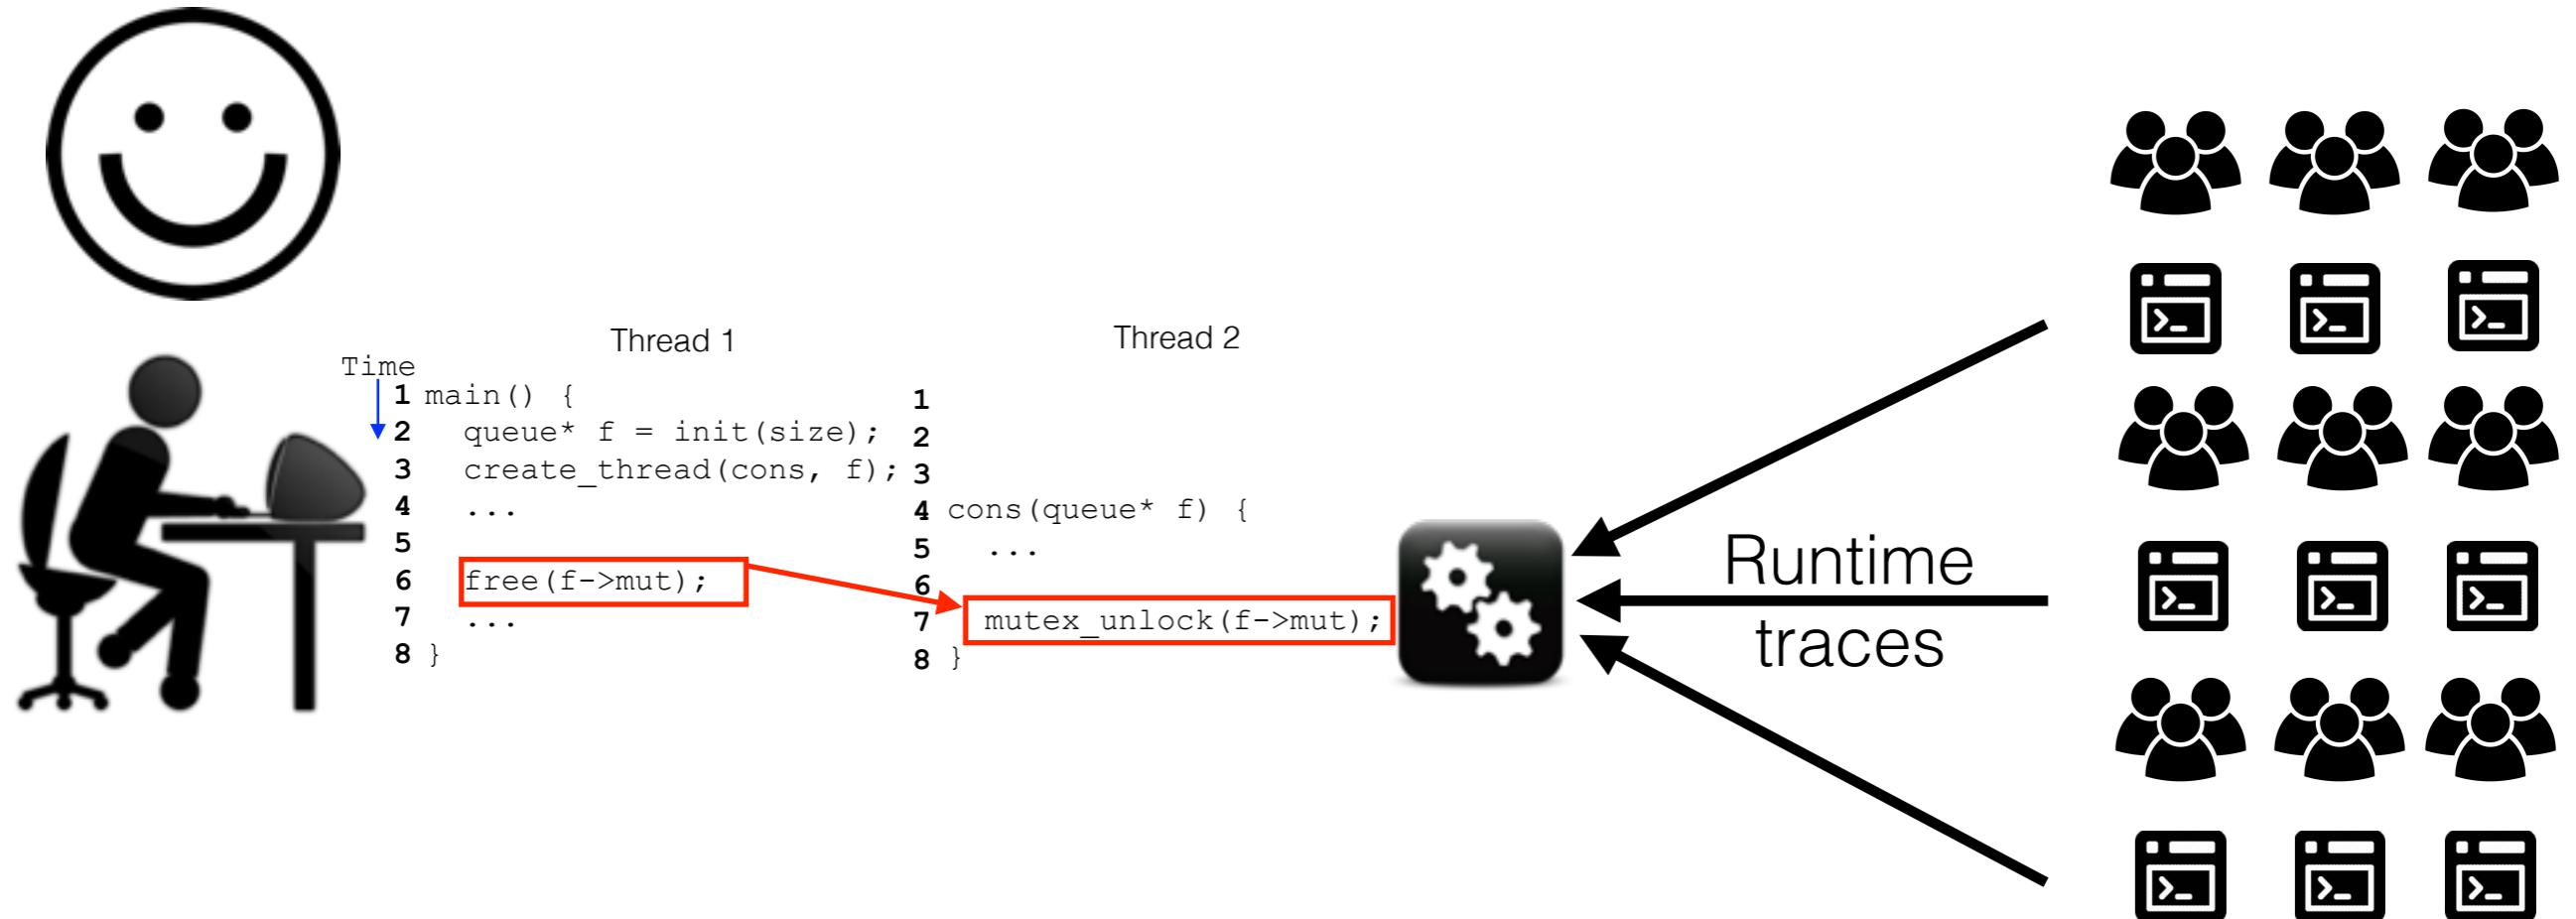
Thread 1	Thread 2
1 main() {	1
2 queue* f = init(size);	2
3 create_thread(cons, f);	3
4 ...	4 cons(queue* f) {
5	5 ...
6 free(f->mut);	6 mutex_unlock(f->mut);
7 ...	7 }
8 }	8 }



Runtime  
traces



# Failure Sketch Usage Model



# Outline

- Challenges
- Design
- Evaluation

# Outline

- Challenges
- Design
- Evaluation

# Challenges of Building Failure Sketches

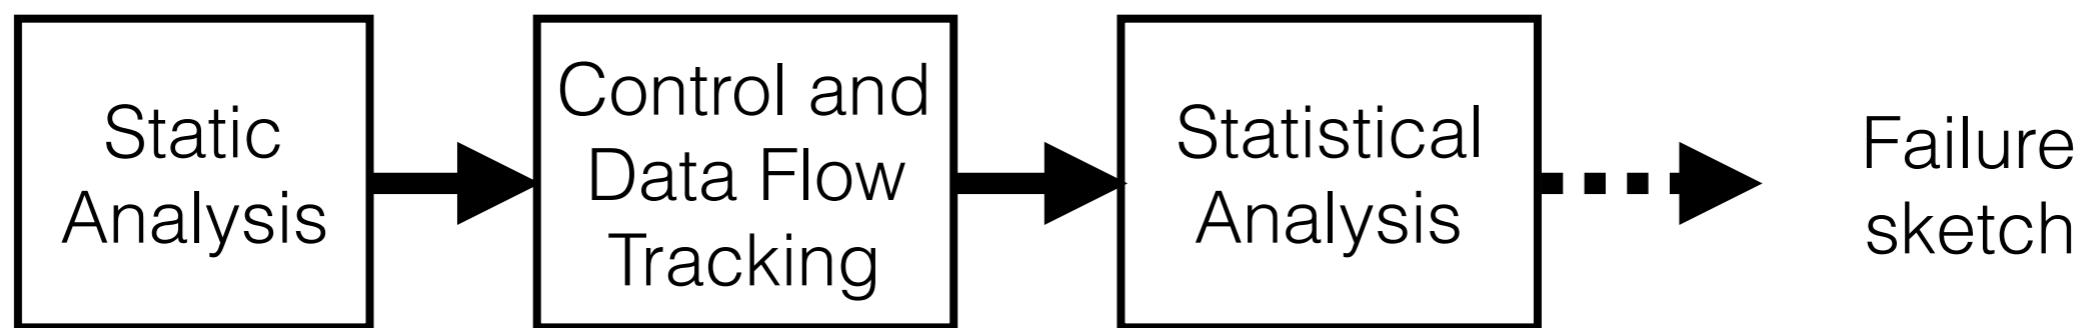
- Accuracy
  - *Exclude all irrelevant information, preserve all relevant one*
- Recurrence
  - *Gathering enough execution information from rare failures*
- Latency
  - *Achieve high accuracy after just a few recurrences*

# Outline

- Challenges
- Design
- Evaluation

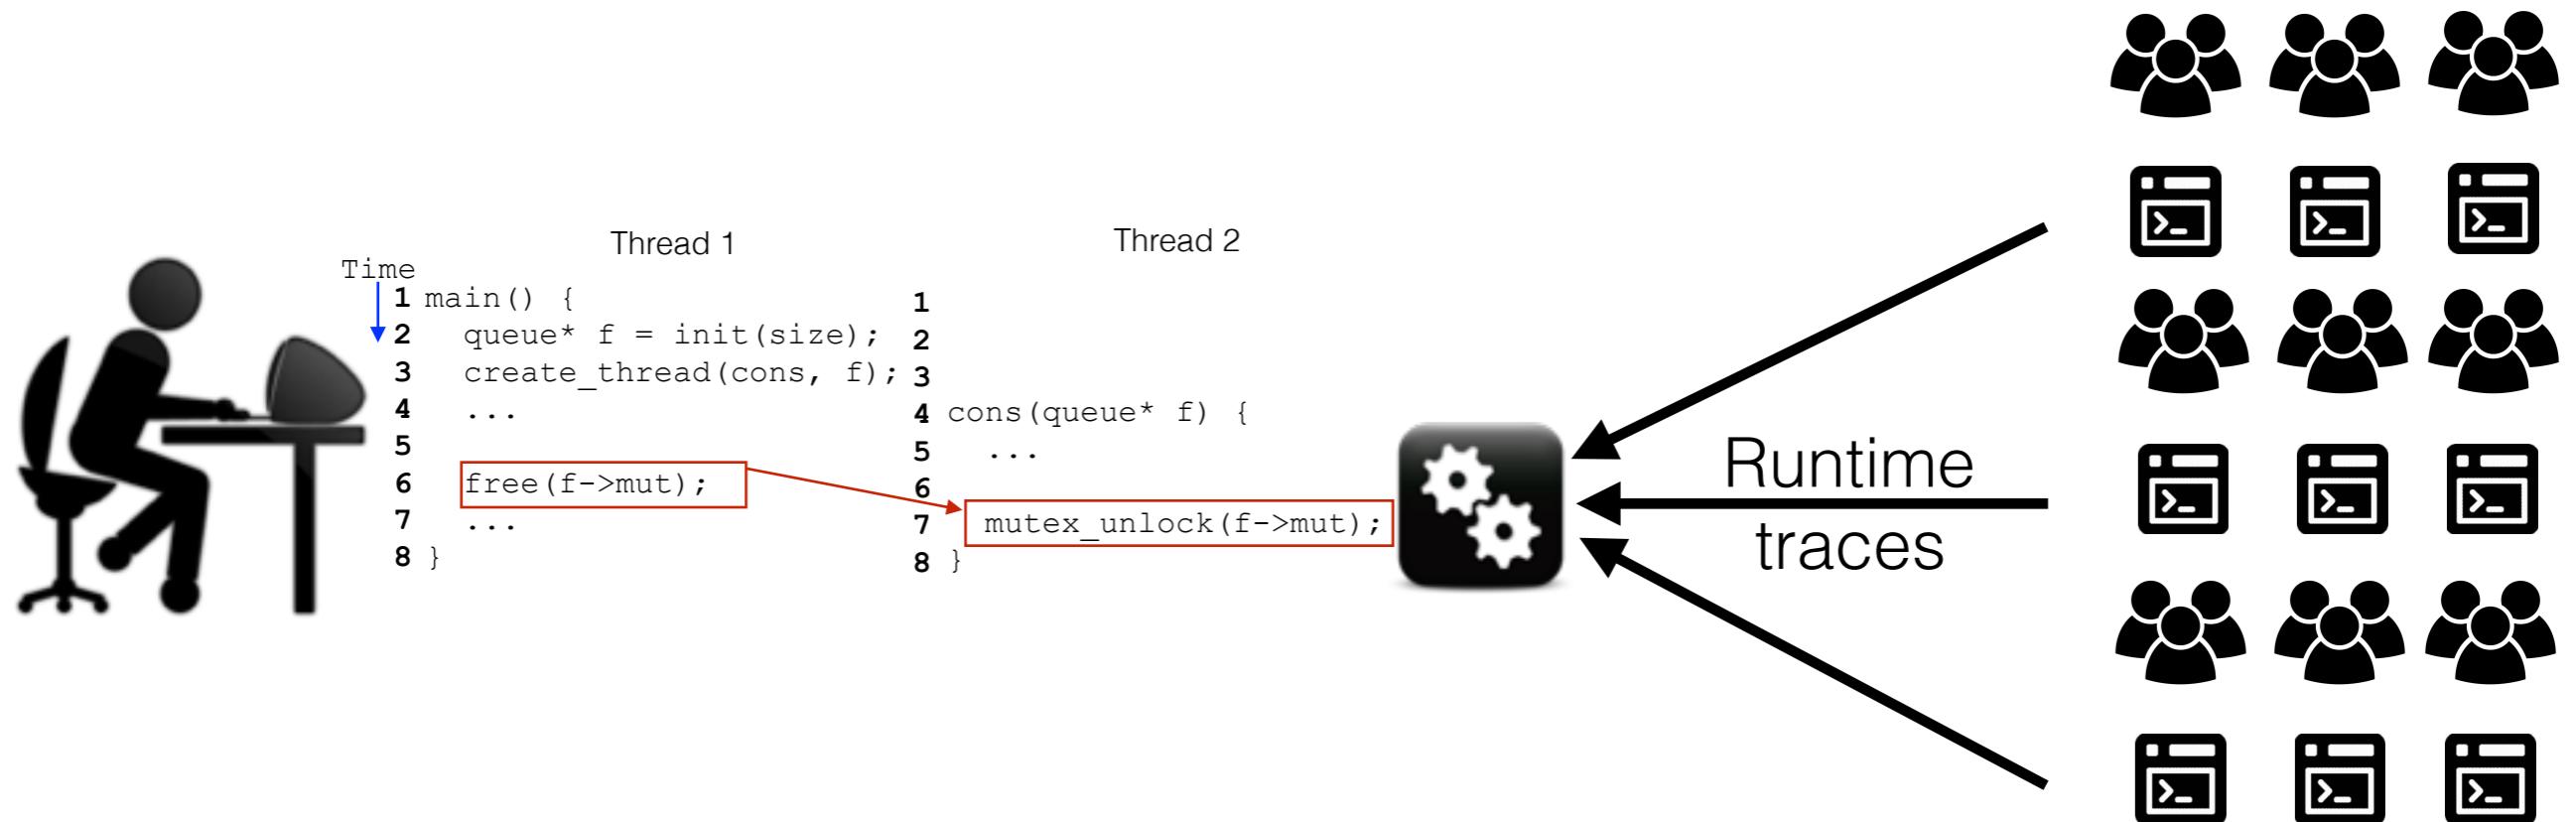
# Outline

- Challenges
- Design

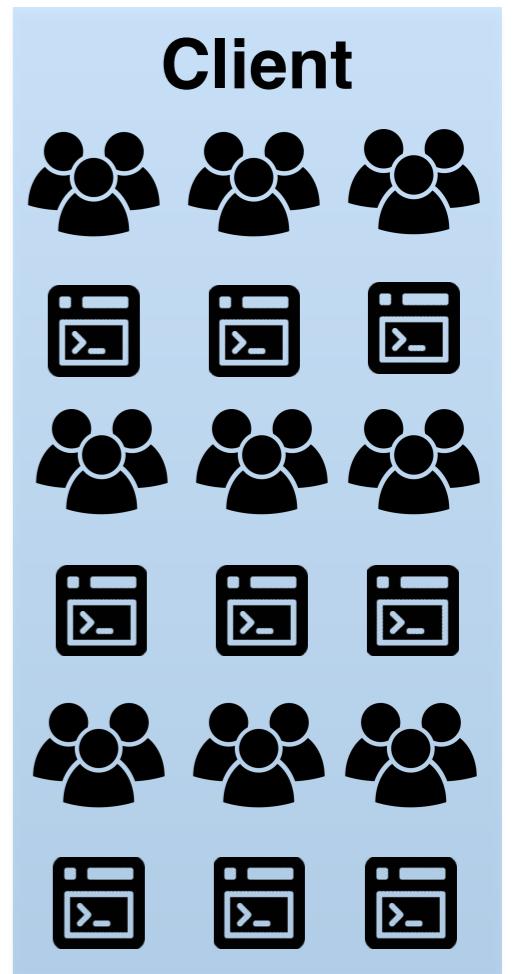
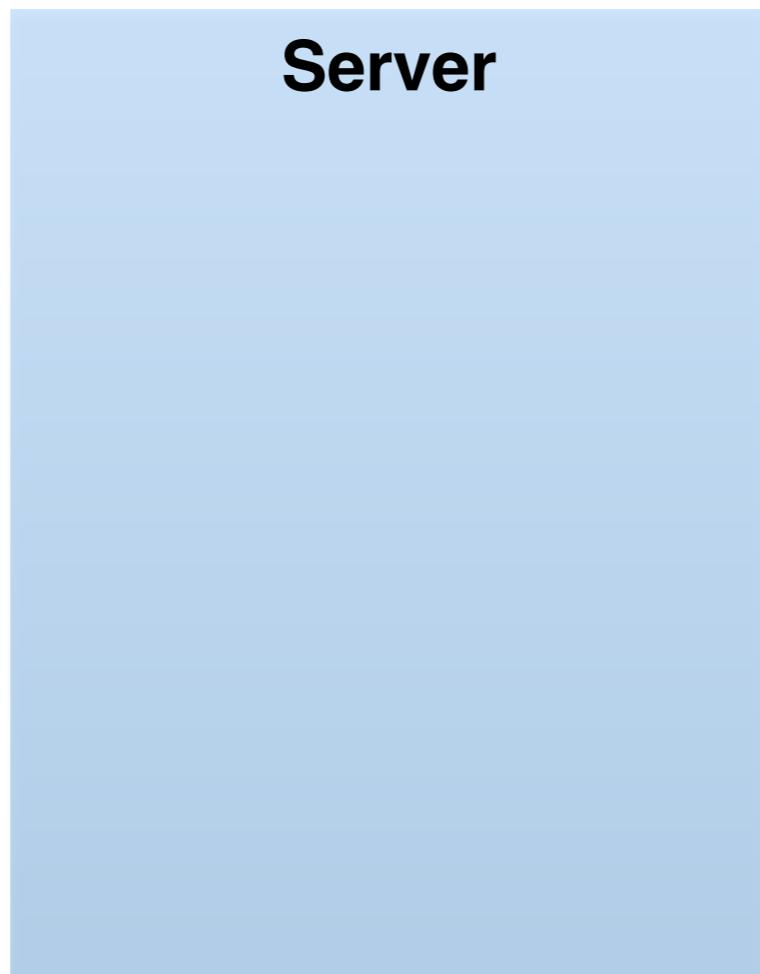


- Evaluation

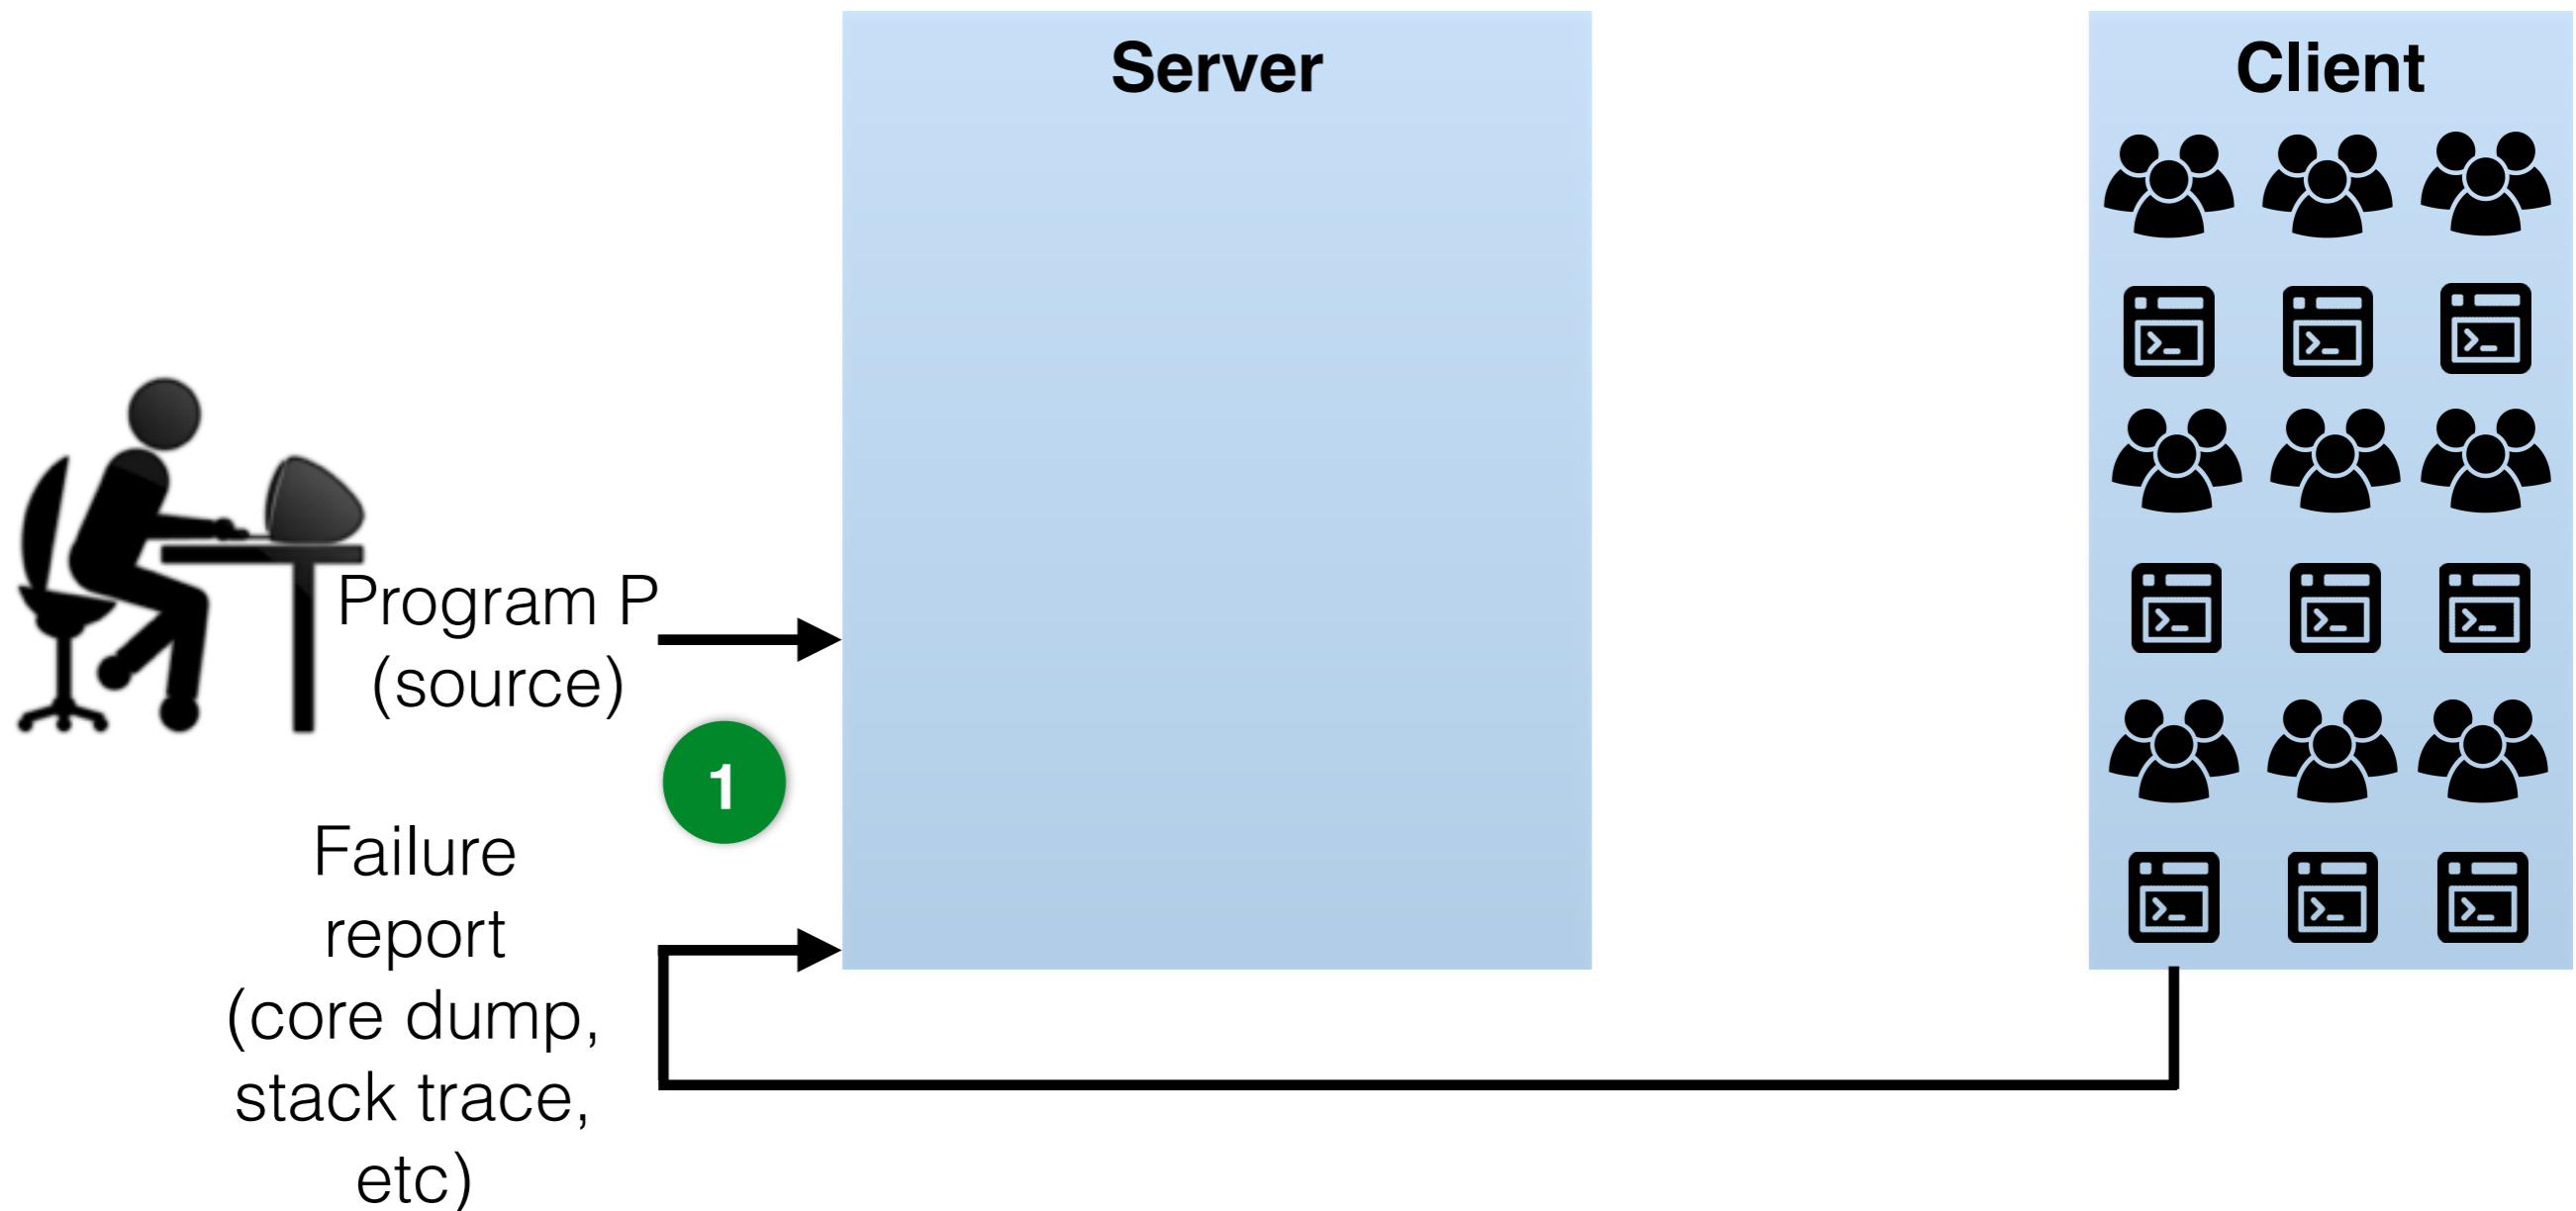
# Gist System Architecture



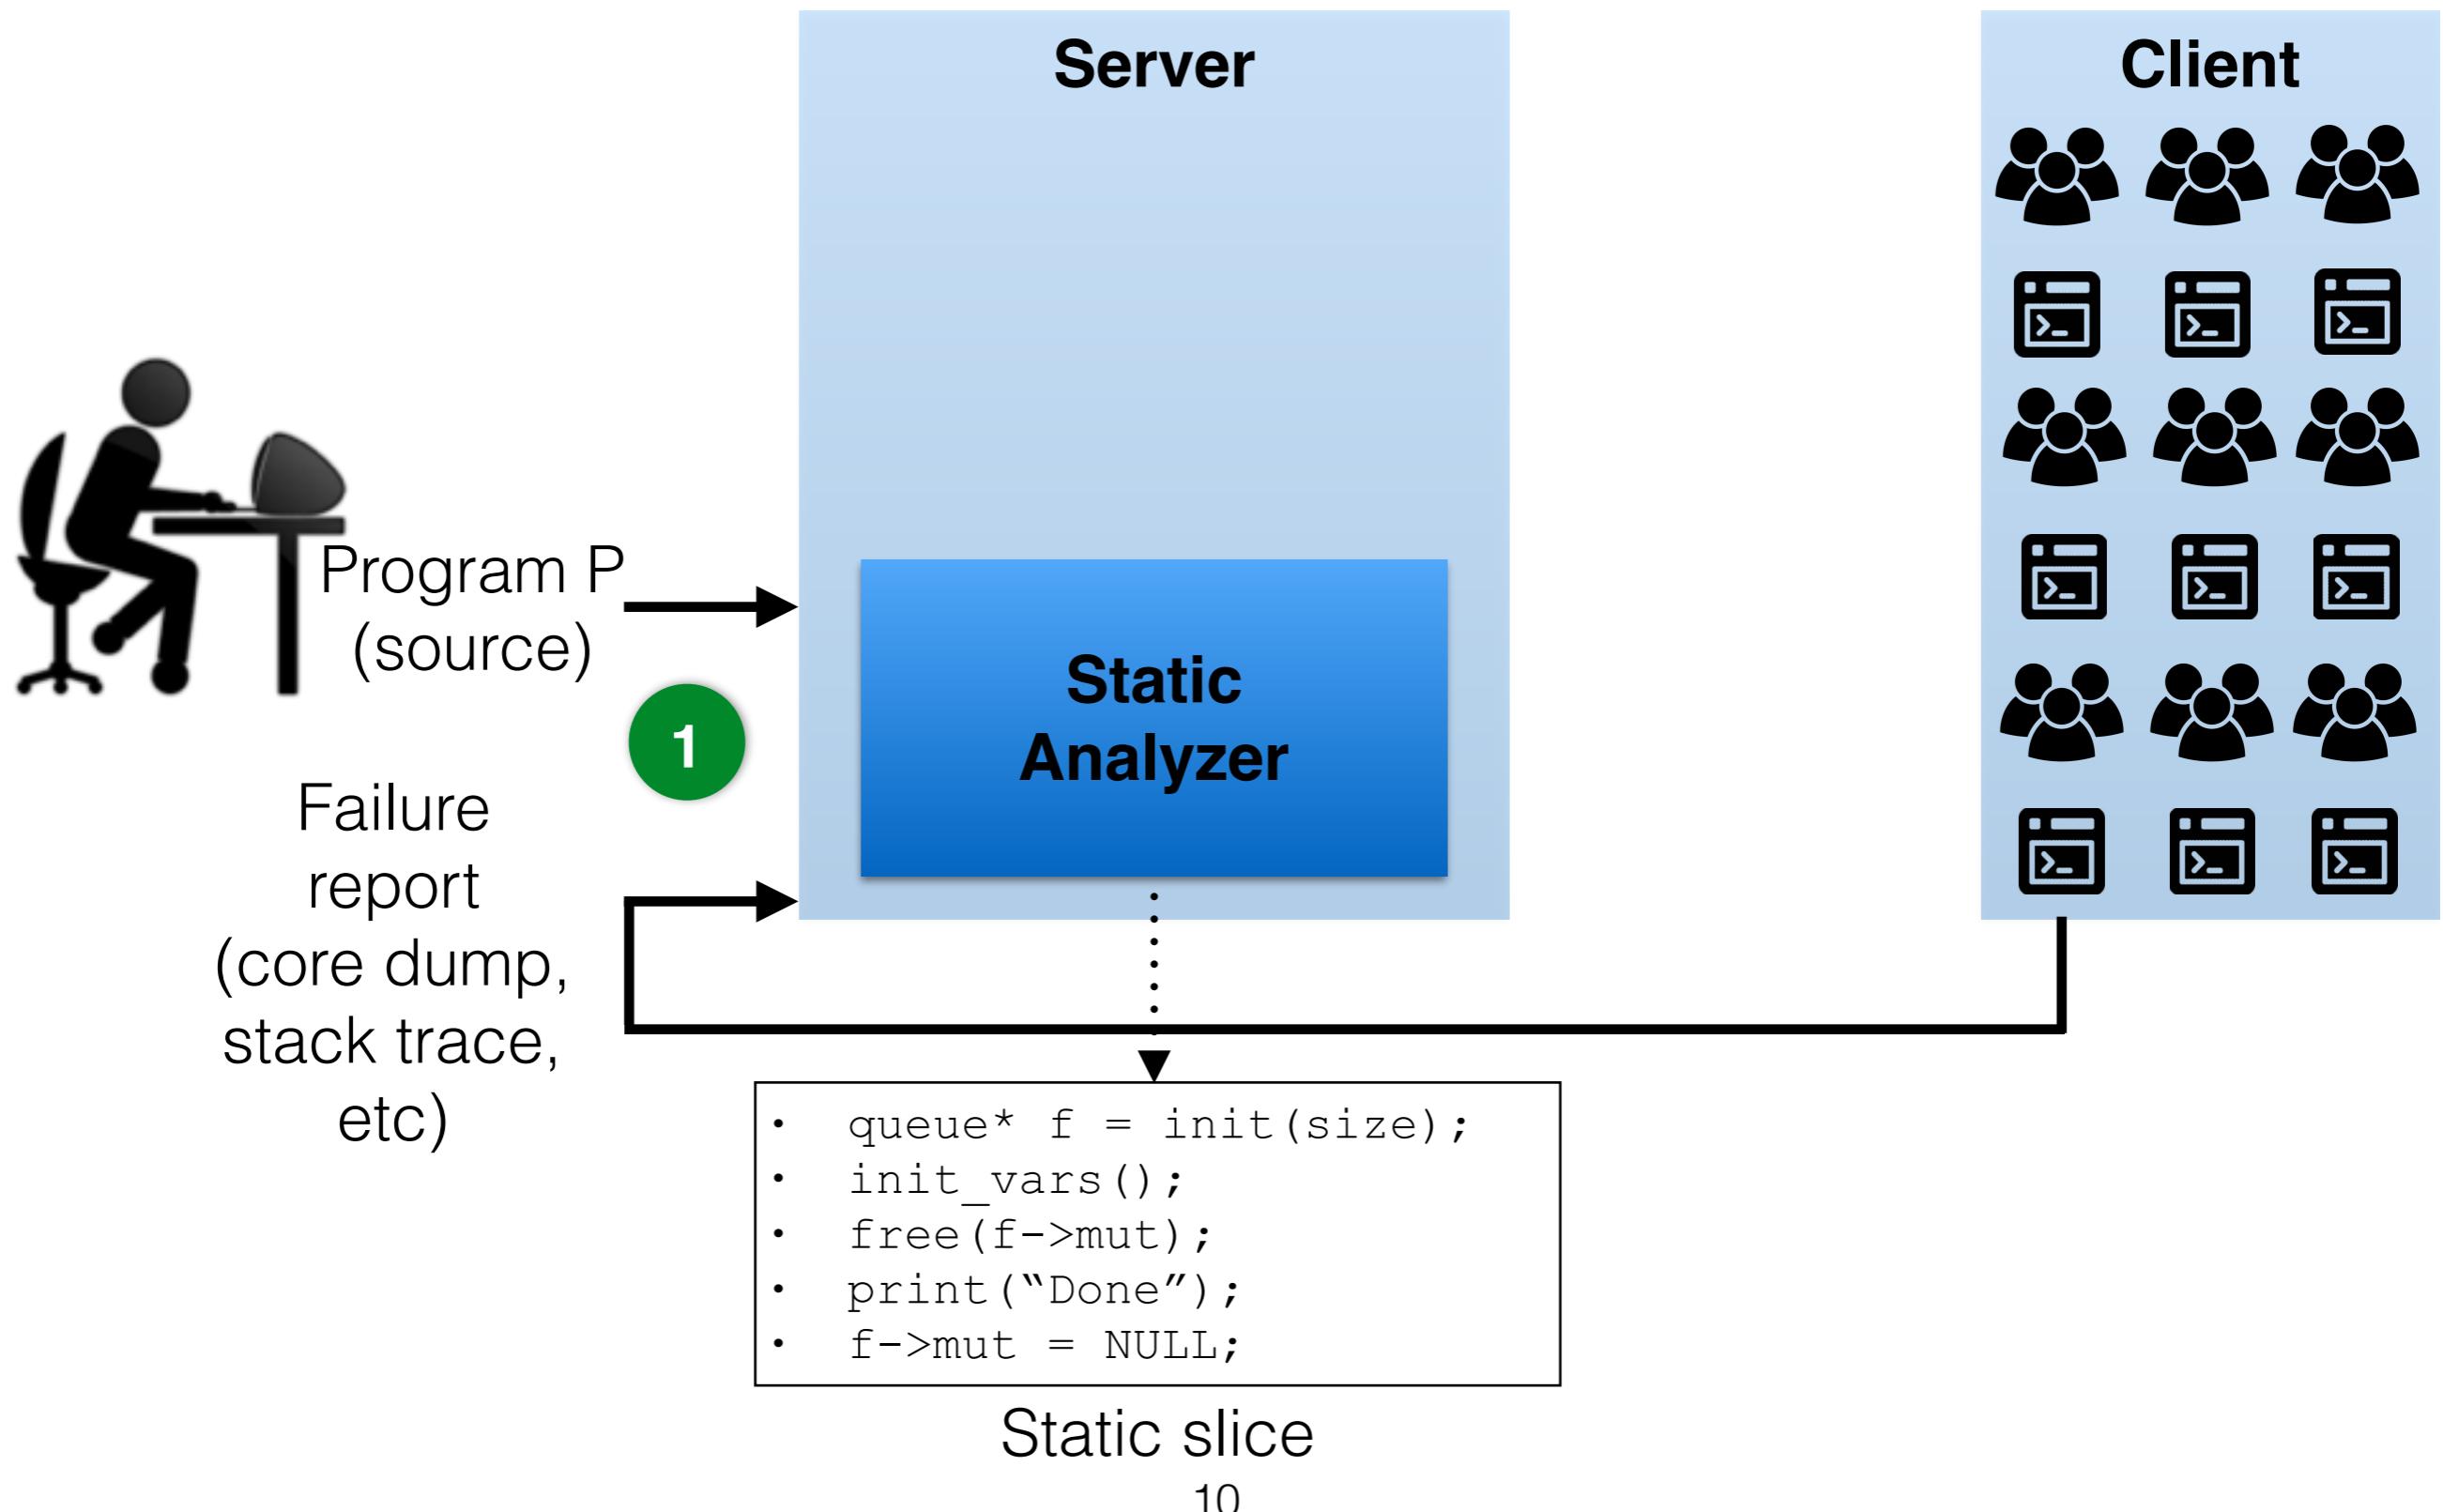
# Gist System Architecture



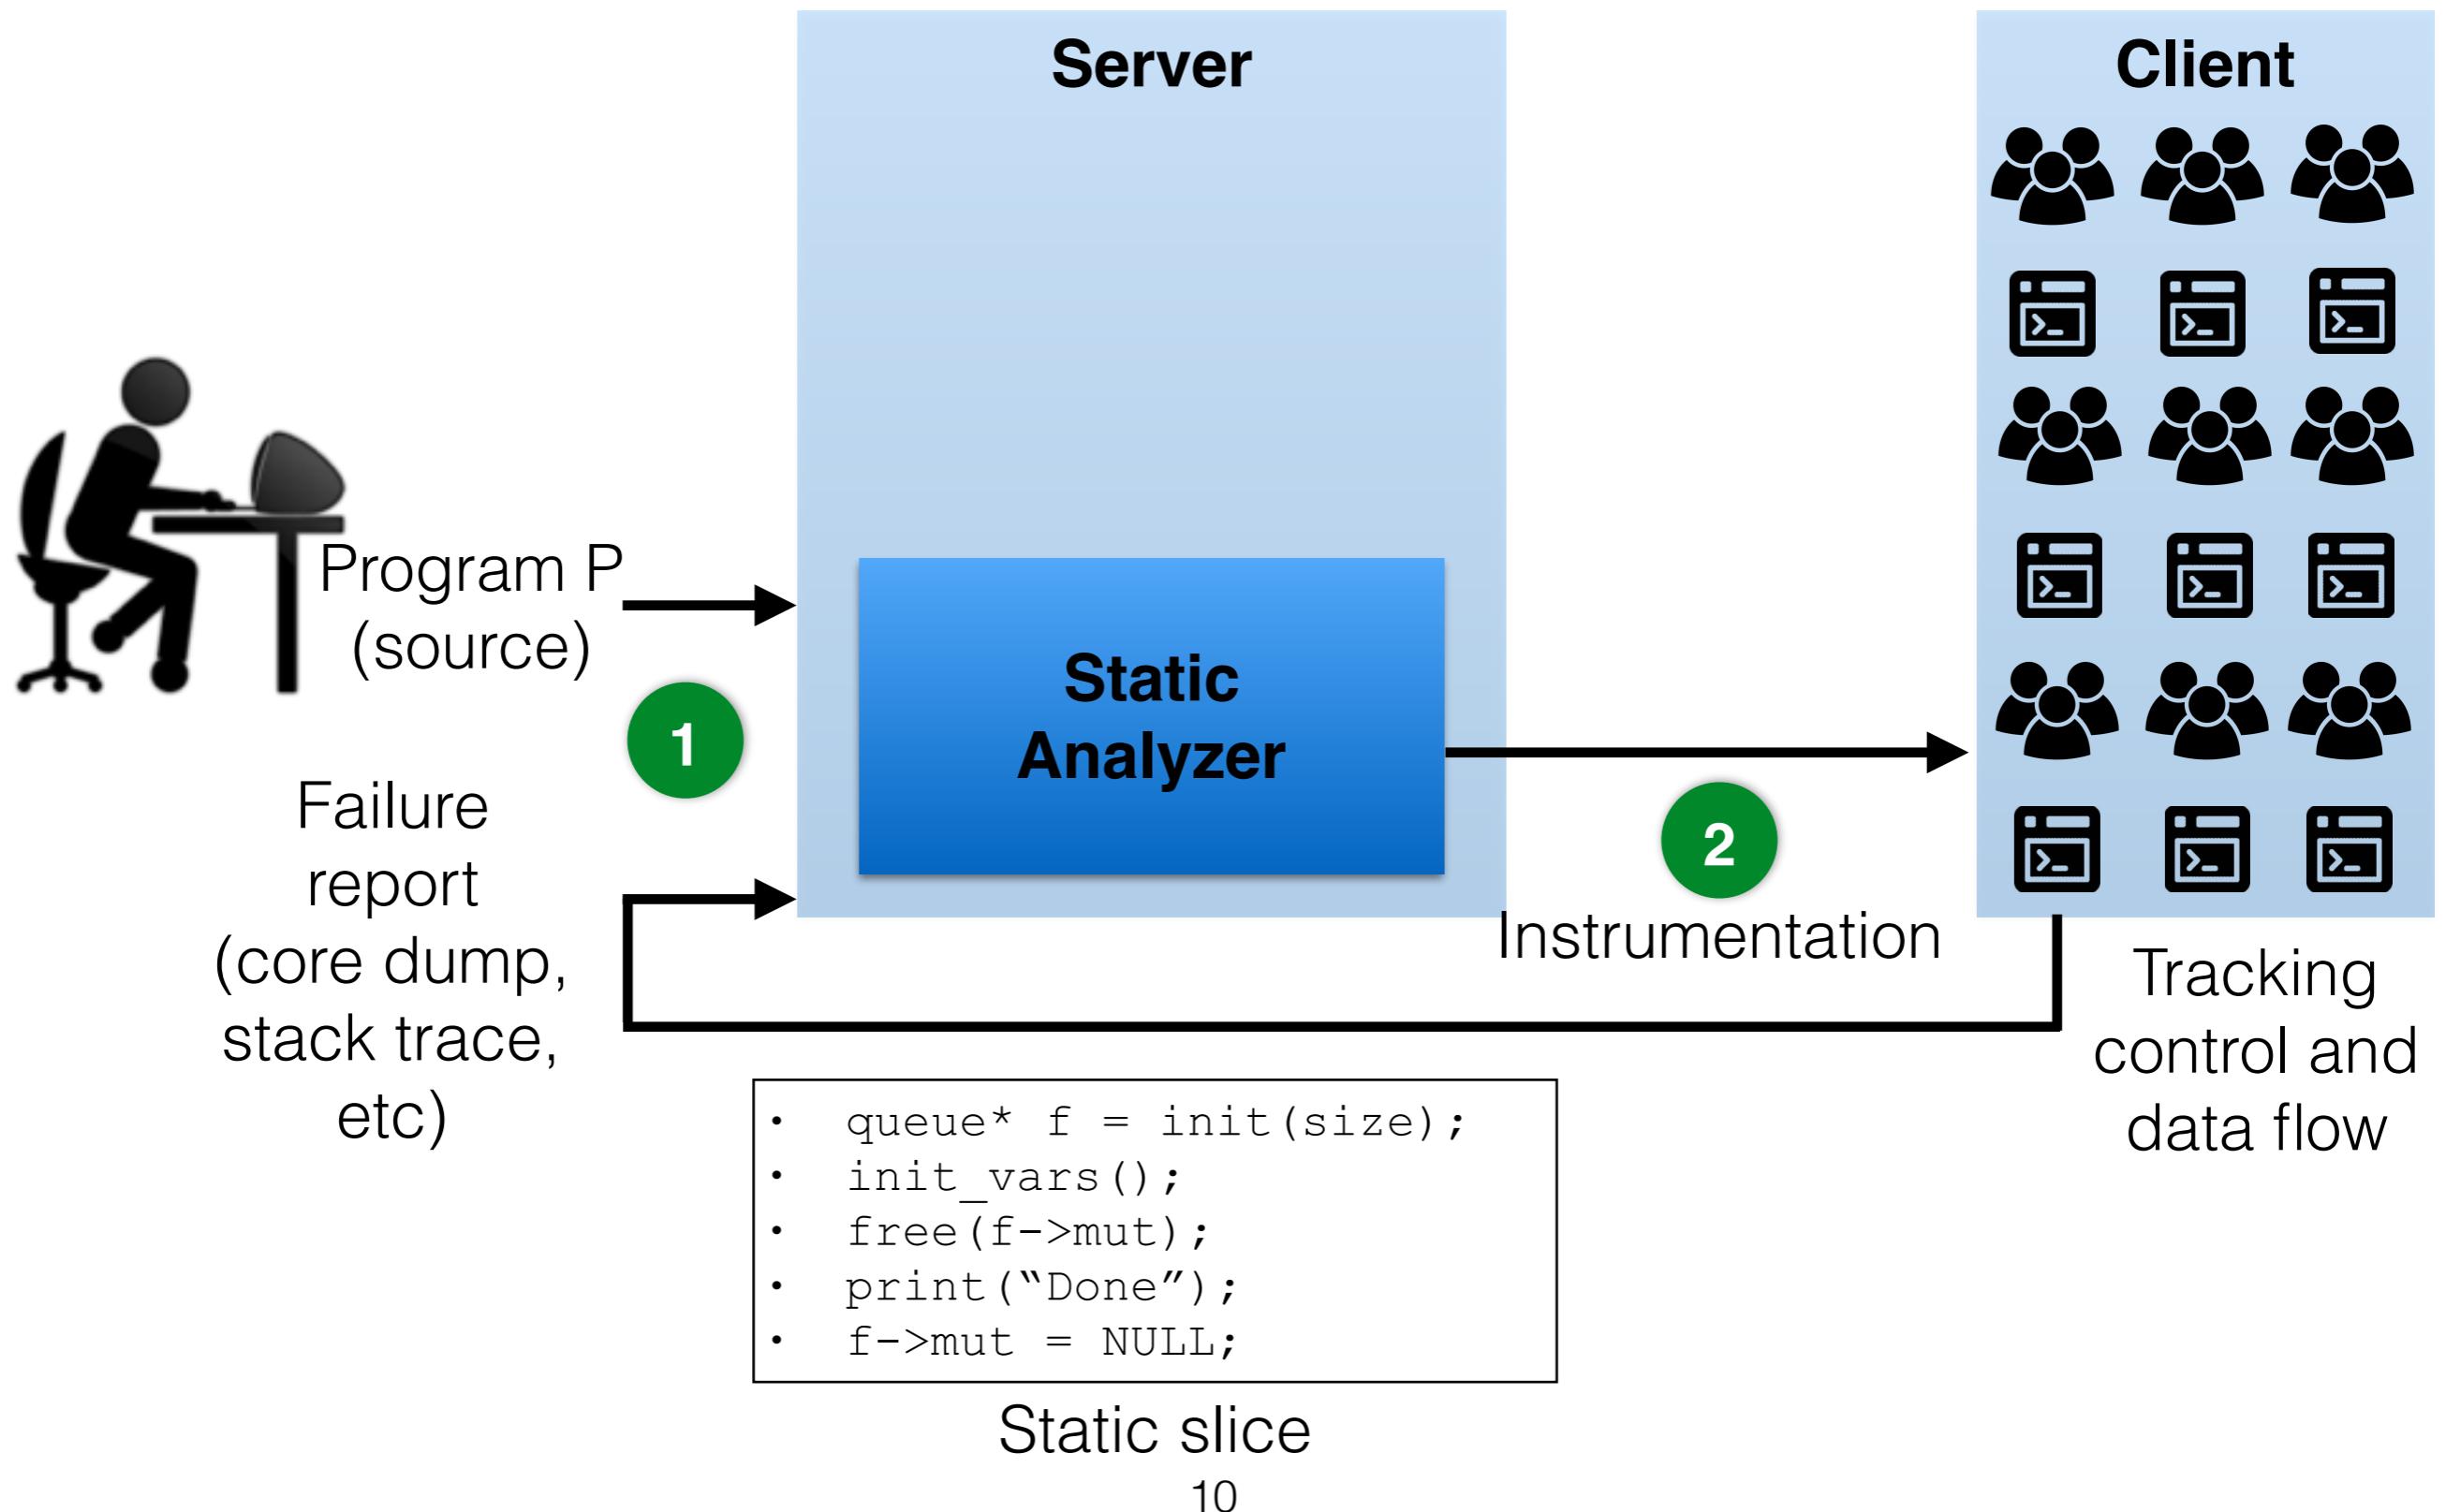
# Gist System Architecture



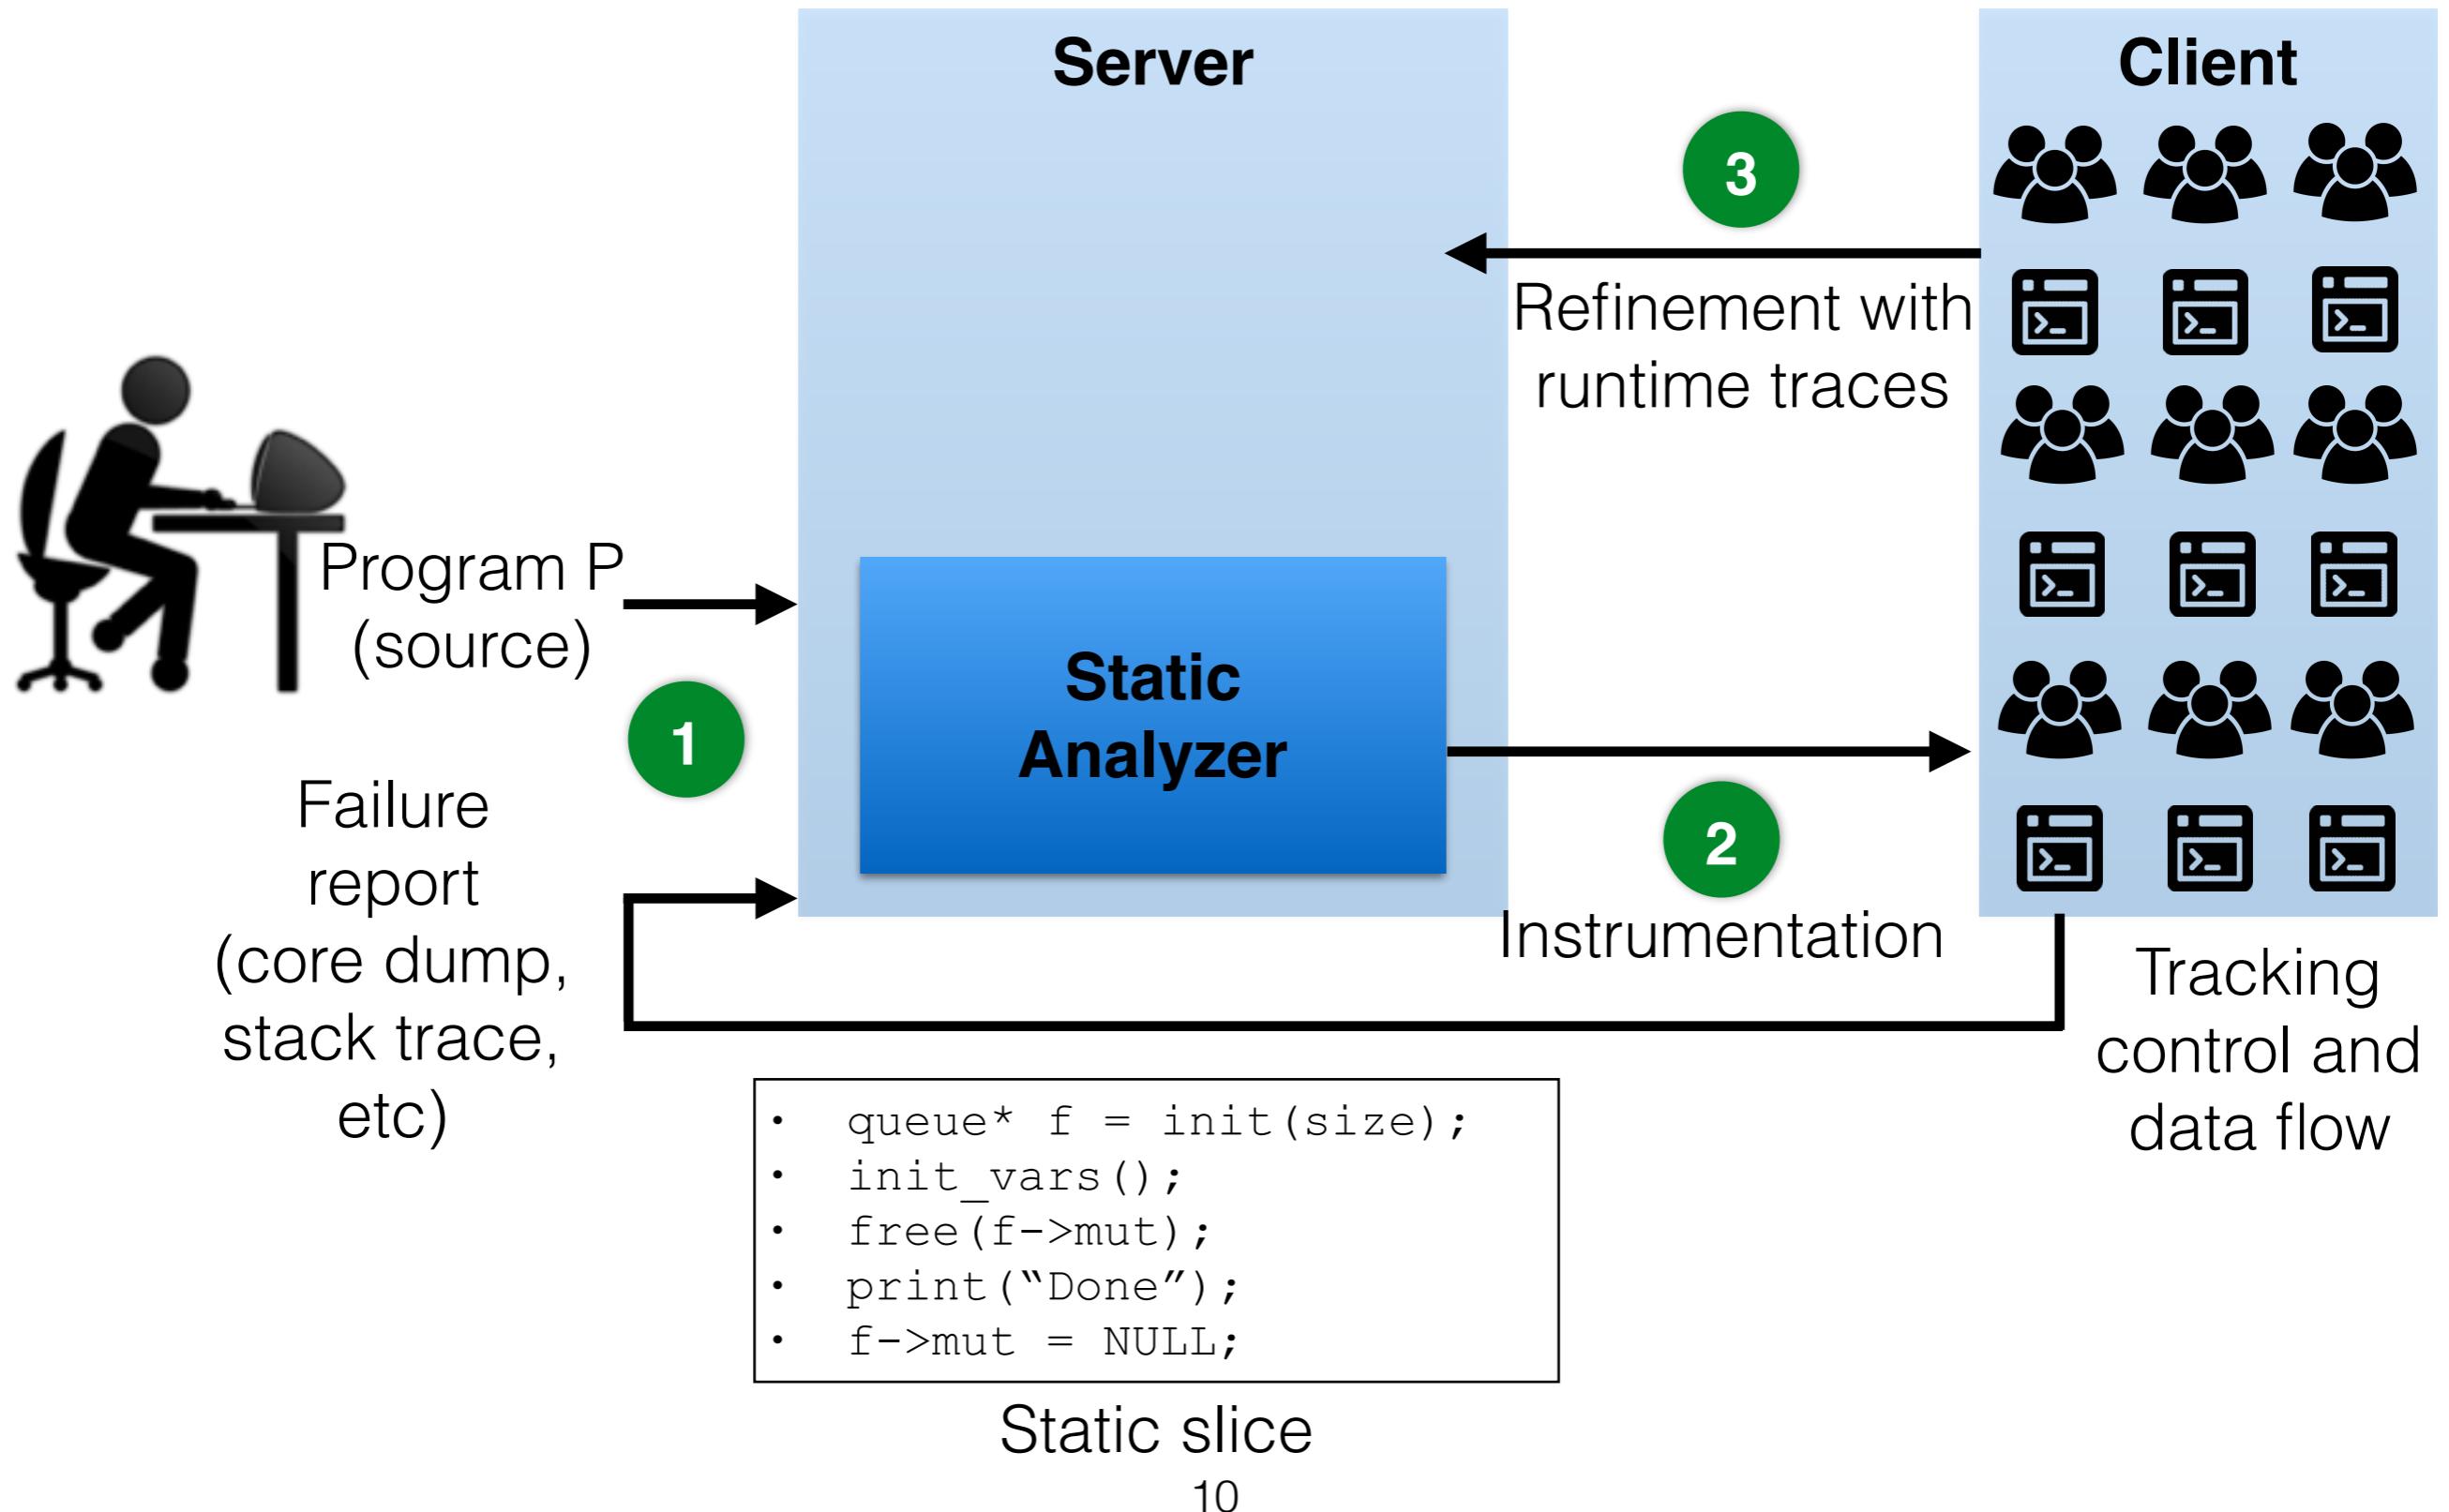
# Gist System Architecture



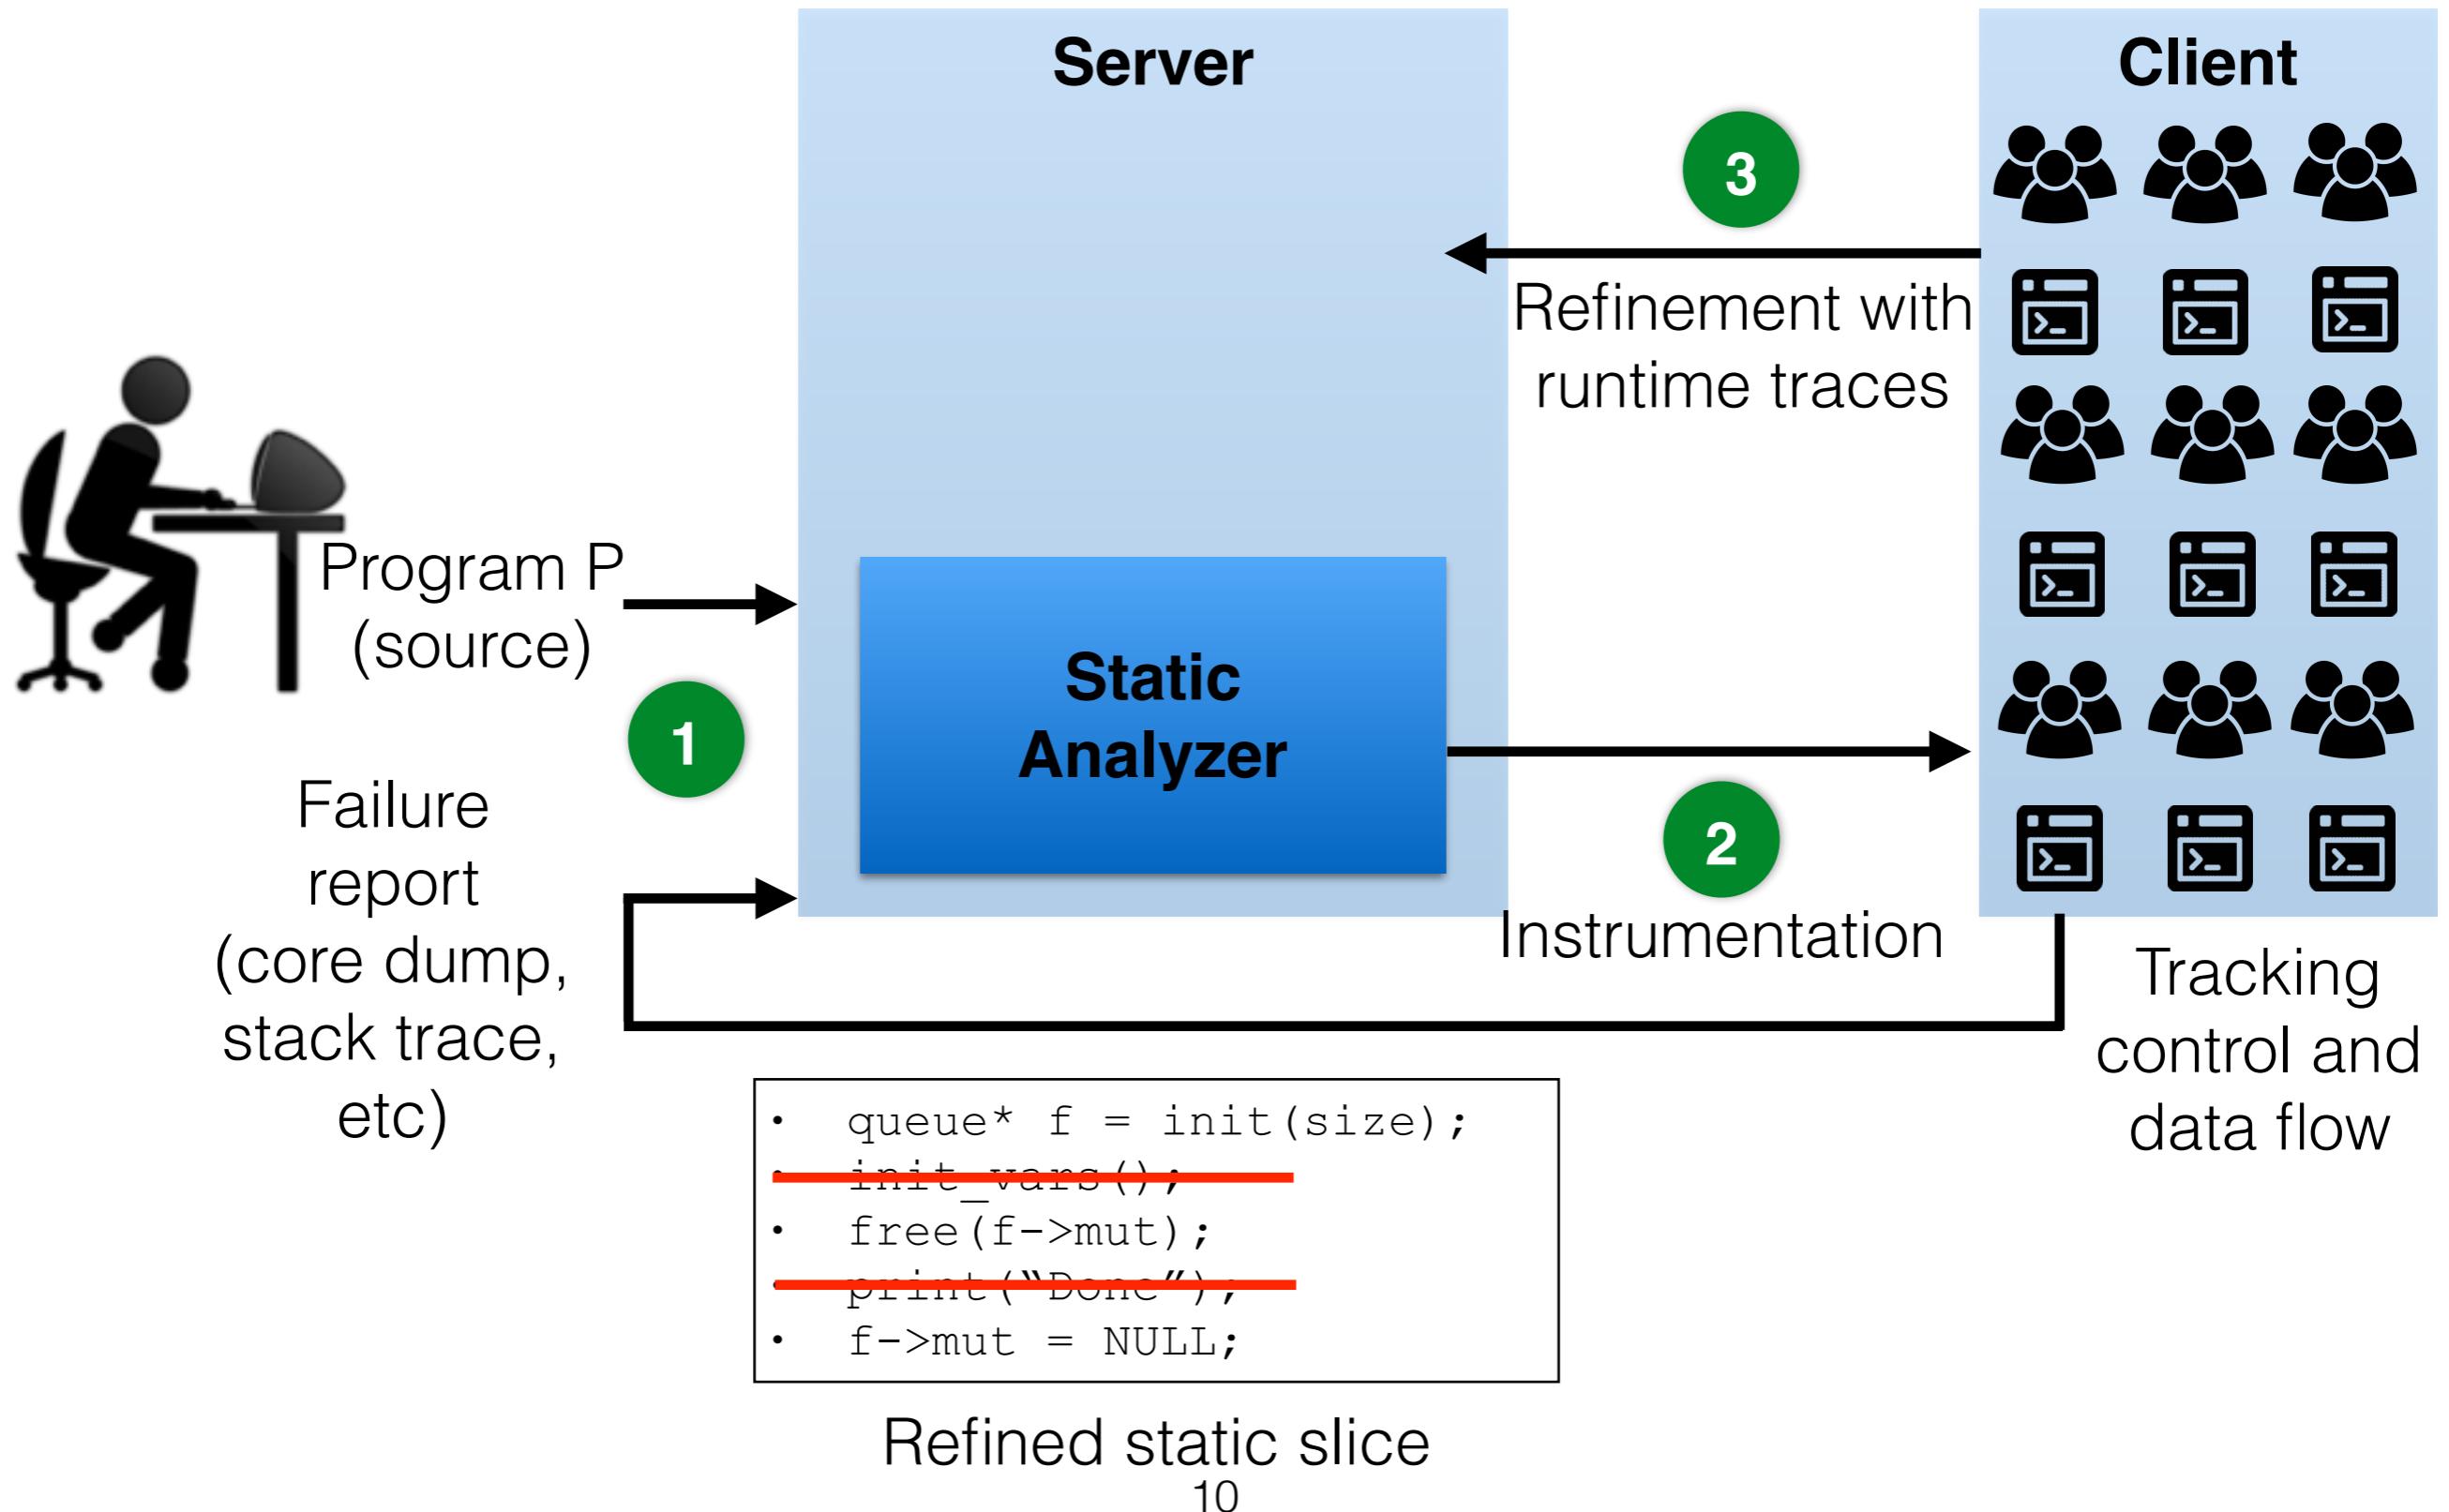
# Gist System Architecture



# Gist System Architecture



# Gist System Architecture



# Gist System Architecture

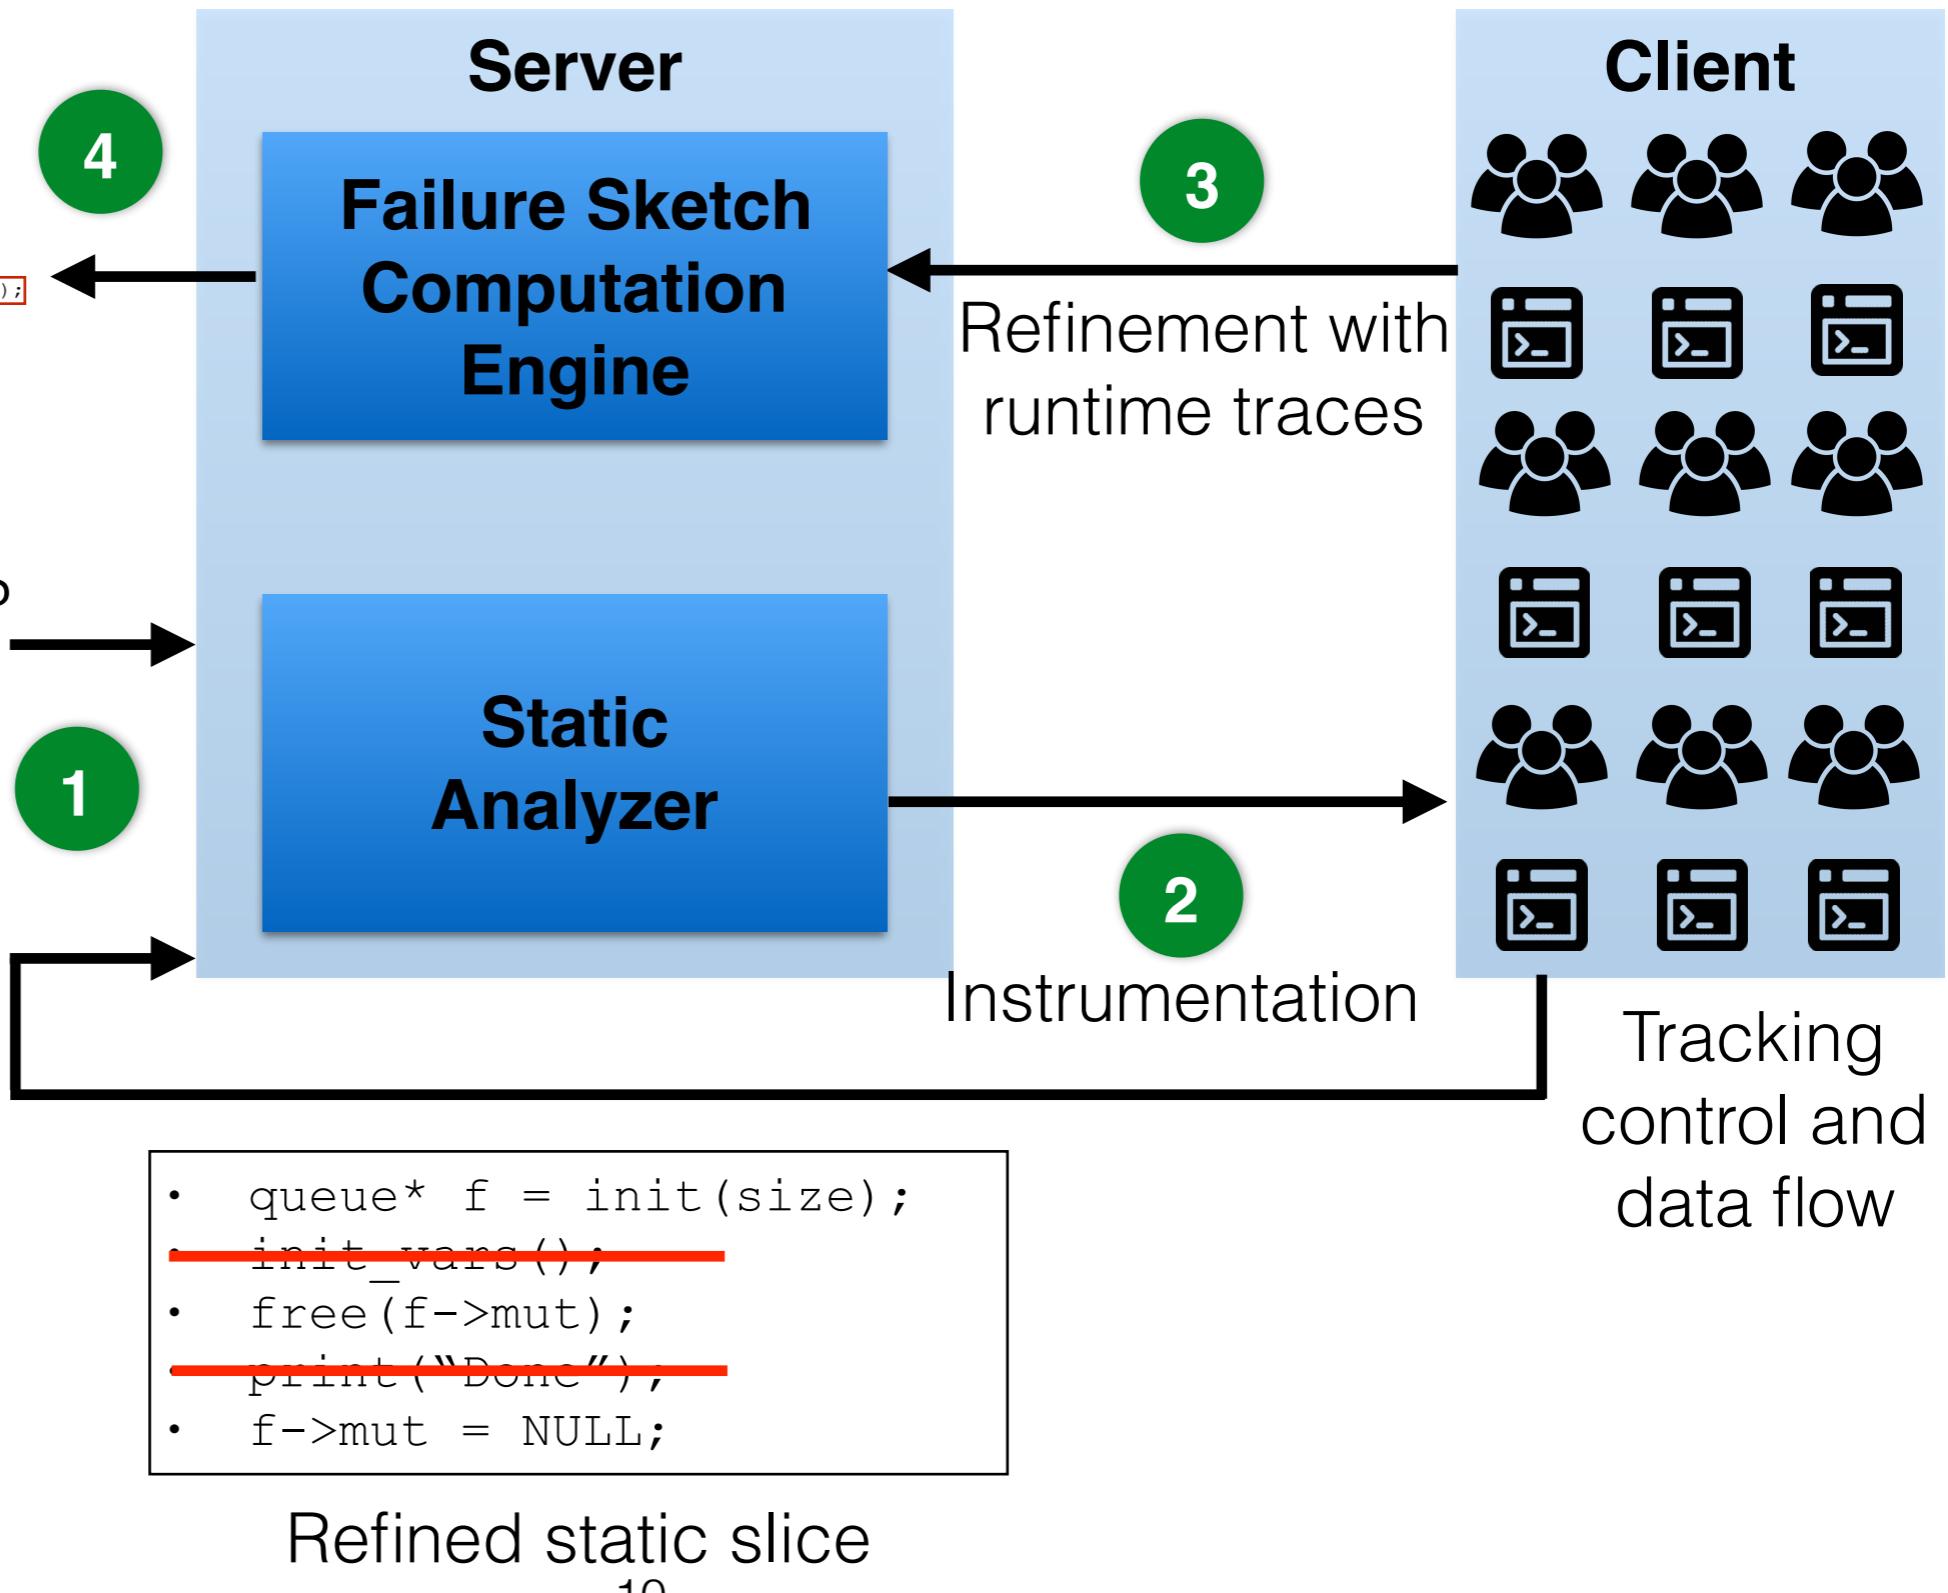
## Failure Sketch

```
Time           Thread 1          Thread 2
1 main() {      1
2   queue* f = init(size);    2
3   create_thread(cons, f);   3
4   ...
5   free(f->mut);          4   cons(queue* f) {
6   f->mut = NULL;          5   ...
7   ...                      6   mutex_unlock(f->mut);
8 }                  7
9 }
```



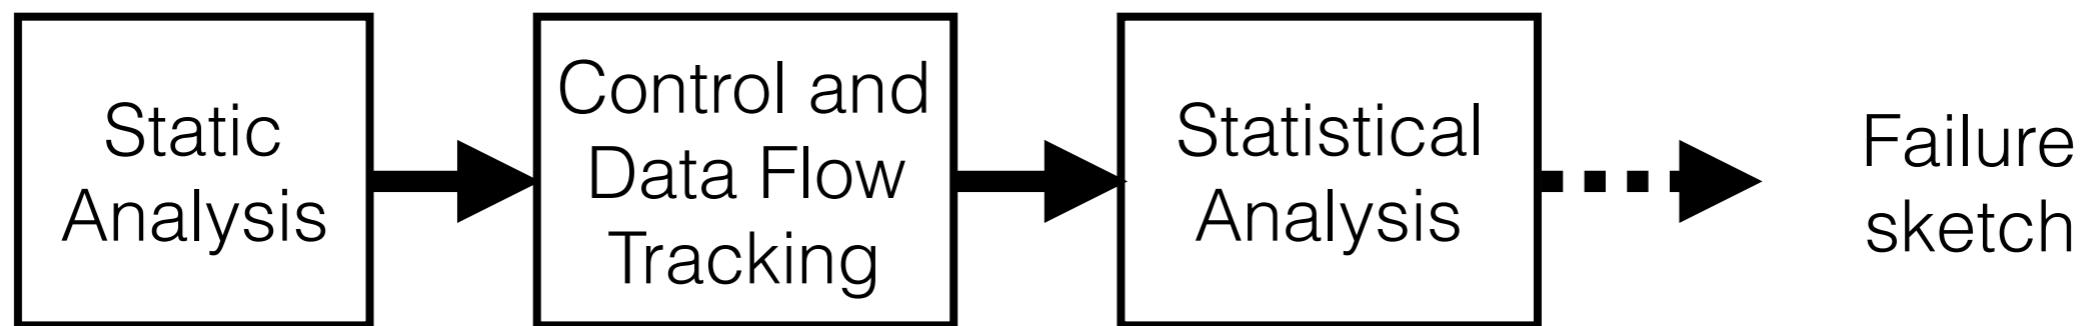
Program P  
(source)

Failure report  
(core dump,  
stack trace,  
etc)



# Outline

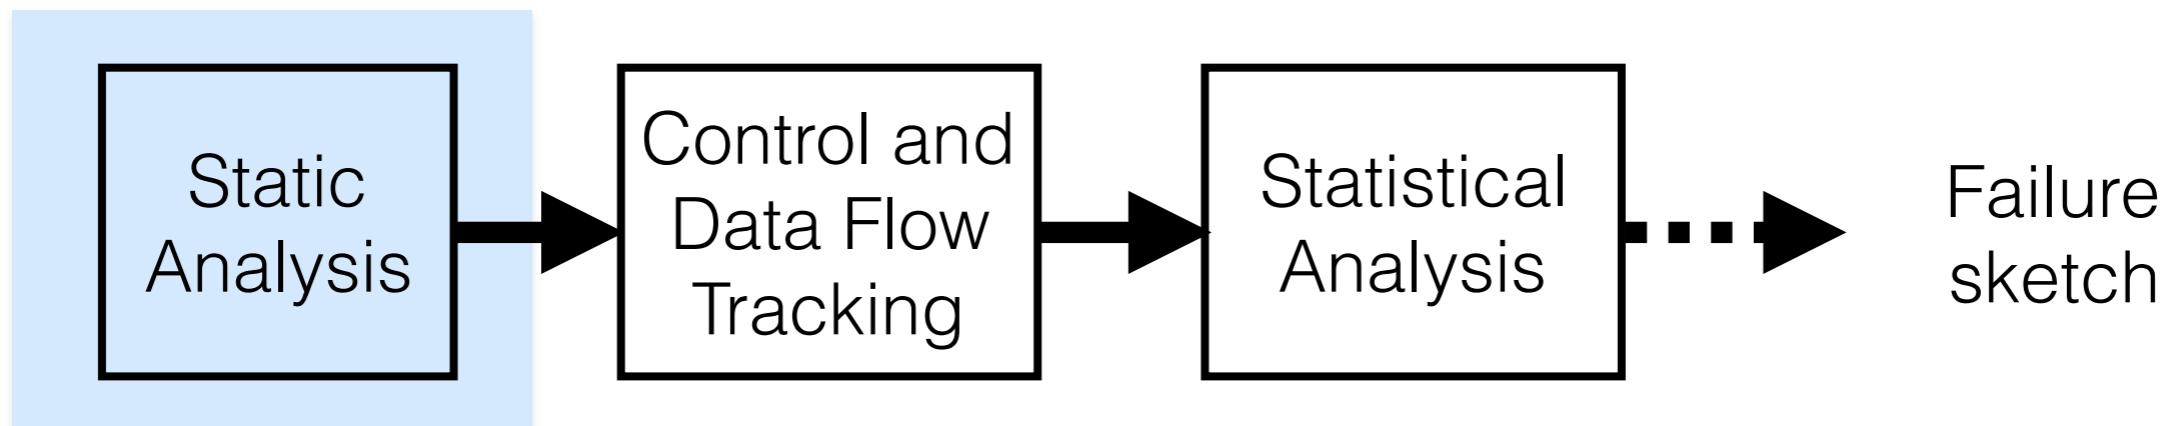
- Challenges
- Design



- Evaluation

# Outline

- Challenges
- Design



- Evaluation

# Static Analysis to Reduce the Overhead

- Computes backward slices
  - *Includes statements with dependencies to the failure*
  - *Excludes all other statements*
- Inter-procedural
  - *Identify dependencies across functions*

# Static Analysis to Reduce the Overhead

- Computes backward slices
  - *Includes statements with dependencies to the failure*
  - *Excludes all other statements*
- Inter-procedural
  - *Identify dependencies across functions*

**Static analysis reduces subsequent runtime tracking (20x)**

# Example

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

# Example

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

# Example

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

# Example

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

# Example

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

# Example

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display size");  
    log("State: %u", s->size);  
}
```



**Segfault**

# Example

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display size");  
    log("State: %u", s->size);  
}
```



Segfault

# Example: Static Backward Slicing

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

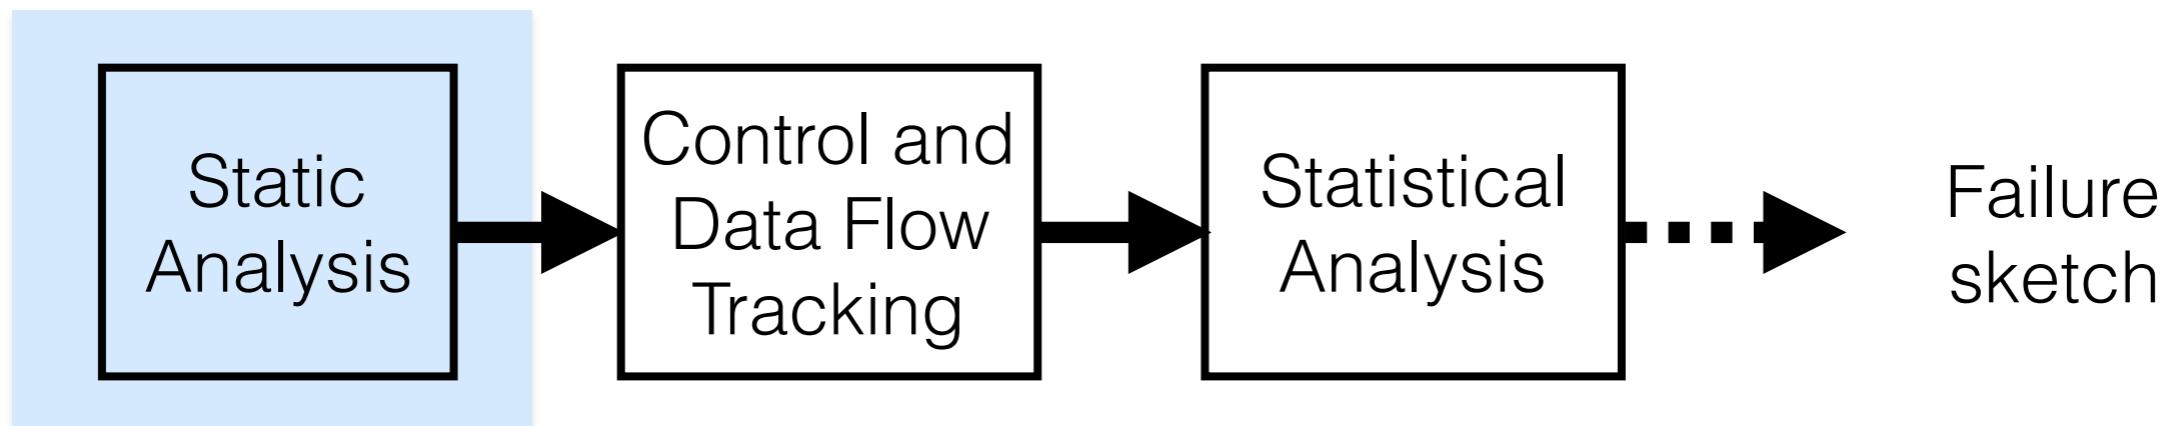
```
void display_size(State* s) {  
    log("Func:display size");  
    log("State: %u", s->size);  
}
```



Segfault

# Outline

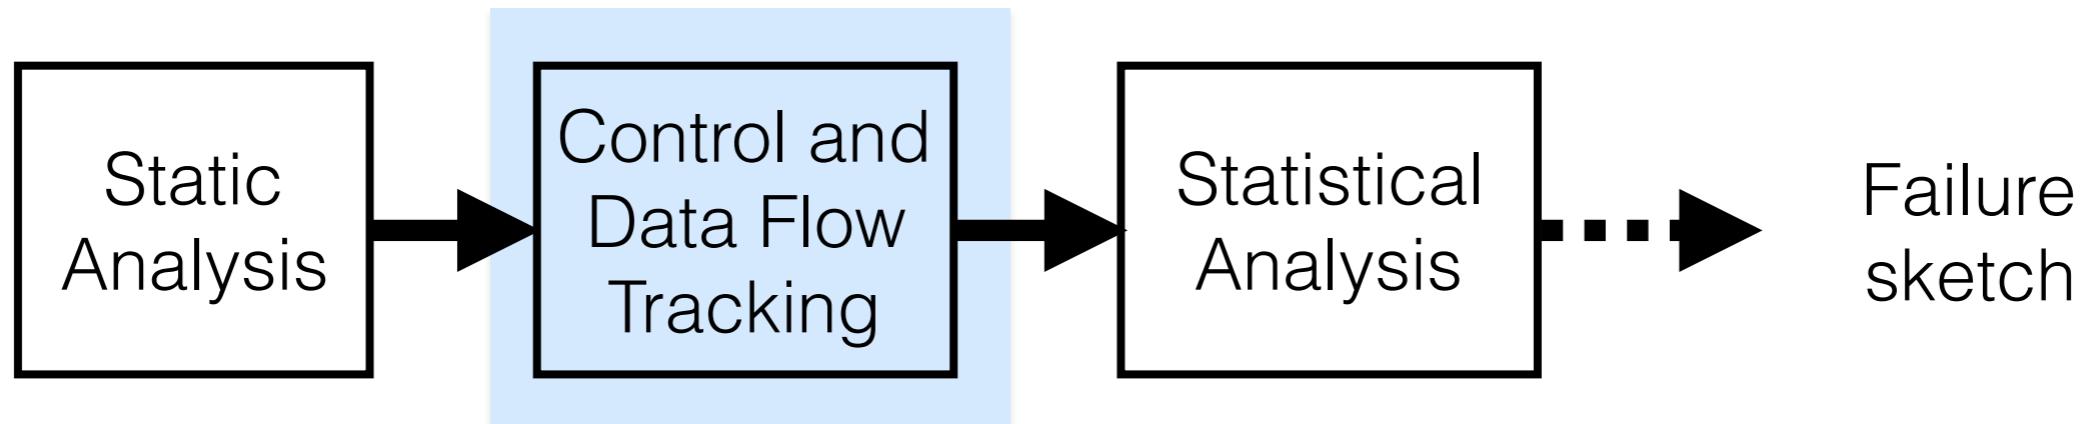
- Challenges
- Design



- Evaluation

# Outline

- Challenges
- Design



- Evaluation

# Low-Overhead Control Flow Tracking

- Software-based tracking is expensive (up to 15x)
- Hardware-based tracking is more efficient
  - *Intel PT: new feature in Intel CPUs (~40%)*
- Gist combines static analysis and hardware-based control flow tracking
  - *Low overhead (~2%)*

# Low-Overhead Control Flow Tracking

- Software-based tracking is expensive (up to 15x)
- Hardware-based tracking is more efficient
  - *Intel PT: new feature in Intel CPUs (~40%)*
- Gist combines static analysis and hardware-based control flow tracking
  - *Low overhead (~2%)*

**Static analysis → Low-overhead control flow tracking**

# Example: Control Flow Tracking (Step 1)

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display size");  
    log("State: %u", s->size);  
}
```



Segfault

# Example: Control Flow Tracking (Step 2)

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display size");  
    log("State: %u", s->size);  
}
```



Segfault

# Example: Control Flow Tracking (Step 2)

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display size");  
    log("State: %u", s->size);  
}
```



Segfault

Static analysis + control flow tracking shorten the sketch

# Data Flow Tracking to Increase Accuracy

# Data Flow Tracking to Increase Accuracy

- Data flow information
  - *Variable values & total order of memory accesses*

# Data Flow Tracking to Increase Accuracy

- Data flow information
  - *Variable values & total order of memory accesses*
- Hardware watchpoints
  - *Allow tracking reads and writes with low overhead*
  - *Allow tracking the total order of accesses*

# Data Flow Tracking to Increase Accuracy

- Data flow information
  - *Variable values & total order of memory accesses*
- Hardware watchpoints
  - *Allow tracking reads and writes with low overhead*
  - *Allow tracking the total order of accesses*
- Monitor multiple clients when run out of watchpoints

# Data Flow Tracking to Increase Accuracy

- Data flow information
  - *Variable values & total order of memory accesses*
- Hardware watchpoints
  - *Allow tracking reads and writes with low overhead*
  - *Allow tracking the total order of accesses*
- Monitor multiple clients when run out of watchpoints

Precise ordering information → High accuracy

# Example: Data Flow Tracking

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display size");  
    log("State: %u", s->size);  
}
```



Segfault

# Example: Data Flow Tracking

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display size");  
    log("State: %u", s->size);  
}
```



Segfault

Watch &s

# Example: Data Flow Tracking

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

```
void display_size(State* s) {  
    log("Func:display size");  
    log("State: %u", s->size);  
}
```



Segfault

Watch &s

# Example: Data Flow Tracking

Thread 1

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

**Success**

Thread 2

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

# Example: Data Flow Tracking

Thread 1

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

**Success**

Thread 2

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

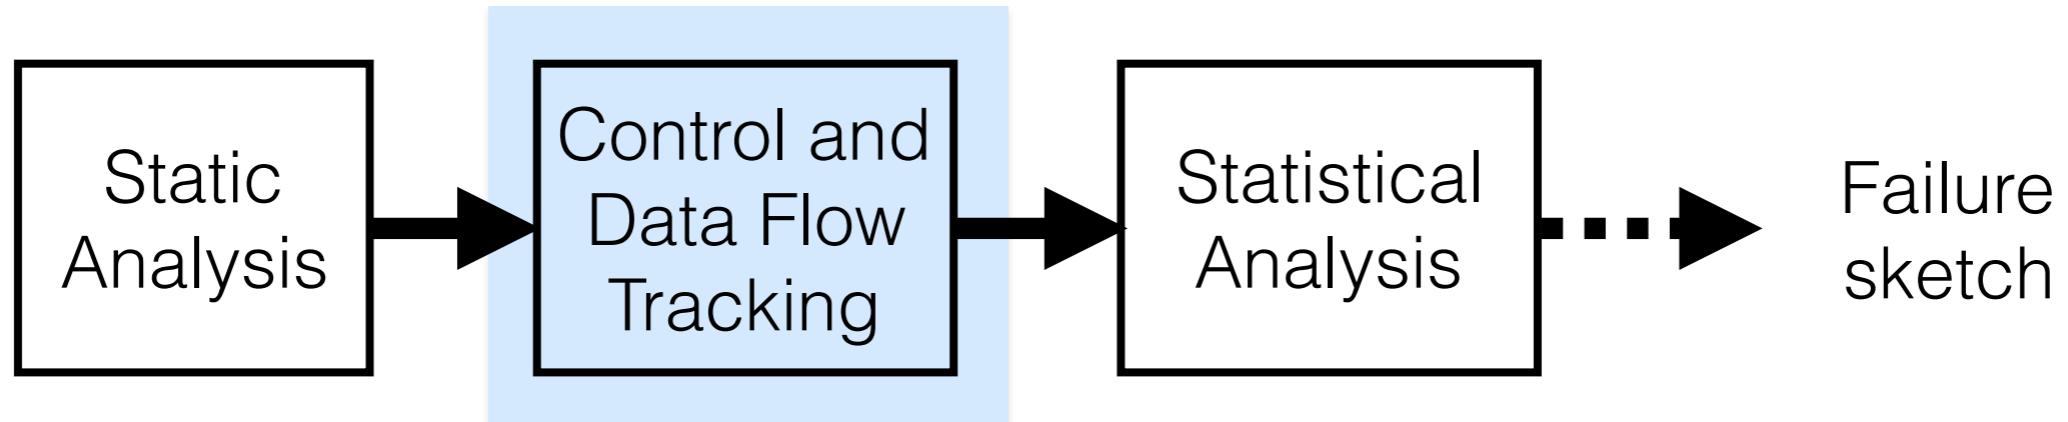
```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Failure**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

# Outline

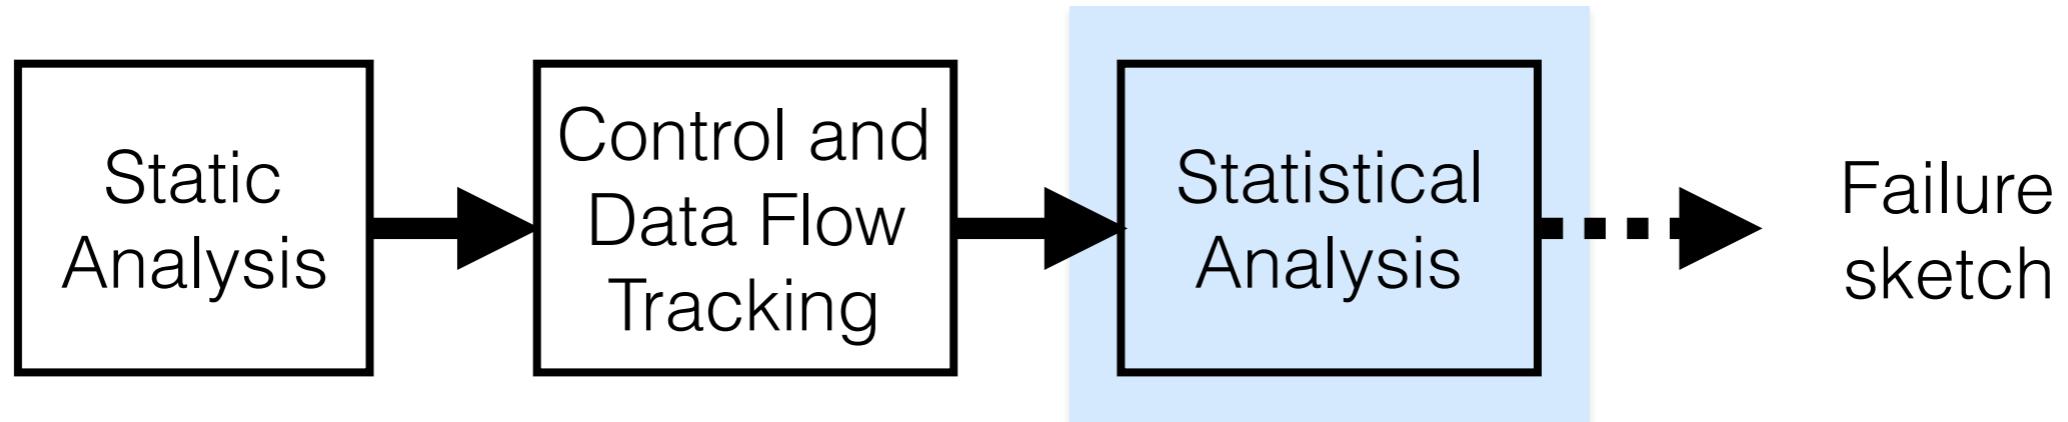
- Challenges
- Design



- Evaluation

# Outline

- Challenges
- Design



- Evaluation

# Statistical Analysis

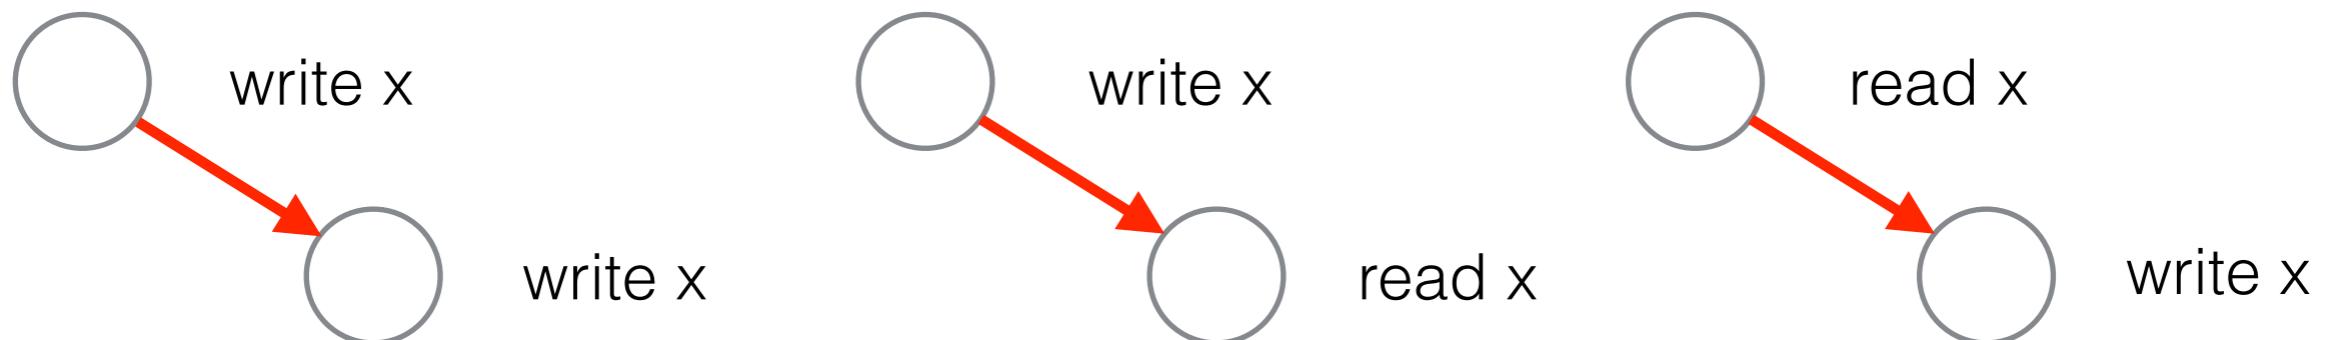
# Statistical Analysis

- Identification of failure predictors<sup>1</sup>
  - *A good predictor portends a failure with high probability (e.g., data races, atomicity violations)*

<sup>1</sup> Liblit, B. et al. Scalable statistical bug isolation. PLDI 2005

# Statistical Analysis

- Identification of failure predictors<sup>1</sup>
  - *A good predictor portends a failure with high probability (e.g., data races, atomicity violations)*
  - *Example: data races*



<sup>1</sup> Liblit, B. et al. Scalable statistical bug isolation. PLDI 2005

# Statistical Analysis

- Identification of failure predictors<sup>1</sup>
  - *A good predictor portends a failure with high probability (e.g., data races, atomicity violations)*
  - *Example: data races*



## Failure predictors across multiple executions

<sup>1</sup> Liblit, B. et al. Scalable statistical bug isolation. PLDI 2005

# Example: Statistical Analysis

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

## Success

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

## Failure

```
void display_size(State* s)  
{  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

# Example: Statistical Analysis

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Failure**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}  
  
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

**Failure**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

# Example: Statistical Analysis

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Failure**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}  
  
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

**Failure**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Failure**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

# Example: Statistical Analysis

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}  
  
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

**Failure**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}  
  
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```

**Failure**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

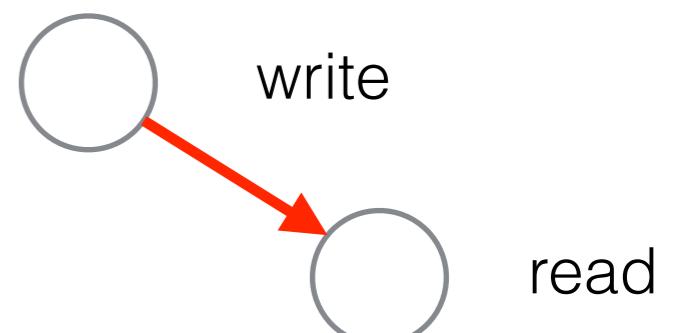
```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}  
  
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Success**

```
void cleanup(State* s) {  
    log("Func:cleanup");  
    if(verbose)  
        log("Cleaning up %p", s);  
    delete s;  
}
```

**Failure**

```
void display_size(State* s) {  
    log("Func:display_size");  
    log("State: %u", s->size);  
}
```



# Example: Statistical Analysis

```
void display_size(State* s) {
    log("Func:display_size");
    log("State: %u", s->size);
}

void cleanup(State* s) {
    log("Func:cleanup");
    if(verbose)
        log("Cleaning up %p", s);
    delete s;
}
```

**Success**

```
void cleanup(State* s) {
    log("Func:cleanup");
    if(verbose)
        log("Cleaning up %p", s);
    delete s;
}

void display_size(State* s) {
    log("Func:display_size");
    log("State: %u", s->size);
}
```

**Failure**

```
void display_size(State* s) {
    log("Func:display_size");
    log("State: %u", s->size);
}

void cleanup(State* s) {
    log("Func:cleanup");
    if(verbose)
        log("Cleaning up %p", s);
    delete s;
}
```

**Success**

```
void cleanup(State* s) {
    log("Func:cleanup");
    if(verbose)
        log("Cleaning up %p", s);
    delete s;
}

void display_size(State* s) {
    log("Func:display_size");
    log("State: %u", s->size);
}
```

**Failure**

```
void display_size(State* s) {
    log("Func:display_size");
    log("State: %u", s->size);
}

void cleanup(State* s) {
    log("Func:cleanup");
    if(verbose)
        log("Cleaning up %p", s);
    delete s;
}
```

**Success**

```
void display_size(State* s) {
    log("Func:display_size");
    log("State: %u", s->size);
}

void cleanup(State* s) {
    log("Func:cleanup");
    if(verbose)
        log("Cleaning up %p", s);
    delete s;
}
```

**Success**

```
void display_size(State* s) {
    log("Func:display_size");
    log("State: %u", s->size);
}

void cleanup(State* s) {
    log("Func:cleanup");
    if(verbose)
        log("Cleaning up %p", s);
    delete s;
}
```

**Success**

```
void display_size(State* s) {
    log("Func:display_size");
    log("State: %u", s->size);
}

void cleanup(State* s) {
    log("Func:cleanup");
    if(verbose)
        log("Cleaning up %p", s);
    delete s;
}
```

**Success**

```
void display_size(State* s) {
    log("Func:display_size");
    log("State: %u", s->size);
}

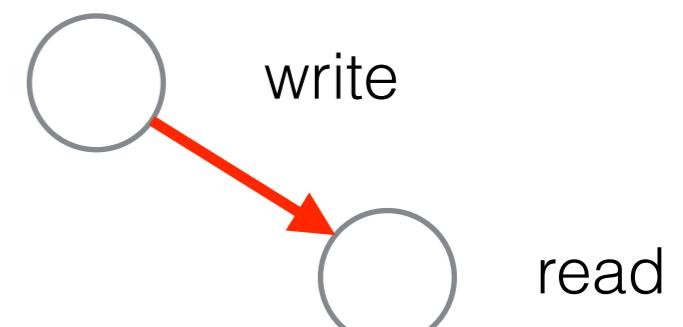
void cleanup(State* s) {
    log("Func:cleanup");
    if(verbose)
        log("Cleaning up %p", s);
    delete s;
}
```

**Success**

```
void cleanup(State* s) {
    log("Func:cleanup");
    if(verbose)
        log("Cleaning up %p", s);
    delete s;
}
```

**Failure**

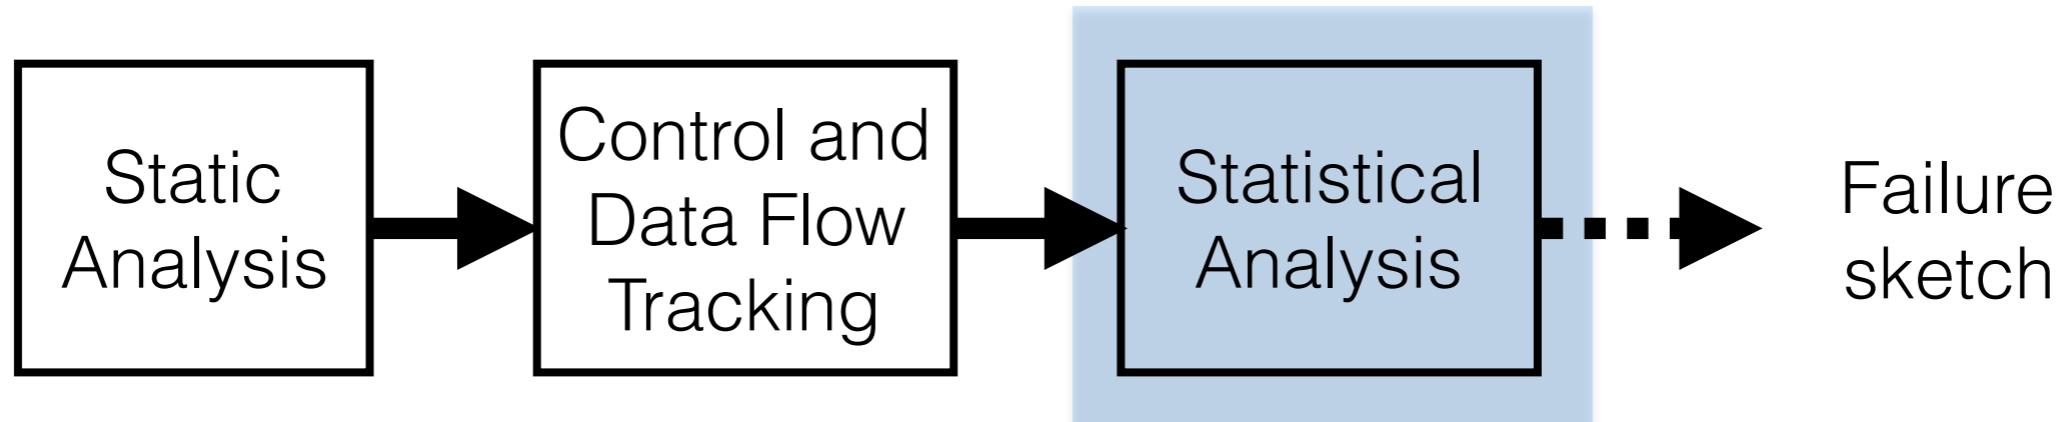
```
void display_size(State* s) {
    log("Func:display_size");
    log("State: %u", s->size);
}
```



## Static analysis + cooperative dynamic analysis

# Outline

- Challenges
- Design



- Evaluation

# Outline

- Challenges
- Design
- Evaluation
  - *Does Gist help developers do root cause diagnosis?*
  - *Is Gist efficient?*
  - *Is Gist accurate?*

# Experimental Setup

- Client side executions are analyzed in the lab
- Real world server and desktop programs



# Do Failure Sketches Help Developers?

- We manually analyzed the usefulness of Gist for 11 failures
- Gist-identified failure predictors point to root causes
  - *Developers eliminated those root causes to fix the bugs*
  - *Average number of statements to look at: 7*

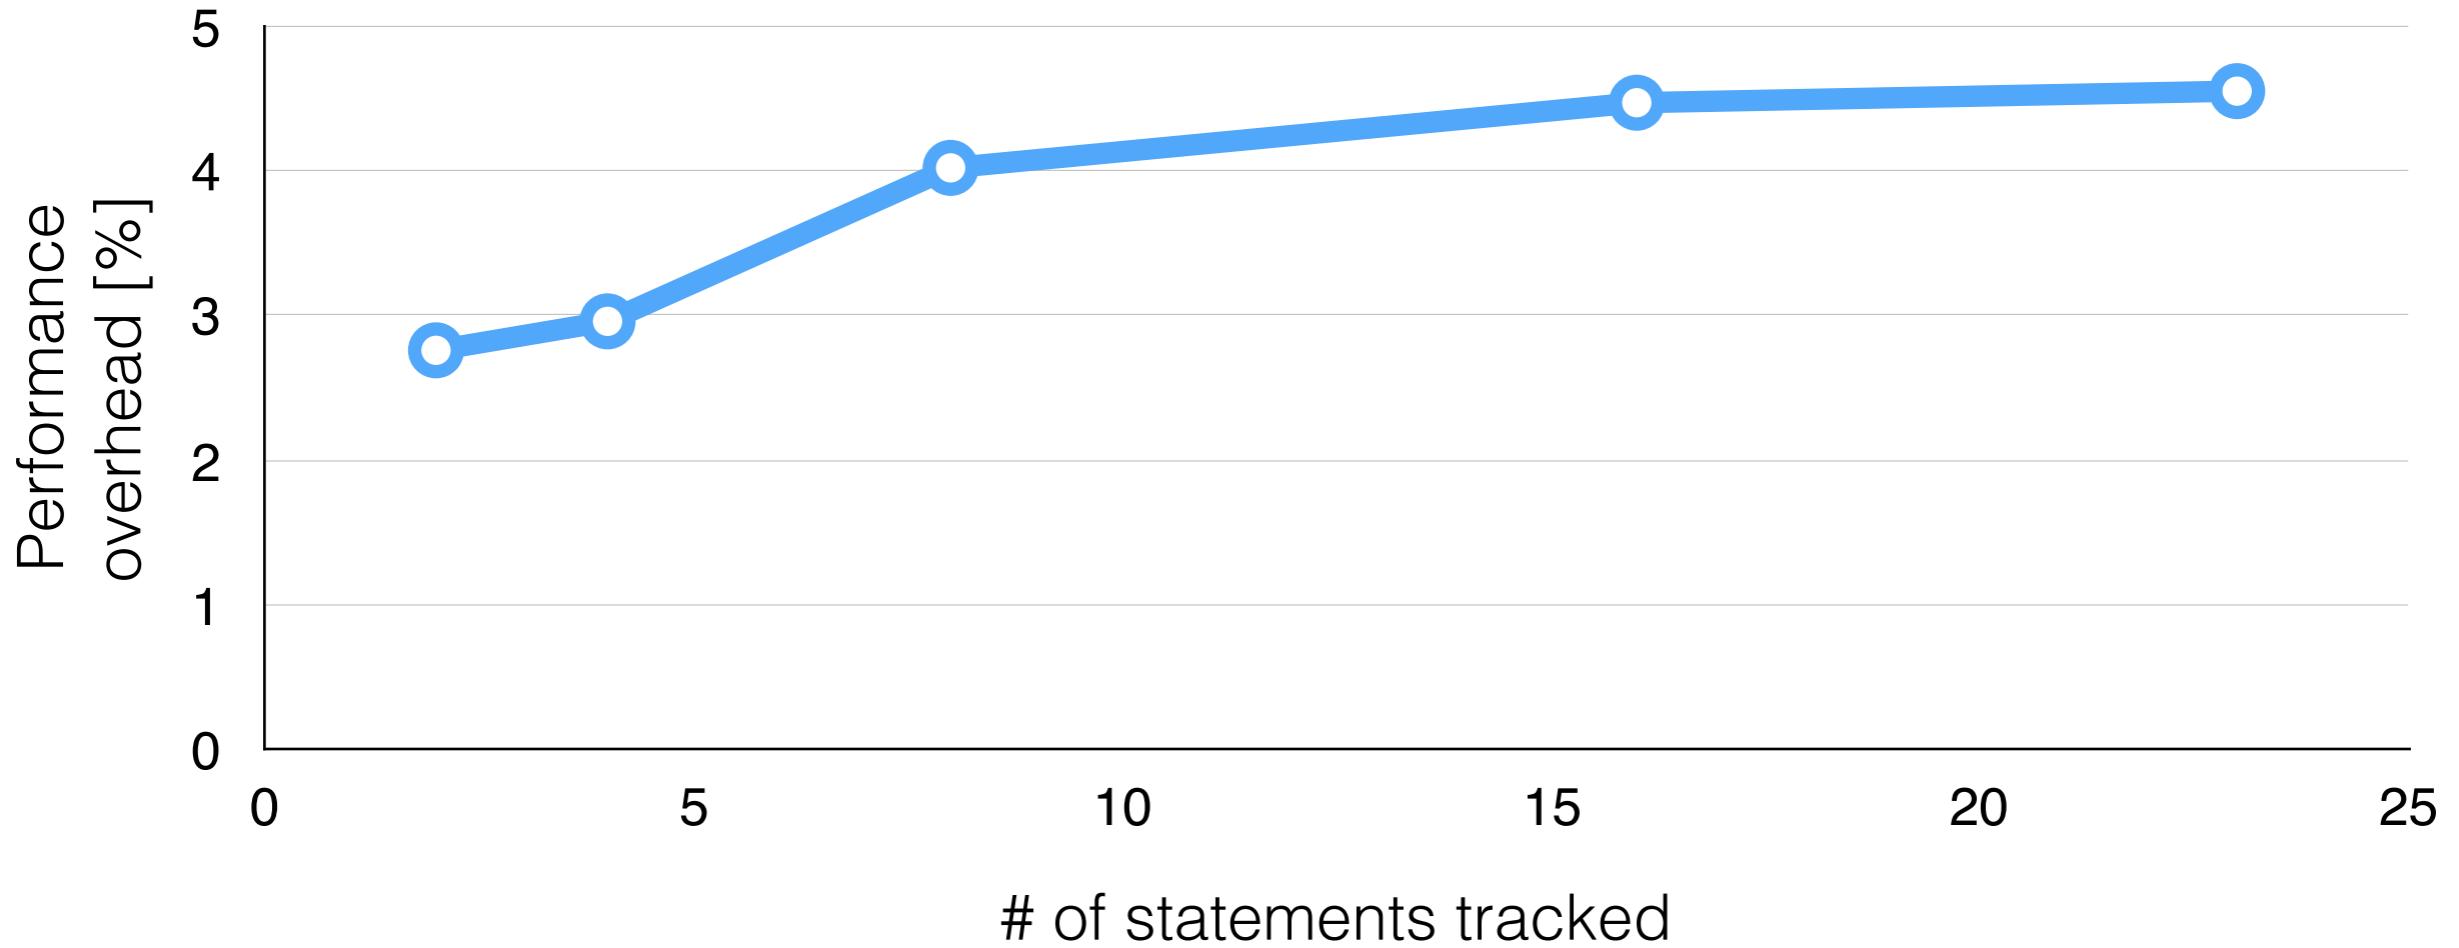
# Do Failure Sketches Help Developers?

- We manually analyzed the usefulness of Gist for 11 failures
- Gist-identified failure predictors point to root causes
  - *Developers eliminated those root causes to fix the bugs*
  - *Average number of statements to look at: 7*

**Gist points developers to the root causes of failures**

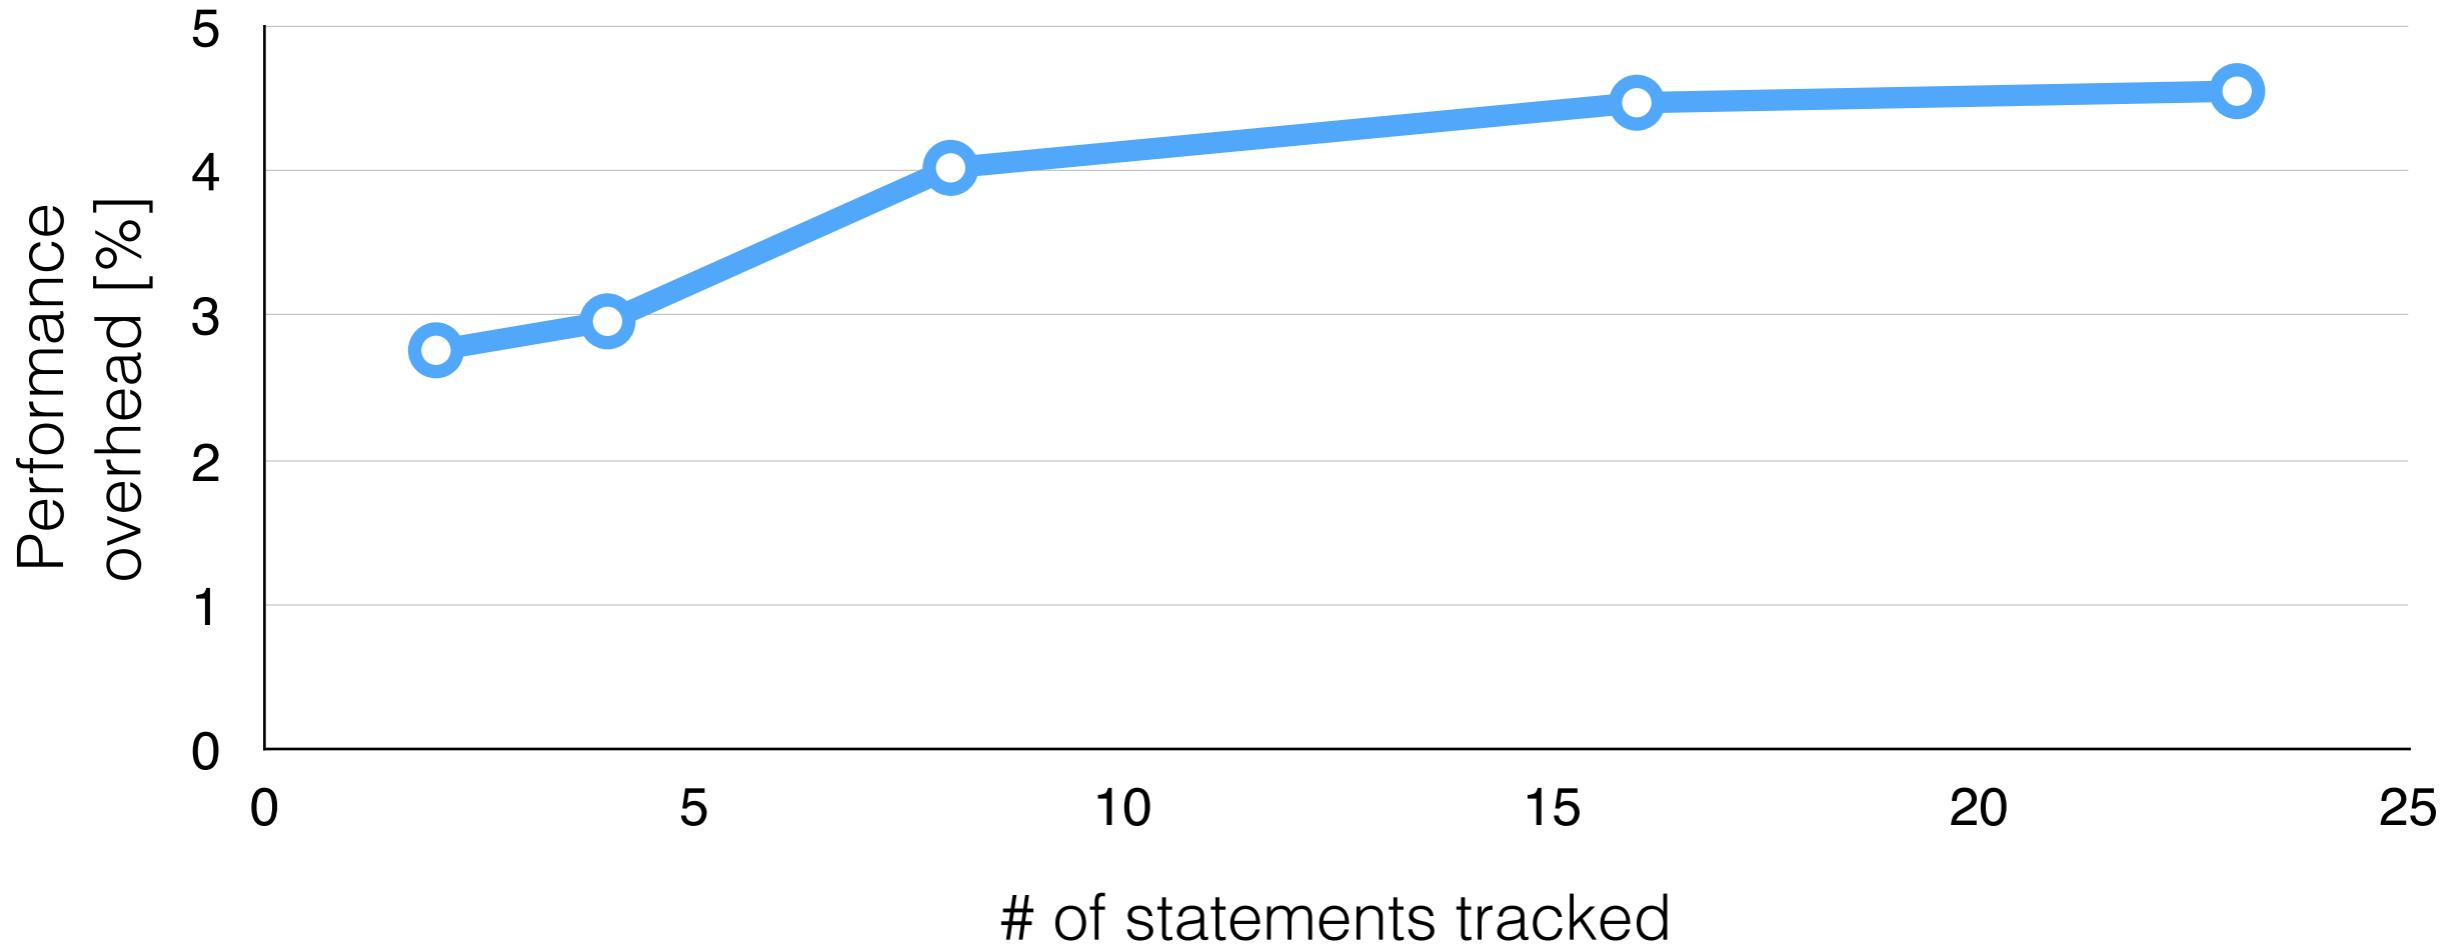
# Efficiency

(Control & data flow tracking)



# Efficiency

(Control & data flow tracking)



**Gist has low average overhead (always below 5%)**

Accuracy [%]

Apache-1

Apache-2

Apache-3

Apache-4

Cppcheck-1

Cppcheck-2

Curl

Transmission

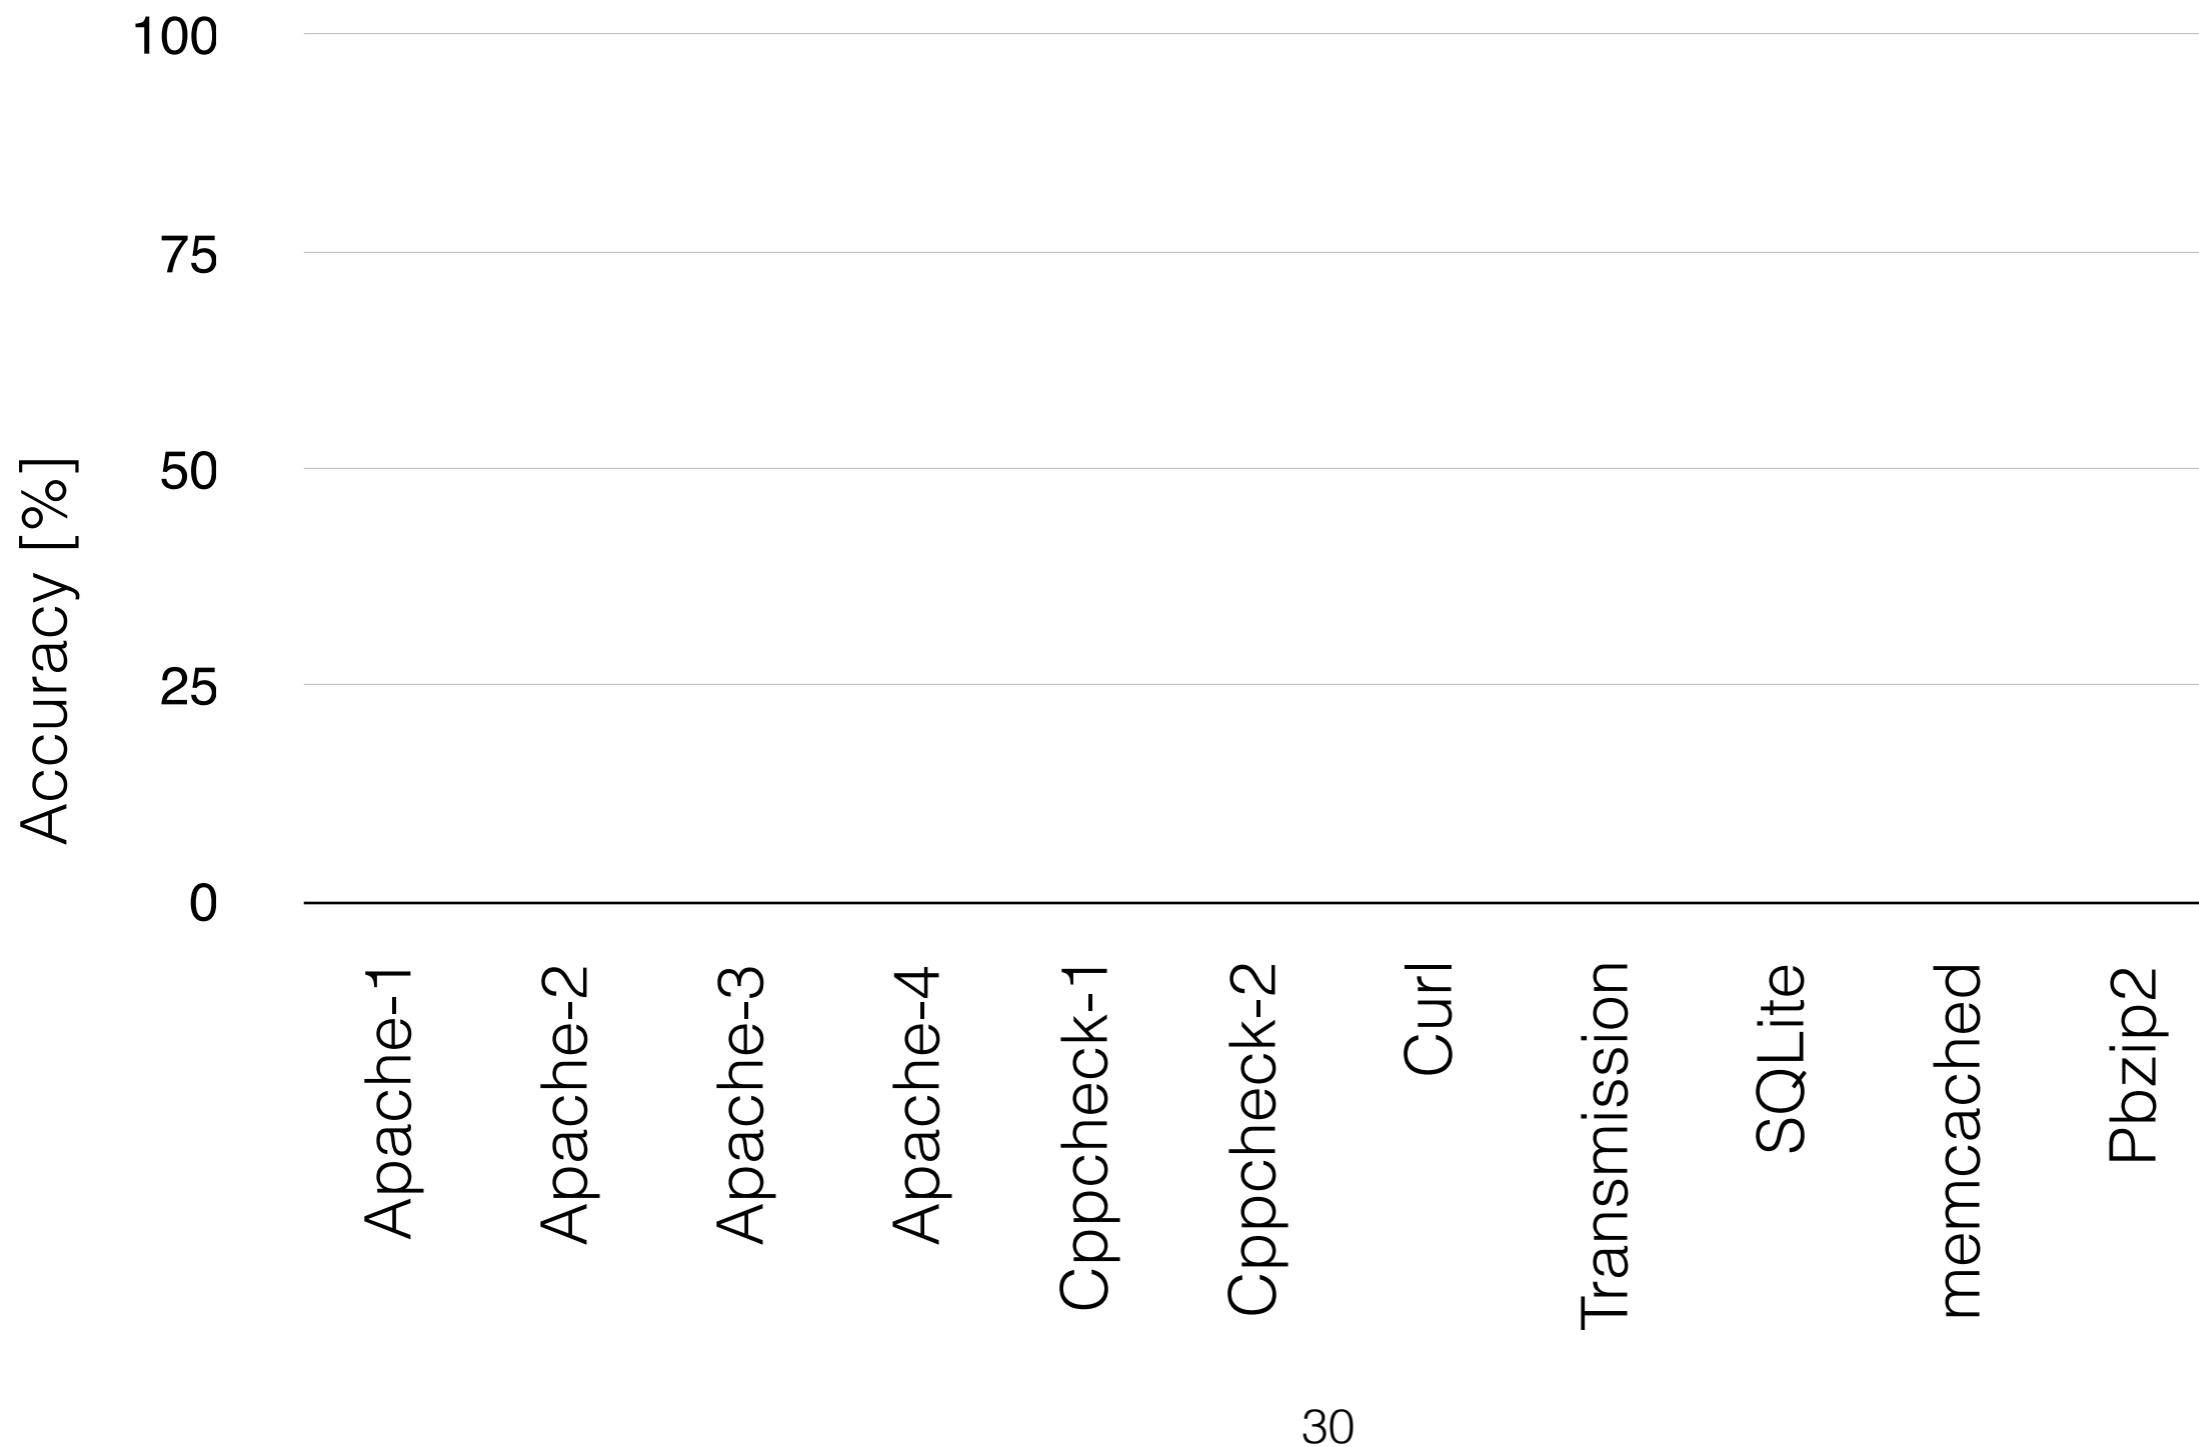
SQLite

memcached

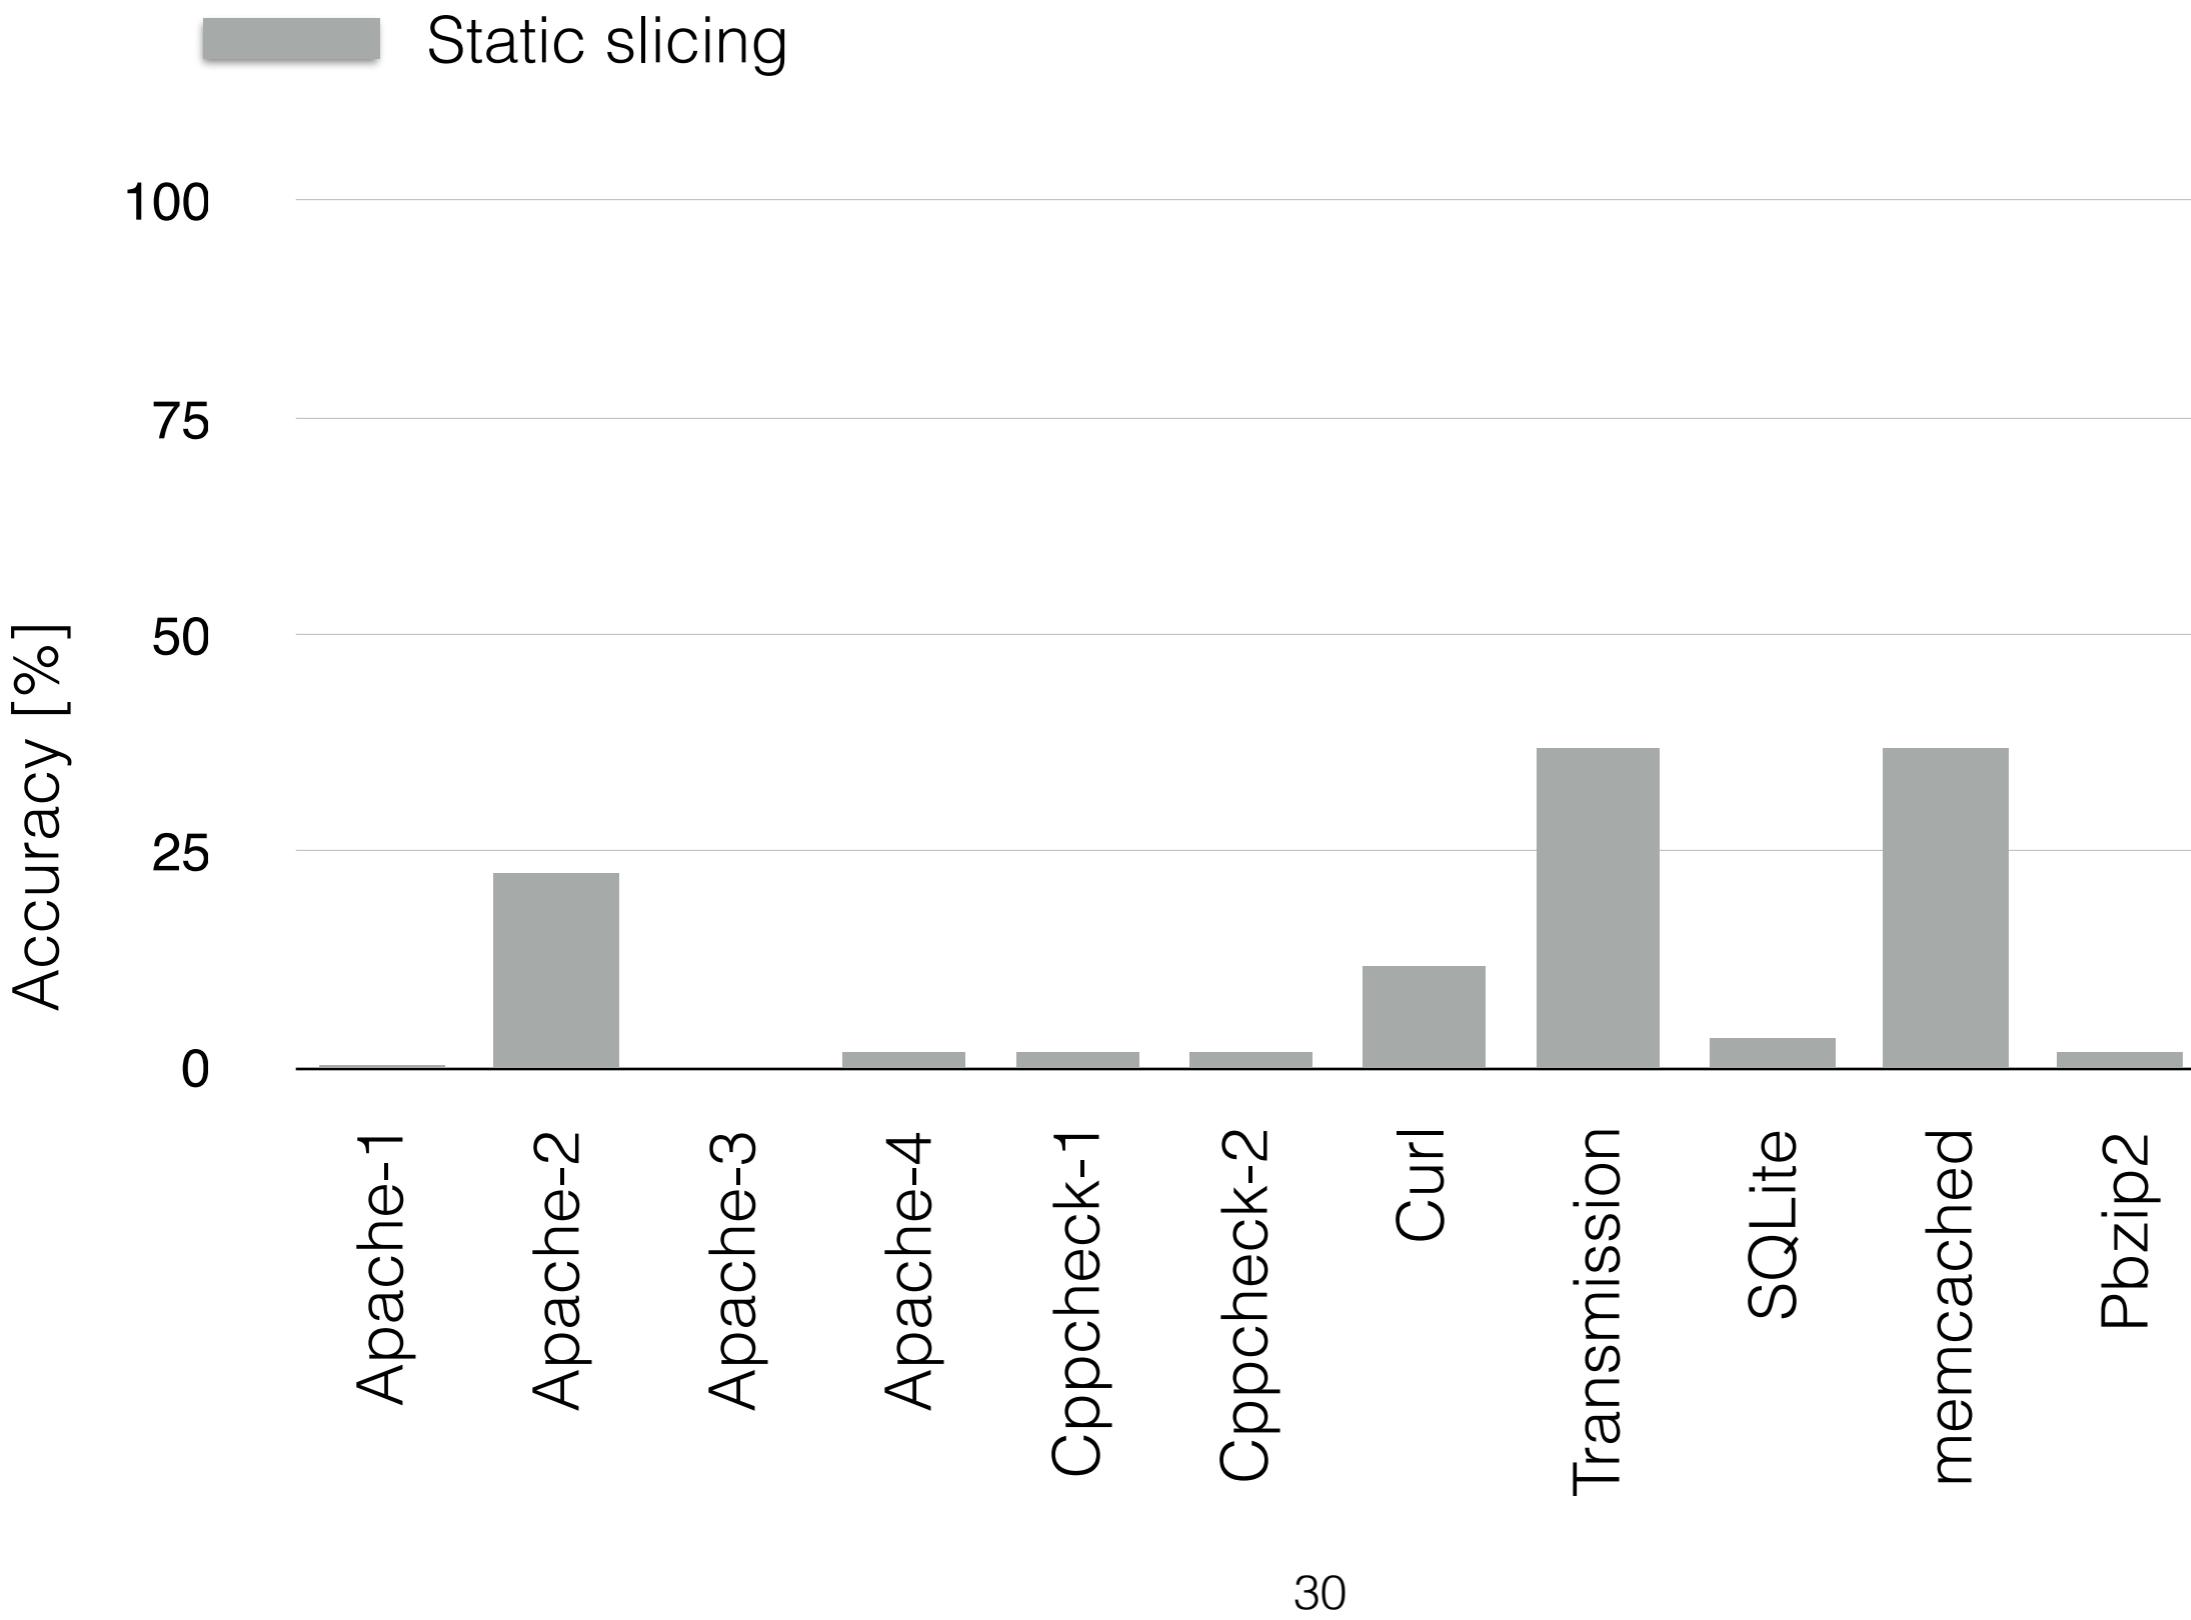
Pbzip2

# Accuracy

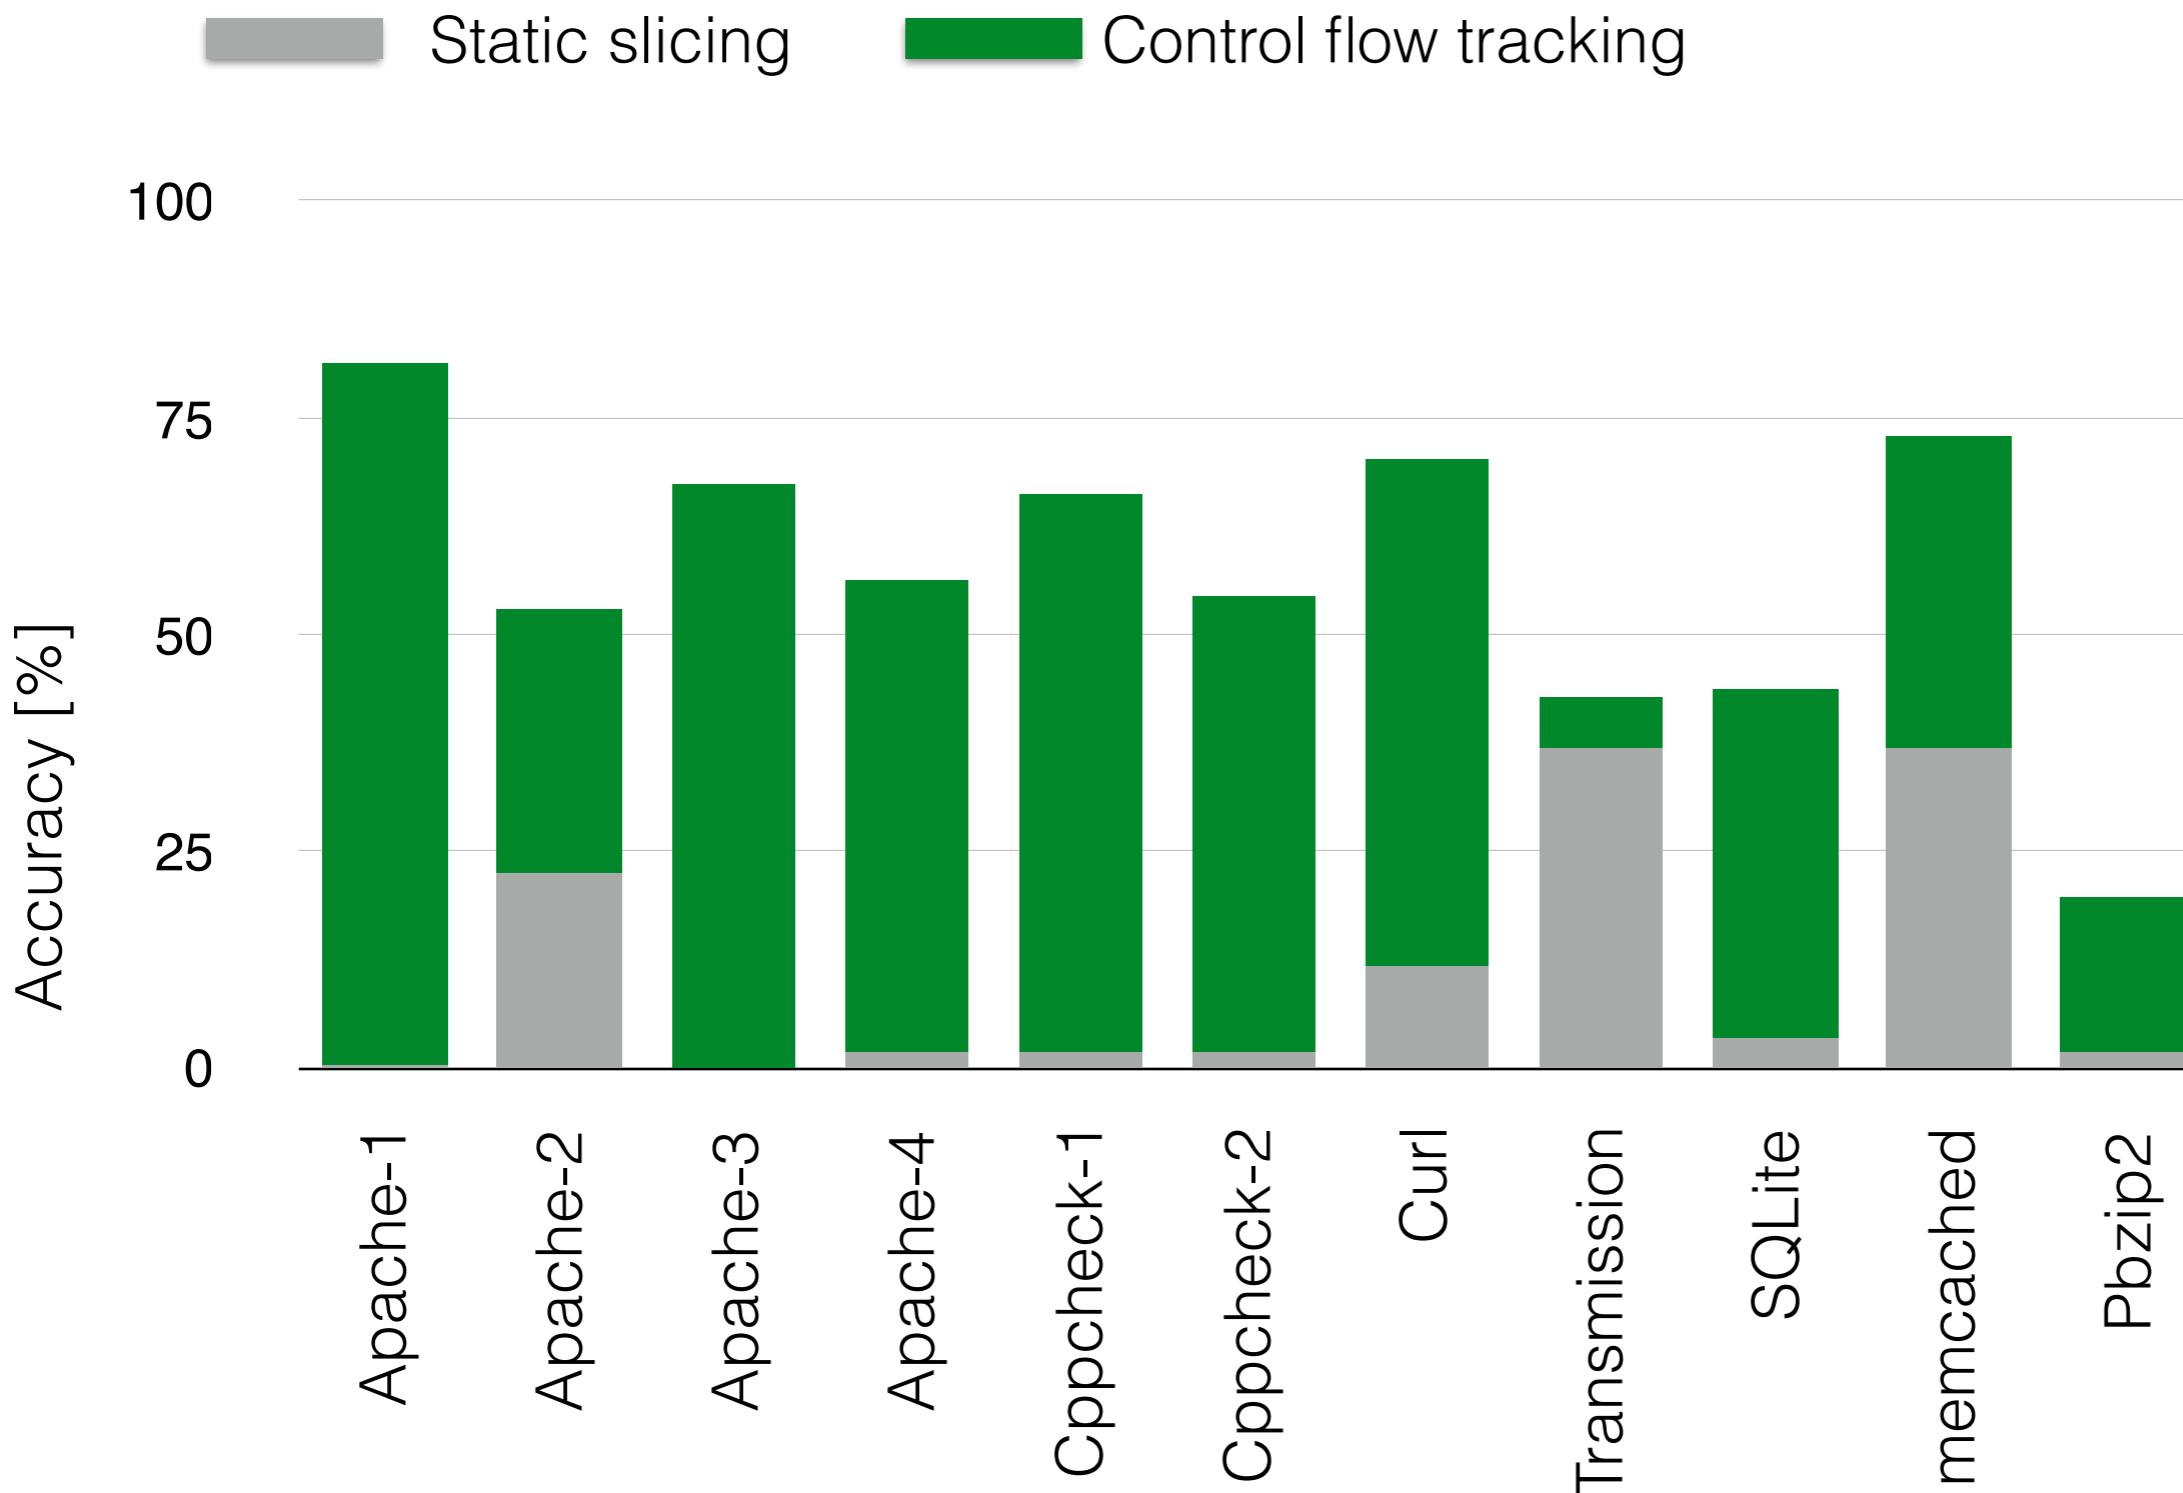
# Accuracy



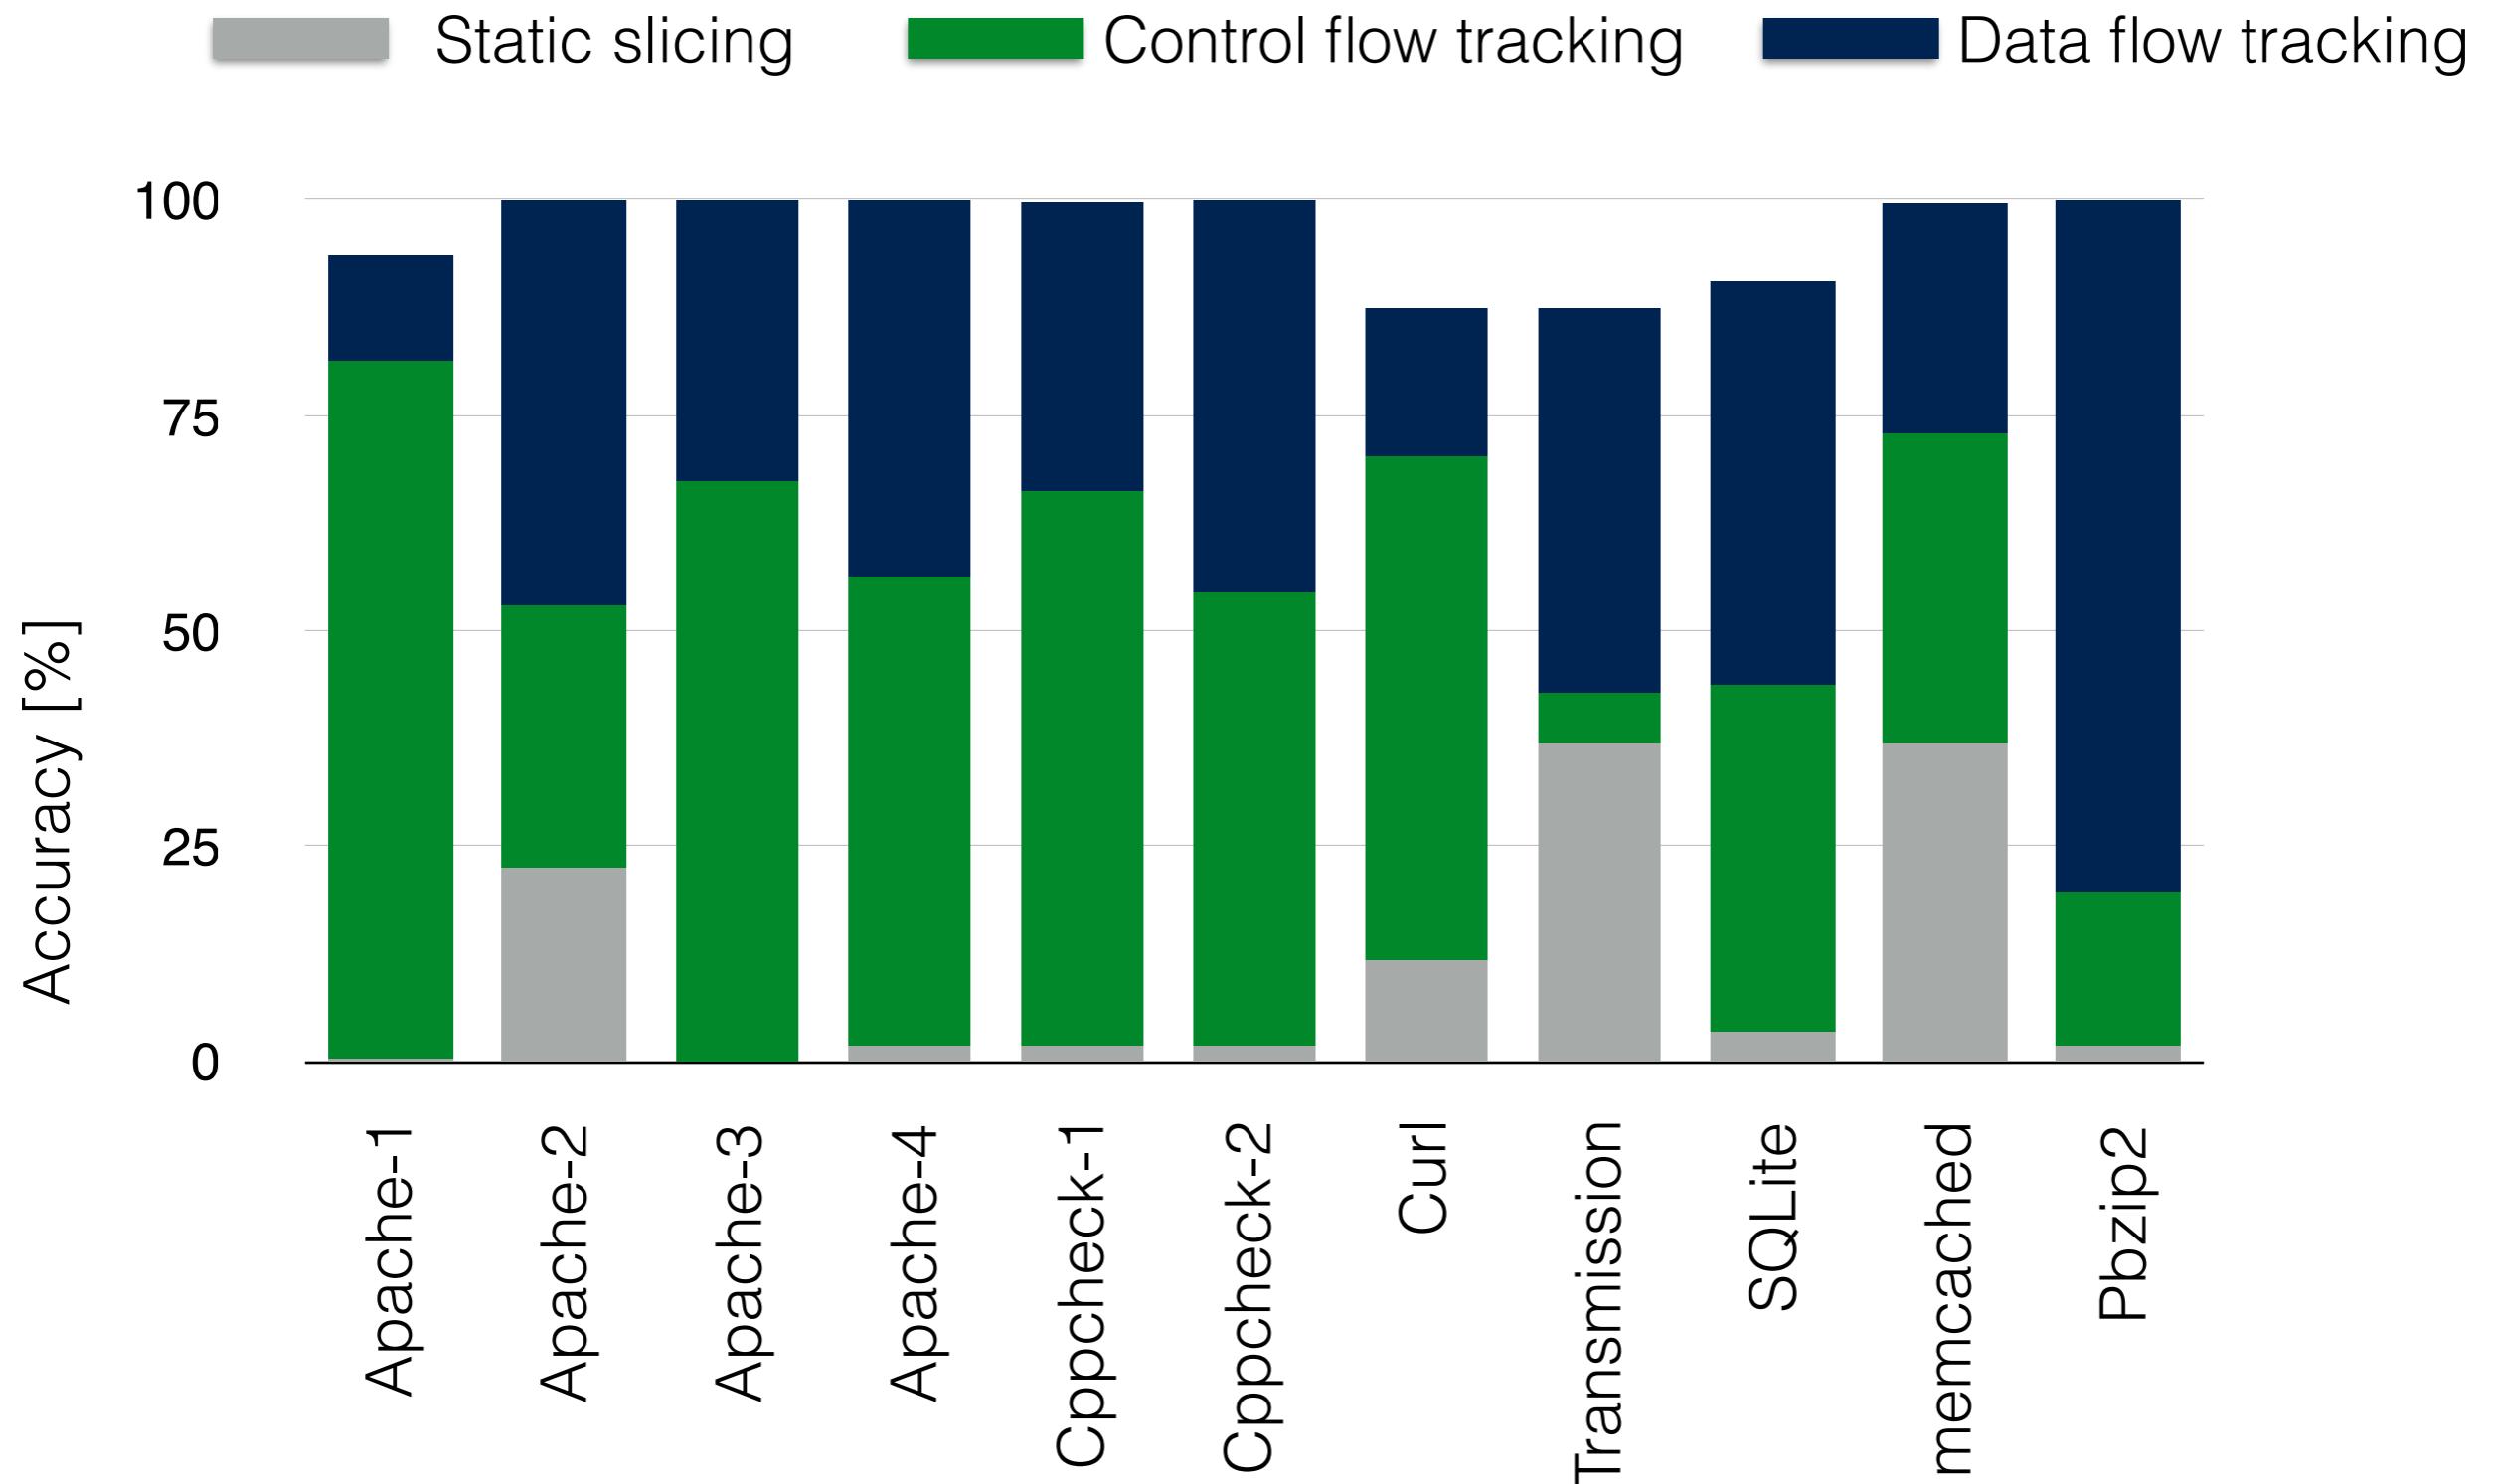
# Accuracy



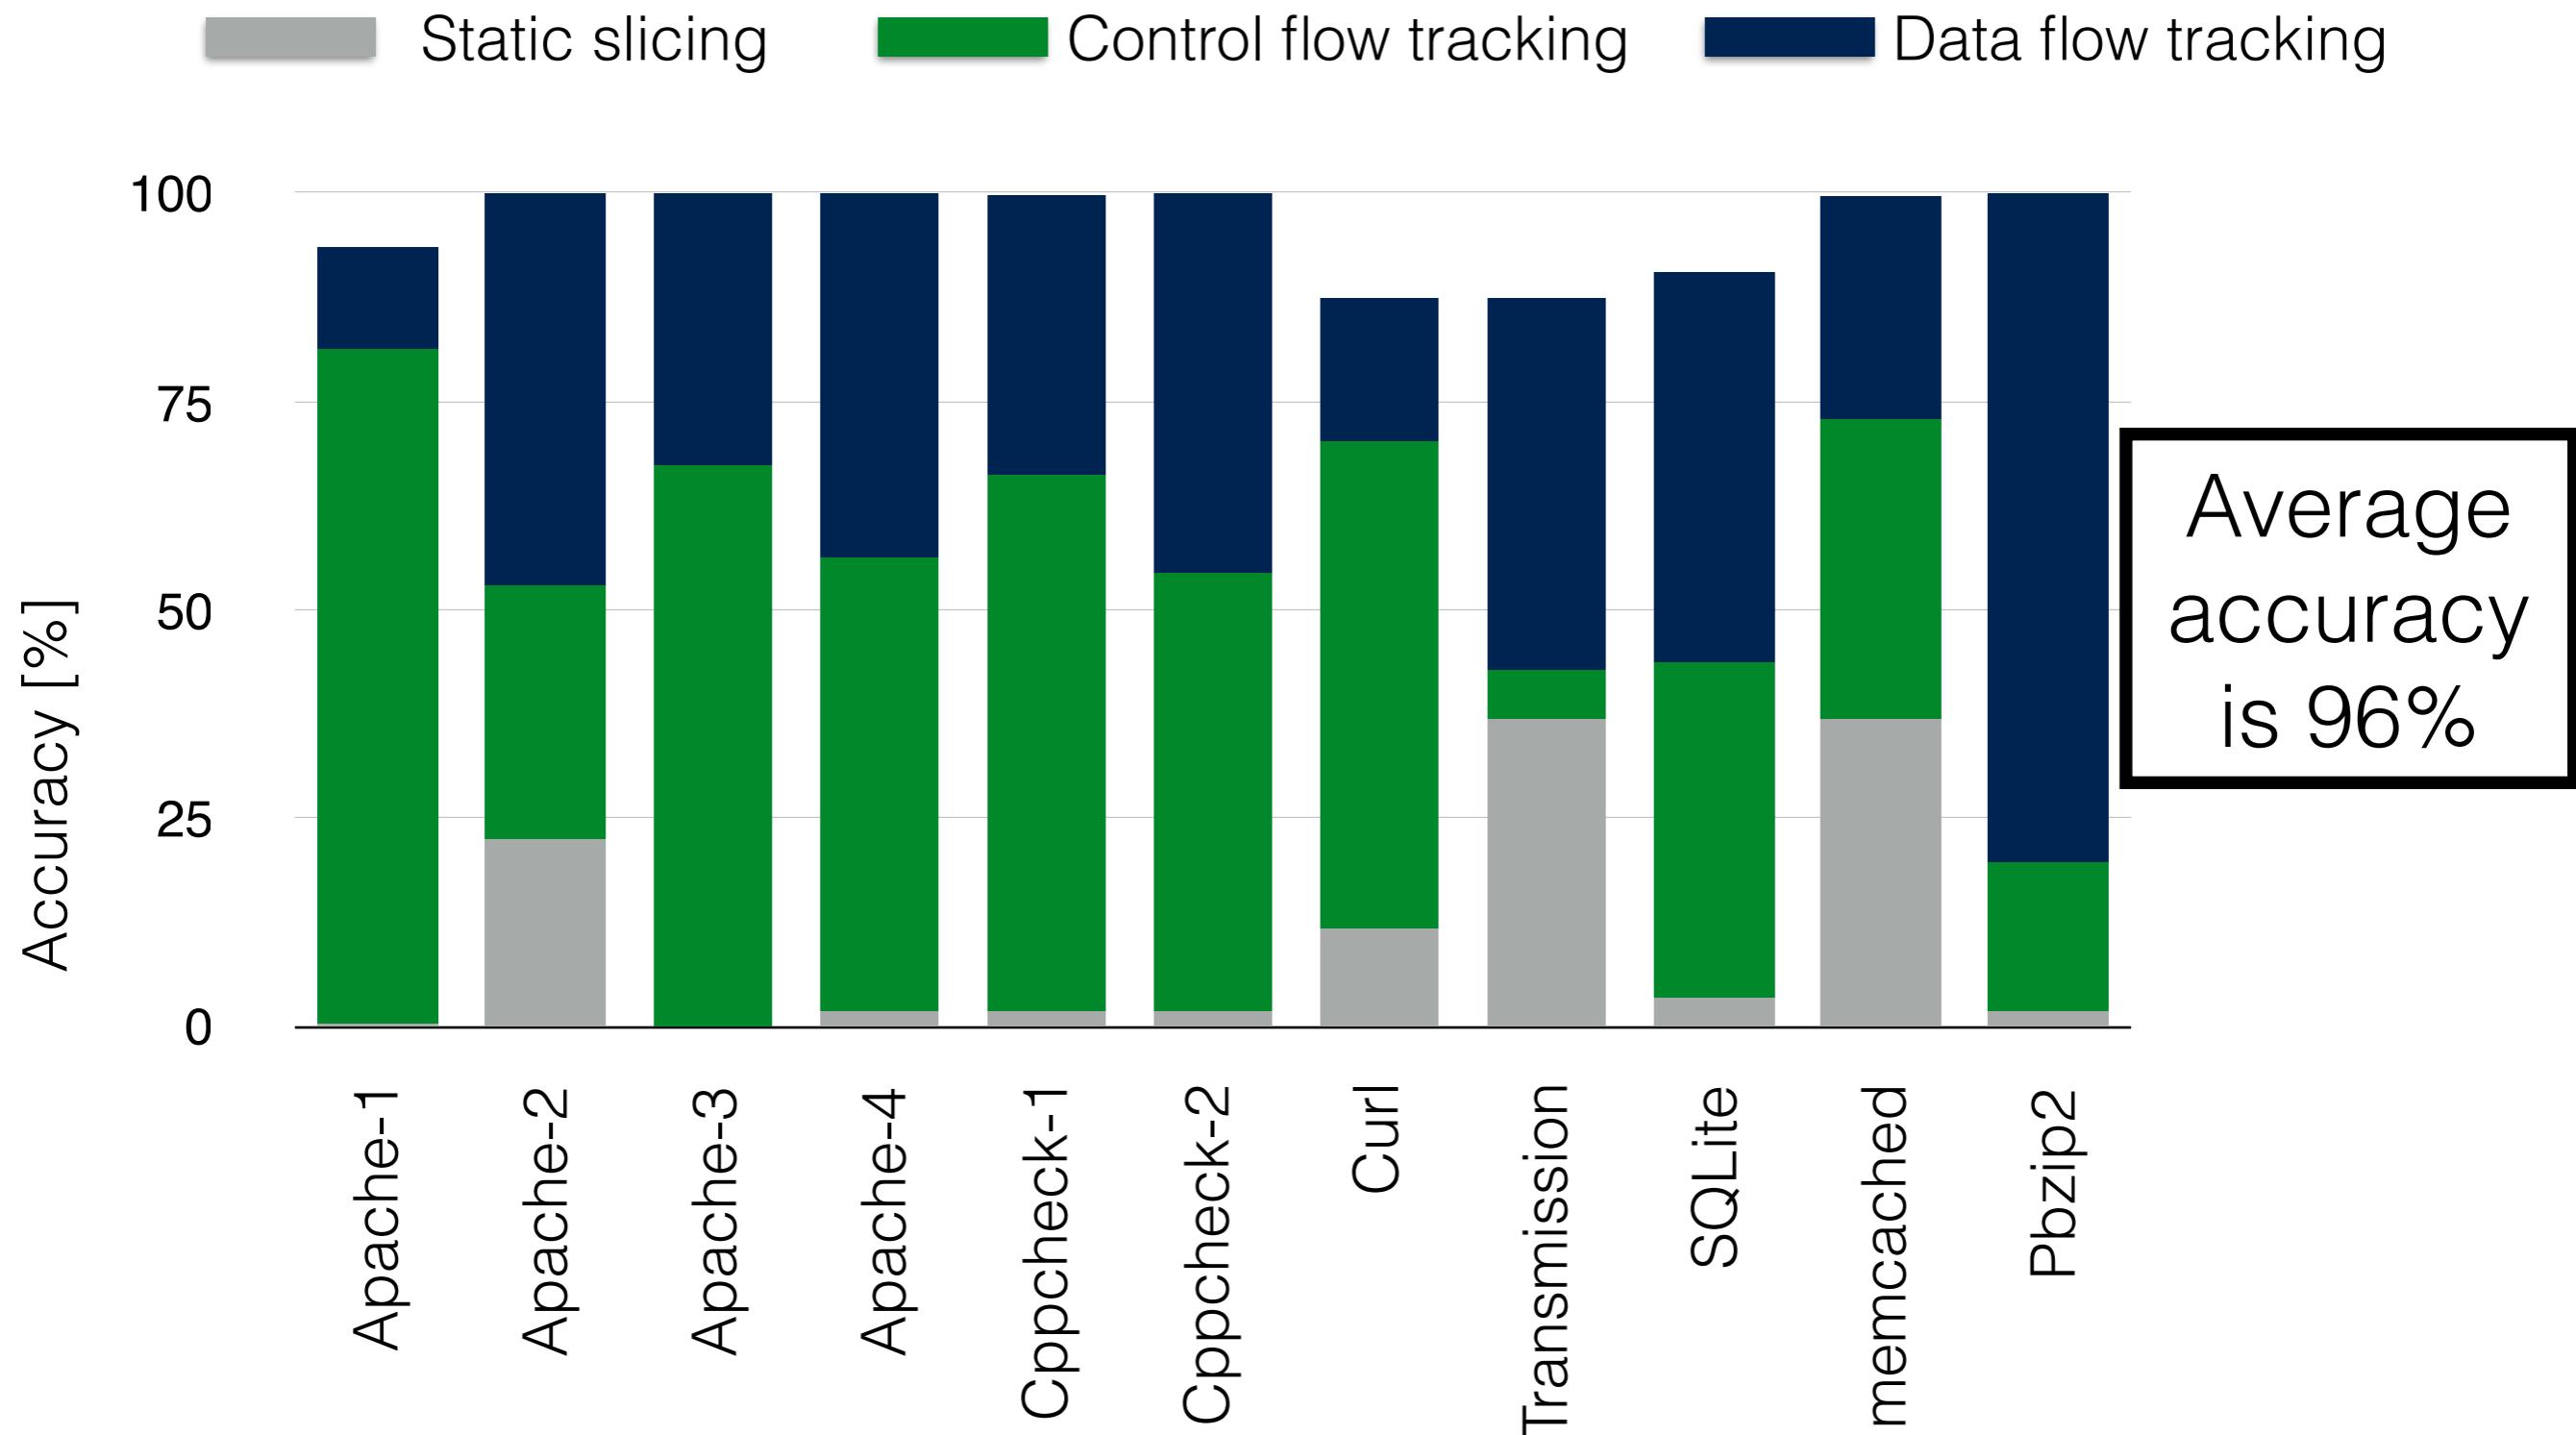
# Accuracy



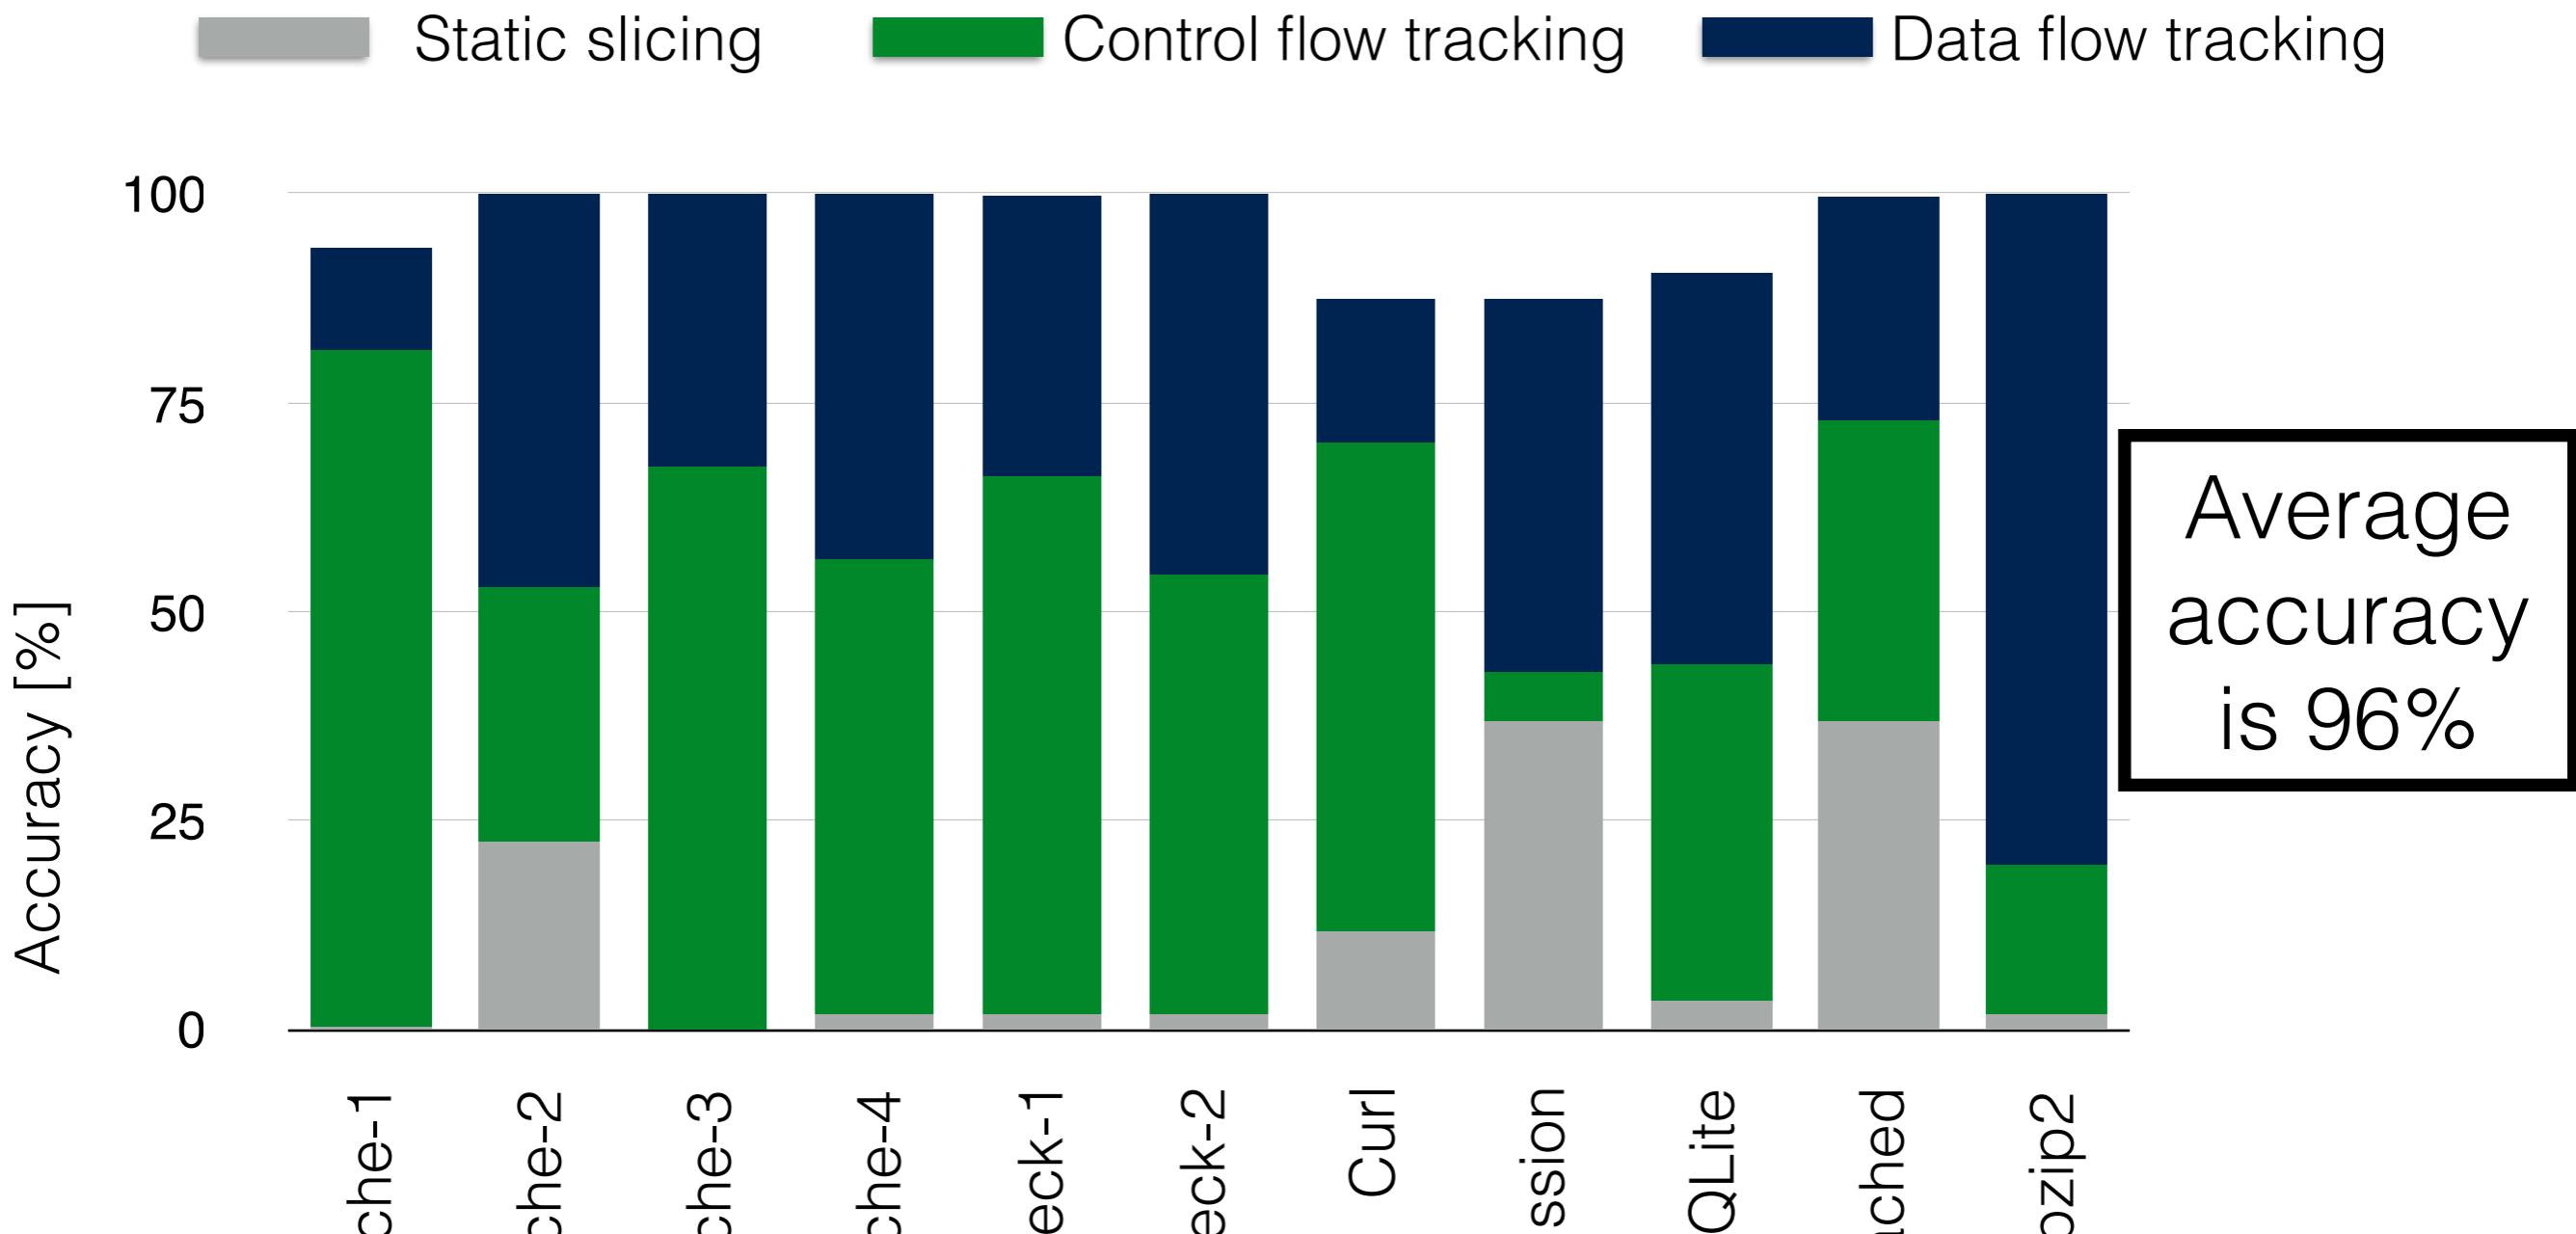
# Accuracy



# Accuracy



# Accuracy



Each technique is needed for accuracy

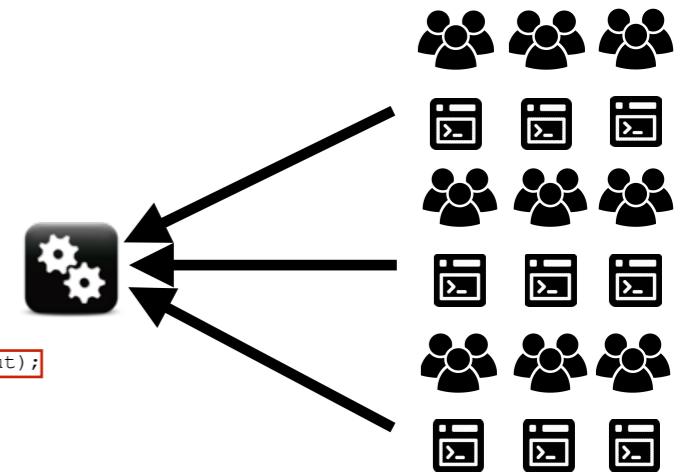
# Conclusion

- Failure sketching
  - *Combination of static and dynamic program analysis*
  - *Failure sketches are summaries explaining failure root causes*
  - *Accurate, efficient, improves developer productivity*



```
Time           Thread 1                                Thread 2
1 main() {          1
2   queue* f = init(size);  2
3   create_thread(cons, f); 3
4   ...
5   free(f->mut);
6   f->mut = NULL;        4   cons(queue* f) {
7   ...                      5
8 }                         6   ...
                           7   mutex_unlock(f->mut);
                           8 }
```

A code snippet showing two threads. Thread 1 contains a main function with several operations, including creating a thread and setting a pointer to NULL. Thread 2 contains a cons function that locks a mutex. A red box highlights the assignment of NULL to the pointer in Thread 1, and a red arrow points from this box to the mutex unlock call in Thread 2, illustrating a race condition.



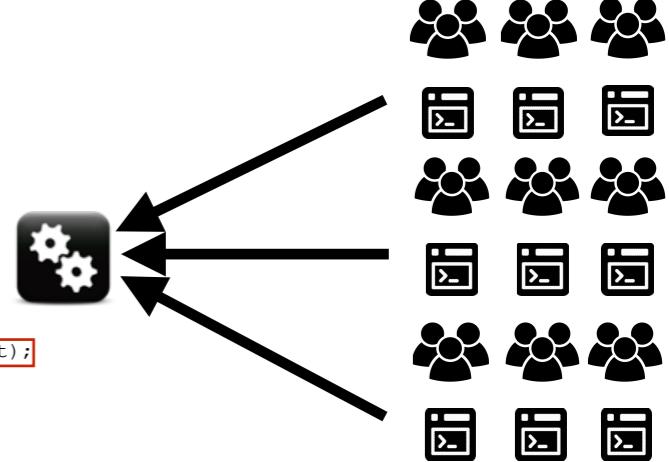
# Conclusion

- Failure sketching
  - *Combination of static and dynamic program analysis*
  - *Failure sketches are summaries explaining failure root causes*
  - *Accurate, efficient, improves developer productivity*

<http://dslab.epfl.ch/proj/gist>



```
Time           Thread 1                                Thread 2
1 main() {          1
2   queue* f = init(size); 2
3   create_thread(cons, f); 3
4   ...                ...
5   free(f->mut);      4   cons(queue* f) {
6   f->mut = NULL;     5   ...
7   ...                6   ...
8 }                  7   mutex_unlock(f->mut);
                     8 }
```



Baris Ben Cristiano Gilles George